

The Continuum Limit of a Space-Filling Layer of Cubes

A Tiling Identity, an Exact Lattice Kernel,
and the Closed-Form Error of the Discrete Layer

28 September 2026

Abstract

A heterogeneous elastic medium discretised into cubic voxels, each carrying a single-site T -matrix and coupled through a lattice Green's tensor, must reproduce the continuous medium as the voxels shrink. For a laterally uniform layer we establish when, how fast, and with what error it does, at normal and at oblique incidence. Five results. (i) Cubes tile the plane, so the lateral form factor of a cube vanishes on the reciprocal lattice; consequently the source-cell-averaged coupling between planes is *exactly* the continuum's plane-wave term, and within a plane the sum over all cells, the voxel's own included, is *exactly* the field of a uniform plate. This tiling identity gives the lattice kernel at normal incidence in closed form, with no lattice sum. (ii) The collocation closure of the single site carries the same self term the kernel removes, so the self term cancels and the discrete system becomes a collocation of the continuum equation with exact cell integrals. (iii) That collocation converges at exactly second order in the voxel size. (iv) Its leading error is closed-form for every discretisation — a midpoint-rule error of the internal field across each voxel, $+(kd)^2/8$ in reflection and $-(kd)^2/24$ in transmission — and the opposite signs show that no scalar correction of a uniform-field single site can remove it. (v) A voxel that carries the first moment of its field, coupled through exact cell-to-cell integrals, removes it and converges at fourth order, with leading error $-(kd)^4/720$ in reflection and transmission alike: one plane of such voxels is 6500 times more accurate than one plane of uniform ones, and the scheme stays fourth order at oblique incidence, in the converted S wave as well as the reflected P. Every result is derived symbolically in Mathematica and checked against an independent numerical implementation.

1 Introduction

Replacing a heterogeneous body by a lattice of small scatterers that interact through the background Green's function is the multiple-scattering description of Foldy and Lax [1, 2]. In electromagnetics its discretised form is the discrete dipole approximation (DDA) [3, 4], whose accuracy has been studied in detail: each dipole's polarisability must account for the local field of the surrounding lattice, which the lattice dispersion relation fixes so that an infinite lattice propagates waves as the continuum does to a given order in kd [5]; the error of the method converges at second order in the dipole spacing when shape errors are absent [6]; and integrating

the Green's tensor over each cell rather than sampling it at the centre improves accuracy, most of all for strong contrasts [7]. The elastic problem adds a structure the scalar and electromagnetic ones lack. A cubic voxel responds to the incident field with a force (through its density contrast) and a moment (through its stiffness contrast), so the state at a voxel has nine components — displacement and strain — and the single site's self-interaction is the elastic depolarisation of a cube, the discrete counterpart of Eshelby's inclusion tensor [8].

This paper settles, for a space-filling layer of cubic voxels, three questions that are usually answered only empirically: whether the discrete Foldy–Lax system converges to the continuous medium, at what order, and with what leading error. The answers are exact rather than measured. The coupling between voxels, averaged over the source cell, is *exactly* the continuum's cell integral at normal incidence, because cubes tile the plane (§5). The collocation closure of the single site carries the same self term, which therefore cancels, so the discrete system is a collocation of the continuum equation (§6). It converges at second order with a leading error known in closed form for every discretisation, whose opposite signs in reflection and transmission show that no correction of a uniform-field single site can remove it (§7). A voxel that carries the first moment of its internal field removes it, and converges at fourth order at normal and oblique incidence (§8); in its spectral form the lattice sum needs no Ewald splitting. Every result is derived symbolically in Mathematica and checked against an independent Python implementation or an exact solution.

2 Setting

The discrete medium. Space is tiled by cubes of side d (half-side $h = d/2$). Cube i carries a contrast $\Delta\lambda, \Delta\mu, \Delta\rho$ from an isotropic background with Lamé parameters λ, μ , density ρ , P and S speeds α, β and P-wave modulus $M_P = \lambda + 2\mu$. Time dependence is $e^{-i\omega t}$; $k_P = \omega/\alpha$, $k_S = \omega/\beta$. Coordinates are ordered (z, x, y) with z down. The field at a voxel is the nine-component state

$$\psi = (u_z, u_x, u_y, \varepsilon_{zz}, \varepsilon_{xx}, \varepsilon_{yy}, 2\varepsilon_{xy}, 2\varepsilon_{zy}, 2\varepsilon_{zx}), \quad (1)$$

displacement followed by engineering strain. A voxel responds with a force and a moment (stress dipole); the 9×9 single-site T -matrix maps the exciting state to them, and the 9×9 lattice Green's tensor $\mathbf{K}(i, j)$ maps a source voxel's force and moment to the state at voxel i . The Foldy–Lax system is

$$\psi_i = \psi_i^0 + \sum_{j \neq i} \mathbf{K}(i, j) \mathbf{T}_j \psi_j, \quad (2)$$

with ψ^0 the incident state. Strain rows of $\mathbf{K}$ are receiver derivatives of the point Green's tensor $\mathbf{G}$; moment columns are derivatives with respect to the receiver coordinate, symmetrised with a factor $\frac{1}{2}$ on the shear components. In that convention the contrast operator is

$$\Delta = \text{blockdiag}(\omega^2 \Delta\rho \mathbf{I}_3, \Delta\mathcal{C}), \quad (3)$$

where $\Delta\mathcal{C}$ is the 6×6 contrast stiffness with $\Delta\lambda + 2\Delta\mu$ and $\Delta\lambda$ in the normal block and $2\Delta\mu$ (not $\Delta\mu$) on the shear diagonal.

The laterally uniform layer. A layer of thickness D is filled with n planes of voxels, $d = D/n$, and illuminated at normal incidence. The field is then laterally uniform, only the zero-lateral-wavenumber (specular) component of the lattice coupling acts, and Eq. (2) becomes exactly a chain of n planes coupled by the specular lattice sums

$$\mathbf{S}(m) = \sum_{\mathbf{R}} \mathbf{G}(\mathbf{R} + m d \hat{\mathbf{z}}), \quad (4)$$

the sum running over the square lattice $\mathbf{R}$ of pitch d , m the plane separation in pitches. The question is whether, and how fast, this chain converges to the continuous layer as $n \rightarrow \infty$ at fixed D .

Implementations. Every result is derived symbolically in Mathematica and compared with an independent numerical implementation or an exact solution. “The package” below is the Python implementation of the voxel scheme (`cubic_scattering`): its cube T -matrix, its lattice kernels and its stratified propagator, which serve both as the object under test and, where validated, as a second implementation.

3 The continuous layer and its thin-layer expansion

At normal incidence a plane force drives a one-dimensional problem per wave type, with modulus M (M_p for P, μ for S), wavenumber k and plane Green’s function $g(z) = \frac{i}{2Mk} e^{ik|z|}$. The layer’s reflection and transmission follow from continuity of displacement and traction at its faces; they agree with the stratified propagator used as the reference in this work to 2.5×10^{-8} , P and S (`ContinuumLimit_Reference.wl`).

For a layer of thickness D with modulus M_1 and density ρ_1 in a background M_0, ρ_0 , the reflection coefficient expands as

$$R = \frac{i\omega D}{2} \sqrt{\frac{M_0}{\rho_0}} \left[\frac{\rho_1 - \rho_0}{M_0} - \rho_0 \left(\frac{1}{M_1} - \frac{1}{M_0} \right) \right] + O(D^2), \quad (5)$$

exact in the contrast (`ContinuumLimit_ThinLayer.wl`): to first order a thin layer is an interface across which the traction jumps by the excess inertia and the displacement by the excess compliance. At oblique incidence the same statement holds with the first-order system $\partial_z \mathbf{b} = \mathbf{A} \mathbf{b}$ for the P–SV state $\mathbf{b} = (u_x, u_z, t_{xz}, t_{zz})$: the layer acts to $O(D)$ as the jump $\mathbf{I} + (\mathbf{A}_1 - \mathbf{A}_0)D$, which reproduces the layer’s $O(D)$ reflection coefficients R_{pp}, R_{ps} to 10^{-12} ; the series truncated at $O(D^k)$ converges at exactly order k ; and a soft thin layer with shear modulus $\mu_1 = D/\eta$ at fixed compliance η tends to the linear-slip interface as $D \rightarrow 0$. These expansions are the target a single plane of voxels must reproduce order by order.

4 The specular sums of the point Green's tensor

By Poisson summation,

$$\mathbf{S}(m) = \frac{1}{d^2} \sum_{\mathbf{g}} \hat{\mathbf{G}}(\mathbf{g}; md), \quad \mathbf{g} = \frac{2\pi}{d}(p, q), \quad (6)$$

with $\hat{\mathbf{G}}$ the two-dimensional transform of the Green's tensor between planes. The $\mathbf{g} = 0$ term is the continuum's plane-wave term; the others are evanescent, decaying as $e^{-2\pi m \sqrt{p^2 + q^2}}$ independently of d . In the strain-from-moment block they carry two powers of $|\mathbf{g}| \propto 1/d$, so their contribution per voxel volume does *not* vanish as $d \rightarrow 0$: at $m = 1$ it is eleven times the continuum term (`ContinuumLimit_SpecularSums.w1`, matched to the package's Ewald lattice sum to 5×10^{-13}). This surviving term is static. Summed over all planes in slab order, the static lattice sum of the point kernel is exactly the depolarisation of a uniform plate — $-1/M_P$ on $\epsilon_{zz} \leftarrow M_{zz}$ and $-1/(2\mu)$ on the normal shears, per unit volume — plus a cubic-symmetric remainder (`ContinuumLimit_StaticLocal.w1`, to 2×10^{-10}). The point kernel therefore reproduces the continuum only if that remainder is cancelled by the voxel's own self term.

5 The tiling identity

Replace the point coupling by the *source-cell average* $\langle \mathbf{G} \rangle(\mathbf{r}) = d^{-3} \int_{\text{cube}} \mathbf{G}(\mathbf{r} - \mathbf{y}) d\mathbf{y}$ — the field at a voxel centre of a uniform polarisation of the source voxel. Each Poisson order in Eq. (6) is then multiplied by the cube's form factor

$$F(\mathbf{g}) = \text{sinc}(g_x h) \text{sinc}(g_y h), \quad (7)$$

which vanishes at every $\mathbf{g} \neq 0$ of the reciprocal lattice since $\text{sinc}(p\pi) = 0$ for integer $p \neq 0$. This is the Fourier form of the fact that cubes tile the plane, $\sum_{\mathbf{R}} \chi(\mathbf{x} - \mathbf{R}) = 1$ for the cube's indicator χ . Two exact statements follow.

Between planes. For $m \neq 0$ every evanescent order is removed and

$$\mathbf{S}_{\text{avg}}(m) = \frac{1}{d^2} \sum_{W=P,S} \hat{\mathbf{G}}_W(0; md) \text{sinc}(k_W h), \quad (8)$$

the continuum's plane-wave term times the vertical form factor. The package's cell-averaged kernel agrees in all 81 entries to 2.8×10^{-16} (`ContinuumLimit_AveragedSums.w1`).

Within a plane. The sum over *all* cells of the plane, the voxel's own included, is the field at the mid-plane of a uniform plate of thickness d :

$$d^3 \mathbf{S}_{\text{avg}}(0) + \mathbf{S}_{\text{self}} = d^3 \mathbf{P}_0, \quad \mathbf{S}_{\text{self}} = \int_{\text{cube}} \mathbf{G}(-\mathbf{y}) d\mathbf{y}, \quad (9)$$

where $\mathbf{P}_0 = d^{-3} \int_{-h}^h \mathbf{G}_{\text{plate}}(-z') dz'$ is elementary: in the one-dimensional transform each entry is $c_0 + A/(Mk_z^2 - \rho\omega^2)$, so it averages to $c_0/d + A(e^{ikh} - 1)/(dMk^2)$. Its only nonzero entries are $u \leftarrow f$ on the diagonal and the three normal strain-from-moment entries, whose local weights $-1/M_p$ and $-1/(2\mu)$ are the plate depolarisation of §4. Statically the identity holds for all 21 independent components of the strain-from-moment tensor to 4×10^{-9} , with every cell integral evaluated analytically (below).

Consequence: an exact kernel with no lattice sum. At normal incidence the same-plane block of the cell-averaged Foldy–Lax kernel is

$$\mathbf{S}_{\text{avg}}(0) = \mathbf{P}_0 - \mathbf{S}_{\text{self}}/d^3, \quad (10)$$

with no Ewald splitting, near shell or asymptotic tail. The route it replaces — Ewald summation with a near shell averaged by quadrature and an $O(d^2)$ Taylor tail beyond it — carries a static strain-from-moment error of 9.4×10^{-4} at its default settings, falling to 2.6×10^{-5} only when both the shell radius and the quadrature order are raised; Eq. (10) is now the package’s kernel at normal incidence.

Cell integrals without quadrature. Write the static (Kelvin) tensor as $G_{bc} = (c_A + c_B) \delta_{bc}/r - c_B \partial_b \partial_c r$ with $c_B = (1/\mu - 1/M_p)/8\pi$ and $c_A = 1/(4\pi\mu) - c_B$. Every cell integral of its derivatives is a cell integral of derivatives of $1/r$ and r . One application of the divergence theorem moves it onto faces $x_p = c_p \pm h$, which never contain the origin, and there the two-dimensional antiderivative in the face coordinates is continuous, so the face integral is a sum over the face’s corners. The distributional term at $r = 0$ is carried by the divergence theorem itself. (A triple antiderivative evaluated at the cube’s corners is *not* usable for repeated derivatives: its arctangent terms jump across the coordinate planes through the origin, and differentiating twice there discards exactly that term.) The self term so obtained reproduces the classical closed form of the cube’s depolarisation to machine precision.

6 The single site

The collocation closure takes the internal field of a voxel as uniform. The voxel’s own field at its centre is then $\mathbf{S}_{\text{self}} \Delta e$ for internal state e , so $e = \psi + \mathbf{S}_{\text{self}} \Delta e$ and

$$\mathbf{T} = d^3 \Delta (\mathbf{I} - \mathbf{S}_{\text{self}} \Delta)^{-1}. \quad (11)$$

With the kernel (10) the self term cancels. Substituting $\psi = (\mathbf{I} - \mathbf{S}_{\text{self}} \Delta)e$ into Eq. (2),

$$e_i = \psi_i^0 + d^3 \sum_j \mathbf{K}_{\text{all}}(i, j) \Delta e_j, \quad \mathbf{K}_{\text{all}}(i, i) = \mathbf{P}_0, \quad \mathbf{K}_{\text{all}}(i, j) = \mathbf{S}_{\text{avg}}(i - j), \quad (12)$$

an identity checked on the matrices to 4×10^{-16} . Every block of $\mathbf{K}_{\text{all}}$ is an *exact* cell integral of the one-dimensional continuum kernel (Eqs. (8), (9)), so Eq. (12) is precisely the collocation of

the continuum Lippmann–Schwinger equation with a strain piecewise constant over voxels. It does not contain $\mathbf{S}_{\text{self}}$ at all.

The package’s cube T -matrix is this closure statically: both blocks agree with Eq. (11) to 2×10^{-5} at $k_s h = 0.01$, the difference falling as $(kh)^2$ (`ContinuumLimit_CubeT.wl`). It departs dynamically by about $0.25(kh)^2$ through its far-field form-factor corrections. Its self integrals are exact in the static and radiation parts but omit the real even-power terms of the Green’s tensor’s expansion ($k^2 r, r^3, \dots$), an error of $0.4\text{--}0.6(kh)^2$ against the analytic self term, whose series agrees with the closed-form elastodynamic tensor to 10^{-14} . By Eq. (12) that error cannot enter the closure scheme.

7 Convergence, and the closed-form error

Second order. For a layer $D = 2\text{ m}$ with the contrast $\Delta\lambda = 2\text{ GPa}$, $\Delta\mu = 1\text{ GPa}$, $\Delta\rho = 100\text{ kg m}^{-3}$ in a background $\alpha = 5000$, $\beta = 3000\text{ m s}^{-1}$, $\rho = 2500\text{ kg m}^{-3}$, illuminated by a P plane force, the reflected displacement of the chain converges to the continuous layer’s at order 2.000 for $n = 4 \rightarrow 16$ at $\omega = 60, 300$ and 600 rad s^{-1} , with the collocation (12) and with the package T -matrix alike (`ContinuumLimit_Chain.wl`). The error depends on $k_s h$ alone.

The leading error in closed form. In Eq. (12) the kernel is integrated exactly; the only approximation is the internal field’s variation across a voxel. At first order in the contrast the discrete scattered field is therefore a midpoint rule applied to the continuum’s Born integral, with the kernel’s cell integral exact. For reflection the integrand carries the two-way phase e^{2ikz} ; for transmission it is flat. With k the background wavenumber of the wave concerned (k_p here), R the reflection coefficient and T^{sc} the scattered part of the transmitted field, for *every* n

$$\frac{R_{\text{disc}}}{R} = \frac{\text{sinc}(kh)}{\text{sinc}(2kh)} = 1 + \frac{(kd)^2}{8} + O(d^4), \quad \frac{T_{\text{disc}}^{\text{sc}}}{T^{\text{sc}}} = \text{sinc}(kh) = 1 - \frac{(kd)^2}{24} + O(d^4), \quad (13)$$

confirmed for $n = 1$ to 16 to 10^{-4} of the error itself (`ContinuumLimit_ErrorConstant.wl`). At the full contrast the reflection constant is unchanged (0.124 against $1/8$); the transmission constant grows to -0.092 .

Why a uniform-field single site stops at second order. The two constants have opposite signs, so no scalar correction of $\mathbf{T}$ — no single $O((kh)^2)$ factor — can cancel both. The package T -matrix illustrates it: its form-factor corrections reduce the reflection error by a factor 4.45 and *increase* the transmission error by 1.7, at every n . The missing term in Eq. (13) is the first moment of the internal field over each voxel, which radiates through the derivative of the kernel. A fourth-order scheme must carry it: a single site with a gradient channel, and a kernel whose cell integrals include the kernel’s derivative.

8 Fourth order: the first-moment voxel

Let each voxel carry the mean *and* the first moment of its internal field. On voxel j , with local coordinate $s \in [-h, h]$ and Legendre functions $\phi_0 = 1$, $\phi_1 = s/h$, the field is $w = \sum_b c_{b,j} \phi_b$, $w = (u, \varepsilon)$ the displacement and strain of the P channel, and the polarisation is $q = (\omega^2 \Delta \rho u, \Delta M \varepsilon)$ with $\Delta M = \Delta \lambda + 2\Delta \mu$. Testing the continuum equation

$$w(z) = w^0(z) + \int \mathbf{K}(z - z') q(z') dz', \quad \mathbf{K} = \begin{pmatrix} g & g' \\ g' & g'' \end{pmatrix}, \quad g'' = -k^2 g - \delta(z)/M_P, \quad (14)$$

with the same two functions on every voxel gives

$$\left(\int_i \phi_a^2 \right) c_{a,i} = \int_i \phi_a w^0 + \sum_{j,b} \left[\iint_j \phi_a(z) \mathbf{K}(z - z') \phi_b(z') dz dz' \right] \Delta_1 c_{b,j}, \quad (15)$$

$\Delta_1 = \text{diag}(\omega^2 \Delta \rho, \Delta M)$. Between voxels the double integral is separable, a product of Legendre moments of $e^{\pm i k z}$; within a voxel it is split at $z' = z$ and the local term contributes $-\int_i \phi_a \phi_b / M_P$. Every entry is closed-form. Keeping only ϕ_0 recovers a second-order scheme; keeping ϕ_0 and ϕ_1 converges at fourth order in both reflection and transmission (`ContinuumLimit_FourthOrder.wl`):

n	collocation	mean only	mean and first moment
1	$1.8 \times 10^{-3} / 1.3 \times 10^{-3}$	$1.2 \times 10^{-3} / 1.8 \times 10^{-3}$	$2.8 \times 10^{-7} / 2.8 \times 10^{-7}$
4	$1.1 \times 10^{-4} / 8.3 \times 10^{-5}$	$7.5 \times 10^{-5} / 1.1 \times 10^{-4}$	$1.1 \times 10^{-9} / 1.1 \times 10^{-9}$
16	$7.0 \times 10^{-6} / 5.2 \times 10^{-6}$	$4.7 \times 10^{-6} / 6.9 \times 10^{-6}$	$4.2 \times 10^{-12} / 4.2 \times 10^{-12}$
order	2 / 2	2 / 2	4 / 4

(relative error of the scattered displacement, reflection / transmission, $\omega = 300 \text{ rad s}^{-1}$; the order is 4.00 at $\omega = 60$ and 600 rad s^{-1} too). A single plane of first-moment voxels represents the 2 m layer 6500 times more accurately than a single plane of uniform ones. The collocation column is the scheme of §6 and reproduces the chain of the full 9×9 system to all printed digits.

The leading error in closed form. Without scattering, the Galerkin solution is the L^2 projection Π of the incident field onto each voxel's basis, and the kernel is integrated over each voxel exactly; at first order in the contrast the error is therefore the kernel's cell integral applied to $\Pi w^0 - w^0$. The field varies as $e^{i k s}$ across a voxel and the kernel toward an outside observer as $e^{i q s}$, $q = +k$ for reflection and $q = -k$ for transmission, so for every scheme and every n

$$\frac{\text{discrete}}{\text{exact}} = \frac{\int_{-h}^h e^{i q s} \Pi[e^{i k s}] ds}{\int_{-h}^h e^{i(q+k)s} ds}, \quad (16)$$

with Π the centre value for collocation, the cell mean for the mean-only voxel and the projection onto $\{1, s/h\}$ for the first-moment voxel (`ContinuumLimit_FourthOrderError.wl`):

	reflection	transmission
collocation	$1 + (kd)^2/8 + O(d^4)$	$1 - (kd)^2/24 + O(d^4)$
mean only	$1 + (kd)^2/12 + O(d^4)$	$1 - (kd)^2/12 + O(d^4)$
first moment	$1 - (kd)^4/720 + O(d^6)$	$1 - (kd)^4/720 + O(d^6)$

The chain of Eq. (15) reproduces all six forms at $n = 1$ to 8 to 10^{-4} of the error, the remainder being the second order in the contrast. The first-moment voxel's leading error is the same in reflection and in transmission: its projection error is orthogonal to everything linear across the voxel, so the leading term no longer depends on the direction in which the field is observed.

In three dimensions. The three-dimensional first-moment voxel carries the nine-component state with Legendre weights $\{1, x/h, y/h, z/h\}$ — 36 unknowns — tested with the same functions and coupled by cell-to-cell double integrals of the Green's tensor. For a laterally uniform layer at normal incidence it reduces *exactly* to Eq. (15) (`ContinuumLimit_Reduction3D.w1`). The lateral first moments are never excited: the lattice and the incident field are symmetric under reflection about every voxel's centre planes, and $x/h, y/h$ are odd there. The mean and the z -moment carry weights that are constant laterally, so their lateral form factor $\text{sinc}(g_x h) \text{sinc}(g_y h)$ vanishes on the reciprocal lattice and the plane sum of the double integrals is the plate-to-plate integral in z alone. Two facts mark where genuinely three-dimensional structure begins. The lateral first moment's form factor, $F_1(g) = i(gh \cos gh - \sin gh)/(gh)^2$, does not vanish on the reciprocal lattice ($F_1 = i(-1)^p/p\pi$), so the gradient channels couple through the evanescent orders. And at oblique incidence the mean's own zeros shift, $\text{sinc}((2\pi p/d + k_x)h) = (-1)^p k_x h/p\pi + O(k_x^2 h^2)$, so the tiling identity holds only to $O(k_{\parallel} h)$.

Oblique incidence. The scheme is built spectrally. By Poisson summation the Bloch coupling between planes i and j at lateral wavenumber $\mathbf{k}_{\parallel}$ is

$$C_{ab}(i, j) = \frac{1}{d^2} \sum_{\mathbf{g}} L_a(\boldsymbol{\kappa}) L_b(-\boldsymbol{\kappa}) \iint \phi_a(z) \hat{\mathbf{G}}(\boldsymbol{\kappa}; z - z') \phi_b(z') dz dz', \quad \boldsymbol{\kappa} = \mathbf{k}_{\parallel} + \mathbf{g}, \quad (17)$$

with L_a the lateral Legendre factors. The whole-space tensor decays as $1/k_z^2$ and the nine-component state adds at most two powers of k_z , so each entry of $\hat{\mathbf{G}}(\boldsymbol{\kappa}; z)$ is a $\delta(z)$ of constant weight plus, per mode, $[i\alpha/2\gamma + i\beta \text{sgn}(z)/2] e^{i\gamma|z|}$, $\gamma^2 = k_W^2 - \kappa^2$: every z integral is closed-form. Double averaging makes the lattice sum converge without Ewald splitting, to three digits at $|p|, |q| \leq 8$. Against the exact layer, for a P wave incident at 20° and 40° (`ContinuumLimit_Oblique.w1`):

n	mean only		mean and first moments	
	R_{pp}	R_{ps}	R_{pp}	R_{ps}
1	9.3×10^{-4}	1.7×10^{-3}	7.2×10^{-8}	4.1×10^{-7}
2	2.3×10^{-4}	4.2×10^{-4}	4.5×10^{-9}	2.5×10^{-8}
8	1.5×10^{-5}	2.7×10^{-5}	1.7×10^{-11}	9.9×10^{-11}
order	2	2	4	4




Figure 1. Relative reflection error of the discrete layer against the number of voxel planes n ($D = 2$ m, $\omega = 300 \text{ rad s}^{-1}$, the contrast of §7). Filled: normal incidence, P reflection — collocation (§6), the mean-only voxel and the first-moment voxel of Eq. (15). Dashed, open markers: P incident at 20° , the converted S wave, mean-only (green) and first-moment (orange) voxels of Eq. (17). Grey: slopes 2 and 4.

(relative error, 20° ; at 40° the orders are again 2 and 4). The mean-only voxel is second order and the first-moment voxel fourth order in both the reflected P and the converted S wave: the breakdown of the tiling identity at $O(k_{\parallel}h)$ and the evanescent coupling of the gradient channels do not cost an order. At $\mathbf{k}_{\parallel} = 0$ the construction reproduces Eq. (15) to all printed digits.

The spectral construction is implemented twice. The second implementation (`scripts/crosscheck_first_moment_voxel.py`) shares only the definition of the scheme: it takes the depth dependence of the transformed kernel from an independent residue construction, obtains the local weight as a numerical large- k_z limit of the inverted Christoffel matrix, evaluates every double integral by quadrature and separates P from S by sampling the upgoing field at two depths. The raw reflected displacements of the two agree to 6×10^{-16} in every case compared — normal and oblique incidence, $n = 1, 2, 4$, with and without the first moments. Figure 1 collects the convergence.

9 Summary

Result	Evidence
Continuous layer, exact and to $O(D^k)$	reference to 2.5×10^{-8} ; series order k
Point-kernel local term survives $d \rightarrow 0$	$11 \times$ the continuum term at $m = 1$
Averaged coupling between planes = continuum $\times$ sinc	2.8×10^{-16} , 81 entries
Tiling identity, Eq. (9)	21/21 components to 4×10^{-9}
Exact kernel, Eq. (10)	replaces a 9.4×10^{-4} static error
Self term cancels, Eq. (12)	4×10^{-16}
Package T = closure statically	2×10^{-5} at $k_S h = 0.01$
Second-order convergence	order 2.000, three frequencies
Closed-form error, Eq. (13)	$n = 1-16$, to 10^{-4} of the error
First-moment voxel, Eq. (15)	order 4.00, three frequencies
First-moment voxel, oblique, Eq. (17)	order 4, P and S, 20° and 40°
First-moment error $-(kd)^4/720$, Eq. (16)	$n = 1-8$, to 10^{-4} of the error
Independent implementation of Eq. (17)	6×10^{-16} , 12 cases

The discrete scheme is exact wherever the voxel is: its coupling is the continuum's cell integral, and its self term cancels. What remains is the assumption that the field inside a voxel is uniform, and that error is known in closed form. A voxel that also carries the first moment of its field removes it and converges at fourth order: in one dimension, exactly in three at normal incidence, and in three at oblique incidence for both the reflected P and the converted S wave; its leading error, $-(kd)^4/720$, no longer depends on the direction of observation.

Scope. The results concern a laterally uniform layer, for which the discrete system is exactly a chain of planes; a laterally heterogeneous medium needs the same voxel in a full three-dimensional solver, which this paper does not build. The numerical evidence is at a moderate, physically typical contrast (a few per cent in density and moduli) and at $k_S h \leq 0.3$, the range in which the single-site T -matrix is validated; the closed forms are first order in the contrast, and the measured departures from them at the full contrast are second order. What is established is the *order* of convergence and its leading constant, not the behaviour at strong contrast, where the single site's own closure, rather than the lattice coupling, may limit the accuracy.

Every numerical value is reproducible from the Mathematica notebooks `Mathematica/ContinuumLimit_X.wl`, for X in Reference, ThinLayer, SpecularSums, StaticLocal, AveragedSums, CubeT, Chain, ErrorConstant, FourthOrder, FourthOrderError, Reduction3D, Oblique and Figure (the figure), and their reference dumps `scripts/continuum_limit_*.py`; the exact kernel is `cubic_scattering/cell_averaged_lattice.py`; the independent implementation of the first-moment voxel is `scripts/crosscheck_first_moment_voxel.py`.

References

- [1] L. L. Foldy, The multiple scattering of waves. I. General theory of isotropic scattering by randomly distributed scatterers, *Phys. Rev.* **67**, 107–119 (1945).
- [2] M. Lax, Multiple scattering of waves, *Rev. Mod. Phys.* **23**, 287–310 (1951).
- [3] E. M. Purcell and C. R. Pennypacker, Scattering and absorption of light by nonspherical dielectric grains, *Astrophys. J.* **186**, 705–714 (1973).
- [4] B. T. Draine and P. J. Flatau, Discrete-dipole approximation for scattering calculations, *J. Opt. Soc. Am. A* **11**, 1491–1499 (1994).
- [5] B. T. Draine and J. Goodman, Beyond Clausius–Mossotti: wave propagation on a polarizable point lattice and the discrete dipole approximation, *Astrophys. J.* **405**, 685–697 (1993).
- [6] M. A. Yurkin, V. P. Maltsev and A. G. Hoekstra, Convergence of the discrete dipole approximation. I. Theoretical analysis, *J. Opt. Soc. Am. A* **23**, 2578–2591 (2006).
- [7] P. C. Chaumet, A. Sentenac and A. Rahmani, Coupled dipole method for scatterers with large permittivity, *Phys. Rev. E* **70**, 036606 (2004).
- [8] J. D. Eshelby, The determination of the elastic field of an ellipsoidal inclusion, and related problems, *Proc. R. Soc. Lond. A* **241**, 376–396 (1957).