

The Continuum Limit of a Space-Filling Layer of Cubes

A Tiling Identity, an Exact Lattice Kernel,
and the Closed-Form Error of the Discrete Layer

T. M. Nestor

Independent researcher, Australia

28 September 2026

Abstract

Voxel schemes for elastic scattering represent a heterogeneous medium as cubic cells, each carrying a single-site T -matrix and coupled through the background Green's tensor. They are built for media with no exact solution, where their error cannot be measured directly and a refinement study can converge smoothly to the wrong limit. A layer of identical voxels exercises such a scheme in full and has an exact answer. Because cubes tile space, it isolates the discretisation error exactly. For this layer we establish whether, how fast and with what error the scheme converges. (i) The cube's lateral form factor vanishes on the reciprocal lattice, so the source-cell-averaged coupling is *exactly* the continuum's, and the lattice kernel at normal incidence is in closed form. (ii) The single-site T -matrix follows from distributional moments of the Green's tensor over the cube. Its uniform-strain closure differs from a sphere's only by a cubic anisotropy of the shear response. Against exact scattering by a sphere it is exact statically, and it departs only by a real, closed-form $(ka)^2$ term. (iii) The closure's self term cancels the kernel's, so the discrete system is a collocation of the continuum equation. (iv) It converges at second order, with leading error $+(kd)^2/8$ in reflection and $-(kd)^2/24$ in transmission. The opposite signs mean that no scalar correction of the single site can remove it. (v) A voxel that also carries the first moment of its internal field converges at fourth order, with leading error $-(kd)^4/720$, for incident P, SV and SH at normal and oblique incidence; with Legendre moments to degree p it converges at order $2p + 2$, with a closed-form constant, and a single plane of such voxels reproduces the thin layer's series in its thickness through $O(D^{2p+2})$. In a randomly stratified layer the orders are unchanged; where the medium varies within a cell, a contrast sampled as constant caps every voxel at second order, and a contrast expanded to the field's degree restores the full order; and the first-moment voxel reaches an accuracy of 10^{-6} with one voxel per model layer. On a voxelised sphere the discretisation error is again second order, and the staircase error of the shape dominates it. Every result is derived symbolically and checked against an independent implementation or an exact solution.

1 Introduction

The problem this paper serves. Seismic waves in the crust are scattered by heterogeneity on every scale. The models that describe it are parameterised in many ways (as layers or blocks, by smooth basis functions, or by interfaces), and a calculation that uses one must first sample it onto a grid. A volume-integral scattering calculation represents the sampled model as cells of constant properties and treats every cell as a scatterer: a voxel carrying a single-site T -matrix and coupled to every other voxel through the Green's tensor of the background, with the multiple-scattering series summed. Where the medium varies within a cell, that sampling is itself an approximation: §9 shows that it caps any voxel at second order, unless the cell's contrast is expanded to the degree of its field. For the media such a scheme is built for, there is no exact solution, and that is the difficulty. Its error cannot be measured against the answer. Comparison with another numerical method mixes the errors of both. And a refinement study can mislead, because a voxel scheme can converge smoothly to the wrong limit. §4 shows how that happens with the most natural choice of coupling, the Green's tensor evaluated between voxel centres: it carries a local term whose size per unit volume does not shrink as the voxels do, and unless the single voxel's response is matched to cancel it, refinement converges to the wrong answer without revealing it.

Why a homogeneous layer. This paper therefore starts from the simplest configuration in which a voxel scheme runs in full (every voxel, every coupling, the whole multiple-scattering series, at normal and at oblique incidence) and the exact answer is nevertheless known: a layer filled with identical voxels. Because cubes tile space, such a layer does not approximate a homogeneous layer; it *is* one. Any difference between its reflection and transmission and those of the continuous layer is therefore error of the discretisation alone, with no modelling error and no shape error behind which to hide. A sphere, the usual test object for volume-integral and discrete-dipole methods [1], cannot give this: its error mixes the discretisation with the staircase approximation of a curved surface by cubes. The layer plays the part of the patch test in finite elements: reproducing a uniform medium correctly is necessary, though not sufficient, for a scheme to be trusted in a heterogeneous one. What the layer establishes about the scheme's two ingredients, the coupling between voxels and the response of a single voxel, are properties of the ingredients, which are the same in any medium: that the averaged coupling is exact, that the self term cancels, and the form of the single-site error. §9 confirms that they carry over to a stratified medium; a laterally heterogeneous one remains to be tested.

What the reader gains. Four things that a heterogeneous test cannot provide. First, a proof, exact at normal incidence, that the voxel scheme converges to the continuum, and the specific ingredient, a coupling averaged over the source voxel, without which it converges to the wrong answer. Second, the leading error in closed form for every discretisation, so that a grid can be chosen for a required accuracy before the calculation is run. Third, a single-site T -matrix derived from first principles, with its error stated and checked against the exact solution for a sphere. Fourth, a voxel that converges at fourth order, and with further Legendre moments at any even order, also in a randomly stratified model (§9): for a reflection accuracy of 10^{-6} it allows voxels 58 times larger in each direction than the standard one (§11).

Why cubic voxels. Cubes are the natural cell for three reasons. They *tile space*, so a uniform region discretised into cubes is still that uniform region, whereas spheres or point scatterers can only approximate it and carry a packing error of their own. Their *regular lattice* makes the interaction between cells translation invariant, so the coupling is a convolution (Toeplitz between planes, Bloch-periodic within them), and the whole system is solved by fast transforms. And a cube small compared with the wavelength has a *single-site response in closed form*, the elastic depolarisation of a cube (§6), so the only approximation a voxel introduces is the internal field it assumes.

Multiple scattering on a lattice. Replacing a heterogeneous body by a lattice of small scatterers that interact through the background Green’s function is the multiple-scattering description of Foldy [2] and Lax [3]. In electromagnetics its discretised form is the discrete dipole approximation [DDA; 4, 5], whose accuracy has been studied in detail: each dipole’s polarisability must account for the local field of the surrounding lattice, which the lattice dispersion relation fixes so that an infinite lattice propagates waves as the continuum does to a given order in kd [6]; the error of the method converges at second order in the dipole spacing when shape errors are absent [1]; and integrating the Green’s tensor over each cell rather than sampling it at the centre improves accuracy, most of all for strong contrasts [7]. The elastic problem adds a structure the scalar and electromagnetic ones lack. A cubic voxel responds to the incident field with a force (through its density contrast) and a moment (through its stiffness contrast), so the state at a voxel has nine components (displacement and strain), and the single site’s self-interaction is the elastic depolarisation of a cube, the discrete counterpart of the inclusion tensor of Eshelby [8].

This paper. The approach originates in the two-and-a-half-dimensional formulation of Nestor [9]: a medium divided into cells that interact through the background Green’s tensor, each carrying a single-site contrast renormalised by its own near field, with the multiple-scattering series summed. This paper extends it to cubic voxels. §2 sets out the discrete medium and §3 the continuous layer it must reproduce. §4 shows why the point coupling fails, and §5 why the source-averaged coupling succeeds exactly. §6 derives the single-site T -matrix from moment integrals, tests it against the exact sphere, and shows that with it the discrete system is a collocation of the continuum equation. §7 gives that collocation’s second-order error in closed form, and §8 the fourth-order voxel, for incident P and S at normal and oblique incidence, and its generalisation to Legendre moments of any degree. §9 repeats the tests on a randomly stratified layer and on a medium that varies within a cell, §10 on a sphere, and §11 discusses what the results imply for voxel schemes in general. Every result is derived symbolically in Mathematica and checked against an independent Python implementation or an exact solution.

2 Setting

The discrete medium. Space is tiled by cubes of side d (half-side $h = d/2$). Cube i carries a contrast $\Delta\lambda, \Delta\mu, \Delta\rho$ from an isotropic background with Lamé parameters λ, μ , density ρ , P and S speeds α, β and P-wave modulus $M_p = \lambda + 2\mu$. Time dependence is $e^{-i\omega t}$; $k_p = \omega/\alpha$, $k_s = \omega/\beta$.

Coordinates are ordered (z, x, y) with z down. The field at a voxel is the nine-component state

$$\psi = (u_z, u_x, u_y, \varepsilon_{zz}, \varepsilon_{xx}, \varepsilon_{yy}, 2\varepsilon_{xy}, 2\varepsilon_{zy}, 2\varepsilon_{zx}), \quad (1)$$

displacement followed by engineering strain. A voxel responds with a force and a moment (stress dipole); the 9×9 single-site T -matrix maps the exciting state to them, and the 9×9 lattice Green's tensor $\mathbf{K}(i, j)$ maps a source voxel's force and moment to the state at voxel i . The Foldy–Lax system is

$$\psi_i = \psi_i^0 + \sum_{j \neq i} \mathbf{K}(i, j) \mathbf{T}_j \psi_j, \quad (2)$$

with ψ^0 the incident state. Strain rows of $\mathbf{K}$ are receiver derivatives of the point Green's tensor $\mathbf{G}$; moment columns are derivatives with respect to the receiver coordinate, symmetrised with a factor $\frac{1}{2}$ on the shear components. In that convention the contrast operator is

$$\Delta = \text{blockdiag}(\omega^2 \Delta \rho \mathbf{I}_3, \Delta \mathcal{C}), \quad (3)$$

where $\Delta \mathcal{C}$ is the 6×6 contrast stiffness with $\Delta \lambda + 2\Delta \mu$ and $\Delta \lambda$ in the normal block and $2\Delta \mu$ (not $\Delta \mu$) on the shear diagonal.

The laterally uniform layer. A layer of thickness D is filled with n planes of voxels, $d = D/n$, and illuminated at normal incidence. The field is then laterally uniform, only the zero-lateral-wavenumber (specular) component of the lattice coupling acts, and Eq. (2) becomes exactly a chain of n planes coupled by the specular lattice sums

$$\mathbf{S}(m) = \sum_{\mathbf{R}} \mathbf{G}(\mathbf{R} + m d \hat{\mathbf{z}}), \quad (4)$$

where the sum runs over the lattice vectors $\mathbf{R} = d(0, p, q)$, p, q integers, that join the centres of the voxels in one plane (a square lattice of pitch d), and m is the separation of the two planes in units of d (Fig. 1a,b). The question is whether, and how fast, this chain converges to the continuous layer as $n \rightarrow \infty$ at fixed D .

Implementations. Every result is derived symbolically in Mathematica and compared with an independent numerical implementation or an exact solution. “The package” below is the Python implementation of the voxel scheme (`cubic_scattering`): its cube T -matrix, its lattice kernels and its stratified propagator, which serve both as the object under test and, where validated, as a second implementation.

3 The continuous layer and its thin-layer expansion

At normal incidence a plane force drives a one-dimensional problem per wave type, with modulus M (M_P for P, μ for S), wavenumber k and plane Green's function $g(z) = \frac{i}{2Mk} e^{ik|z|}$. The layer's reflection and transmission follow from continuity of displacement and traction at its faces; they




Figure 1. The lattice and its sums. (a) Plan view of one plane of voxels: their centres form a square lattice of pitch d , and a lattice vector $\mathbf{R} = d(0, p, q)$ joins the centre of one voxel (red) to another. (b) Cross-section (z down), between planes: the voxel of a receiver plane (red) is coupled to every voxel of a source plane m planes above it, one for each lattice vector $\mathbf{R}$; the sum over $\mathbf{R}$ is the specular lattice sum of Eq. (4). (c) Within a plane: the sum over every cell of the plane, the receiver's own cell (outlined) included, is the field of a uniform plate of thickness d , the tiling identity of Eq. (11).

Alt text: Three schematic panels. On the left, a plan view of a five-by-five grid of square voxels with a dot at each centre; the centre voxel is highlighted, an arrow labelled $\mathbf{R}$ runs from its centre to another voxel centre, and the spacing between neighbouring centres is marked d . At the top right, a cross-section through three rows of voxels: arrows from the centre of every voxel in the top row converge on a highlighted voxel two rows below, with the plane separation $m d$ marked. Below it, a single row of voxels with one outlined as the receiver's own cell is shown equal to a continuous shaded band of the same thickness d , a uniform plate.

agree with the stratified propagator used as the reference in this work to 2.5×10^{-8} , $\mathbf{P}$ and $\mathbf{S}$ (`ContinuumLimit_Reference.wl`).

For a layer of thickness D with modulus M_1 and density ρ_1 in a background M_0, ρ_0 , write $Z_j = \sqrt{M_j \rho_j}$ for the impedances, $\zeta = Z_1/Z_0$ for their ratio and $k_1 = \omega \sqrt{\rho_1/M_1}$ for the layer's wavenumber. The reflection coefficient expands as

$$R = \frac{ik_1 D}{2}(\zeta - \zeta^{-1}) + \frac{(k_1 D)^2}{4}(\zeta^{-2} - \zeta^2) - \frac{i(k_1 D)^3}{24}(\zeta - \zeta^{-1})(3\zeta^{-2} + 2 + 3\zeta^2) + \frac{(k_1 D)^4}{48}(\zeta^2 - \zeta^{-2})(3\zeta^2 - 2 + 3\zeta^{-2}) + O(D^5), \quad (5)$$

exact in the contrast (`ContinuumLimit_ThinLayer.wl`, which checks each coefficient symbolically and the truncated series against the exact layer). The first term is

$$\frac{i\omega D}{2} \sqrt{\frac{M_0}{\rho_0}} \left[\frac{\rho_1 - \rho_0}{M_0} - \rho_0 \left(\frac{1}{M_1} - \frac{1}{M_0} \right) \right] : \quad (6)$$

to first order a thin layer is an interface across which the traction jumps by the excess inertia and the displacement by the excess compliance. At oblique incidence the same statement holds with the first-order system $\partial_z \mathbf{b} = \mathbf{A} \mathbf{b}$ for the P–SV state $\mathbf{b} = (u_x, u_z, t_{xz}, t_{zz})$: the layer acts to $O(D)$ as the jump $\mathbf{I} + (\mathbf{A}_1 - \mathbf{A}_0)D$, which reproduces the layer’s $O(D)$ reflection coefficients R_{PP}, R_{PS} to 10^{-12} ; the series truncated at $O(D^k)$ converges at exactly order k ; and a soft thin layer with shear modulus $\mu_1 = D/\eta$ at fixed compliance η tends to the linear-slip interface as $D \rightarrow 0$. These expansions are the target a single plane of voxels must reproduce order by order. A plane of voxels that carry Legendre moments of their internal field to degree p reproduces Eq. (5) term by term through $O(D^{2p+2})$ and first differs at D^{2p+3} (§8): the mean-only voxel reproduces the equivalent interface and the D^2 term, the first-moment voxel all four terms shown.

4 The specular sums of the point Green’s tensor

Why transform the sum. The specular sum (4) adds, at a receiver, the field of every voxel of a plane m planes away. In that form it is a poor object to study. The Green’s tensor decays slowly with distance, so the sum converges slowly, and it hides the one question that matters here: which part of it is the continuum, and which part is the discretisation. The Poisson summation formula [10] answers both at once, and is the standard tool for lattice sums of wave fields [11]. For a function f on the plane and the square lattice of pitch d it states

$$\sum_{\mathbf{R}} f(\mathbf{R}) = \frac{1}{d^2} \sum_{\mathbf{g}} \hat{f}(\mathbf{g}), \quad \hat{f}(\mathbf{g}) = \int f(\mathbf{x}) e^{-i\mathbf{g} \cdot \mathbf{x}} d^2x, \quad \mathbf{g} = \frac{2\pi}{d}(p, q), \quad (7)$$

with p, q integers: a sum of samples of f on the lattice equals a sum of samples of its Fourier transform on the reciprocal lattice, whose pitch is $2\pi/d$ (Fig. 2). Applied to the specular sum,

$$\mathbf{S}(m) = \frac{1}{d^2} \sum_{\mathbf{g}} \hat{\mathbf{G}}(\mathbf{g}; md), \quad (8)$$

with $\hat{\mathbf{G}}(\mathbf{g}; z)$ the two-dimensional (lateral) transform of the Green’s tensor between planes a distance z apart. Each term has a physical meaning. Transformed laterally, the field of a plane of sources is a superposition of plane waves, one for each lateral wavenumber, and the lattice selects the wavenumbers $\mathbf{g}$. The order $\mathbf{g} = 0$ is the laterally uniform plane wave that a continuous layer radiates: it is the continuum’s term. Every other order has $|\mathbf{g}| \geq 2\pi/d$, far beyond the wavenumber of the field when the voxels are small compared with the wavelength, so it is evanescent and decays away from the plane as $e^{-|\mathbf{g}|md} = e^{-2\pi m \sqrt{p^2+q^2}}$, independently of d . The transform therefore splits the lattice sum exactly into the continuum and a discretisation correction, and the correction converges after a few orders.

The local term that survives. In the strain-from-moment block the evanescent orders carry two powers of $|\mathbf{g}| \propto 1/d$, so their contribution per voxel volume does *not* vanish as $d \rightarrow 0$: at $m = 1$ it is eleven times the continuum term (`ContinuumLimit_SpecularSums.wl`, matched to the package’s Ewald lattice sum to 5×10^{-13}). This surviving term is static. Summed over all




Figure 2. The Poisson summation formula (7) applied to a plane of voxels, along one lateral axis. (a) Real space: the field of each voxel of a source plane at a receiver one plane away, sampled at the lattice points; the samples decay only as $1/|\mathbf{R}|$, so their sum converges slowly. (b) The same sum after Poisson summation, on a logarithmic scale: samples of the lateral Fourier transform at the reciprocal-lattice orders $\mathbf{g} = 2\pi\mathbf{p}/d$. The order $\mathbf{g} = 0$ is the continuum's plane wave; every other order is evanescent and, at $m = 1$, smaller than its neighbour by $e^{2\pi} \approx 535$. Magnitudes are schematic: the static decay $e^{-|\mathbf{g}|md}$ is shown.

Alt text: Two panels. On the left, a smooth bell-shaped curve that decays slowly on both sides is sampled by vertical stems at integer positions from minus six to six, the stems remaining clearly visible far from the centre. On the right, on a logarithmic vertical axis from one down to ten to the minus nine, a stem of height one at zero is labelled as the continuum, and stems at plus and minus one, two and three fall by about three powers of ten per step and are labelled as evanescent orders.

planes in slab order, the static lattice sum of the point kernel is exactly the depolarisation of a uniform plate ($-1/M_p$ on $\varepsilon_{zz} \leftarrow M_{zz}$ and $-1/(2\mu)$ on the normal shears, per unit volume) plus a cubic-symmetric remainder (ContinuumLimit_StaticLocal.w1, to 2×10^{-10}). The point kernel therefore reproduces the continuum only if that remainder is cancelled by the voxel's own self term.

5 The tiling identity

Replace the point coupling by the *source-cell average* $\langle \mathbf{G} \rangle(\mathbf{r}) = d^{-3} \int_{\text{cube}} \mathbf{G}(\mathbf{r} - \mathbf{y}) d\mathbf{y}$, the field at a voxel centre of a uniform polarisation of the source voxel. Each Poisson order in Eq. (8) is then multiplied by the cube's form factor

$$F(\mathbf{g}) = \text{sinc}(g_x h) \text{sinc}(g_y h), \quad (9)$$

which vanishes at every $\mathbf{g} \neq 0$ of the reciprocal lattice since $\text{sinc}(p\pi) = 0$ for integer $p \neq 0$. This is the Fourier form of the fact that cubes tile the plane, $\sum_{\mathbf{R}} \chi(\mathbf{x} - \mathbf{R}) = 1$ for the cube's indicator χ : a partition of unity (Fig. 3). Two exact statements follow.

Between planes. For $m \neq 0$ every evanescent order is removed and

$$\mathbf{S}_{\text{avg}}(m) = \frac{1}{d^2} \sum_{W=P,S} \hat{\mathbf{G}}_W(0; md) \text{sinc}(k_W h), \quad (10)$$

the continuum's plane-wave term times the vertical form factor. The package's cell-averaged kernel agrees in all 81 entries to 2.8×10^{-16} (ContinuumLimit_AveragedSums.w1).




Figure 3. Why cubes make the coupling exact. (a) Cubes tile the plane: the indicator functions χ of the translated cells sum to exactly one everywhere, a partition of unity. (b) The same fact in Fourier space: the cube’s lateral form factor $F(g) = \text{sinc}(gh)$ (shown along one axis) vanishes at every reciprocal-lattice order $g_p = 2\pi p/d$ with $p \neq 0$, so of all the orders in the lattice sum (8) only $g = 0$, the continuum’s plane wave, survives the source-cell average.

Alt text: Two panels. On the left, adjacent box functions of unit height and unit width, alternately shaded, lie side by side along an axis marked in voxel widths, and a thick horizontal line at height one marks their sum. On the right, the sinc curve, equal to one at the origin and oscillating with decaying amplitude, is plotted against wavenumber in units of $2\pi/d$; dots mark the reciprocal-lattice orders, one at height one at the origin and all the others on the axis at the zeros of the curve.

Within a plane. The sum over *all* cells of the plane, the voxel’s own included, is the field at the mid-plane of a uniform plate of thickness d (Fig. 1c):

$$d^3 \mathbf{S}_{\text{avg}}(0) + \mathbf{S}_{\text{self}} = d^3 \mathbf{P}_0, \quad \mathbf{S}_{\text{self}} = \int_{\text{cube}} \mathbf{G}(-\mathbf{y}) d\mathbf{y}, \quad (11)$$

where $\mathbf{P}_0 = d^{-3} \int_{-h}^h \mathbf{G}_{\text{plate}}(-z') dz'$ is elementary: in the one-dimensional transform each entry is $c_0 + A/(Mk_z^2 - \rho\omega^2)$, so it averages to $c_0/d + A(e^{ikh} - 1)/(dMk^2)$. Its only nonzero entries are $u \leftarrow f$ on the diagonal and the three normal strain-from-moment entries, whose local weights $-1/M_p$ and $-1/(2\mu)$ are the plate depolarisation of §4. Statically the identity holds for all 21 independent components of the strain-from-moment tensor to 4×10^{-9} , with every cell integral evaluated analytically (below).

Consequence: an exact kernel with no lattice sum. At normal incidence the same-plane block of the cell-averaged Foldy–Lax kernel is

$$\mathbf{S}_{\text{avg}}(0) = \mathbf{P}_0 - \mathbf{S}_{\text{self}}/d^3, \quad (12)$$

with no Ewald splitting, near shell or asymptotic tail. The route it replaces (Ewald summation with a near shell averaged by quadrature and an $O(d^2)$ Taylor tail beyond it) carries a static strain-from-moment error of 9.4×10^{-4} at its default settings, falling to 2.6×10^{-5} only when both the shell radius and the quadrature order are raised; Eq. (12) is now the package’s kernel at normal incidence.

Cell integrals without quadrature. Write the static (Kelvin) tensor as $G_{bc} = (c_A + c_B) \delta_{bc}/r - c_B \partial_b \partial_c r$ with $c_B = (1/\mu - 1/M_p)/8\pi$ and $c_A = 1/(4\pi\mu) - c_B$. Every cell integral of its derivatives is

a cell integral of derivatives of $1/r$ and r . One application of the divergence theorem moves it onto faces $x_p = c_p \pm h$, which never contain the origin, and there the two-dimensional antiderivative in the face coordinates is continuous, so the face integral is a sum over the face's corners. The distributional term at $r = 0$ is carried by the divergence theorem itself. (A triple antiderivative evaluated at the cube's corners is *not* usable for repeated derivatives: its arctangent terms jump across the coordinate planes through the origin, and differentiating twice there discards exactly that term.) The self term so obtained reproduces the classical closed form of the cube's depolarisation to machine precision.

6 The single site: the cube T -matrix from moment integrals

Every error of a voxel scheme is made at the single site (§11), so the single site is derived here rather than imported. The classical route imports it. For an ellipsoid, Eshelby's uniformity theorem makes the internal strain under a uniform load uniform, with a depolarisation tensor in closed form [8, 12, 13]; its dynamic generalisation is known for spheres and ellipsoids [14, 15]. For a cube or a general polyhedron the static field of a uniform eigenstrain is also in closed form [16–19], but the strain it produces is not uniform, so there is no exact uniform-strain inclusion to import. In scattering theory the same object appears as the self-consistent or effective-field closure of the volume integral equation [20–22]: a first-order (Born) scatterer [23] renormalised by the geometric series of its own near field, the step that turns a contrast into an effective contrast. Nestor [9] introduced the single site in this form, summing the series of the cell's self integral, with the singular near field evaluated over an infinitesimal cell. The moments below evaluate the same self integral over the finite cube, as a distribution. The elastic T -matrix of Waterman [24] represents a scatterer in spherical multipoles, the natural basis for a sphere; a voxel in a lattice needs one compatible with the lattice, and the moments over the cube provide it. The construction below makes the approximation of a voxel explicit rather than implicit: it derives the cube's T -matrix from the moments of the elastodynamic Green's tensor over the cube, and identifies the uniform-strain closure as one truncation of an exact hierarchy.

6.1 The closed hierarchy

Let the voxel be $V = [-h, h]^3$ with contrast $\delta c_{ijkl} = \Delta\lambda \delta_{ij}\delta_{kl} + \Delta\mu(\delta_{ik}\delta_{jl} + \delta_{il}\delta_{jk})$ and $\Delta\rho$. Expand the internal displacement about the centre to second degree and project the Lippmann–Schwinger equation onto 1 , x_p and x_px_q . The unknowns are $u_i(0)$, $\partial_p u_i(0)$ and $\partial_p \partial_q u_i(0)$, and the result is the closed system

$$(\delta_{ij} - \omega^2 \Delta\rho G_{ij})u_j + (N_{in,k}^r \delta c_{nksj} - \frac{1}{2}\omega^2 \Delta\rho K_{ij}^{rs})\partial_r \partial_s u_j = u_i^0, \quad (13)$$

$$(\delta_{pr}\delta_{ij} + \omega^2 \Delta\rho N_{ij,p}^r - M_{in,pk} \delta c_{nkrj})\partial_r u_j = \partial_p u_i^0, \quad (14)$$

$$-\omega^2 \Delta\rho M_{ij,pq} u_j + (\delta_{qr}\delta_{ps}\delta_{ij} + Q_{in,pqk}^r \delta c_{nksj} - \frac{1}{2}\omega^2 \Delta\rho P_{ij,pq}^{rs})\partial_r \partial_s u_j = \partial_p \partial_q u_i^0, \quad (15)$$

all fields evaluated at the centre. Its coefficients are six moments of the background Green's tensor $\mathbf{G}$ over the voxel,

$$\begin{aligned} \mathbf{G}_{ij} &= \int_V G_{ij} dV, & \mathbf{K}_{ij}^{rs} &= \int_V G_{ij} x_r x_s dV, & \mathbf{N}_{in,k}^r &= \int_V \partial'_k G_{in} x_r dV, \\ \mathbf{M}_{in,pk} &= \int_V \partial'_p \partial'_k G_{in} dV, & \mathbf{P}_{ij,pq}^{rs} &= \int_V \partial'_q \partial'_p G_{ij} x_r x_s dV, & \mathbf{Q}_{in,pqk}^r &= \int_V \partial'_q \partial'_p \partial'_k G_{in} x_r dV, \end{aligned} \quad (16)$$

with derivatives taken with respect to the source point, $\partial' = -\partial$. That sign reaches $\mathbf{N}$ and $\mathbf{Q}$ only, and no symmetry, parity or scaling check can detect it, so it is fixed by convention. The moments are graded by (D, W) , the number of derivatives and the degree of the weight: $\mathbf{G}(0, 0)$, $\mathbf{K}(0, 2)$, $\mathbf{N}(1, 1)$, $\mathbf{M}(2, 0)$, $\mathbf{P}(2, 2)$, $\mathbf{Q}(3, 1)$.

6.2 The moments as distributions

Write the elastodynamic tensor as

$$G_{in}(\mathbf{r}) = \frac{1}{4\pi\mu} \left[\delta_{in} g_S(r) + \frac{1}{k_S^2} \partial_i \partial_n (g_S(r) - g_P(r)) \right], \quad g_c(r) = \frac{e^{ik_c r}}{r} \quad (17)$$

[25, 26]. Expanding g_c in powers of r reduces every moment in Eq. (16) to a finite sum of scalar moments

$$E[m; d_1 \dots d_D; w_1 \dots w_W] = \langle \partial_{d_1} \dots \partial_{d_D} r^m, x_{w_1} \dots x_{w_W} \mathbf{1}_V \rangle. \quad (18)$$

For $m = -1$ and $D \geq 2$ the classical integrand is not absolutely integrable. $\partial^2(1/r)$ decays as r^{-3} , its angular average vanishes, and the value of the integral depends on the shape of the region excluded around the origin. This is the source of the δ -function term of the inclusion problem, which a numerical quadrature misses. Without it the depolarisation is wrong; for a sphere it vanishes altogether (§6.3). It is not added afterwards here. The derivatives are distributional from the start, and each is moved onto the voxel's faces by pairing with the test function $w \mathbf{1}_V$, using $\partial_q \mathbf{1}_V = -n_q \delta_{\partial V}$:

$$\langle \partial_q T, w \mathbf{1}_V \rangle = \oint_{\partial V} n_q w T dA - \langle T, (\partial_q w) \mathbf{1}_V \rangle. \quad (19)$$

The surface term is evaluated on the faces, a distance h from the origin, where every derivative of r^m is smooth. The recursion lowers D and W together and ends on convergent integrals. The face integrals are done in closed form by the corner sums of §5. The sharp check is the sum rule

$$\sum_p E[-1; p, p;] = \nabla^2(1/r) \text{ paired with } \mathbf{1}_V = -4\pi, \quad (20)$$

which the face route satisfies exactly and a spherical excision returns as 0. The closed forms contain logarithms of the units $(2 \pm \sqrt{3})^n$ of the Pell equation $x^2 - 3y^2 = 1$. Left unreduced, these cancel catastrophically: for example, $\log(72010600134783751 - 41575339372323900\sqrt{3})$ is a cancellation of 35 digits. They are therefore reduced to $n \log(2 - \sqrt{3})$, which is exact. Because every power of r is paired this way, the dynamic moments keep the even-power terms of the expansion ($k^2 r$, r^3 , ...) as well as the odd ones. A route that integrates only the polynomial

(odd-power) terms of $\mathbf{G}$ misses those terms, at a real $O((kh)^2)$ cost (§6.4).

6.3 Parity, and the nine-component block

The cube is centrosymmetric. A moment that couples a field quantity of degree a to one of degree b therefore vanishes unless $a + b$ is even. u , ∂u and $\partial\partial u$ have degrees 0, 1 and 2, so Eq. (14) contains the first gradient alone: ∂u is determined by its own block and couples to neither u nor $\partial\partial u$. The leading scattering by a stiffness contrast is the stress dipole $\int_V \delta c : \nabla u \, dV$, which needs only ∂u . The leading far field of the stiffness contrast is therefore fixed by Eq. (14), and *no treatment of the second-gradient block (15) can change it*. What Eq. (15) contributes reaches the far field only through the $O((kh)^2)$ coupling in Eq. (13). That is the order, and the origin (the variation of the field across the voxel), of the collocation error of §7, which the first-moment voxel of §8 removes.

With $\Delta\rho = 0$, Eq. (14) is the 9×9 operator $\delta_{pr}\delta_{ij} - M_{in,pk}\delta c_{nkrj}$, built from $\mathbf{M}$ alone. Under O_h the nine gradient components decompose as $A_{1g} \oplus E_g \oplus T_{1g} \oplus T_{2g}$, and the operator is diagonal in that basis. Its eigenvalues are the inverse strain concentrations:

$$A_{1g} \text{ (dilatation)} : 1 + \frac{\Delta\lambda + \frac{2}{3}\Delta\mu}{\lambda + 2\mu}, \quad (21)$$

$$T_{2g} \text{ (off-diagonal shear)} : 1 + 2\Delta\mu S_{\text{off}}, \quad S_{\text{off}} = \frac{\pi(\lambda + 2\mu) - \sqrt{3}(\lambda + \mu)}{3\pi\mu(\lambda + 2\mu)}, \quad (22)$$

$$E_g \text{ (diagonal shear)} : 1 + 2\Delta\mu S_{\text{diag}}, \quad S_{\text{diag}} = \frac{3\sqrt{3}(\lambda + \mu) + 2\pi\mu}{6\pi\mu(\lambda + 2\mu)}. \quad (23)$$

The density channel is the displacement block of Eq. (13) with the second gradients dropped, an amplification $(1 - \omega^2\Delta\rho\mathbf{G})^{-1}$ with $\mathbf{G}$ the cube integral of $\mathbf{G}$.

The cube against the sphere. For a sphere, Eshelby's tensor gives a dilatational channel $1 + \Delta K/(\lambda + 2\mu)$, with $\Delta K = \Delta\lambda + \frac{2}{3}\Delta\mu$, and a single deviatoric channel $1 + 2\Delta\mu S_{\text{sph}}$, with $2\mu S_{\text{sph}} = 2(4 - 5\nu)/[15(1 - \nu)]$. The cube's dilatational channel, Eq. (21), is the sphere's exactly: its eigenvalue involves $\mathbf{M}$ only through the contraction $\int_V \partial_i \partial_n G_{in} \, dV$, and the static Navier equation makes $\partial_i \partial_n G_{in} = -\delta(\mathbf{r})/(\lambda + 2\mu)$, whose integral does not depend on the shape. The cube's shear channels differ from each other. Their isotropic average, however, is the sphere's, also exactly:

$$\frac{1}{5}(3S_{\text{off}} + 2S_{\text{diag}}) = S_{\text{sph}}, \quad S_{\text{diag}} - S_{\text{off}} = \frac{(\lambda + \mu)(5\sqrt{3} - 2\pi)}{6\pi\mu(\lambda + 2\mu)} > 0. \quad (24)$$

The $\sqrt{3}$ terms cancel in the average. What a cube adds to a sphere at this order is therefore a pure cubic anisotropy. It splits the deviatoric channel about the sphere's value, weighted 3 : 2 by the dimensions of T_{2g} and E_g , and the diagonal shear is the more depolarised. At the background of §7 ($\nu = 7/32$), $2\mu S_{\text{off}} = 0.4314$, $2\mu S_{\text{diag}} = 0.5929$ and the sphere's value is 0.4960.

Where Eshelby enters. Equations (21)–(23) follow from the moment integrals and the group theory alone. No ellipsoidal result is used, and no δ term is added by hand. What corresponds to Eshelby is the *truncation*. Keeping ∂u and discarding $\partial\partial u$ is exactly the uniform-internal-strain ansatz, and by the uniformity theorem that ansatz is exact if and only if the inclusion is an ellipsoid. For a sphere the construction is therefore exact in the static limit, and for a cube it is an approximation. The hierarchy (13)–(15) states that approximation as the omitted blocks, rather than as an unquantified shape correction. A polynomial expansion about the centre cannot represent the edge and corner concentrations of an isolated cube’s internal field at any finite degree. A voxel in a uniform layer has no such concentrations: its neighbours have the same properties, so its internal field is the continuum’s, which is smooth. That is why a low-degree internal field suffices in a lattice (§§7–8) although it is only an approximation for an isolated cube.

Against the exact sphere. These statements are tested against the exact elastic Mie solution for a sphere, with no Eshelby formula in between. The concentration of Mie order n is $E_n = a_n/a_n^{\text{Born}}$ (orders 0, 1, 2 are the dilatational, density and shear channels), extrapolated to $ka \rightarrow 0$. The formulation is applied to a ball of radius a , with its moment $\mathbf{M}$ evaluated by the surface route of Eq. (19) on $|\mathbf{r}| = a$. Its dilatational and shear channels equal Mie’s to 2×10^{-10} , and the density channel is 1. The cube’s closed forms meet the same numbers. Its dilatational channel is Mie’s E_0 to 2×10^{-10} , and its isotropic shear average, Eq. (24), is the depolarisation read from Mie’s E_2 to 10^{-8} . The same moment taken as a principal value, over the ball with a small sphere excluded around the origin, has zero depolarisation in every channel. The whole static response of the sphere is the δ term.

At finite frequency the closure uses the dynamic ball moments, which are in closed form. It departs from Mie at second order in ka and first order in the contrast, and the departure is closed-form too: the ball’s moments are elementary, and Mie’s low-frequency series with the contrast kept exact gives E_n to relative order $(ka)^3$. To first order in the contrast,

$$\frac{E_n^{\text{closure}}}{E_n^{\text{Mie}}} - 1 = \begin{cases} -\frac{(k_p a)^2}{10} \frac{\Delta\lambda}{\lambda + 2\mu}, & n = 0, \\ \frac{2(k_s a)^2 + (k_p a)^2}{30} \frac{\Delta\rho}{\rho}, & n = 1, \\ -\frac{3(k_s a)^2 + 2(k_p a)^2 \beta^2 / \alpha^2}{75} \frac{\Delta\mu}{\mu}, & n = 2, \end{cases} \quad (25)$$

with no $O(ka)$ term and no $O((ka)^3)$ term. The imaginary $(ka)^3$ terms, the radiation damping, agree exactly in all three channels. So the closure’s only error at this order is the real $(ka)^2$ term, the uniform-field ansatz’s own. That is the order the hierarchy assigns to the omitted second gradients (§6.3), and it has the same origin, the variation of the field across the scatterer, as the collocation error of §7. Eq. (25) is derived symbolically (`SphereClosureDynamic.wl`). It agrees with the coefficient extracted from the numerical Mie solve and an independent implementation of the closure to 5×10^{-4} , which is the size of the next order in the contrast.

6.4 The collocation closure

The collocation closure takes the internal field of a voxel as uniform, the truncation of §6.3 for the density and stiffness channels together. The voxel's own field at its centre is then $\mathbf{S}_{\text{self}}\Delta \mathbf{e}$ for internal state $\mathbf{e}$, where $\mathbf{S}_{\text{self}}$ is the nine-component form of the moments $\mathbf{G}$ and $\mathbf{M}$. By parity it is block-diagonal. So $\mathbf{e} = \boldsymbol{\psi} + \mathbf{S}_{\text{self}}\Delta \mathbf{e}$ and

$$\mathbf{T} = d^3 \Delta (\mathbf{I} - \mathbf{S}_{\text{self}}\Delta)^{-1}. \quad (26)$$

With the kernel (12) the self term cancels. Substituting $\boldsymbol{\psi} = (\mathbf{I} - \mathbf{S}_{\text{self}}\Delta)\mathbf{e}$ into Eq. (2),

$$e_i = \psi_i^0 + d^3 \sum_j \mathbf{K}_{\text{all}}(i, j) \Delta e_j, \quad \mathbf{K}_{\text{all}}(i, i) = \mathbf{P}_0, \quad \mathbf{K}_{\text{all}}(i, j) = \mathbf{S}_{\text{avg}}(i - j), \quad (27)$$

an identity checked on the matrices to 4×10^{-16} . Every block of $\mathbf{K}_{\text{all}}$ is an *exact* cell integral of the one-dimensional continuum kernel (Eqs. (10), (11)), so Eq. (27) is precisely the collocation of the continuum Lippmann–Schwinger equation with a strain piecewise constant over voxels. It does not contain $\mathbf{S}_{\text{self}}$ at all.

The package's cube T -matrix is this closure statically: both blocks agree with Eq. (26) to 2×10^{-5} at $k_S h = 0.01$, the difference falling as $(kh)^2$ (`ContinuumLimit_CubeT.wl`). It departs dynamically by about $0.25(kh)^2$ through its far-field form-factor corrections. Its self integrals are exact in the static and radiation parts but omit the real even-power terms of the Green's tensor's expansion ($k^2 r, r^3, \dots$), an error of $0.4\text{--}0.6(kh)^2$ against the analytic self term, whose series agrees with the closed-form elastodynamic tensor to 10^{-14} . By Eq. (27) that error cannot enter the closure scheme.

7 Convergence, and the closed-form error

Second order. For a layer $D = 2$ m with the contrast $\Delta\lambda = 2$ GPa, $\Delta\mu = 1$ GPa, $\Delta\rho = 100 \text{ kg m}^{-3}$ in a background $\alpha = 5000$, $\beta = 3000 \text{ ms}^{-1}$, $\rho = 2500 \text{ kg m}^{-3}$, illuminated by a P plane force, the reflected displacement of the chain converges to the continuous layer's at order 2.000 for $n = 4 \rightarrow 16$ at $\omega = 60, 300$ and 600 rad s^{-1} , with the collocation (27) and with the package T -matrix alike (`ContinuumLimit_Chain.wl`). The error depends on $k_S h$ alone.

The leading error in closed form. In Eq. (27) the kernel is integrated exactly; the only approximation is the internal field's variation across a voxel. At first order in the contrast the discrete scattered field is therefore a midpoint rule applied to the continuum's Born integral, with the kernel's cell integral exact. For reflection the integrand carries the two-way phase e^{2ikz} ; for transmission it is flat. With k the background wavenumber of the wave concerned (k_P for an incident P wave, k_S for an incident S wave), R the reflection coefficient and T^{sc} the scattered part of the transmitted field, for *every* n

$$\frac{R_{\text{disc}}}{R} = \frac{\text{sinc}(kh)}{\text{sinc}(2kh)} = 1 + \frac{(kd)^2}{8} + O(d^4), \quad \frac{T_{\text{disc}}^{\text{sc}}}{T^{\text{sc}}} = \text{sinc}(kh) = 1 - \frac{(kd)^2}{24} + O(d^4), \quad (28)$$

confirmed for $n = 1$ to 16 to 10^{-4} of the error itself (`ContinuumLimit_ErrorConstant.wl`); for an incident S wave, with $k = k_S$, the same forms and those of Eq. (31) hold to 4×10^{-4} (`ContinuumLimit_IncidentS.wl`, §8). At the full contrast the reflection constant is unchanged (0.124 against 1/8); the transmission constant grows to -0.092 .

Why a uniform-field single site stops at second order. The two constants have opposite signs, so no scalar correction of $\mathbf{T}$, no single $O((kh)^2)$ factor, can cancel both. The package *T*-matrix illustrates it: its form-factor corrections reduce the reflection error by a factor 4.45 and *increase* the transmission error by 1.7, at every n . The missing term in Eq. (28) is the first moment of the internal field over each voxel, which radiates through the derivative of the kernel. A fourth-order scheme must carry it: a single site with a gradient channel, and a kernel whose cell integrals include the kernel's derivative.

8 Fourth order and beyond: Legendre moments of the voxel field

Let each voxel carry the mean *and* the first moment of its internal field. On voxel j , with local coordinate $s \in [-h, h]$ and Legendre functions $\phi_b(s) = P_b(s/h)$ ($\phi_0 = 1$, $\phi_1 = s/h$, $\phi_2 = (3s^2/h^2 - 1)/2$, ...), the field is $w = \sum_b c_{b,j} \phi_b$, $w = (u, \varepsilon)$ the displacement and strain of the P channel, and the polarisation is $q = (\omega^2 \Delta \rho u, \Delta M \varepsilon)$ with $\Delta M = \Delta \lambda + 2\Delta \mu$. Testing the continuum equation

$$w(z) = w^0(z) + \int \mathbf{K}(z - z') q(z') dz', \quad \mathbf{K} = \begin{pmatrix} g & g' \\ g' & g'' \end{pmatrix}, \quad g'' = -k^2 g - \delta(z)/M_P, \quad (29)$$

with the same functions on every voxel gives

$$\left(\int_i \phi_a^2 dz \right) c_{a,i} = \int_i \phi_a w^0 dz + \sum_{j,b} \left[\iint_j \phi_a(z) \mathbf{K}(z - z') \phi_b(z') dz dz' \right] \Delta_1 c_{b,j}, \quad (30)$$

$\Delta_1 = \text{diag}(\omega^2 \Delta \rho, \Delta M)$. Between voxels the double integral is separable, a product of Legendre moments of $e^{\pm i k z}$; within a voxel it is split at $z' = z$ and the local term contributes $-\int_i \phi_a \phi_b dz / M_P$. Every entry is closed-form. Keeping only ϕ_0 recovers a second-order scheme; keeping ϕ_0 and ϕ_1 converges at fourth order in both reflection and transmission (`ContinuumLimit_FourthOrder.wl`; Table 1). A single plane of first-moment voxels represents the 2 m layer 6500 times more accurately

Table 1. Relative error of the scattered displacement of the layer, reflection / transmission, at $\omega = 300 \text{ rad s}^{-1}$, against the number of voxel planes n : collocation, the mean-only voxel and the first-moment voxel of Eq. (30). The order is 4.00 at $\omega = 60$ and 600 rad s^{-1} too.

n	collocation	mean only	mean and first moment
1	$1.8 \times 10^{-3} / 1.3 \times 10^{-3}$	$1.2 \times 10^{-3} / 1.8 \times 10^{-3}$	$2.8 \times 10^{-7} / 2.8 \times 10^{-7}$
4	$1.1 \times 10^{-4} / 8.3 \times 10^{-5}$	$7.5 \times 10^{-5} / 1.1 \times 10^{-4}$	$1.1 \times 10^{-9} / 1.1 \times 10^{-9}$
16	$7.0 \times 10^{-6} / 5.2 \times 10^{-6}$	$4.7 \times 10^{-6} / 6.9 \times 10^{-6}$	$4.2 \times 10^{-12} / 4.2 \times 10^{-12}$
order	2 / 2	2 / 2	4 / 4

than a single plane of uniform ones. The collocation column is the scheme of §6 and reproduces the chain of the full 9×9 system to all printed digits.

Higher moments. Carrying the Legendre moments to degree p changes nothing in the construction: the moments and the same-cell double integrals of $P_b(s/h)$ against exponentials are again closed-form, and for $p \leq 1$ the general-degree solver reproduces the chain above exactly. The second-moment voxel ($p = 2$) converges at sixth order and the third-moment voxel ($p = 3$) at eighth, in reflection and transmission (Table 2; `ContinuumLimit_SecondMoment.wl`); the second-moment voxel's order is 6.00 at $\omega = 60, 300$ and 600 rad s^{-1} . The errors fall to 10^{-22} , so they are computed in 40-digit arithmetic.

Table 2. Relative error of the scattered displacement of the layer, reflection / transmission, at $\omega = 300 \text{ rad s}^{-1}$, for voxels carrying Legendre moments to degree p . The order is the mean over the halvings $n = 2 \rightarrow 4 \rightarrow 8$.

n	first moment ($p = 1$)	second moment ($p = 2$)	third moment ($p = 3$)
1	$2.8 \times 10^{-7} / 2.8 \times 10^{-7}$	$2.8 \times 10^{-11} / 4.2 \times 10^{-11}$	$1.5 \times 10^{-15} / 1.5 \times 10^{-15}$
2	$1.7 \times 10^{-8} / 1.7 \times 10^{-8}$	$4.4 \times 10^{-13} / 6.5 \times 10^{-13}$	$6.0 \times 10^{-18} / 6.0 \times 10^{-18}$
4	$1.1 \times 10^{-9} / 1.1 \times 10^{-9}$	$6.9 \times 10^{-15} / 1.0 \times 10^{-14}$	$2.4 \times 10^{-20} / 2.4 \times 10^{-20}$
8	$6.7 \times 10^{-11} / 6.7 \times 10^{-11}$	$1.1 \times 10^{-16} / 1.6 \times 10^{-16}$	$9.2 \times 10^{-23} / 9.2 \times 10^{-23}$
order	4 / 4	6 / 6	8 / 8

The leading error in closed form. Without scattering, the Galerkin solution is the L^2 projection Π of the incident field onto each voxel's basis, and the kernel is integrated over each voxel exactly; at first order in the contrast the error is therefore the kernel's cell integral applied to $\Pi w^0 - w^0$. The field varies as e^{iks} across a voxel and the kernel toward an outside observer as e^{iqs} , $q = +k$ for reflection and $q = -k$ for transmission, so for every scheme and every n

$$\frac{\text{discrete}}{\text{exact}} = \frac{\int_{-h}^h e^{iqs} \Pi[e^{iks}] ds}{\int_{-h}^h e^{i(q+k)s} ds}, \quad (31)$$

with Π the centre value for collocation, the cell mean for the mean-only voxel and the projection onto $\{P_0, \dots, P_p\}$ for the voxel of degree p (`ContinuumLimit_FourthOrderError.wl`, `ContinuumLimit_SecondMoment.wl`). The results are in Table 3. The chain of Eq. (30) reproduces the forms for collocation and degrees 0 and 1 at $n = 1$ to 8 to 10^{-4} of the error, the remainder being the second order in the contrast, and those for degrees 2 and 3 to 4×10^{-7} at a contrast reduced by 10^{-6} . The constant for degree p follows from the Legendre expansion $e^{iks} = \sum_j i^j (2j+1) j_j(kh) P_j(s/h)$. The projection drops the terms $j > p$, led by $j = p+1$, whose coefficient is $O((kh)^{p+1})$; its moment against e^{iqs} is $2h i^{p+1} j_{p+1}(qh) = O((qh)^{p+1})$. With $j_n(x) \approx x^n / (2n+1)!!$ and $h = d/2$,

$$c_p = \frac{2p+3}{4^{p+1} [(2p+3)!!]^2} = \frac{[(p+1)!]^2}{(2p+2)! (2p+3)!}, \quad c_0 = \frac{1}{12}, \quad c_1 = \frac{1}{720}, \quad c_2 = \frac{1}{100800}, \quad (32)$$

Table 3. The leading error of the discrete layer in closed form, Eq. (31): discrete over exact scattered field at first order in the contrast, for every number of voxel planes n .

	reflection	transmission
collocation	$1 + (kd)^2/8 + O(d^4)$	$1 - (kd)^2/24 + O(d^4)$
mean only	$1 + (kd)^2/12 + O(d^4)$	$1 - (kd)^2/12 + O(d^4)$
first moment	$1 - (kd)^4/720 + O(d^6)$	$1 - (kd)^4/720 + O(d^6)$
second moment	$1 + (kd)^6/100800 + O(d^8)$	$1 - (kd)^6/100800 + O(d^8)$
third moment	$1 - (kd)^8/25401600 + O(d^{10})$	$1 - (kd)^8/25401600 + O(d^{10})$
degree p	$1 + (-1)^p c_p (kd)^{2p+2}$	$1 - c_p (kd)^{2p+2}$

and the leading error carries $(q/k)^{p+1}$: the same sign in reflection and transmission for odd p , opposite signs for even p . The first-moment voxel's leading error therefore does not depend on the direction in which the field is observed. Eq. (32) is checked symbolically for $p = 0$ to 5 (every lower power of kd vanishes) in Mathematica and, independently, in sympy.

One plane of voxels and the thin-layer series. Refining at fixed thickness is one limit; thinning a single plane is the other. For one plane of voxels ($n = 1$, $d = D$) the series in D of the reflection coefficient, computed exactly at the full contrast, coincides with that of the exact layer, Eq. (5), term by term through D^{2p+2} and first differs at D^{2p+3} , for $p = 0$ to 3 and for a layer both stiffer and softer than its surroundings (`ContinuumLimit_SecondMoment.wl`). Each further Legendre moment adds two terms of the thin-layer series.

In three dimensions. The three-dimensional first-moment voxel carries the nine-component state with Legendre weights $\{1, x/h, y/h, z/h\}$, 36 unknowns in all, tested with the same functions and coupled by cell-to-cell double integrals of the Green's tensor. For a laterally uniform layer at normal incidence it reduces *exactly* to Eq. (30) (`ContinuumLimit_Reduction3D.wl`). The lateral first moments are never excited: the lattice and the incident field are symmetric under reflection about every voxel's centre planes, and $x/h, y/h$ are odd there. The mean and the z -moment carry weights that are constant laterally, so their lateral form factor $\text{sinc}(g_x h) \text{sinc}(g_y h)$ vanishes on the reciprocal lattice and the plane sum of the double integrals is the plate-to-plate integral in z alone. Two facts mark where genuinely three-dimensional structure begins. The lateral first moment's form factor, $F_1(g) = i(g h \cos g h - \sin g h)/(g h)^2$, does not vanish on the reciprocal lattice ($F_1 = i(-1)^p/p\pi$), so the gradient channels couple through the evanescent orders. And at oblique incidence the mean's own zeros shift, $\text{sinc}((2\pi p/d + k_x)h) = (-1)^p k_x h/p\pi + O(k_x^2 h^2)$, so the tiling identity holds only to $O(k_{\parallel} h)$.

Oblique incidence. The scheme is built spectrally. By Poisson summation the Bloch coupling between planes i and j at lateral wavenumber $\mathbf{k}_{\parallel}$ is

$$C_{ab}(i, j) = \frac{1}{d^2} \sum_{\mathbf{g}} L_a(\boldsymbol{\kappa}) L_b(-\boldsymbol{\kappa}) \iint \phi_a(\mathbf{z}) \hat{\mathbf{G}}(\boldsymbol{\kappa}; \mathbf{z} - \mathbf{z}') \phi_b(\mathbf{z}') d\mathbf{z} d\mathbf{z}', \quad \boldsymbol{\kappa} = \mathbf{k}_{\parallel} + \mathbf{g}, \quad (33)$$

with L_a the lateral Legendre factors. The whole-space tensor decays as $1/k_z^2$ and the nine-component state adds at most two powers of k_z , so each entry of $\hat{\mathbf{G}}(\boldsymbol{\kappa}; z)$ is a $\delta(z)$ of constant weight plus, per mode, $[\alpha/2\gamma + i\beta \operatorname{sgn}(z)/2] e^{i\gamma|z|}$, $\gamma^2 = k_W^2 - \kappa^2$: every z integral is closed-form. Double averaging makes the lattice sum converge without Ewald splitting, to three digits at $|p|, |q| \leq 8$. Against the exact layer, for a P wave incident at 20° and 40° (`ContinuumLimit_Oblique.wl`; Table 4). The mean-only voxel is second order and the first-moment voxel fourth order in both

Table 4. Relative error of the reflected P (R_{PP}) and converted S (R_{PS}) waves for a P wave incident at 20° , with the mean-only and first-moment voxels of Eq. (33). At 40° the orders are again 2 and 4.

n	mean only		mean and first moments	
	R_{PP}	R_{PS}	R_{PP}	R_{PS}
1	9.3×10^{-4}	1.7×10^{-3}	7.2×10^{-8}	4.1×10^{-7}
2	2.3×10^{-4}	4.2×10^{-4}	4.5×10^{-9}	2.5×10^{-8}
8	1.5×10^{-5}	2.7×10^{-5}	1.7×10^{-11}	9.9×10^{-11}
order	2	2	4	4

the reflected P and the converted S wave: the breakdown of the tiling identity at $O(k_{\parallel}h)$ and the evanescent coupling of the gradient channels do not cost an order. At $\mathbf{k}_{\parallel} = 0$ the construction reproduces Eq. (30) to all printed digits.

Incident S. Nothing in the construction is specific to P. For a normally incident S wave the layer is the one-dimensional problem of Eq. (30) with $M_P \rightarrow \mu$, $k_P \rightarrow k_S$ and $\Delta M \rightarrow \Delta\mu$; the collocation and mean-only schemes converge at order 2.00 and the first-moment voxel at 4.00, in reflection and transmission, and the full 9×9 construction with an incident SH wave reproduces that one-dimensional result to all printed digits. At oblique incidence, against the exact layer (`ContinuumLimit_IncidentS.wl`), the orders are those of Table 5. The orders, and with $k = k_S$

Table 5. Observed order of convergence ($n = 2 \rightarrow 8$, $\omega = 300 \text{ rad s}^{-1}$) for incident SV and SH waves, with the mean-only and first-moment voxels. The SV angles lie below the critical angle $\arcsin(\beta/\alpha) = 36.9^\circ$.

	mean only	mean and first moments
SV at 20° : R_{SS} / R_{SP}	2.00 / 2.00	4.00 / 3.99
SV at 30° : R_{SS} / R_{SP}	1.98 / 2.00	3.97 / 4.00
SH at 20° and 40° : R_{SH}	2.00	4.00

the closed-form constants of Eqs. (28) and (31), are those of the incident P wave.

The spectral construction is implemented twice. The second implementation (`scripts/crosscheck_first_moment_voxel.py`) shares only the definition of the scheme: it takes the depth dependence of the transformed kernel from an independent residue construction, obtains the local weight as a numerical large- k_z limit of the inverted Christoffel matrix, evaluates every double integral by quadrature and separates P from S by sampling the upgoing field at two depths. The raw reflected displacements of the two agree to 9×10^{-16} in all 36 cases compared: incident P, SV and SH, normal incidence and 20° , $n = 1, 2, 4$, with and without the first moments. The one-dimensional scheme of any degree is also implemented twice. The second implementation (`scripts/crosscheck_second_moment_voxel.py`) evaluates every integral by Gauss quadrature,

split at the receiver within a voxel, and the exact layer from its own interface system, in double precision. Its errors agree with the 40-digit ones within its round-off floor (3×10^{-14} in reflection, 1.6×10^{-13} in transmission); at $\omega = 1200 \text{ rad s}^{-1}$ it measures orders 4.02 / 3.99 and 6.02 / 6.00; for one plane as the layer thins its differences from the exact layer fall as $D^{3.00}$, $D^{5.00}$ and $D^{7.00}$ for $p = 0, 1$ and 2 ; and it rederives c_p for $p \leq 5$. Figure 5(a,b) shows the convergence, with the orders fitted to it.

9 A heterogeneous layer

A heterogeneous medium is not a layer of identical cells. The simplest heterogeneous case with an exact answer is a stratified model built from layers of constant properties, with real interfaces between them: the form of many layered crustal velocity models. The last test takes such a model. A layer $D = 4 \text{ m}$ thick is divided into eight model layers of 0.5 m , and in each one $\Delta\lambda$, $\Delta\mu$ and $\Delta\rho$ are drawn uniformly within $\pm 15\%$ of the background (the values are listed in `ContinuumLimit_Heterogeneous.wl`). The result is a random stratification with physical contrasts. The voxels refine the model's layering, with m planes of voxels in each model layer, so no voxel straddles a model interface: the natural discretisation when the model's interfaces are known. At $\omega = 1500 \text{ rad s}^{-1}$ the layer is $k_p D = 1.2$ and $k_s D = 2.0$ thick, so the field reverberates between the model interfaces. The exact answer is the product of the model layers' propagators. Its same-type reflection coefficients agree with those of the reflection–transmission recursion of Kennett [27], computed independently, to 10^{-14} . The schemes are those of §§7–8, with the contrast now changing from plane to plane; nothing else changes. The errors are in Table 6 and Fig. 5(c).

Table 6. Relative error against the exact stratified layer, for m planes of voxels per model layer (§9). At normal incidence the order is measured from $m = 4$ to 16, where it is also 2.00 for collocation and 4.00 in transmission; at 20° it is measured from $m = 2$ to 4.

		$m = 1$	$m = 2$	$m = 4$	order
P, normal, R	mean only	1.9×10^{-3}	4.7×10^{-4}	1.2×10^{-4}	2.00
	first moment	8.5×10^{-7}	5.3×10^{-8}	3.3×10^{-9}	4.00
S, normal, R	mean only	5.3×10^{-3}	1.3×10^{-3}	3.3×10^{-4}	2.00
	first moment	6.2×10^{-6}	3.9×10^{-7}	2.4×10^{-8}	4.00
P, 20° , R_{pp}	mean only	1.4×10^{-3}	3.6×10^{-4}	9.0×10^{-5}	2.00
	first moments	3.3×10^{-7}	2.1×10^{-8}	1.3×10^{-9}	4.00
P, 20° , R_{ps}	mean only	2.8×10^{-3}	6.9×10^{-4}	1.7×10^{-4}	2.00
	first moments	1.4×10^{-6}	8.8×10^{-8}	5.5×10^{-9}	4.00

The interfaces cost nothing. Every order is the one measured for the uniform layer. At a model interface the strain of the true field jumps. A voxel basis that is discontinuous from voxel to voxel represents that jump exactly when the interface lies on a voxel boundary, so the interface adds no error of its own. The only requirement is that voxel boundaries coincide with the model's interfaces. A voxel that straddled an interface would carry a contrast that is not uniform within it, a case not examined here.

What fourth order buys. For a reflection error of 10^{-6} at normal incidence (P), the mean-only voxel needs 347 planes of voxels, about 43 per model layer. The first-moment voxel needs 8, one per model layer: the model's own layering, which already gives 8.5×10^{-7} . Counting unknowns (9 per voxel against 36), that is 3123 against 288 per vertical column of voxels, a factor of eleven. In three dimensions, with the lateral grid refined as the vertical one is, the number of voxels falls by $(347/8)^3 \approx 8 \times 10^4$ and the number of unknowns by about 2×10^4 . The same model has been solved independently in Python, with a contrast for each plane: the raw reflected displacements agree with the Mathematica scheme to 1.1×10^{-15} at normal incidence and at 20° , with and without the first moments. Lateral heterogeneity needs the full three-dimensional solver; it is the first extension discussed below.

A medium that varies within a cell. The stratified model is exact on its own grid: every cell is homogeneous. A smoothly varying medium is not, and sampling it as cells of constant properties is an approximation of its own. Let the contrast of cell j be its Legendre projection to degree r , $\Delta(z) = \Delta_0 \sum_{c \leq r} f_c^{(j)} P_c(s/h)$, with the voxel's field of degree p as in §8. The product of contrast and field is again a Legendre sum, $P_c P_b = \sum_e A_{cbe} P_e$, with $A_{cbe} = (2e+1) \begin{pmatrix} c & b & e \\ 0 & 0 & 0 \end{pmatrix}^2$ the square of a Wigner $3j$ symbol, so the coupling of Eq. (30) becomes

$$\iint_j \phi_a(z) \mathbf{K}(z-z') \Delta(z') \phi_b(z') dz dz' = \sum_{c \leq r} \sum_{e \leq p+r} f_c^{(j)} A_{cbe} \left[\iint_j \phi_a(z) \mathbf{K}(z-z') \phi_e(z') dz dz' \right] \Delta_0 : (34)$$

the same kernel moments, now to degree $p+r$, and nothing else new. Against the exact graded layer, a high-precision solution of $(M(z)u')' + \omega^2 \rho(z)u = 0$, the order of convergence is

$$\min(2p+2, 2r+2), \quad (35)$$

or $2p+2$ when degree r represents the profile exactly (Table 7; `ContinuumLimit_GradedContrast.wl`). A contrast that is constant in each cell therefore caps every voxel at second order, however rich its field basis, and for the smooth profile the voxel of degree 2 with a linear contrast has the error of the voxel of degree 1 to within 0.06%. At first order in the contrast this is what one expects: the contrast's projection error is orthogonal to the polynomials of degree r , so against the smooth product of the incident and adjoint fields it contributes at $O(h^{2r+2})$. To keep order $2p+2$ the contrast must be represented to the degree of the field. The closed forms of the construction are tested

Table 7. A contrast that varies within the cells: the field of each voxel to Legendre degree p , its contrast to degree r . Relative reflection error at $n = 16$ planes and the order ($n = 4 \rightarrow 8 \rightarrow 16$), $\omega = 300 \text{ rad s}^{-1}$, for the profiles $1 + (2z/D - 1)/2$ (linear) and $1 + \frac{1}{2} \sin(2\pi z/D)$ (smooth) times the contrast of §7. Predicted: Eq. (35), and $2p+2$ for the linear profile once $r \geq 1$.

p	r	linear profile		smooth profile	
		error	order	error	order
0	0	7.7×10^{-5}	2.00		
1	0	7.6×10^{-5}	2.00	5.5×10^{-5}	1.95
1	1	6.4×10^{-11}	4.00	1.4×10^{-7}	3.95
2	1	2.3×10^{-17}	6.00	1.4×10^{-7}	3.95
2	2			1.6×10^{-10}	5.98

one by one by an independent implementation (`scripts/crosscheck_graded_contrast.py`): the Legendre moments of the kernel against $2hi^n j_n(ch)$ and 30-digit quadrature (to 2×10^{-16}), the linearisation coefficients against the $3j$ symbols (exactly), and the same-cell double integrals against 30-digit quadrature (1.4×10^{-16}). It then assembles the scheme without linearisation, evaluating the projected contrast at the quadrature nodes, against its own 30-digit solution of the graded layer; its errors agree with those of Table 7 to 10^{-3} , and at $\omega = 1200 \text{ rad s}^{-1}$ it measures the orders of Eq. (35) again.

10 A sphere: what the layer predicts in three dimensions

The layer and the stratified layer are laterally uniform. A sphere is not, and it has an exact solution for elastic scattering [28, 29]: it is the usual benchmark for volume-integral and discrete-dipole methods. This section solves it with the full three-dimensional voxel scheme, to test what the layer predicts for a laterally heterogeneous body. The sphere is voxelised by the cubes of an n_{sub}^3 grid whose centres lie inside it, and the Foldy–Lax system is solved with the uniform-field voxel (the second-order scheme of §§6–7; the first-moment voxel of §8 has not been built for finite three-dimensional scatterers) and the cell-averaged coupling, by an FFT-accelerated iterative solver. The single site is the package’s cube T -matrix, which differs from the collocation closure only through its far-field corrections (§6.4) and has the same order. A P wave is incident; the contrast is that of §7. The error is the largest difference between the scattered far field and the exact one over nine scattering angles from 11° to 169° , relative to the largest exact value, at a distance of 5×10^4 radii where the far field has converged (`scripts/pilot_sphere_voxel_vs_mie.py`).

The raw error does not converge. Table 8 shows the error for grids of 4, 6 and 8 voxels across. It does not fall steadily: at $k_s a = 0.5$ it is 6.1%, then 17%, then 2.4%. The staircase changes with every grid, and with it the volume and the shape of the scatterer. Rescaling by the ratio of the sphere’s volume to the staircase’s removes the largest swings but leaves an error of about 2% at $k_s a = 0.5$, and 8 to 11% at $k_s a = 1.0$, that does not decrease. A refinement study of the voxelised sphere measures the staircase, not the scheme.

Table 8. A sphere voxelised on grids of n_{sub} voxels across, against exact elastic scattering: relative far-field error, raw and rescaled by the ratio of the sphere’s volume to the staircase’s. It does not converge with refinement, because the staircase changes with the grid.

n_{sub}	voxels	$k_s a = 0.5$		$k_s a = 1.0$	
		raw	volume-rescaled	raw	volume-rescaled
4	32	6.1×10^{-2}	1.8×10^{-2}	1.19×10^{-1}	7.7×10^{-2}
6	136	1.71×10^{-1}	2.8×10^{-2}	6.8×10^{-2}	1.12×10^{-1}
8	280	2.4×10^{-2}	2.1×10^{-2}	4.6×10^{-2}	8.6×10^{-2}

Refining a fixed shape. The staircase is itself a union of cubes, so it has its own exact solution, and the voxel scheme converges to it when each voxel is split into m^3 sub-voxels of the same shape (Fig. 4). The change of the far field under that refinement is the discretisation




Figure 4. Refining a fixed shape. (a) A cross-section through the $n_{\text{sub}} = 8$ voxelisation of a sphere (red circle): the voxels whose centres lie inside it. (b) The same staircase with every voxel split into $2 \times 2 \times 2$ sub-voxels ($m = 2$): the shape, and so its exact solution, is unchanged, and only the discretisation is refined.

Alt text: Two panels, each a square grid of shaded cells approximating a circle drawn in red; the shaded cells form a stepped outline that crosses the circle. In the right panel every shaded cell is divided into four smaller cells by thin lines, while the stepped outline is exactly the same as on the left.

error alone; the difference between the converged staircase field and the sphere's is the shape error alone. Table 9 refines the $n_{\text{sub}} = 8$ staircase to $m = 2$ and 3, up to 7560 voxels (`scripts/pilot_sphere_staircase_refinement.py`). The discretisation error converges at order 1.99 at $k_s a = 0.5$ and 2.06 at $k_s a = 1.0$ (Fig. 5d): the second order of the uniform-field voxel on the layer, now on a laterally heterogeneous body. At the original grid it is 5.4×10^{-4} and 2.7×10^{-3} , larger by a factor of about five when the frequency doubles, close to the factor four of its $(kh)^2$ scaling. The shape error, from the field extrapolated to $m \rightarrow \infty$, is 2.36×10^{-2} and 4.83×10^{-2} : 44 and 18 times the discretisation error. On a sphere the staircase dominates the error of the scheme, which is why the layer, where the staircase is exact, is the test that measures the scheme itself.

Table 9. The $n_{\text{sub}} = 8$ staircase with every voxel split into m^3 sub-voxels: far-field error against the exact sphere, and the change from the previous level. The change is the discretisation error; its order is measured from the two changes, and the field extrapolated to $m \rightarrow \infty$ gives the shape error.

m	voxels	$k_s a = 0.5$		$k_s a = 1.0$	
		vs sphere	change	vs sphere	change
1	280	2.403×10^{-2}		4.604×10^{-2}	
2	2240	2.369×10^{-2}	4.04×10^{-4}	4.779×10^{-2}	2.05×10^{-3}
3	7560	2.364×10^{-2}	7.51×10^{-5}	4.811×10^{-2}	3.65×10^{-4}
order of the discretisation error		1.99		2.06	
shape error ($m \rightarrow \infty$)		2.36×10^{-2}		4.83×10^{-2}	

11 Discussion

Figure 5 collects the convergence of every test, with the order fitted to each by least squares (`scripts/plot_convergence_orders.py`).

Where the error of a voxel scheme lives. The results separate the two approximations a voxel scheme makes. The *coupling* between voxels makes none: averaged over the source cell it is the continuum’s cell integral exactly, between planes by the form factor’s zeros and within a plane by the tiling identity, and the self term it omits is exactly the one the single site carries. Every error therefore comes from the *single site*, from the internal field it assumes. For a uniform field that error is the midpoint rule’s, known in closed form; for a linear field it is the next term of the same expansion. The converse is the practical warning. A kernel that is not the cell integral (the point Green’s tensor, or a cell average that does not match the single site’s closure) leaves a local term that does not vanish as the voxels shrink (§4): the evanescent lattice orders carry a static contribution of fixed size per voxel volume. A scheme built that way converges to the wrong limit, and refinement cannot reveal it, because the error is scale-invariant.

What the single site is. The moment hierarchy of §6 turns the uniform-field single site from a modelling choice into a derived object with a stated error. Its static part is exact for a sphere, and for a cube it differs from the sphere’s only by the cubic split of Eq. (24). Its dynamic part departs from exact scattering by a sphere only through the real $(ka)^2$ term of Eq. (25); the radiation damping is exact. That term and the collocation error of Eq. (28) have the same origin, the variation of the field across the scatterer, and the same remedy, a richer internal field. In a lattice the remedy is cheap: the field a voxel must represent is the smooth continuum field, not the field of an isolated cube, with its concentrations at the edges and corners.

Relation to the discrete dipole approximation. The same partition appears in the electromagnetic literature as the choice of polarisability. The Clausius–Mossotti relation, the lattice dispersion relation [6] and the integrated Green tensor [7] all adjust the single site so that the lattice reproduces the continuum. For space-filling cubes the adjustment is neither needed nor a free parameter: the tiling identity fixes the local field exactly, the self term of the collocation closure is the one that makes it cancel, and what is left is a closed-form error rather than an empirical one. The second-order convergence reported for the DDA by Yurkin et al. [1] corresponds to the uniform-field voxel’s, and the opposite-signed constants of Eq. (28) explain why a polarisability correction can improve one observable only at the expense of another, as the package’s own form-factor corrections do in reflection and transmission. Going beyond second order requires a richer internal field, which is the point of higher-order bases for volume integral equations; the first-moment voxel is the lowest such basis, and in its spectral form the Galerkin double average makes the lattice sum converge without Ewald splitting.

The cost of fourth order. A first-moment voxel carries four times as many unknowns as a uniform one (36 against 9 in three dimensions). The closed forms show what that buys. To

reach an error ε in reflection, the uniform voxel needs $kd = \sqrt{8\varepsilon}$ and the first-moment voxel $kd = (720\varepsilon)^{1/4}$; at $\varepsilon = 10^{-6}$ that is $kd = 2.8 \times 10^{-3}$ against 0.16, voxels 58 times larger in each direction. In a three-dimensional model the voxel count falls by $58^3 \approx 2 \times 10^5$ and the number of unknowns by about 5×10^4 (these are the Born-order constants; at the contrast studied the first-moment voxel's measured errors lie within 5% of them, while the uniform voxel's transmission constant grows with the contrast, so the comparison is if anything conservative). The stratified model of §9 measures it directly: 347 planes of uniform-field voxels against 8 first-moment voxels for the same accuracy. Higher moments continue the trend in one dimension: by Eq. (32) the second-moment voxel reaches $\varepsilon = 10^{-6}$ at $kd = 0.68$ and the third-moment voxel at $kd = 1.5$, where a voxel is already a quarter of a wavelength. A three-dimensional voxel of degree p has not been built; with the Legendre polynomials of total degree at most p it would carry 10 ($p = 2$) or 20 ($p = 3$) weights per state component, against 4 for the first-moment voxel. The gain is largest precisely where it matters: when the model varies on scales that are coarse compared with the accuracy the wavefield must have.

Limitations and outlook. The results concern laterally uniform media, a homogeneous layer and a randomly stratified one, for which the discrete system is exactly a chain of planes; a laterally heterogeneous medium needs the first-moment voxel in a full three-dimensional solver, where the gradient channels couple through the evanescent lattice orders (§8) and the tiling identity holds only to $O(k_{\parallel}h)$. The sphere of §10 shows that the second order of the uniform-field voxel carries over to a laterally heterogeneous body; the fourth order has still to be shown there. A contrast that varies within a cell, which §9 shows is needed beyond second order in a smoothly varying medium, has a further consequence in three dimensions: its linear part is odd about the cell's centre, so it couples the single site's uniform-field and first-moment blocks, which parity keeps apart for a homogeneous cube (§6.3). A graded voxel therefore needs at least the first-moment block; it has not been built. The numerical evidence is at moderate, physically typical contrasts (up to 15% in density and moduli) and at $k_s h \leq 0.3$, the range in which the single-site T -matrix is validated; the closed forms are first order in the contrast, and the departures from them at the full contrast are of second order. What is established is the *order* of convergence and its leading constant, not the behaviour at strong contrast, where the single site's own closure may limit the accuracy before the lattice coupling does. The single site's dynamic error is in closed form for a sphere only; for a cube, for which no exact solution exists, it is known through the layer (§§7–8). Two extensions follow directly. The first is the three-dimensional solver itself. The second is a stratified *background*, a layered reference medium, as distinct from the stratified contrast of §9. There the lattice coupling then carries the reverberation of the layering, which is not translation invariant in depth, and whether a tiling identity holds for it is the natural next question.

12 Conclusions

A lattice of cubic voxels reproduces a continuous elastic medium in a precise sense. Because cubes tile the plane, the source-cell-averaged coupling between voxels is the continuum's cell integral exactly, and at normal incidence the lattice kernel is available in closed form with no

lattice sum. The single site is derived rather than assumed. The moments of the Green's tensor over the cube, taken as distributions, give a closed hierarchy whose first-gradient truncation is the uniform-strain closure. That closure is exact for a sphere in the static limit, and it departs from exact scattering by a sphere only by a real $(ka)^2$ term, with the radiation damping exact. For a cube it differs from the sphere only by a cubic anisotropy of the shear response. The collocation closure carries the same self term as the lattice kernel, which cancels, so the discrete system is a collocation of the continuum equation. It converges at second order, with a leading error that is a midpoint-rule error of the internal field, known in closed form for every discretisation, and opposite in sign in reflection and transmission. A voxel that also carries the first moment of its internal field converges at fourth order, at normal incidence (where the three-dimensional scheme reduces exactly to one dimension) and at oblique incidence, with a leading error, $-(kd)^4/720$, that no longer depends on the direction of observation. With Legendre moments to degree p it converges at order $2p + 2$ with the closed-form constant of Eq. (32), and a single plane of such voxels reproduces the thin-layer series term by term through $O(D^{2p+2})$. Where the medium varies within a cell, a contrast sampled as constant caps every voxel at second order; with the contrast projected to degree r the order is $\min(2p + 2, 2r + 2)$, and the coupling needs nothing beyond the same kernel moments to degree $p + r$. The orders survive the interfaces of a randomly stratified model, provided voxel boundaries coincide with its interfaces, and there the first-moment voxel reaches an accuracy of 10^{-6} with one voxel per model layer, where the uniform-field voxel needs about forty. On a voxelised sphere the discretisation error is again second order, and the staircase approximation of the surface dominates it, by factors of 18 to 44. The results, and the evidence for each, are collected in Table 10.

Every numerical value is reproducible from the Mathematica notebooks `Mathematica/ContinuumLimit_X.wl`, for X in Reference, ThinLayer, SpecularSums, StaticLocal, AveragedSums, CubeT, Chain, ErrorConstant, FourthOrder, FourthOrderError, Reduction3D, Oblique, IncidentS, Heterogeneous, SecondMoment, GradedContrast and Figure (data for the figure), and their reference dumps `scripts/continuum_limit_*.py`; the moment integrals are `Mathematica/CubeMomentCore.wl` (gate `CubeMomentCoreTest.wl`), `CubeA22Block.wl`, `CubeMomentArchiveCheck.wl` and `CubeShearSplit.wl`, with the Python half of the last in `scripts/gate_cube_shear_split.py`; the stratified reference against Kennett's recursion is `scripts/gate_heterogeneous_reference_vs_kennett.py`; the voxelised sphere is `scripts/pilot_sphere_voxel_vs_mie.py` and `scripts/pilot_sphere_staircase_refinement.py`; the test against exact Mie is `scripts/gate_sphere_closure_vs_mie.py`, with the closed forms in `Mathematica/SphereClosureDynamic.wl`; the exact kernel is `cubic_scattering/cell_averaged_lattice.py`; the independent implementation of the first-moment voxel is `scripts/crosscheck_first_moment_voxel.py`, and of the voxel of any degree `scripts/crosscheck_second_moment_voxel.py`, and of the graded cell `scripts/crosscheck_graded_contrast.py`; the convergence figure is drawn by `scripts/plot_convergence_orders.py`. All are archived at <https://doi.org/10.5281/zenodo.23003241> and developed at <https://github.com/tmnestor/MultipleScatteringCalculations>

Acknowledgments

This work began thirty years ago in the author's doctoral research at the Research School of Earth Sciences of the Australian National University, Canberra [9].

Table 10. The results of the paper, and the evidence for each.

Result	Evidence
Continuous layer, exact and to $O(D^k)$	reference to 2.5×10^{-8} ; series order k
Thin-layer series, Eq. (5)	each coefficient symbolically, to D^4
Point-kernel local term survives $d \rightarrow 0$	11× the continuum term at $m = 1$
Averaged coupling between planes = continuum× sinc	2.8×10^{-16} , 81 entries
Tiling identity, Eq. (11)	21/21 components to 4×10^{-9}
Exact kernel, Eq. (12)	replaces a 9.4×10^{-4} static error
Moments as distributions, Eq. (19)	sum rule (20) exact; M, Q equal an independent archive, 81/81
Channels (21)–(23)	a production code by another route, 2×10^{-15}
Cube vs sphere, Eq. (24)	symbolic; and numerically at five ν , 4×10^{-16}
Ball closure = exact Mie, static	2×10^{-10} ; cube’s bulk and shear average 10^{-8}
Ball closure vs Mie, dynamic, Eq. (25)	symbolic; numerical Mie 5×10^{-4} ; $(ka)^3$ terms exact
Self term cancels, Eq. (27)	4×10^{-16}
Package T = closure statically	2×10^{-5} at $k_s h = 0.01$
Second-order convergence	order 2.000, three frequencies
Closed-form error, Eq. (28)	$n = 1$ –16, to 10^{-4} of the error
First-moment voxel, Eq. (30)	order 4.00, three frequencies
First-moment voxel, oblique, Eq. (33)	order 4, P and S, 20° and 40°
Incident SV and SH, all schemes	orders 2 and 4; constants with k_s to 4×10^{-4}
First-moment error $-(kd)^4/720$, Eq. (31)	$n = 1$ –8, to 10^{-4} of the error
Voxel of degree p : order $2p + 2$	6.00 ($p = 2$, three frequencies), 8.00 ($p = 3$)
Constant c_p , Eq. (32)	symbolic, $p = 0$ –5, twice; chain to 4×10^{-7}
One plane of degree p = thin-layer series to D^{2p+2}	exact series, $p = 0$ –3, two layers; $D^{3.00}$, $D^{5.00}$, $D^{7.00}$
Independent implementation, any degree	within round-off; orders 4.02, 6.02
Independent implementation of Eq. (33)	1.1×10^{-15} , 40 cases: P, SV, SH, stratified
Random stratified layer, §9	orders 2 and 4, P and S, normal and 20°
Stratified reference against Kennett’s recursion	10^{-14}
Contrast varying within a cell, Eq. (35)	8 cases, linear and smooth profiles; closed forms tested one by one to 2×10^{-16}
Sphere, fixed staircase refined, §10	order 1.99 and 2.06; shape error 18–44×

References

- [1] M. A. Yurkin, V. P. Maltsev, and A. G. Hoekstra. Convergence of the discrete dipole approximation. I. Theoretical analysis. *J. Opt. Soc. Am. A*, 23:2578–2591, 2006.
- [2] L. L. Foldy. The multiple scattering of waves. I. General theory of isotropic scattering by randomly distributed scatterers. *Phys. Rev.*, 67:107–119, 1945.
- [3] M. Lax. Multiple scattering of waves. *Rev. Mod. Phys.*, 23:287–310, 1951.
- [4] E. M. Purcell and C. R. Pennypacker. Scattering and absorption of light by nonspherical dielectric grains. *Astrophys. J.*, 186:705–714, 1973.
- [5] B. T. Draine and P. J. Flatau. Discrete-dipole approximation for scattering calculations. *J. Opt. Soc. Am. A*, 11:1491–1499, 1994.
- [6] B. T. Draine and J. Goodman. Beyond Clausius–Mossotti: wave propagation on a polarizable point lattice and the discrete dipole approximation. *Astrophys. J.*, 405:685–697, 1993.
- [7] P. C. Chaumet, A. Sentenac, and A. Rahmani. Coupled dipole method for scatterers with large permittivity. *Phys. Rev. E*, 70:036606, 2004.
- [8] J. D. Eshelby. The determination of the elastic field of an ellipsoidal inclusion, and related problems. *Proc. R. Soc. Lond. A*, 241:376–396, 1957.
- [9] T. M. Nestor. *Seismic Wave Propagation in Complex Media: Summing the Multiple Scattering Series*. PhD thesis, Research School of Earth Sciences, Australian National University, Canberra, ACT, Australia, 1996.
- [10] E. M. Stein and G. Weiss. *Introduction to Fourier Analysis on Euclidean Spaces*. Princeton University Press, Princeton, NJ, 1971.
- [11] C. M. Linton. Lattice sums for the Helmholtz equation. *SIAM Rev.*, 52(4):630–674, 2010. doi: 10.1137/09075130X.
- [12] J. D. Eshelby. The elastic field outside an ellipsoidal inclusion. *Proc. R. Soc. Lond. A*, 252: 561–569, 1959. doi: 10.1098/rspa.1959.0173.
- [13] T. Mura. *Micromechanics of Defects in Solids*. Martinus Nijhoff, Dordrecht, 2nd edition, 1987. doi: 10.1007/978-94-009-3489-4.
- [14] Y. Mikata and S. Nemat-Nasser. Elastic field due to a dynamically transforming spherical inclusion. *J. Appl. Mech.*, 57:845–849, 1990. doi: 10.1115/1.2897650.
- [15] T. M. Michelitsch, H. Gao, and V. M. Levin. Dynamic Eshelby tensor and potentials for ellipsoidal inclusions. *Proc. R. Soc. Lond. A*, 459:863–890, 2003. doi: 10.1098/rspa.2002.1054.
- [16] Y. P. Chiu. On the stress field due to initial strains in a cuboid surrounded by an infinite elastic space. *J. Appl. Mech.*, 44:587–590, 1977. doi: 10.1115/1.3424140.

- [17] J. K. Lee and W. C. Johnson. Calculation of the elastic strain field of a cuboidal precipitate in an anisotropic matrix. *Phys. Status Solidi A*, 46:267–272, 1978. doi: 10.1002/pssa.2210460134.
- [18] G. J. Rodin. Eshelby’s inclusion problem for polygons and polyhedra. *J. Mech. Phys. Solids*, 44:1977–1995, 1996. doi: 10.1016/S0022-5096(96)00066-X.
- [19] H. Nozaki and M. Taya. Elastic fields in a polyhedral inclusion with uniform eigenstrains and related problems. *J. Appl. Mech.*, 68:441–452, 2001. doi: 10.1115/1.1362670.
- [20] A. K. Mal and L. Knopoff. Elastic wave velocities in two-component systems. *J. Inst. Math. Appl.*, 3:376–387, 1967. doi: 10.1093/imamat/3.4.376.
- [21] J. E. Gubernatis, E. Domany, and J. A. Krumhansl. Formal aspects of the theory of the scattering of ultrasound by flaws in elastic materials. *J. Appl. Phys.*, 48:2804–2811, 1977. doi: 10.1063/1.324141.
- [22] S. K. Kanaun and V. M. Levin. *Self-Consistent Methods for Composites, Vol. 2: Wave Propagation in Heterogeneous Materials*. Solid Mechanics and Its Applications. Springer, Dordrecht, 2008. doi: 10.1007/978-1-4020-6968-0.
- [23] R.-S. Wu and K. Aki. Scattering characteristics of elastic waves by an elastic heterogeneity. *Geophysics*, 50:582–595, 1985. doi: 10.1190/1.1441934.
- [24] P. C. Waterman. Matrix theory of elastic wave scattering. *J. Acoust. Soc. Am.*, 60:567–580, 1976. doi: 10.1121/1.381130.
- [25] V. D. Kupradze. *Dynamical Problems in Elasticity*, volume 3 of *Progress in Solid Mechanics*. North-Holland, Amsterdam, 1963.
- [26] K. Aki and P. G. Richards. *Quantitative Seismology*. University Science Books, Sausalito, CA, 2nd edition, 2002.
- [27] B. L. N. Kennett. *Seismic Wave Propagation in Stratified Media*. Cambridge University Press, Cambridge, 1983.
- [28] C. F. Ying and R. Truell. Scattering of a plane longitudinal wave by a spherical obstacle in an isotropically elastic solid. *J. Appl. Phys.*, 27(9):1086–1097, 1956. doi: 10.1063/1.1722545.
- [29] Y.-H. Pao and C.-C. Mow. *Diffraction of Elastic Waves and Dynamic Stress Concentrations*. Crane Russak, New York, 1973.




Figure 5. Convergence of every test, with least-squares regression lines. Each line fits $\log e = c - p \log n$ to the filled points; open points are pre-asymptotic and are not fitted. The legend gives the order p and its standard error. Colour marks the scheme: collocation (§6), the mean-only voxel and the voxels carrying Legendre moments to degree $p = 1, 2$ and 3 (Eqs. (30) and (33)). (a) The uniform layer, P reflection at normal incidence ($D = 2$ m, $\omega = 300$ rad s $^{-1}$, the contrast of §7), fitted from $n = 4$; the second- and third-moment voxels, computed in 40-digit arithmetic, fitted from $n = 2$. (b) The uniform layer at 20°, fitted from $n = 2$: the S wave converted from incident P (solid) and reflected from incident SV (dashed). (c) The stratified layer of §9, reflection at normal incidence for P (solid) and S (dashed), against the number m of voxel planes per model layer, fitted from $m = 4$. (d) The voxelised sphere of §10: the far-field difference between the $n_{\text{sub}} = 8$ staircase with every voxel split into m^3 sub-voxels and the staircase field extrapolated to $m \rightarrow \infty$, fitted over $m = 1, 2, 3$. That limit is extrapolated from the same three levels, so panel (d) checks consistency; its fitted orders differ slightly from those of Table 9 (1.99 and 2.06), which are measured from the changes between levels.

Alt text: Four log-log panels of relative error against refinement, each showing data points and fitted straight lines with the fitted order in the legend. (a) Uniform layer at normal incidence, 1 to 32 voxel planes: the collocation and mean-only lines have order 2.000, the first-moment line order 4.000, falling from about 3×10^{-7} to 2×10^{-13} ; the second-moment line has order 6.000 and the third-moment line order 8.000, over 1 to 8 planes, reaching about 10^{-22} . (b) Uniform layer at 20 degrees, 1 to 8 planes: P-to-S and SV-to-S lines, order 2.00 for the mean-only voxel and 4.00 for the first-moment voxel. (c) Stratified layer, 1 to 16 planes per model layer: P and S reflection, order 2.000 for collocation and mean-only, 4.000 for first moment. (d) Voxelised sphere, 1 to 3 sub-voxels per voxel edge: two lines of order about 2, for $k_{sa} = 0.5$ and 1.