

Highlights

The continuum limit of a space-filling layer of cubes: a tiling identity, an exact lattice kernel and the closed-form error of the discrete layer

T. M. Nestor

- Voxel coupling averaged over the source cell equals the continuum's exactly
- The cube's single-site T-matrix follows from distributional moment integrals
- The voxel scheme's leading error is known in closed form for every grid
- A first-moment voxel converges at fourth order for P, SV and SH waves
- In a random stratified layer, one voxel per model cell reaches 1e-6 accuracy

The continuum limit of a space-filling layer of cubes: a tiling identity, an exact lattice kernel and the closed-form error of the discrete layer

T. M. Nestor

Independent researcher, Australia

Abstract

We analyse the convergence of voxel discretisations of the elastic volume integral equation, in which a medium is divided into cubic cells, each represented by a single-site T -matrix and coupled through the background Green's tensor: the elastic counterpart of the discrete dipole approximation. Because cubes tile space, a layer of identical voxels is exactly a homogeneous layer, so it isolates the discretisation error, free of modelling and shape error. We show that (i) the cube's lateral form factor vanishes on the reciprocal lattice, so the source-cell-averaged coupling equals the continuum's and the lattice kernel at normal incidence is in closed form; (ii) the single-site T -matrix follows from distributional moments of the Green's tensor and, against exact scattering by a sphere, is exact statically and departs only by a real, closed-form $(ka)^2$ term; (iii) its self term cancels, so the scheme is a collocation of the continuum equation, of second order with leading error $+(kd)^2/8$ in reflection and $-(kd)^2/24$ in transmission, which no scalar correction can remove; (iv) a Galerkin voxel carrying the first moment of its internal field converges at fourth order, with leading error $-(kd)^4/720$, for incident P, SV

and SH waves at normal and oblique incidence, and its spectral lattice sum needs no Ewald splitting. In a randomly stratified layer the orders are unchanged, and the fourth-order voxel reaches an accuracy of 10^{-6} with one voxel per model cell. Every result is derived symbolically and verified by an independent implementation or an exact solution.

Keywords: Elastic wave scattering, Volume integral equation, Foldy–Lax multiple scattering, Discrete dipole approximation, Galerkin method, Convergence analysis

1. Introduction

The problem this paper serves.. Seismic waves in the crust are scattered by heterogeneity on every scale, and the models that describe it — tomographic images, full-waveform inversions, geological models — arrive as grids of cells of constant properties. A scattering calculation that uses such a model directly treats every cell as a scatterer: a voxel carrying a single-site T -matrix and coupled to every other voxel through the Green’s tensor of the background, with the multiple-scattering series summed. For the media such a scheme is built for, there is no exact solution, and that is the difficulty. Its error cannot be measured against the answer. Comparison with another numerical method mixes the errors of both. And a refinement study can mislead, because a voxel scheme can converge smoothly to the wrong limit. §4 shows how that happens with the most natural choice of coupling, the Green’s tensor evaluated between voxel centres: it carries a local term whose size per unit volume does not shrink as the voxels do, and unless the single voxel’s response is matched to cancel it, refinement converges to the wrong

17 answer without revealing it.

18 *Why a homogeneous layer..* This paper therefore studies the one configura-
19 tion in which a voxel scheme runs in full — every voxel, every coupling, the
20 whole multiple-scattering series, at normal and at oblique incidence — and
21 the exact answer is nevertheless known: a laterally uniform layer filled with
22 identical voxels. Because cubes tile space, such a layer does not approximate
23 a homogeneous layer; it *is* one. Any difference between its reflection and
24 transmission and those of the continuous layer is therefore error of the dis-
25 cretisation alone, with no modelling error and no shape error behind which
26 to hide. A sphere, the usual test object for volume-integral and discrete-
27 dipole methods [1], cannot give this: its error mixes the discretisation with
28 the staircase approximation of a curved surface by cubes. The layer plays the
29 part of the patch test in finite elements. A scheme that cannot reproduce a
30 uniform medium correctly cannot be trusted in a heterogeneous one. What
31 the layer reveals about the scheme’s two ingredients — the coupling between
32 voxels and the response of a single voxel — does not depend on the medium,
33 because the same two ingredients are used wherever the scheme is applied.

34 *What the reader gains..* Four things that a heterogeneous test cannot pro-
35 vide. First, a proof, exact at normal incidence, that the voxel scheme con-
36 verges to the continuum, and the specific ingredient — a coupling averaged
37 over the source voxel — without which it converges to the wrong answer. Sec-
38 ond, the leading error in closed form for every discretisation, so that a grid
39 can be chosen for a required accuracy before the calculation is run. Third,
40 a single-site T -matrix derived from first principles, with its error stated and
41 checked against the exact solution for a sphere. Fourth, a voxel that converges

42 at fourth order, also in a randomly stratified model (§9): for a reflection ac-
43 curacy of 10^{-6} it allows voxels 58 times larger in each direction than the
44 standard one (§10).

45 *Why cubic voxels..* Cubes are the natural cell for three reasons. They *tile*
46 *space*, so a uniform region discretised into cubes is still that uniform region,
47 whereas spheres or point scatterers can only approximate it and carry a pack-
48 ing error of their own. Their *regular lattice* makes the interaction between
49 cells translation invariant, so the coupling is a convolution — Toeplitz be-
50 tween planes, Bloch-periodic within them — and the whole system is solved
51 by fast transforms. And a cube small compared with the wavelength has a
52 *single-site response in closed form*, the elastic depolarisation of a cube (§6),
53 so the only approximation a voxel introduces is the internal field it assumes.

54 *Multiple scattering on a lattice..* Replacing a heterogeneous body by a lattice
55 of small scatterers that interact through the background Green’s function
56 is the multiple-scattering description of Foldy [2] and Lax [3]. In electro-
57 magnetics its discretised form is the discrete dipole approximation [DDA;
58 4, 5], whose accuracy has been studied in detail: each dipole’s polarisability
59 must account for the local field of the surrounding lattice, which the lattice
60 dispersion relation fixes so that an infinite lattice propagates waves as the
61 continuum does to a given order in kd [6]; the error of the method converges
62 at second order in the dipole spacing when shape errors are absent [1]; and
63 integrating the Green’s tensor over each cell rather than sampling it at the
64 centre improves accuracy, most of all for strong contrasts [7]. The elastic
65 problem adds a structure the scalar and electromagnetic ones lack. A cubic

66 voxel responds to the incident field with a force (through its density con-
67 trast) and a moment (through its stiffness contrast), so the state at a voxel
68 has nine components — displacement and strain — and the single site’s self-
69 interaction is the elastic depolarisation of a cube, the discrete counterpart of
70 the inclusion tensor of Eshelby [8].

71 *This paper..* The approach originates in the two-and-a-half-dimensional for-
72 mulation of Nestor [9]: a medium divided into cells that interact through the
73 background Green’s tensor, each carrying a single-site contrast renormalised
74 by its own near field, with the multiple-scattering series summed. This pa-
75 per extends it to cubic voxels. §2 sets out the discrete medium and §3 the
76 continuous layer it must reproduce. §4 shows why the point coupling fails,
77 and §5 why the source-averaged coupling succeeds exactly. §6 derives the
78 single-site T -matrix from moment integrals, tests it against the exact sphere,
79 and shows that with it the discrete system is a collocation of the continuum
80 equation. §7 gives that collocation’s second-order error in closed form, and
81 §8 the fourth-order voxel, for incident P and S at normal and oblique inci-
82 dence. §9 repeats the tests on a randomly stratified layer, and §10 discusses
83 what the results imply for voxel schemes in general. Every result is derived
84 symbolically in Mathematica and checked against an independent Python
85 implementation or an exact solution.

86 2. Setting

87 *The discrete medium..* Space is tiled by cubes of side d (half-side $h = d/2$).
88 Cube i carries a contrast $\Delta\lambda, \Delta\mu, \Delta\rho$ from an isotropic background with
89 Lamé parameters λ, μ , density ρ , P and S speeds α, β and P-wave modulus

90 $M_P = \lambda + 2\mu$. Time dependence is $e^{-i\omega t}$; $k_P = \omega/\alpha$, $k_S = \omega/\beta$. Coordinates
 91 are ordered (z, x, y) with z down. The field at a voxel is the nine-component
 92 state

$$\psi = (u_z, u_x, u_y, \varepsilon_{zz}, \varepsilon_{xx}, \varepsilon_{yy}, 2\varepsilon_{xy}, 2\varepsilon_{zy}, 2\varepsilon_{zx}), \quad (1)$$

93 displacement followed by engineering strain. A voxel responds with a force
 94 and a moment (stress dipole); the 9×9 single-site T -matrix maps the exciting
 95 state to them, and the 9×9 lattice Green's tensor $\mathbf{K}(i, j)$ maps a source
 96 voxel's force and moment to the state at voxel i . The Foldy–Lax system is

$$\psi_i = \psi_i^0 + \sum_{j \neq i} \mathbf{K}(i, j) \mathbf{T}_j \psi_j, \quad (2)$$

97 with ψ^0 the incident state. Strain rows of $\mathbf{K}$ are receiver derivatives of the
 98 point Green's tensor $\mathbf{G}$; moment columns are derivatives with respect to the
 99 receiver coordinate, symmetrised with a factor $\frac{1}{2}$ on the shear components.
 100 In that convention the contrast operator is

$$\Delta = \text{blockdiag}(\omega^2 \Delta \rho \mathbf{I}_3, \Delta \mathcal{C}), \quad (3)$$

101 where $\Delta \mathcal{C}$ is the 6×6 contrast stiffness with $\Delta \lambda + 2\Delta \mu$ and $\Delta \lambda$ in the normal
 102 block and $2\Delta \mu$ (not $\Delta \mu$) on the shear diagonal.

103 *The laterally uniform layer.* A layer of thickness D is filled with n planes
 104 of voxels, $d = D/n$, and illuminated at normal incidence. The field is then
 105 laterally uniform, only the zero-lateral-wavenumber (specular) component of
 106 the lattice coupling acts, and Eq. (2) becomes exactly a chain of n planes
 107 coupled by the specular lattice sums

$$\mathbf{S}(m) = \sum_{\mathbf{R}} \mathbf{G}(\mathbf{R} + md\hat{\mathbf{z}}), \quad (4)$$

the sum running over the square lattice $\mathbf{R}$ of pitch d , m the plane separation in pitches. The question is whether, and how fast, this chain converges to the continuous layer as $n \rightarrow \infty$ at fixed D .

Implementations.. Every result is derived symbolically in Mathematica and compared with an independent numerical implementation or an exact solution. “The package” below is the Python implementation of the voxel scheme (`cubic_scattering`): its cube T -matrix, its lattice kernels and its stratified propagator, which serve both as the object under test and, where validated, as a second implementation.

3. The continuous layer and its thin-layer expansion

At normal incidence a plane force drives a one-dimensional problem per wave type, with modulus M (M_P for P, μ for S), wavenumber k and plane Green’s function $g(z) = \frac{i}{2Mk} e^{ik|z|}$. The layer’s reflection and transmission follow from continuity of displacement and traction at its faces; they agree with the stratified propagator used as the reference in this work to 2.5×10^{-8} , P and S (`ContinuumLimit_Reference.wl`).

For a layer of thickness D with modulus M_1 and density ρ_1 in a background M_0, ρ_0 , the reflection coefficient expands as

$$R = \frac{i\omega D}{2} \sqrt{\frac{M_0}{\rho_0}} \left[\frac{\rho_1 - \rho_0}{M_0} - \rho_0 \left(\frac{1}{M_1} - \frac{1}{M_0} \right) \right] + O(D^2), \quad (5)$$

exact in the contrast (`ContinuumLimit_ThinLayer.wl`): to first order a thin layer is an interface across which the traction jumps by the excess inertia and the displacement by the excess compliance. At oblique incidence the same statement holds with the first-order system $\partial_z \mathbf{b} = \mathbf{A} \mathbf{b}$ for the P–SV state

130 $\mathbf{b} = (u_x, u_z, t_{xz}, t_{zz})$: the layer acts to $O(D)$ as the jump $\mathbf{I} + (\mathbf{A}_1 - \mathbf{A}_0)D$,
 131 which reproduces the layer's $O(D)$ reflection coefficients R_{PP}, R_{PS} to 10^{-12} ;
 132 the series truncated at $O(D^k)$ converges at exactly order k ; and a soft thin
 133 layer with shear modulus $\mu_1 = D/\eta$ at fixed compliance η tends to the linear-
 134 slip interface as $D \rightarrow 0$. These expansions are the target a single plane of
 135 voxels must reproduce order by order.

136 4. The specular sums of the point Green's tensor

137 By Poisson summation,

$$\mathbf{S}(m) = \frac{1}{d^2} \sum_{\mathbf{g}} \hat{\mathbf{G}}(\mathbf{g}; md), \quad \mathbf{g} = \frac{2\pi}{d}(p, q), \quad (6)$$

138 with $\hat{\mathbf{G}}$ the two-dimensional transform of the Green's tensor between planes.
 139 The $\mathbf{g} = 0$ term is the continuum's plane-wave term; the others are evanes-
 140 cent, decaying as $e^{-2\pi m \sqrt{p^2 + q^2}}$ independently of d . In the strain-from-moment
 141 block they carry two powers of $|\mathbf{g}| \propto 1/d$, so their contribution per voxel vol-
 142 ume does *not* vanish as $d \rightarrow 0$: at $m = 1$ it is eleven times the continuum
 143 term (`ContinuumLimit_SpecularSums.wl`, matched to the package's Ewald
 144 lattice sum to 5×10^{-13}). This surviving term is static. Summed over all
 145 planes in slab order, the static lattice sum of the point kernel is exactly the
 146 depolarisation of a uniform plate — $-1/M_P$ on $\varepsilon_{zz} \leftarrow M_{zz}$ and $-1/(2\mu)$
 147 on the normal shears, per unit volume — plus a cubic-symmetric remainder
 148 (`ContinuumLimit_StaticLocal.wl`, to 2×10^{-10}). The point kernel there-
 149 fore reproduces the continuum only if that remainder is cancelled by the
 150 voxel's own self term.

151 5. The tiling identity

152 Replace the point coupling by the *source-cell average* $\langle \mathbf{G} \rangle(\mathbf{r}) = d^{-3} \int_{\text{cube}} \mathbf{G}(\mathbf{r} -$
 153 $\mathbf{y}) d\mathbf{y}$ — the field at a voxel centre of a uniform polarisation of the source
 154 voxel. Each Poisson order in Eq. (6) is then multiplied by the cube’s form
 155 factor

$$F(\mathbf{g}) = \text{sinc}(g_x h) \text{sinc}(g_y h), \quad (7)$$

156 which vanishes at every $\mathbf{g} \neq 0$ of the reciprocal lattice since $\text{sinc}(p\pi) = 0$ for
 157 integer $p \neq 0$. This is the Fourier form of the fact that cubes tile the plane,
 158 $\sum_{\mathbf{R}} \chi(\mathbf{x} - \mathbf{R}) = 1$ for the cube’s indicator χ . Two exact statements follow.

159 *Between planes..* For $m \neq 0$ every evanescent order is removed and

$$\mathbf{S}_{\text{avg}}(m) = \frac{1}{d^2} \sum_{W=P,S} \hat{\mathbf{G}}_W(0; md) \text{sinc}(k_W h), \quad (8)$$

160 the continuum’s plane-wave term times the vertical form factor. The pack-
 161 age’s cell-averaged kernel agrees in all 81 entries to 2.8×10^{-16} (`ContinuumLimit_`
 162 `AveragedSums.wl`).

163 *Within a plane..* The sum over *all* cells of the plane, the voxel’s own included,
 164 is the field at the mid-plane of a uniform plate of thickness d :

$$d^3 \mathbf{S}_{\text{avg}}(0) + \mathbf{S}_{\text{self}} = d^3 \mathbf{P}_0, \quad \mathbf{S}_{\text{self}} = \int_{\text{cube}} \mathbf{G}(-\mathbf{y}) d\mathbf{y}, \quad (9)$$

165 where $\mathbf{P}_0 = d^{-3} \int_{-h}^h \mathbf{G}_{\text{plate}}(-z') dz'$ is elementary: in the one-dimensional
 166 transform each entry is $c_0 + A/(Mk_z^2 - \rho\omega^2)$, so it averages to $c_0/d + A(e^{ikh} -$
 167 $1)/(dMk^2)$. Its only nonzero entries are $u \leftarrow f$ on the diagonal and the
 168 three normal strain-from-moment entries, whose local weights $-1/M_P$ and
 169 $-1/(2\mu)$ are the plate depolarisation of §4. Statically the identity holds for

170 all 21 independent components of the strain-from-moment tensor to 4×10^{-9} ,
 171 with every cell integral evaluated analytically (below).

172 *Consequence: an exact kernel with no lattice sum..* At normal incidence the
 173 same-plane block of the cell-averaged Foldy–Lax kernel is

$$\mathbf{S}_{\text{avg}}(0) = \mathbf{P}_0 - \mathbf{S}_{\text{self}}/d^3, \quad (10)$$

174 with no Ewald splitting, near shell or asymptotic tail. The route it replaces
 175 — Ewald summation with a near shell averaged by quadrature and an $O(d^2)$
 176 Taylor tail beyond it — carries a static strain-from-moment error of 9.4×10^{-4}
 177 at its default settings, falling to 2.6×10^{-5} only when both the shell radius
 178 and the quadrature order are raised; Eq. (10) is now the package’s kernel at
 179 normal incidence.

180 *Cell integrals without quadrature..* Write the static (Kelvin) tensor as $G_{bc} =$
 181 $(c_A + c_B) \delta_{bc}/r - c_B \partial_b \partial_c r$ with $c_B = (1/\mu - 1/M_P)/8\pi$ and $c_A = 1/(4\pi\mu) - c_B$.
 182 Every cell integral of its derivatives is a cell integral of derivatives of $1/r$ and
 183 r . One application of the divergence theorem moves it onto faces $x_p = c_p \pm h$,
 184 which never contain the origin, and there the two-dimensional antiderivative
 185 in the face coordinates is continuous, so the face integral is a sum over the
 186 face’s corners. The distributional term at $r = 0$ is carried by the divergence
 187 theorem itself. (A triple antiderivative evaluated at the cube’s corners is
 188 *not* usable for repeated derivatives: its arctangent terms jump across the
 189 coordinate planes through the origin, and differentiating twice there discards
 190 exactly that term.) The self term so obtained reproduces the classical closed
 191 form of the cube’s depolarisation to machine precision.

192 6. The single site: the cube T -matrix from moment integrals

193 Every error of a voxel scheme is made at the single site (§10), so the single
194 site is derived here rather than imported. The classical route imports it. For
195 an ellipsoid, Eshelby’s uniformity theorem makes the internal strain under a
196 uniform load uniform, with a depolarisation tensor in closed form [8, 10, 11];
197 its dynamic generalisation is known for spheres and ellipsoids [12, 13]. For
198 a cube or a general polyhedron the static field of a uniform eigenstrain is
199 also in closed form [14, 15, 16, 17], but the strain it produces is not uniform,
200 so there is no exact uniform-strain inclusion to import. In scattering theory
201 the same object appears as the self-consistent or effective-field closure of
202 the volume integral equation [18, 19, 20]: a first-order (Born) scatterer [21]
203 renormalised by the geometric series of its own near field, the step that turns
204 a contrast into an effective contrast. Nestor [9] introduced the single site
205 in this form, summing the series of the cell’s self integral, with the singular
206 near field evaluated over an infinitesimal cell. The moments below evaluate
207 the same self integral over the finite cube, as a distribution. The elastic T -
208 matrix of Waterman [22] represents a scatterer in spherical multipoles, the
209 natural basis for a sphere; a voxel in a lattice needs one compatible with the
210 lattice, and the moments over the cube provide it. The construction below
211 makes the approximation of a voxel explicit rather than implicit: it derives
212 the cube’s T -matrix from the moments of the elastodynamic Green’s tensor
213 over the cube, and identifies the uniform-strain closure as one truncation of
214 an exact hierarchy.

215 *6.1. The closed hierarchy*

216 Let the voxel be $V = [-h, h]^3$ with contrast $\delta c_{ijkl} = \Delta\lambda \delta_{ij}\delta_{kl} + \Delta\mu(\delta_{ik}\delta_{jl} +$
 217 $\delta_{il}\delta_{jk})$ and $\Delta\rho$. Expand the internal displacement about the centre to second
 218 degree and project the Lippmann–Schwinger equation onto 1, x_p and x_px_q .
 219 The unknowns are $u_i(0)$, $\partial_p u_i(0)$ and $\partial_p \partial_q u_i(0)$, and the result is the closed
 220 system

$$(\delta_{ij} - \omega^2 \Delta\rho \mathbf{G}_{ij}) u_j + (\mathbf{N}_{in,k}^r \delta c_{nksj} - \frac{1}{2} \omega^2 \Delta\rho \mathbf{K}_{ij}^{rs}) \partial_r \partial_s u_j = u_i^0, \quad (11)$$

$$(\delta_{pr} \delta_{ij} + \omega^2 \Delta\rho \mathbf{N}_{ij,p}^r - \mathbf{M}_{in,pk} \delta c_{nkrj}) \partial_r u_j = \partial_p u_i^0, \quad (12)$$

$$-\omega^2 \Delta\rho \mathbf{M}_{ij,pq} u_j + (\delta_{qr} \delta_{ps} \delta_{ij} + \mathbf{Q}_{in,pqk}^r \delta c_{nksj} - \frac{1}{2} \omega^2 \Delta\rho \mathbf{P}_{ij,pq}^{rs}) \partial_r \partial_s u_j = \partial_p \partial_q u_i^0, \quad (13)$$

221 all fields evaluated at the centre. Its coefficients are six moments of the
 222 background Green’s tensor $\mathbf{G}$ over the voxel,

$$\begin{aligned} \mathbf{G}_{ij} &= \int_V G_{ij} \, dV, & \mathbf{K}_{ij}^{rs} &= \int_V G_{ij} x_r x_s \, dV, & \mathbf{N}_{in,k}^r &= \int_V \partial'_k G_{in} x_r \, dV, \\ \mathbf{M}_{in,pk} &= \int_V \partial'_p \partial'_k G_{in} \, dV, & \mathbf{P}_{ij,pq}^{rs} &= \int_V \partial'_q \partial'_p G_{ij} x_r x_s \, dV, & \mathbf{Q}_{in,pqk}^r &= \int_V \partial'_q \partial'_p \partial'_k G_{in} x_r \, dV, \end{aligned} \quad (14)$$

223 with derivatives taken with respect to the source point, $\partial' = -\partial$. That sign
 224 reaches $\mathbf{N}$ and $\mathbf{Q}$ only, and no symmetry, parity or scaling check can detect
 225 it, so it is fixed by convention. The moments are graded by (D, W) , the
 226 number of derivatives and the degree of the weight: $\mathbf{G}(0, 0)$, $\mathbf{K}(0, 2)$, $\mathbf{N}(1, 1)$,
 227 $\mathbf{M}(2, 0)$, $\mathbf{P}(2, 2)$, $\mathbf{Q}(3, 1)$.

228 *6.2. The moments as distributions*

229 Write the elastodynamic tensor as

$$G_{in}(\mathbf{r}) = \frac{1}{4\pi\mu} \left[\delta_{in} g_S(r) + \frac{1}{k_S^2} \partial_i \partial_n (g_S(r) - g_P(r)) \right], \quad g_c(r) = \frac{e^{ik_c r}}{r} \quad (15)$$

230 [23, 24]. Expanding g_c in powers of r reduces every moment in Eq. (14) to a
231 finite sum of scalar moments

$$E[m; d_1 \dots d_D; w_1 \dots w_W] = \langle \partial_{d_1} \dots \partial_{d_D} r^m, x_{w_1} \dots x_{w_W} \mathbf{1}_V \rangle. \quad (16)$$

232 For $m = -1$ and $D \geq 2$ the classical integrand is not absolutely integrable.
233 $\partial^2(1/r)$ decays as r^{-3} , its angular average vanishes, and the value of the inte-
234 gral depends on the shape of the region excluded around the origin. This is
235 the source of the δ -function term of the inclusion problem, which a numerical
236 quadrature misses. Without it the depolarisation is wrong; for a sphere it
237 vanishes altogether (§6.3). It is not added afterwards here. The derivatives
238 are distributional from the start, and each is moved onto the voxel's faces by
239 pairing with the test function $w\mathbf{1}_V$, using $\partial_q \mathbf{1}_V = -n_q \delta_{\partial V}$:

$$\langle \partial_q T, w\mathbf{1}_V \rangle = \oint_{\partial V} n_q w T \, dA - \langle T, (\partial_q w) \mathbf{1}_V \rangle. \quad (17)$$

240 The surface term is evaluated on the faces, a distance h from the origin, where
241 every derivative of r^m is smooth. The recursion lowers D and W together
242 and ends on convergent integrals. The face integrals are done in closed form
243 by the corner sums of §5. The sharp check is the sum rule

$$\sum_p E[-1; p, p;] = \nabla^2(1/r) \text{ paired with } \mathbf{1}_V = -4\pi, \quad (18)$$

244 which the face route satisfies exactly and a spherical excision returns as
245 0. The closed forms contain logarithms of the units $(2 \pm \sqrt{3})^n$ of the Pell

246 equation $x^2 - 3y^2 = 1$. Left unreduced, these cancel catastrophically: for
 247 example, $\log(72010600134783751 - 41575339372323900\sqrt{3})$ is a cancellation
 248 of 35 digits. They are therefore reduced to $n \log(2 - \sqrt{3})$, which is exact.
 249 Because every power of r is paired this way, the dynamic moments keep the
 250 even-power terms of the expansion (k^2r , r^3 , ...) as well as the odd ones.
 251 A route that integrates only the polynomial (odd-power) terms of $\mathbf{G}$ misses
 252 those terms, at a real $O((kh)^2)$ cost (§6.4).

253 6.3. Parity, and the nine-component block

254 The cube is centrosymmetric. A moment that couples a field quantity of
 255 degree a to one of degree b therefore vanishes unless $a + b$ is even. u , ∂u and
 256 $\partial\partial u$ have degrees 0, 1 and 2, so Eq. (12) contains the first gradient alone: ∂u is
 257 determined by its own block and couples to neither u nor $\partial\partial u$. The leading
 258 scattering by a stiffness contrast is the stress dipole $\int_V \delta c : \nabla u \, dV$, which
 259 needs only ∂u . The leading far field of the stiffness contrast is therefore fixed
 260 by Eq. (12), and *no treatment of the second-gradient block (13) can change*
 261 *it*. What Eq. (13) contributes reaches the far field only through the $O((kh)^2)$
 262 coupling in Eq. (11). That is the order, and the origin (the variation of the
 263 field across the voxel), of the collocation error of §7, which the first-moment
 264 voxel of §8 removes.

265 With $\Delta\rho = 0$, Eq. (12) is the 9×9 operator $\delta_{pr}\delta_{ij} - \mathbf{M}_{in,pk}\delta c_{nkrj}$, built
 266 from $\mathbf{M}$ alone. Under O_h the nine gradient components decompose as $A_{1g} \oplus$
 267 $E_g \oplus T_{1g} \oplus T_{2g}$, and the operator is diagonal in that basis. Its eigenvalues are

268 the inverse strain concentrations:

$$A_{1g} \text{ (dilatation)} : 1 + \frac{\Delta\lambda + \frac{2}{3}\Delta\mu}{\lambda + 2\mu}, \quad (19)$$

$$T_{2g} \text{ (off-diagonal shear)} : 1 + 2\Delta\mu S_{\text{off}}, \quad S_{\text{off}} = \frac{\pi(\lambda + 2\mu) - \sqrt{3}(\lambda + \mu)}{3\pi\mu(\lambda + 2\mu)}, \quad (20)$$

$$E_g \text{ (diagonal shear)} : 1 + 2\Delta\mu S_{\text{diag}}, \quad S_{\text{diag}} = \frac{3\sqrt{3}(\lambda + \mu) + 2\pi\mu}{6\pi\mu(\lambda + 2\mu)}. \quad (21)$$

269 The density channel is the displacement block of Eq. (11) with the second
 270 gradients dropped, an amplification $(1 - \omega^2 \Delta\rho \mathbf{G})^{-1}$ with $\mathbf{G}$ the cube integral
 271 of $\mathbf{G}$.

272 *The cube against the sphere..* For a sphere, Eshelby's tensor gives a dilata-
 273 tional channel $1 + \Delta K / (\lambda + 2\mu)$, with $\Delta K = \Delta\lambda + \frac{2}{3}\Delta\mu$, and a single deviatoric
 274 channel $1 + 2\Delta\mu S_{\text{sph}}$, with $2\mu S_{\text{sph}} = 2(4 - 5\nu)/[15(1 - \nu)]$. The cube's di-
 275 latational channel, Eq. (19), is the sphere's exactly: its eigenvalue involves
 276 $\mathbf{M}$ only through the contraction $\int_V \partial_i \partial_n G_{in} dV$, and the static Navier equa-
 277 tion makes $\partial_i \partial_n G_{in} = -\delta(\mathbf{r})/(\lambda + 2\mu)$, whose integral does not depend on
 278 the shape. The cube's shear channels differ from each other. Their isotropic
 279 average, however, is the sphere's, also exactly:

$$\frac{1}{5}(3S_{\text{off}} + 2S_{\text{diag}}) = S_{\text{sph}}, \quad S_{\text{diag}} - S_{\text{off}} = \frac{(\lambda + \mu)(5\sqrt{3} - 2\pi)}{6\pi\mu(\lambda + 2\mu)} > 0. \quad (22)$$

280 The $\sqrt{3}$ terms cancel in the average. What a cube adds to a sphere at this
 281 order is therefore a pure cubic anisotropy. It splits the deviatoric channel
 282 about the sphere's value, weighted 3 : 2 by the dimensions of T_{2g} and E_g ,
 283 and the diagonal shear is the more depolarised. At the background of §7

284 $(\nu = 7/32)$, $2\mu S_{\text{off}} = 0.4314$, $2\mu S_{\text{diag}} = 0.5929$ and the sphere's value is
 285 0.4960.

286 *Where Eshelby enters..* Equations (19)–(21) follow from the moment inte-
 287 grals and the group theory alone. No ellipsoidal result is used, and no δ term
 288 is added by hand. What corresponds to Eshelby is the *truncation*. Keeping
 289 ∂u and discarding $\partial\partial u$ is exactly the uniform-internal-strain ansatz, and by
 290 the uniformity theorem that ansatz is exact if and only if the inclusion is
 291 an ellipsoid. For a sphere the construction is therefore exact in the static
 292 limit, and for a cube it is an approximation. The hierarchy (11)–(13) states
 293 that approximation as the omitted blocks, rather than as an unquantified
 294 shape correction. A polynomial expansion about the centre cannot represent
 295 the edge and corner concentrations of an isolated cube's internal field at any
 296 finite degree. A voxel in a uniform layer has no such concentrations: its
 297 neighbours have the same properties, so its internal field is the continuum's,
 298 which is smooth. That is why a low-degree internal field suffices in a lattice
 299 (§§7–8) although it is only an approximation for an isolated cube.

300 *Against the exact sphere..* These statements are tested against the exact
 301 elastic Mie solution for a sphere, with no Eshelby formula in between. The
 302 concentration of Mie order n is $E_n = a_n/a_n^{\text{Born}}$ (orders 0, 1, 2 are the dilata-
 303 tional, density and shear channels), extrapolated to $ka \rightarrow 0$. The formulation
 304 is applied to a ball of radius a , with its moment $\mathbf{M}$ evaluated by the surface
 305 route of Eq. (17) on $|\mathbf{r}| = a$. Its dilatational and shear channels equal Mie's
 306 to 2×10^{-10} , and the density channel is 1. The cube's closed forms meet
 307 the same numbers. Its dilatational channel is Mie's E_0 to 2×10^{-10} , and its
 308 isotropic shear average, Eq. (22), is the depolarisation read from Mie's E_2

309 to 10^{-8} . The same moment taken as a principal value, over the ball with
 310 a small sphere excluded around the origin, has zero depolarisation in every
 311 channel. The whole static response of the sphere is the δ term.

312 At finite frequency the closure uses the dynamic ball moments, which are
 313 in closed form. It departs from Mie at second order in ka and first order in
 314 the contrast, and the departure is closed-form too: the ball's moments are
 315 elementary, and Mie's low-frequency series with the contrast kept exact gives
 316 E_n to relative order $(ka)^3$. To first order in the contrast,

$$\frac{E_n^{\text{closure}}}{E_n^{\text{Mie}}} - 1 = \begin{cases} -\frac{(k_P a)^2}{10} \frac{\Delta \lambda}{\lambda + 2\mu}, & n = 0, \\ \frac{2(k_S a)^2 + (k_P a)^2}{30} \frac{\Delta \rho}{\rho}, & n = 1, \\ -\frac{3(k_S a)^2 + 2(k_P a)^2 \beta^2 / \alpha^2}{75} \frac{\Delta \mu}{\mu}, & n = 2, \end{cases} \quad (23)$$

317 with no $O(ka)$ term and no $O((ka)^3)$ term. The imaginary $(ka)^3$ terms,
 318 the radiation damping, agree exactly in all three channels. So the closure's
 319 only error at this order is the real $(ka)^2$ term, the uniform-field ansatz's
 320 own. That is the order the hierarchy assigns to the omitted second gradi-
 321 ents (§6.3), and it has the same origin, the variation of the field across the
 322 scatterer, as the collocation error of §7. Eq. (23) is derived symbolically
 323 (`SphereClosureDynamic.wl`). It agrees with the coefficient extracted from
 324 the numerical Mie solve and an independent implementation of the closure
 325 to 5×10^{-4} , which is the size of the next order in the contrast.

326 6.4. The collocation closure

327 The collocation closure takes the internal field of a voxel as uniform, the
 328 truncation of §6.3 for the density and stiffness channels together. The voxel's

own field at its centre is then $\mathbf{S}_{\text{self}}\Delta e$ for internal state e , where $\mathbf{S}_{\text{self}}$ is the
 nine-component form of the moments $\mathbf{G}$ and $\mathbf{M}$. By parity it is block-diagonal.
 So $e = \psi + \mathbf{S}_{\text{self}}\Delta e$ and

$$\mathbf{T} = d^3 \Delta (\mathbf{I} - \mathbf{S}_{\text{self}}\Delta)^{-1}. \quad (24)$$

With the kernel (10) the self term cancels. Substituting $\psi = (\mathbf{I} - \mathbf{S}_{\text{self}}\Delta)e$
 into Eq. (2),

$$e_i = \psi_i^0 + d^3 \sum_j \mathbf{K}_{\text{all}}(i, j) \Delta e_j, \quad \mathbf{K}_{\text{all}}(i, i) = \mathbf{P}_0, \quad \mathbf{K}_{\text{all}}(i, j) = \mathbf{S}_{\text{avg}}(i - j), \quad (25)$$

an identity checked on the matrices to 4×10^{-16} . Every block of $\mathbf{K}_{\text{all}}$ is an
exact cell integral of the one-dimensional continuum kernel (Eqs. (8), (9)), so
 Eq. (25) is precisely the collocation of the continuum Lippmann–Schwinger
 equation with a strain piecewise constant over voxels. It does not contain
 $\mathbf{S}_{\text{self}}$ at all.

The package’s cube T -matrix is this closure statically: both blocks agree
 with Eq. (24) to 2×10^{-5} at $k_S h = 0.01$, the difference falling as $(kh)^2$
 (`ContinuumLimit_CubeT.w1`). It departs dynamically by about $0.25(kh)^2$
 through its far-field form-factor corrections. Its self integrals are exact in the
 static and radiation parts but omit the real even-power terms of the Green’s
 tensor’s expansion ($k^2 r, r^3, \dots$), an error of 0.4–0.6 $(kh)^2$ against the analytic
 self term, whose series agrees with the closed-form elastodynamic tensor to
 10^{-14} . By Eq. (25) that error cannot enter the closure scheme.

347 7. Convergence, and the closed-form error

348 *Second order..* For a layer $D = 2\text{ m}$ with the contrast $\Delta\lambda = 2\text{ GPa}$, $\Delta\mu =$
 349 1 GPa , $\Delta\rho = 100\text{ kg m}^{-3}$ in a background $\alpha = 5000$, $\beta = 3000\text{ m s}^{-1}$, $\rho =$
 350 2500 kg m^{-3} , illuminated by a P plane force, the reflected displacement of
 351 the chain converges to the continuous layer's at order 2.000 for $n = 4 \rightarrow$
 352 16 at $\omega = 60, 300$ and 600 rad s^{-1} , with the collocation (25) and with the
 353 package *T*-matrix alike (`ContinuumLimit_Chain.wl`). The error depends on
 354 $k_S h$ alone.

355 *The leading error in closed form..* In Eq. (25) the kernel is integrated exactly;
 356 the only approximation is the internal field's variation across a voxel. At first
 357 order in the contrast the discrete scattered field is therefore a midpoint rule
 358 applied to the continuum's Born integral, with the kernel's cell integral exact.
 359 For reflection the integrand carries the two-way phase e^{2ikz} ; for transmission
 360 it is flat. With k the background wavenumber of the wave concerned (k_P for
 361 an incident P wave, k_S for an incident S wave), R the reflection coefficient
 362 and T^{sc} the scattered part of the transmitted field, for *every* n

$$\frac{R_{\text{disc}}}{R} = \frac{\text{sinc}(kh)}{\text{sinc}(2kh)} = 1 + \frac{(kd)^2}{8} + O(d^4), \quad \frac{T_{\text{disc}}^{\text{sc}}}{T^{\text{sc}}} = \text{sinc}(kh) = 1 - \frac{(kd)^2}{24} + O(d^4), \quad (26)$$

363 confirmed for $n = 1$ to 16 to 10^{-4} of the error itself (`ContinuumLimit_`
 364 `ErrorConstant.wl`); for an incident S wave, with $k = k_S$, the same forms
 365 and those of Eq. (29) hold to 4×10^{-4} (`ContinuumLimit_IncidentS.wl`, §8).
 366 At the full contrast the reflection constant is unchanged (0.124 against 1/8);
 367 the transmission constant grows to -0.092 .

368 *Why a uniform-field single site stops at second order..* The two constants
 369 have opposite signs, so no scalar correction of $\mathbf{T}$ — no single $O((kh)^2)$ fac-
 370 tor — can cancel both. The package T -matrix illustrates it: its form-factor
 371 corrections reduce the reflection error by a factor 4.45 and *increase* the trans-
 372 mission error by 1.7, at every n . The missing term in Eq. (26) is the first
 373 moment of the internal field over each voxel, which radiates through the
 374 derivative of the kernel. A fourth-order scheme must carry it: a single site
 375 with a gradient channel, and a kernel whose cell integrals include the kernel's
 376 derivative.

377 8. Fourth order: the first-moment voxel

378 Let each voxel carry the mean *and* the first moment of its internal field.
 379 On voxel j , with local coordinate $s \in [-h, h]$ and Legendre functions $\phi_0 = 1$,
 380 $\phi_1 = s/h$, the field is $w = \sum_b c_{b,j} \phi_b$, $w = (u, \varepsilon)$ the displacement and
 381 strain of the P channel, and the polarisation is $q = (\omega^2 \Delta \rho u, \Delta M \varepsilon)$ with
 382 $\Delta M = \Delta \lambda + 2\Delta \mu$. Testing the continuum equation

$$w(z) = w^0(z) + \int \mathbf{K}(z-z') q(z') dz', \quad \mathbf{K} = \begin{pmatrix} g & g' \\ g' & g'' \end{pmatrix}, \quad g'' = -k^2 g - \delta(z)/M_P, \quad (27)$$

383 with the same two functions on every voxel gives

$$\left(\int_i \phi_a^2 \right) c_{a,i} = \int_i \phi_a w^0 + \sum_{j,b} \left[\int_i \int_j \phi_a(z) \mathbf{K}(z-z') \phi_b(z') dz dz' \right] \Delta_1 c_{b,j}, \quad (28)$$

384 $\Delta_1 = \text{diag}(\omega^2 \Delta \rho, \Delta M)$. Between voxels the double integral is separable, a
 385 product of Legendre moments of $e^{\pm i k z}$; within a voxel it is split at $z' = z$ and
 386 the local term contributes $-\int_i \phi_a \phi_b / M_P$. Every entry is closed-form. Keeping

only ϕ_0 recovers a second-order scheme; keeping ϕ_0 and ϕ_1 converges at fourth order in both reflection and transmission (`ContinuumLimit_FourthOrder.w1`; Table 1). A single plane of first-moment voxels represents the 2 m layer

Table 1: Relative error of the scattered displacement of the layer, reflection / transmission, at $\omega = 300 \text{ rad s}^{-1}$, against the number of voxel planes n : collocation, the mean-only voxel and the first-moment voxel of Eq. (28). The order is 4.00 at $\omega = 60$ and 600 rad s^{-1} too.

n	collocation	mean only	mean and first moment
1	$1.8 \times 10^{-3} / 1.3 \times 10^{-3}$	$1.2 \times 10^{-3} / 1.8 \times 10^{-3}$	$2.8 \times 10^{-7} / 2.8 \times 10^{-7}$
4	$1.1 \times 10^{-4} / 8.3 \times 10^{-5}$	$7.5 \times 10^{-5} / 1.1 \times 10^{-4}$	$1.1 \times 10^{-9} / 1.1 \times 10^{-9}$
16	$7.0 \times 10^{-6} / 5.2 \times 10^{-6}$	$4.7 \times 10^{-6} / 6.9 \times 10^{-6}$	$4.2 \times 10^{-12} / 4.2 \times 10^{-12}$
order	2 / 2	2 / 2	4 / 4

6500 times more accurately than a single plane of uniform ones. The collocation column is the scheme of §6 and reproduces the chain of the full 9×9 system to all printed digits.

The leading error in closed form.. Without scattering, the Galerkin solution is the L^2 projection Π of the incident field onto each voxel's basis, and the kernel is integrated over each voxel exactly; at first order in the contrast the error is therefore the kernel's cell integral applied to $\Pi w^0 - w^0$. The field varies as e^{iks} across a voxel and the kernel toward an outside observer as e^{iqs} , $q = +k$ for reflection and $q = -k$ for transmission, so for every scheme and every n

$$\frac{\text{discrete}}{\text{exact}} = \frac{\int_{-h}^h e^{iqs} \Pi[e^{iks}] \, ds}{\int_{-h}^h e^{i(q+k)s} \, ds}, \quad (29)$$

400 with Π the centre value for collocation, the cell mean for the mean-only voxel
 401 and the projection onto $\{1, s/h\}$ for the first-moment voxel (`ContinuumLimit_`
`FourthOrderError.wl`). The results are in Table 2. The chain of Eq. (28)

Table 2: The leading error of the discrete layer in closed form, Eq. (29): discrete over exact scattered field at first order in the contrast, for every number of voxel planes n .

	reflection	transmission
collocation	$1 + (kd)^2/8 + O(d^4)$	$1 - (kd)^2/24 + O(d^4)$
mean only	$1 + (kd)^2/12 + O(d^4)$	$1 - (kd)^2/12 + O(d^4)$
first moment	$1 - (kd)^4/720 + O(d^6)$	$1 - (kd)^4/720 + O(d^6)$

402
 403 reproduces all six forms at $n = 1$ to 8 to 10^{-4} of the error, the remainder
 404 being the second order in the contrast. The first-moment voxel's leading
 405 error is the same in reflection and in transmission: its projection error is or-
 406 thogonal to everything linear across the voxel, so the leading term no longer
 407 depends on the direction in which the field is observed.

408 *In three dimensions..* The three-dimensional first-moment voxel carries the
 409 nine-component state with Legendre weights $\{1, x/h, y/h, z/h\}$ — 36 un-
 410 knowns — tested with the same functions and coupled by cell-to-cell double
 411 integrals of the Green's tensor. For a laterally uniform layer at normal in-
 412 cidence it reduces *exactly* to Eq. (28) (`ContinuumLimit_Reduction3D.wl`).
 413 The lateral first moments are never excited: the lattice and the incident
 414 field are symmetric under reflection about every voxel's centre planes, and
 415 $x/h, y/h$ are odd there. The mean and the z -moment carry weights that
 416 are constant laterally, so their lateral form factor $\text{sinc}(g_x h) \text{sinc}(g_y h)$ van-
 417 ishes on the reciprocal lattice and the plane sum of the double integrals

418 is the plate-to-plate integral in z alone. Two facts mark where genuinely
 419 three-dimensional structure begins. The lateral first moment's form fac-
 420 tor, $F_1(g) = i(gh \cos gh - \sin gh)/(gh)^2$, does not vanish on the recipro-
 421 cal lattice ($F_1 = i(-1)^p/p\pi$), so the gradient channels couple through the
 422 evanescent orders. And at oblique incidence the mean's own zeros shift,
 423 $\text{sinc}((2\pi p/d + k_x)h) = (-1)^p k_x h / p\pi + O(k_x^2 h^2)$, so the tiling identity holds
 424 only to $O(k_{\parallel} h)$.

425 *Oblique incidence..* The scheme is built spectrally. By Poisson summation
 426 the Bloch coupling between planes i and j at lateral wavenumber $\mathbf{k}_{\parallel}$ is

$$C_{ab}(i, j) = \frac{1}{d^2} \sum_{\mathbf{g}} L_a(\boldsymbol{\kappa}) L_b(-\boldsymbol{\kappa}) \iint \phi_a(z) \hat{\mathbf{G}}(\boldsymbol{\kappa}; z - z') \phi_b(z') dz dz', \quad \boldsymbol{\kappa} = \mathbf{k}_{\parallel} + \mathbf{g}, \quad (30)$$

427 with L_a the lateral Legendre factors. The whole-space tensor decays as
 428 $1/k_z^2$ and the nine-component state adds at most two powers of k_z , so each
 429 entry of $\hat{\mathbf{G}}(\boldsymbol{\kappa}; z)$ is a $\delta(z)$ of constant weight plus, per mode, $[i\alpha/2\gamma +$
 430 $i\beta \text{sgn}(z)/2] e^{i\gamma|z|}$, $\gamma^2 = k_W^2 - \kappa^2$: every z integral is closed-form. Double
 431 averaging makes the lattice sum converge without Ewald splitting, to three
 432 digits at $|p|, |q| \leq 8$. Against the exact layer, for a P wave incident at 20°
 433 and 40° (`ContinuumLimit_Oblique.wl`; Table 3). The mean-only voxel is
 434 second order and the first-moment voxel fourth order in both the reflected
 435 P and the converted S wave: the breakdown of the tiling identity at $O(k_{\parallel} h)$
 436 and the evanescent coupling of the gradient channels do not cost an order.
 437 At $\mathbf{k}_{\parallel} = 0$ the construction reproduces Eq. (28) to all printed digits.

438 *Incident S..* Nothing in the construction is specific to P. For a normally
 439 incident S wave the layer is the one-dimensional problem of Eq. (28) with

Table 3: Relative error of the reflected P (R_{PP}) and converted S (R_{PS}) waves for a P wave incident at 20° , with the mean-only and first-moment voxels of Eq. (30). At 40° the orders are again 2 and 4.

	mean only		mean and first moments	
n	R_{PP}	R_{PS}	R_{PP}	R_{PS}
1	9.3×10^{-4}	1.7×10^{-3}	7.2×10^{-8}	4.1×10^{-7}
2	2.3×10^{-4}	4.2×10^{-4}	4.5×10^{-9}	2.5×10^{-8}
8	1.5×10^{-5}	2.7×10^{-5}	1.7×10^{-11}	9.9×10^{-11}
order	2	2	4	4

440 $M_P \rightarrow \mu$, $k_P \rightarrow k_S$ and $\Delta M \rightarrow \Delta \mu$; the collocation and mean-only schemes
441 converge at order 2.00 and the first-moment voxel at 4.00, in reflection and
442 transmission, and the full 9×9 construction with an incident SH wave re-
443 produces that one-dimensional result to all printed digits. At oblique inci-
444 dence, against the exact layer (`ContinuumLimit_IncidentS.wl`), the orders
are those of Table 4. The orders, and with $k = k_S$ the closed-form constants

Table 4: Observed order of convergence ($n = 2 \rightarrow 8$, $\omega = 300 \text{ rad s}^{-1}$) for incident SV and SH waves, with the mean-only and first-moment voxels. The SV angles lie below the critical angle $\arcsin(\beta/\alpha) = 36.9^\circ$.

	mean only	mean and first moments
SV at 20° : R_{SS} / R_{SP}	2.00 / 2.00	4.00 / 3.99
SV at 30° : R_{SS} / R_{SP}	1.98 / 2.00	3.97 / 4.00
SH at 20° and 40° : R_{SH}	2.00	4.00

445

446 of Eqs. (26) and (29), are those of the incident P wave.

447 The spectral construction is implemented twice. The second implementa-
448 tion (`scripts/crosscheck_first_moment_voxel.py`) shares only the defi-
449 nition of the scheme: it takes the depth dependence of the transformed ker-
450 nel from an independent residue construction, obtains the local weight as a
451 numerical large- k_z limit of the inverted Christoffel matrix, evaluates every
452 double integral by quadrature and separates P from S by sampling the up-
453 going field at two depths. The raw reflected displacements of the two agree
454 to 9×10^{-16} in all 36 cases compared — incident P, SV and SH, normal
455 incidence and 20° , $n = 1, 2, 4$, with and without the first moments. Figure 1
456 collects the convergence.

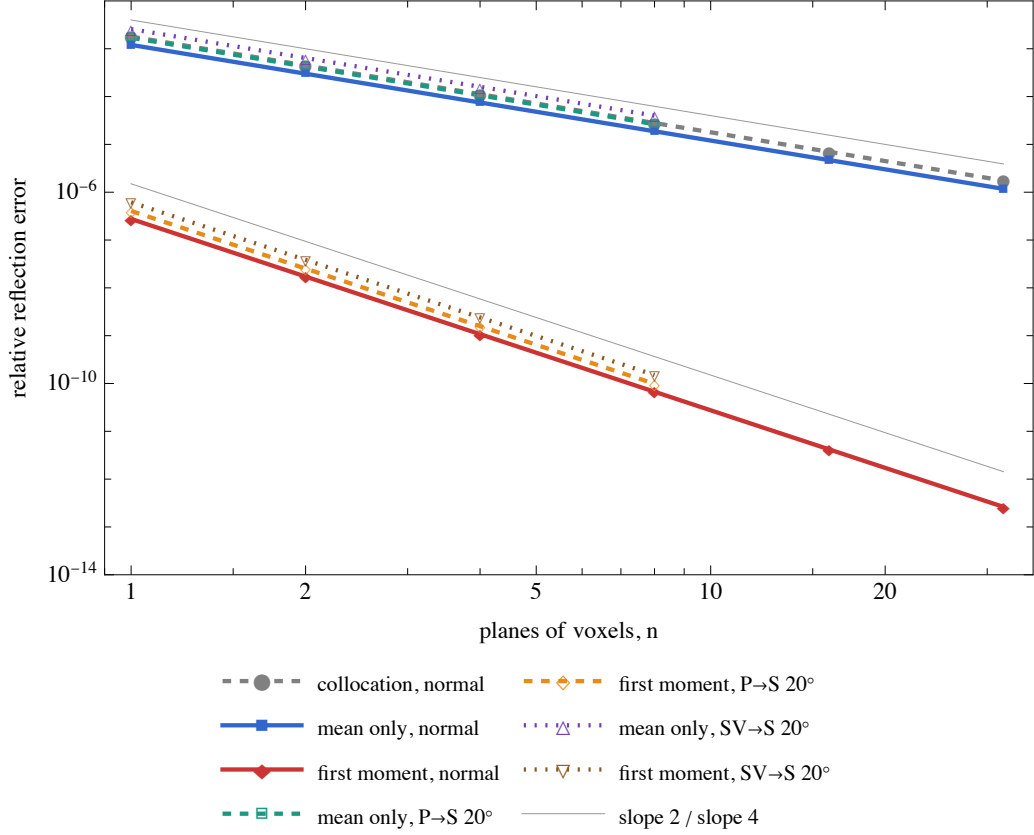


Figure 1: Relative reflection error of the discrete layer against the number of voxel planes n ($D = 2\text{ m}$, $\omega = 300\text{ rad s}^{-1}$, the contrast of §7). Filled: normal incidence, P reflection — collocation (§6), the mean-only voxel and the first-moment voxel of Eq. (28). Dashed, open squares and diamonds: P incident at 20° , the converted S wave, mean-only (green) and first-moment (orange) voxels of Eq. (30). Dotted, open triangles: SV incident at 20° , the reflected S wave, mean-only (purple) and first-moment (brown). Grey: slopes 2 and 4.

Alt text: A log-log plot of relative reflection error (vertical axis, 10^{-14} to 10^{-2}) against the number of voxel planes (horizontal axis, 1 to 32). The collocation and mean-only curves fall as straight lines of slope -2 from about 10^{-3} at one plane; the first-moment curves fall as straight lines of slope -4 from about 10^{-6} at one plane to about 10^{-13} at 32 planes. The oblique P-to-S and incident-SV curves lie alongside the normal-incidence curves of the same scheme, and grey guide lines show slopes 2 and 4.

457 9. A heterogeneous layer

458 A gridded Earth model is not a layer of identical cells. It is a set of model
459 cells of constant properties with real interfaces between them. The last test
460 takes such a model. A layer $D = 4$ m thick is divided into eight model cells of
461 0.5 m, and in each one $\Delta\lambda$, $\Delta\mu$ and $\Delta\rho$ are drawn uniformly within $\pm 15\%$ of
462 the background (the values are listed in `ContinuumLimit_Heterogeneous.`
463 `w1`). The result is a random stratification with physical contrasts. The voxels
464 refine the model grid, with m planes of voxels in each model cell, so no voxel
465 straddles a model interface; that is how a gridded model is discretised. At
466 $\omega = 1500 \text{ rad s}^{-1}$ the layer is $k_P D = 1.2$ and $k_S D = 2.0$ thick, so the field
467 reverberates between the model interfaces. The exact answer is the product
468 of the model cells' propagators. Its same-type reflection coefficients agree
469 with those of the reflection–transmission recursion of Kennett [25], computed
470 independently, to 10^{-14} . The schemes are those of §§7–8, with the contrast
471 now changing from plane to plane; nothing else changes. The errors are in
472 Table 5.

473 *The interfaces cost nothing..* Every order is the one measured for the uniform
474 layer. At a model interface the strain of the true field jumps. A voxel basis
475 that is discontinuous from voxel to voxel represents that jump exactly when
476 the interface lies on a voxel boundary, so the interface adds no error of its
477 own. The only requirement is that the voxels resolve the model grid. A voxel
478 that straddled an interface would carry a contrast that is not uniform within
479 it, a case not examined here.

Table 5: Relative error against the exact stratified layer, for m planes of voxels per model cell (§9). At normal incidence the order is measured from $m = 4$ to 16, where it is also 2.00 for collocation and 4.00 in transmission; at 20° it is measured from $m = 2$ to 4.

		$m = 1$	$m = 2$	$m = 4$	order
P, normal, R	mean only	1.9×10^{-3}	4.7×10^{-4}	1.2×10^{-4}	2.00
	first moment	8.5×10^{-7}	5.3×10^{-8}	3.3×10^{-9}	4.00
S, normal, R	mean only	5.3×10^{-3}	1.3×10^{-3}	3.3×10^{-4}	2.00
	first moment	6.2×10^{-6}	3.9×10^{-7}	2.4×10^{-8}	4.00
P, 20° , R_{PP}	mean only	1.4×10^{-3}	3.6×10^{-4}	9.0×10^{-5}	2.00
	first moments	3.3×10^{-7}	2.1×10^{-8}	1.3×10^{-9}	4.00
P, 20° , R_{PS}	mean only	2.8×10^{-3}	6.9×10^{-4}	1.7×10^{-4}	2.00
	first moments	1.4×10^{-6}	8.8×10^{-8}	5.5×10^{-9}	4.00

480 *What fourth order buys..* For a reflection error of 10^{-6} at normal incidence
481 (P), the mean-only voxel needs 347 planes of voxels, about 43 per model cell.
482 The first-moment voxel needs 8, one per model cell: the model grid itself,
483 which already gives 8.5×10^{-7} . Counting unknowns (9 per voxel against 36),
484 that is 3123 against 288 per vertical column of voxels, a factor of eleven.
485 In three dimensions, with the lateral grid refined as the vertical one is, the
486 number of voxels falls by $(347/8)^3 \approx 8 \times 10^4$ and the number of unknowns by
487 about 2×10^4 . The same model has been solved independently in Python,
488 with a contrast for each plane: the raw reflected displacements agree with
489 the Mathematica scheme to 1.1×10^{-15} at normal incidence and at 20° , with
490 and without the first moments. Lateral heterogeneity needs the full three-
491 dimensional solver; it is the first extension discussed below.

492 10. Discussion

493 *Where the error of a voxel scheme lives..* The results separate the two ap-
 494 proximations a voxel scheme makes. The *coupling* between voxels makes
 495 none: averaged over the source cell it is the continuum’s cell integral exactly,
 496 between planes by the form factor’s zeros and within a plane by the tiling
 497 identity, and the self term it omits is exactly the one the single site carries.
 498 Every error therefore comes from the *single site*, from the internal field it
 499 assumes. For a uniform field that error is the midpoint rule’s, known in
 500 closed form; for a linear field it is the next term of the same expansion. The
 501 converse is the practical warning. A kernel that is not the cell integral —
 502 the point Green’s tensor, or a cell average that does not match the single
 503 site’s closure — leaves a local term that does not vanish as the voxels shrink
 504 (§4): the evanescent lattice orders carry a static contribution of fixed size
 505 per voxel volume. A scheme built that way converges to the wrong limit,
 506 and refinement cannot reveal it, because the error is scale-invariant.

507 *What the single site is..* The moment hierarchy of §6 turns the uniform-field
 508 single site from a modelling choice into a derived object with a stated error.
 509 Its static part is exact for a sphere, and for a cube it differs from the sphere’s
 510 only by the cubic split of Eq. (22). Its dynamic part departs from exact
 511 scattering by a sphere only through the real $(ka)^2$ term of Eq. (23); the
 512 radiation damping is exact. That term and the collocation error of Eq. (26)
 513 have the same origin, the variation of the field across the scatterer, and the
 514 same remedy, a richer internal field. In a lattice the remedy is cheap: the
 515 field a voxel must represent is the smooth continuum field, not the field of
 516 an isolated cube, with its concentrations at the edges and corners.

517 *Relation to the discrete dipole approximation..* The same partition appears in
 518 the electromagnetic literature as the choice of polarisability. The Clausius–
 519 Mossotti relation, the lattice dispersion relation [6] and the integrated Green
 520 tensor [7] all adjust the single site so that the lattice reproduces the con-
 521 tinuum. For space-filling cubes the adjustment is neither needed nor a free
 522 parameter: the tiling identity fixes the local field exactly, the self term of the
 523 collocation closure is the one that makes it cancel, and what is left is a closed-
 524 form error rather than an empirical one. The second-order convergence re-
 525 ported for the DDA by Yurkin et al. [1] corresponds to the uniform-field
 526 voxel’s, and the opposite-signed constants of Eq. (26) explain why a polaris-
 527 ability correction can improve one observable only at the expense of another,
 528 as the package’s own form-factor corrections do in reflection and transmis-
 529 sion. Going beyond second order requires a richer internal field, which is the
 530 point of higher-order bases for volume integral equations; the first-moment
 531 voxel is the lowest such basis, and in its spectral form the Galerkin double
 532 average makes the lattice sum converge without Ewald splitting.

533 *The cost of fourth order..* A first-moment voxel carries four times as many
 534 unknowns as a uniform one (36 against 9 in three dimensions). The closed
 535 forms show what that buys. To reach an error ε in reflection, the uniform
 536 voxel needs $kd = \sqrt{8\varepsilon}$ and the first-moment voxel $kd = (720\varepsilon)^{1/4}$; at $\varepsilon = 10^{-6}$
 537 that is $kd = 2.8 \times 10^{-3}$ against 0.16, voxels 58 times larger in each direction.
 538 In a three-dimensional model the voxel count falls by $58^3 \approx 2 \times 10^5$ and the
 539 number of unknowns by about 5×10^4 (these are the Born-order constants;
 540 at the contrast studied the first-moment voxel’s measured errors lie within
 541 5% of them, while the uniform voxel’s transmission constant grows with the

contrast, so the comparison is if anything conservative). The stratified model of §9 measures it directly: 347 planes of uniform-field voxels against 8 first-moment voxels for the same accuracy. The gain is largest precisely where it matters: when the grid of a geological model is coarse compared with the accuracy the wavefield must have.

Limitations and outlook.. The results concern laterally uniform media — a homogeneous layer and a randomly stratified one — for which the discrete system is exactly a chain of planes; a laterally heterogeneous medium needs the first-moment voxel in a full three-dimensional solver, where the gradient channels couple through the evanescent lattice orders (§8) and the tiling identity holds only to $O(k_{\parallel}h)$. The numerical evidence is at moderate, physically typical contrasts (up to 15% in density and moduli) and at $k_S h \leq 0.3$, the range in which the single-site T -matrix is validated; the closed forms are first order in the contrast, and the departures from them at the full contrast are of second order. What is established is the *order* of convergence and its leading constant, not the behaviour at strong contrast, where the single site’s own closure may limit the accuracy before the lattice coupling does. The single site’s dynamic error is in closed form for a sphere only; for a cube, for which no exact solution exists, it is known through the layer (§§7–8). Two extensions follow directly. The first is the three-dimensional solver itself. The second is a stratified *background* — a layered reference medium, as distinct from the stratified contrast of §9: the lattice coupling then carries the reverberation of the layering, which is not translation invariant in depth, and whether a tiling identity holds for it is the natural next question.

566 11. Conclusions

567 A lattice of cubic voxels reproduces a continuous elastic medium in a
568 precise sense. Because cubes tile the plane, the source-cell-averaged cou-
569 pling between voxels is the continuum’s cell integral exactly, and at normal
570 incidence the lattice kernel is available in closed form with no lattice sum.
571 The single site is derived rather than assumed. The moments of the Green’s
572 tensor over the cube, taken as distributions, give a closed hierarchy whose
573 first-gradient truncation is the uniform-strain closure. That closure is exact
574 for a sphere in the static limit, and it departs from exact scattering by a
575 sphere only by a real $(ka)^2$ term, with the radiation damping exact. For
576 a cube it differs from the sphere only by a cubic anisotropy of the shear
577 response. The collocation closure carries the same self term as the lattice
578 kernel, which cancels, so the discrete system is a collocation of the contin-
579 uum equation. It converges at second order, with a leading error that is a
580 midpoint-rule error of the internal field, known in closed form for every dis-
581 cretisation, and opposite in sign in reflection and transmission. A voxel that
582 also carries the first moment of its internal field converges at fourth order —
583 at normal incidence, where the three-dimensional scheme reduces exactly to
584 one dimension, and at oblique incidence — with a leading error, $-(kd)^4/720$,
585 that no longer depends on the direction of observation. The orders survive
586 the interfaces of a randomly stratified model, provided the voxels resolve its
587 grid, and there the first-moment voxel reaches an accuracy of 10^{-6} with one
588 voxel per model cell, where the uniform-field voxel needs about forty. The
589 results, and the evidence for each, are collected in Table 6.

Table 6: The results of the paper, and the evidence for each.

Result	Evidence
Continuous layer, exact and to $O(D^k)$	reference to 2.5×10^{-8} ; series order k
Point-kernel local term survives $d \rightarrow 0$	$11 \times$ the continuum term at $m = 1$
Averaged coupling between planes = continuum $\times$ sinc	2.8×10^{-16} , 81 entries
Tiling identity, Eq. (9)	21/21 components to 4×10^{-9}
Exact kernel, Eq. (10)	replaces a 9.4×10^{-4} static error
Moments as distributions, Eq. (17)	sum rule (18) exact; M, Q equal an independent archive, 81/81
Channels (19)–(21)	a production code by another route, 2×10^{-15}
Cube vs sphere, Eq. (22)	symbolic; and numerically at five ν , 4×10^{-16}
Ball closure = exact Mie, static	2×10^{-10} ; cube’s bulk and shear average 10^{-8}
Ball closure vs Mie, dynamic, Eq. (23)	symbolic; numerical Mie 5×10^{-4} ; $(ka)^3$ terms exact
Self term cancels, Eq. (25)	4×10^{-16}
Package T = closure statically	2×10^{-5} at $k_S h = 0.01$
Second-order convergence	order 2.000, three frequencies
Closed-form error, Eq. (26)	$n = 1$ –16, to 10^{-4} of the error
First-moment voxel, Eq. (28)	order 4.00, three frequencies
First-moment voxel, oblique, Eq. (30)	order 4, P and S, 20° and 40°
Incident SV and SH, all schemes	orders 2 and 4; constants with k_S to 4×10^{-4}

590 **Acknowledgments**

591 This work began thirty years ago in the author’s doctoral research at
592 the Research School of Earth Sciences of the Australian National University,
593 Canberra [9].

594 **Declaration of competing interest**

595 The author declares no known competing financial interests or personal
596 relationships that could have appeared to influence the work reported in this
597 paper.

598 **Data availability**

599 The Mathematica notebooks and Python scripts that derive and check ev-
600 ery result, and the package they test, are archived at [https://doi.org/10.](https://doi.org/10.5281/zenodo.23003242)
601 [5281/zenodo.23003242](https://doi.org/10.5281/zenodo.23003242) and developed at [https://github.com/tmnestor/](https://github.com/tmnestor/MultipleScatteringCalculations)
602 [MultipleScatteringCalculations](https://github.com/tmnestor/MultipleScatteringCalculations).

603 **References**

- 604 [1] M. A. Yurkin, V. P. Maltsev, A. G. Hoekstra, Convergence of the dis-
605 crete dipole approximation. I. Theoretical analysis, J. Opt. Soc. Am. A
606 23 (2006) 2578–2591.
- 607 [2] L. L. Foldy, The multiple scattering of waves. I. General theory of
608 isotropic scattering by randomly distributed scatterers, Phys. Rev. 67
609 (1945) 107–119.

- 610 [3] M. Lax, Multiple scattering of waves, *Rev. Mod. Phys.* 23 (1951)
611 287–310.
- 612 [4] E. M. Purcell, C. R. Pennypacker, Scattering and absorption of light by
613 nonspherical dielectric grains, *Astrophys. J.* 186 (1973) 705–714.
- 614 [5] B. T. Draine, P. J. Flatau, Discrete-dipole approximation for scattering
615 calculations, *J. Opt. Soc. Am. A* 11 (1994) 1491–1499.
- 616 [6] B. T. Draine, J. Goodman, Beyond Clausius–Mossotti: wave propaga-
617 tion on a polarizable point lattice and the discrete dipole approximation,
618 *Astrophys. J.* 405 (1993) 685–697.
- 619 [7] P. C. Chaumet, A. Sentenac, A. Rahmani, Coupled dipole method for
620 scatterers with large permittivity, *Phys. Rev. E* 70 (2004) 036606.
- 621 [8] J. D. Eshelby, The determination of the elastic field of an ellipsoidal
622 inclusion, and related problems, *Proc. R. Soc. Lond. A* 241 (1957)
623 376–396.
- 624 [9] T. M. Nestor, *Seismic Wave Propagation in Complex Media: Sum-*
625 *ming the Multiple Scattering Series*, Ph.D. thesis, Research School of
626 Earth Sciences, Australian National University, Canberra, ACT, Aus-
627 tralia, 1996.
- 628 [10] J. D. Eshelby, The elastic field outside an ellipsoidal inclusion, *Proc. R.*
629 *Soc. Lond. A* 252 (1959) 561–569. doi:10.1098/rspa.1959.0173.
- 630 [11] T. Mura, *Micromechanics of Defects in Solids*, 2nd ed., Martinus Nijhoff,
631 Dordrecht, 1987. doi:10.1007/978-94-009-3489-4.

- [12] Y. Mikata, S. Nemat-Nasser, Elastic field due to a dynamically transforming spherical inclusion, *J. Appl. Mech.* 57 (1990) 845–849. doi:10.1115/1.2897650.
- [13] T. M. Michelitsch, H. Gao, V. M. Levin, Dynamic Eshelby tensor and potentials for ellipsoidal inclusions, *Proc. R. Soc. Lond. A* 459 (2003) 863–890. doi:10.1098/rspa.2002.1054.
- [14] Y. P. Chiu, On the stress field due to initial strains in a cuboid surrounded by an infinite elastic space, *J. Appl. Mech.* 44 (1977) 587–590. doi:10.1115/1.3424140.
- [15] J. K. Lee, W. C. Johnson, Calculation of the elastic strain field of a cuboidal precipitate in an anisotropic matrix, *Phys. Status Solidi A* 46 (1978) 267–272. doi:10.1002/pssa.2210460134.
- [16] G. J. Rodin, Eshelby’s inclusion problem for polygons and polyhedra, *J. Mech. Phys. Solids* 44 (1996) 1977–1995. doi:10.1016/S0022-5096(96)00066-X.
- [17] H. Nozaki, M. Taya, Elastic fields in a polyhedral inclusion with uniform eigenstrains and related problems, *J. Appl. Mech.* 68 (2001) 441–452. doi:10.1115/1.1362670.
- [18] A. K. Mal, L. Knopoff, Elastic wave velocities in two-component systems, *J. Inst. Math. Appl.* 3 (1967) 376–387. doi:10.1093/imamat/3.4.376.
- [19] J. E. Gubernatis, E. Domany, J. A. Krumhansl, Formal aspects of the

- theory of the scattering of ultrasound by flaws in elastic materials, J. Appl. Phys. 48 (1977) 2804–2811. doi:10.1063/1.324141.
- [20] S. K. Kanaun, V. M. Levin, Self-Consistent Methods for Composites, Vol. 2: Wave Propagation in Heterogeneous Materials, Solid Mechanics and Its Applications, Springer, Dordrecht, 2008. doi:10.1007/978-1-4020-6968-0.
- [21] R.-S. Wu, K. Aki, Scattering characteristics of elastic waves by an elastic heterogeneity, Geophysics 50 (1985) 582–595. doi:10.1190/1.1441934.
- [22] P. C. Waterman, Matrix theory of elastic wave scattering, J. Acoust. Soc. Am. 60 (1976) 567–580. doi:10.1121/1.381130.
- [23] V. D. Kupradze, Dynamical Problems in Elasticity, volume 3 of *Progress in Solid Mechanics*, North-Holland, Amsterdam, 1963.
- [24] K. Aki, P. G. Richards, Quantitative Seismology, 2nd ed., University Science Books, Sausalito, CA, 2002.
- [25] B. L. N. Kennett, Seismic Wave Propagation in Stratified Media, Cambridge University Press, Cambridge, 1983.