

# The First-Order Contrast Operator

The elastodynamic matrix–vector wave equation as a route to the single-site  $T$ -matrix: derivation, validation, and what it actually buys

17 September 2026

## Abstract

The elastodynamic matrix–vector wave equation  $\partial_3 \mathbf{q} = \mathcal{A} \mathbf{q} + \mathbf{d}$  expresses the wave field as a first-order recursion in depth. This note asks whether the cube moment method can be applied to it directly, to obtain a single-site  $T$ -matrix without passing through the Navier equation. The answer is yes, and the construction is not new: the thesis already carries it in  $2\frac{1}{2}$ -D. What is new here is the 3-D Cartesian operator, and a measurement of what the route is worth.

Three results. *First*, the operator matrix  $\mathcal{A}$  is derived from the stiffness tensor and reproduces the thesis system matrix (Akdef) under one constant similarity transform; the paper’s  $\nu_1, \nu_2$  are the thesis’s  $\zeta, \chi$ . A fourth, arrived at late and by the wrong road, is that the thesis’s own spectral representation of  $\mathbf{A}$  — analytic eigenvectors, a piecewise branch, and an inverse obtained from the symplectic structure rather than by inversion — removes three numerical hazards outright that the operator route had to discover and repair one at a time. *Second*, the test space is *derived* rather than chosen:  $\mathbf{J}_6$  is antisymmetric and  $\Delta \mathcal{A}^t \mathbf{J}_6 = -\mathbf{J}_6 \Delta \mathcal{A}$ , so  $\mathbf{J}_6 \Delta \mathcal{A}$  is symmetric and the only natural projection is test =  $\mathbf{J}_6 \times$  trial. This dissolves an apparent mismatch of ten conditions against nine unknowns. *Third*, and least expected, the resulting coupling agrees with the ordinary Navier weak form at Born order but differs at  $O(\Delta c^2)$  entirely within the traction channel — and that difference splits into a genuine  $O_h$ -invariant resummation *and* a spurious preference for the  $x_3$  axis, in comparable amounts. A single-site score against an independent arbiter in a homogeneous background would mix the two and settle neither. The directional preference ceases to be an artefact once the background is layered, which is where the route was wanted.

The  $T$ -matrix itself is assembled in §9 and the comparison is reported channel by channel. The sharpest statement is not a norm:  $e_{12}$  and  $e_{13}$  are equivalent under  $O_h$  and so must amplify identically, and in the first-order  $T$  they split by  $8.4 \times 10^{-4}$  while the ordinary route splits them by exactly zero. That split is an arbiter-free *lower bound* on the first-order error,  $\geq \delta/2$ ; it bounds nothing about the ordinary route, which can be symmetrically wrong, so no ordering of the two follows from it.

The ordering is settled elsewhere, in the one geometry where both formulations apply and an external reference exists. A periodic patch of space-filling cubes *is* a homogeneous layer, so a reflectivity recursion gives its reflection coefficient exactly; running both routes through the same lattice sum and the same solve, with only  $T_0$  substituted, the first-order  $T$  is the more accurate of the two, by a margin that grows with contrast and that *grows further as the shared lattice error is refined away* — which is what distinguishes a real advantage from an accidental cancellation (§13).

Every claim below is backed by an executable check; the scripts report 25/25, 44/44, 14/14, 17/17, 8/8, 7/7 and 5/5, and a separate gate binds the numerical propagator to the derived  $\mathcal{A}$ . Section 15 states plainly what is *not* yet established.

## Contents

|          |                                                                                             |           |
|----------|---------------------------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Conventions and objects</b>                                                              | <b>4</b>  |
| 1.1      | Indices, time, and units . . . . .                                                          | 4         |
| 1.2      | The field vector . . . . .                                                                  | 4         |
| 1.3      | The metric $\mathbf{J}_6$ . . . . .                                                         | 5         |
| <b>2</b> | <b>The first-order system for an isotropic medium</b>                                       | <b>5</b>  |
| 2.1      | From the stiffness tensor to the stiffness matrices . . . . .                               | 5         |
| 2.2      | Eliminating the in-plane tractions . . . . .                                                | 5         |
| 2.3      | The operator matrices . . . . .                                                             | 6         |
| <b>3</b> | <b><math>\mathcal{A}</math> against the thesis, and against the published operator form</b> | <b>6</b>  |
| 3.1      | Arbiter A: the published isotropic operator matrices . . . . .                              | 7         |
| 3.2      | The anchor: the thesis system matrix (Akdef) . . . . .                                      | 7         |
| 3.3      | Structural checks . . . . .                                                                 | 7         |
| <b>4</b> | <b>The contrast operator</b>                                                                | <b>8</b>  |
| 4.1      | Definition and the Lippmann–Schwinger equation . . . . .                                    | 8         |
| 4.2      | The foundation check . . . . .                                                              | 8         |
| 4.3      | $\Delta\mathcal{A}$ is not linear in the contrast . . . . .                                 | 8         |
| <b>5</b> | <b>The augmentation</b>                                                                     | <b>9</b>  |
| 5.1      | Why a derivative must be moved . . . . .                                                    | 9         |
| 5.2      | The augmented dimension . . . . .                                                           | 9         |
| <b>6</b> | <b>The test space, derived</b>                                                              | <b>10</b> |
| 6.1      | The argument . . . . .                                                                      | 10        |
| 6.2      | The ten-versus-nine mismatch dissolves . . . . .                                            | 10        |
| 6.3      | The effective coupling matrix . . . . .                                                     | 10        |
| <b>7</b> | <b>What the first-order coupling actually buys</b>                                          | <b>11</b> |
| 7.1      | The trial spaces coincide . . . . .                                                         | 11        |
| 7.2      | The density channel cannot differ . . . . .                                                 | 11        |
| 7.3      | The units the two couplings are written in . . . . .                                        | 11        |
| 7.4      | Born agreement, and an $O(\Delta c^2)$ difference . . . . .                                 | 12        |
| 7.5      | Half of it is a broken symmetry . . . . .                                                   | 12        |
| 7.6      | What follows, and what does not . . . . .                                                   | 13        |

|           |                                                                         |           |
|-----------|-------------------------------------------------------------------------|-----------|
| <b>8</b>  | <b>The propagator, and the bridge to it</b>                             | <b>13</b> |
| 8.1       | It already exists . . . . .                                             | 13        |
| 8.2       | The bridge identity . . . . .                                           | 13        |
| 8.3       | Reporting residuals against the arithmetic floor . . . . .              | 14        |
| <b>9</b>  | <b>The single-site <math>T</math>-matrix</b>                            | <b>14</b> |
| 9.1       | What is new and what is borrowed . . . . .                              | 14        |
| 9.2       | The self moments, derived . . . . .                                     | 15        |
| 9.2.1     | The static Green's tensor . . . . .                                     | 15        |
| 9.2.2     | The moment as a distribution, and why no delta is inserted . . . . .    | 15        |
| 9.2.3     | The face integrals, and one of them by pure geometry . . . . .          | 16        |
| 9.2.4     | The closed forms . . . . .                                              | 16        |
| 9.2.5     | The check that certifies the delta content . . . . .                    | 17        |
| 9.3       | The tabulated route agrees . . . . .                                    | 17        |
| 9.4       | Two implementations of the coupling . . . . .                           | 17        |
| 9.5       | Calibration: every normalisation fixed against existing code . . . . .  | 17        |
| 9.6       | The first-order $T$ , channel by channel . . . . .                      | 18        |
| 9.7       | The broken symmetry, now visible in $T$ itself . . . . .                | 18        |
| 9.8       | What this does and does not settle . . . . .                            | 19        |
| <b>10</b> | <b>The propagator moment from the first-order system</b>                | <b>19</b> |
| 10.1      | Step 1: $\Gamma$ , built from $\mathcal{A}$ alone . . . . .             | 20        |
| 10.2      | Step 2: the vertical integral, in closed form . . . . .                 | 21        |
| 10.2.1    | The same integrals as a series, and why they are needed . . . . .       | 21        |
| 10.3      | Step 3: the lateral integral . . . . .                                  | 22        |
| 10.3.1    | The scaling that makes the expansion uniform . . . . .                  | 22        |
| 10.3.2    | Which entries fail to decay . . . . .                                   | 23        |
| 10.3.3    | The asymptote in closed form . . . . .                                  | 24        |
| 10.3.4    | The lateral integral, and where it must be evaluated . . . . .          | 24        |
| 10.3.5    | The principal-value piece vanishes . . . . .                            | 25        |
| 10.3.6    | Why the transferred form converges . . . . .                            | 26        |
| 10.3.7    | The tail, in closed form . . . . .                                      | 26        |
| 10.3.8    | The remainder, and the test that closes the construction . . . . .      | 26        |
| 10.4      | Step 4: how the augmentation composes with the Schwinger form . . . . . | 27        |
| <b>11</b> | <b>The thesis spectral representation, and what it removes</b>          | <b>31</b> |

|                                                                                  |           |
|----------------------------------------------------------------------------------|-----------|
| 11.1 What the thesis gives . . . . .                                             | 31        |
| 11.2 What it removes . . . . .                                                   | 32        |
| <b>12 A layer in a whole space: the external arbiter</b>                         | <b>32</b> |
| 12.1 Why a layer is the right test . . . . .                                     | 32        |
| 12.2 The closed form, by continuity . . . . .                                    | 33        |
| 12.3 The first-order route . . . . .                                             | 33        |
| 12.4 What it gives . . . . .                                                     | 34        |
| <b>13 Which formulation is more accurate</b>                                     | <b>34</b> |
| 13.1 Substituting at the one place both routes speak the same language . . . . . | 35        |
| 13.2 The measurement . . . . .                                                   | 35        |
| 13.3 Why this is not cancellation . . . . .                                      | 35        |
| 13.4 Which channel carries it . . . . .                                          | 36        |
| 13.5 What the $x_3$ preference is worth, bracketed . . . . .                     | 36        |
| <b>14 Summary of checks</b>                                                      | <b>37</b> |
| <b>15 What is not established</b>                                                | <b>39</b> |
| <b>16 Where the route belongs</b>                                                | <b>40</b> |

## 1 Conventions and objects

### 1.1 Indices, time, and units

The positive  $x_3$  axis points downward;  $x_1, x_2$  are lateral. Greek indices  $\alpha, \beta$  run over the lateral pair  $\{1, 2\}$  only, Latin indices over  $\{1, 2, 3\}$ . The time convention is  $e^{-i\omega t}$  throughout, so  $\partial_t \rightarrow -i\omega$  and the particle velocity is  $\mathbf{v} = -i\omega \mathbf{u}$ . The lateral transform is  $e^{+i(k_1 x_1 + k_2 x_2)}$ , hence  $\partial_\alpha \rightarrow ik_\alpha$ .

Numerical work uses seismic units: km,  $\text{km s}^{-1}$ ,  $\text{g cm}^{-3}$ , GPa. These are consistent, since  $\text{GPa} = \text{g cm}^{-3} \cdot (\text{km s}^{-1})^2$ .

Two orderings appear in the sources and must not be conflated. The paper orders components  $(1, 2, 3)$ ; the thesis orders them  $(z, x, y) = (x_3, x_1, x_2)$ . The map between them is written explicitly in §3.2.

### 1.2 The field vector

The matrix–vector wave equation is

$$\partial_3 \mathbf{q} = \mathcal{A} \mathbf{q} + \mathbf{d}, \quad \mathbf{q} = \begin{pmatrix} \mathbf{q}_1 \\ \mathbf{q}_2 \end{pmatrix}, \quad \mathcal{A} = \begin{pmatrix} \mathcal{A}_{11} & \mathcal{A}_{12} \\ \mathcal{A}_{21} & \mathcal{A}_{22} \end{pmatrix}, \quad (1)$$

with the elastodynamic choice

$$\mathbf{q}_1 = -\boldsymbol{\tau}_3, \quad \mathbf{q}_2 = \mathbf{v}, \quad (\boldsymbol{\tau}_j)_i = \tau_{ij}, \quad (2)$$

so that  $\mathbf{q} = (-\tau_{13}, -\tau_{23}, -\tau_{33}, v_1, v_2, v_3)^t$ . This is not an arbitrary pairing:  $-\boldsymbol{\tau}_3$  and  $\mathbf{v}$  are conjugate in the vertical power flux,  $j = \frac{1}{4}(-\boldsymbol{\tau}_3^\dagger \mathbf{v} - \mathbf{v}^\dagger \boldsymbol{\tau}_3)$ , and that is what makes the metric of §1.3 the natural one.

### 1.3 The metric $\mathbf{J}_6$

The symplectic metric is

$$\mathbf{J}_6 = \begin{pmatrix} \mathbf{0} & \mathbf{I}_3 \\ -\mathbf{I}_3 & \mathbf{0} \end{pmatrix}, \quad \mathbf{J}_6^t = -\mathbf{J}_6, \quad \mathbf{J}_6^2 = -\mathbf{I}_6. \quad (3)$$

**A naming warning.** This project, following the thesis, calls this matrix  $\mathbf{J}_6$ ; it is already defined as such in the package. The paper calls the *same* matrix  $\mathbf{N}$ , and uses  $\mathbf{J}$  for  $\text{diag}(\mathbf{I}_3, -\mathbf{I}_3)$ , a different object. Translate on the way in. Every occurrence of  $\mathbf{N}$  in the paper's equations (16)–(30) corresponds to  $\mathbf{J}_6$  here.

The operator matrix is *quasi-Hamiltonian* with respect to it,

$$\mathcal{A}^t \mathbf{J}_6 = -\mathbf{J}_6 \mathcal{A}, \quad (4)$$

which is the paper's eq. (27) and the thesis's ATdef. Everything in §6 follows from (4) together with the antisymmetry of  $\mathbf{J}_6$ , and from nothing else.

## 2 The first-order system for an isotropic medium

### 2.1 From the stiffness tensor to the stiffness matrices

For an isotropic solid,

$$c_{ijkl} = \lambda \delta_{ij} \delta_{kl} + \mu (\delta_{ik} \delta_{jl} + \delta_{il} \delta_{jk}), \quad \rho_{ij} = \rho \delta_{ij}, \quad (5)$$

with  $\lambda = \lambda(\mathbf{x}, \omega)$ ,  $\mu = \mu(\mathbf{x}, \omega)$ ,  $\rho = \rho(\mathbf{x}, \omega)$  and  $K_c = \lambda + 2\mu$ . The stiffness matrices are  $(\mathbf{C}_{jl})_{ik} = c_{ijkl}$ . In particular

$$\mathbf{C}_{33} = \text{diag}(\mu, \mu, K_c), \quad \mathbf{C}_{33}^{-1} = \text{diag}(\mu^{-1}, \mu^{-1}, K_c^{-1}), \quad (6)$$

and  $\mathbf{C}_{jl}^t = \mathbf{C}_{lj}$ , which follows from  $c_{ijkl} = c_{klij}$ . All nine blocks are reproduced from (5) and checked entry by entry.

### 2.2 Eliminating the in-plane tractions

The components  $\boldsymbol{\tau}_1, \boldsymbol{\tau}_2$  do not appear in (2) and must be removed. Doing so introduces the reduced stiffness

$$\mathbf{U}_{\alpha l} = \mathbf{C}_{\alpha l} - \mathbf{C}_{\alpha 3} \mathbf{C}_{33}^{-1} \mathbf{C}_{3l}, \quad (7)$$

with the two structural properties  $\mathbf{U}_{\alpha 3} = \mathbf{0}$  and  $\mathbf{U}_{\alpha\beta}^t = \mathbf{U}_{\beta\alpha}$ . Both are checked rather than assumed, because it is they — not the algebra that produced them — that make the symmetry relations of §3.3 hold.

The isotropic  $\mathbf{U}_{\alpha\beta}$  is where two constants are *born*:

$$v_1 \equiv (\mathbf{U}_{11})_{11} = \frac{4\mu(\lambda + \mu)}{\lambda + 2\mu}, \quad v_2 \equiv (\mathbf{U}_{12})_{12} = \frac{2\mu\lambda}{\lambda + 2\mu}, \quad v_1 - v_2 = 2\mu. \quad (8)$$

They are read off  $\mathbf{U}$  and only then compared with the published definitions; they are not independent quantities. Every  $\mathbf{U}_{\alpha\beta}$  has a null third row and column, which is the statement that  $\tau_{33}$  has been fully eliminated.

### 2.3 The operator matrices

The general construction gives

$$\mathcal{A}_{11} = -\partial_\alpha \mathbf{C}_{\alpha 3} \mathbf{C}_{33}^{-1}, \quad \mathcal{A}_{12} = i\omega\rho - \frac{1}{i\omega} \partial_\alpha \mathbf{U}_{\alpha\beta} \partial_\beta, \quad \mathcal{A}_{21} = i\omega \mathbf{C}_{33}^{-1}, \quad \mathcal{A}_{22} = -\mathbf{C}_{33}^{-1} \mathbf{C}_{3\beta} \partial_\beta, \quad (9)$$

and for the isotropic medium these evaluate to

$$\mathcal{A}_{11} = \begin{pmatrix} 0 & 0 & -\partial_1 \frac{\lambda}{\lambda+2\mu} \\ 0 & 0 & -\partial_2 \frac{\lambda}{\lambda+2\mu} \\ -\partial_1 & -\partial_2 & 0 \end{pmatrix}, \quad \mathcal{A}_{22} = \begin{pmatrix} 0 & 0 & -\partial_1 \\ 0 & 0 & -\partial_2 \\ -\frac{\lambda}{\lambda+2\mu} \partial_1 & -\frac{\lambda}{\lambda+2\mu} \partial_2 & 0 \end{pmatrix}, \quad (10)$$

$$\mathcal{A}_{12} = \begin{pmatrix} i\omega\rho - \frac{1}{i\omega}(\partial_1 v_1 \partial_1 + \partial_2 v_2 \partial_2) & -\frac{1}{i\omega}(\partial_2 \mu \partial_1 + \partial_1 v_2 \partial_2) & 0 \\ -\frac{1}{i\omega}(\partial_2 v_2 \partial_1 + \partial_1 \mu \partial_2) & i\omega\rho - \frac{1}{i\omega}(\partial_1 \mu \partial_1 + \partial_2 v_1 \partial_2) & 0 \\ 0 & 0 & i\omega\rho \end{pmatrix}, \quad (11)$$

with  $\mathcal{A}_{21} = i\omega \text{diag}(\mu^{-1}, \mu^{-1}, K_c^{-1})$ .

**The ordering is the content.** Compare the (1, 3) entry of  $\mathcal{A}_{11}$  with the (3, 1) entry of  $\mathcal{A}_{22}$ . In the first the derivative stands to the *left* of the modulus and differentiates the product; in the second it stands to the *right* and differentiates the field alone. For a laterally invariant medium the distinction collapses and is invisible. For a laterally varying one it is the physics. This is why the checks in §3.1 compare *operators* — by their action on generic fields — rather than expressions, and why they are accompanied by a negative control that corrupts precisely this ordering.

## 3 $\mathcal{A}$ against the thesis, and against the published operator form

**A note on where this starts.** An earlier version of this note presented the published operator matrices as the foundation and the thesis system matrix as an arbiter for them. That is the wrong way round. The thesis carries  $\mathbf{A}(k_x, k_y)$  in §2.1, equation (Akdef), and the purpose of this work is to extend and test the thesis, not to check the thesis against a later paper. What the operator form genuinely contributes is narrower: the thesis has the system only in the laterally *invariant* case, where  $\partial_\alpha \rightarrow ik_\alpha$  has already

been taken, and a cube requires the coefficients to vary laterally. So the operator form is a *generalisation* of (Akdef), and it is that generalisation — not the system — that is new here. The consequences of getting this backwards were not only presentational; they are collected in §11.

### 3.1 Arbiter A: the published isotropic operator matrices

The blocks (10)–(11) are transcribed separately and compared against those built from (9). The comparison is made by acting on a vector of generic symbolic functions, so it tests the derivative ordering as well as the coefficients. All four match.

A check that cannot fail is worth nothing, so a corrupted  $\mathcal{A}_{22}$  — with the derivative wrongly placed outside the modulus — is also tested, and is rejected.

### 3.2 The anchor: the thesis system matrix (Akdef)

The thesis carries a  $6 \times 6$  Fourier-domain system matrix  $\mathbf{A}(k_x, k_y)$  for a laterally invariant medium, in a *different* field vector

$$\mathbf{b} = (u_3, u_1, u_2, \tau_{33}, \tau_{31}, \tau_{32})^t, \quad (12)$$

i.e. displacement–stress rather than velocity–traction, ordered  $(z, x, y)$  rather than  $(1, 2, 3)$ , and normalised so that  $\rho\omega^2$  appears where the paper has  $i\omega\rho$ . It uses the shorthand

$$\gamma = \frac{\lambda}{\lambda + 2\mu}, \quad a = \frac{1}{\lambda + 2\mu}, \quad b = \frac{1}{\mu}, \quad \zeta = \frac{4\mu(\lambda + \mu)}{\lambda + 2\mu}, \quad \chi = \frac{2\mu\lambda}{\lambda + 2\mu}. \quad (13)$$

The thesis attaches a warning to these that must travel with them: the shorthand  $\gamma, a, b, \zeta, \chi$  is *local to* (Akdef), and is *not* the  $\gamma, \zeta, \chi$  of the List of Symbols, which denote the Chapter 5 material-contrast ratios. Since this note uses  $\zeta$  and  $\chi$  throughout as the in-plane stiffnesses that survive the elimination of  $\varepsilon_{33}$ , the collision is worth stating once rather than inheriting silently. With that said,  $\zeta - \chi \equiv 2\mu$ , which is the thesis’s own footnote and the cleanest check on the pair.

**Comparing (8) with these:**  $\zeta \equiv \nu_1$  and  $\chi \equiv \nu_2$ . The thesis’s in-plane stiffnesses and the paper’s auxiliary moduli are the same two quantities, born in the same elimination (7). This is verified, not asserted.

Setting  $\lambda, \mu, \rho$  constant and acting on  $e^{i(k_1 x_1 + k_2 x_2)}$  reduces  $\mathcal{A}$  to a matrix. With the constant map  $\mathbf{b} = \mathbf{S}\mathbf{q}$  built from  $\mathbf{u} = \mathbf{v}/(-i\omega)$ ,  $\boldsymbol{\tau}_3 = -\mathbf{q}_{1:3}$  and the permutation  $(1, 2, 3) \rightarrow (3, 1, 2)$ , the identity

$$\mathbf{A}_{\text{thesis}} = \mathbf{S} \mathbf{A}_{\text{paper}} \mathbf{S}^{-1} \quad (14)$$

holds entry by entry. Since the two matrices were derived independently and published separately, this is evidence rather than bookkeeping.

### 3.3 Structural checks

Three further properties are verified on the reduced  $6 \times 6$ :

1. the operator symmetry relations  $\mathcal{A}_{11}^t = -\mathcal{A}_{22}$ ,  $\mathcal{A}_{12}^t = \mathcal{A}_{12}$ ,  $\mathcal{A}_{21}^t = \mathcal{A}_{21}$ , where the transpose is the formal adjoint under integration by parts ( $\partial_\alpha^t = -\partial_\alpha$ );
2. the quasi-Hamiltonian relation (4), which in the wavenumber domain reads  $\mathbf{A}^t(-\mathbf{k})\mathbf{J}_6 = -\mathbf{J}_6\mathbf{A}(\mathbf{k})$ ;
3. the characteristic polynomial factorises as  $(s^2 + k_{z,p}^2)(s^2 + k_{z,s}^2)^2$ , recovering  $\alpha^2 = (\lambda + 2\mu)/\rho$  and  $\beta^2 = \mu/\rho$  from a construction in which no wave speed was ever written down.

Item 3 is the one that binds the algebra to physics: nothing upstream knows the wave speeds.

## 4 The contrast operator

### 4.1 Definition and the Lippmann–Schwinger equation

Inside the scatterer the medium is  $(\lambda + \Delta\lambda, \mu + \Delta\mu, \rho + \Delta\rho)$ . Writing  $\mathcal{A}^{\text{tot}} = \mathcal{A} + \Delta\mathcal{A}$ , the total field solves the background system driven by its own contrast,

$$(\partial_3 - \mathcal{A})\mathbf{q} = \Delta\mathcal{A}\mathbf{q} + \mathbf{d}, \quad (15)$$

which is the Lippmann–Schwinger equation of the first-order system.

### 4.2 The foundation check

Before (15) means anything,  $\mathbf{q}$  must actually satisfy (1). Building  $\mathbf{q}$  from a *generic* displacement field in a *laterally varying* medium and forming  $\partial_3\mathbf{q} - \mathcal{A}\mathbf{q}$  gives two halves with different strengths:

- rows 4–6 vanish **identically**, for any  $\mathbf{u}$  whatever. They are the constitutive relation rearranged; no equation of motion is used. Failure here would mean the elimination (7) is wrong.
- rows 1–3 equal **exactly** the Navier body force  $f_i = -\rho\omega^2 u_i - \partial_j\tau_{ij}$ . Anything left over would be a defect in  $\mathcal{A}_{11}$  or  $\mathcal{A}_{12}$ .

Both hold. This is a considerably stronger statement than the plane-wave and eigenvalue checks of §3.3, because  $\mathbf{u}$  is generic and the moduli vary laterally, so every ordering term in (10) is exercised.

### 4.3 $\Delta\mathcal{A}$ is not linear in the contrast

The Navier potential is exactly linear in  $(\Delta\lambda, \Delta\mu, \Delta\rho)$ .  $\Delta\mathcal{A}$  is not, because the elimination inverted  $\mathbf{C}_{33}$ :

$$(\Delta\mathcal{A}_{21})_{11} = i\omega \left( \frac{1}{\mu + \Delta\mu} - \frac{1}{\mu} \right) = -\frac{i\omega \Delta\mu}{\mu(\mu + \Delta\mu)}, \quad (16)$$

whose ratio to its linearisation  $-i\omega\Delta\mu/\mu^2$  is exactly  $\mu/(\mu + \Delta\mu)$ . At the project’s validated moderate contrast ( $\alpha = 5$ ,  $\beta = 3$ ,  $\rho = 2.5$ , so  $\lambda = 17.5$ ,  $\mu = 22.5$  GPa;  $\Delta\lambda = 2$ ,  $\Delta\mu = 1$  GPa):

| quantity                                      | value   |
|-----------------------------------------------|---------|
| $\Delta\mu/\mu$                               | 0.04444 |
| exact / linear for $\Delta b = \Delta(1/\mu)$ | 0.95745 |
| exact / linear for $\Delta a = \Delta(1/K_c)$ | 0.93985 |

The first-order form has already summed part of the local response exactly. §7 measures what that is worth, and finds the answer is not the obvious one.

## 5 The augmentation

### 5.1 Why a derivative must be moved

Inside the scatterer the contrasts are constant, so every coefficient of  $\Delta\mathcal{A}$  carries the indicator  $\mathbf{1}_V$ . A derivative standing to the *left* of such a coefficient differentiates the indicator and produces a surface delta on the cube faces. Integration by parts over all space moves that derivative onto the test function instead, where it meets something smooth:

$$\int \phi_i (-\partial_\alpha (m \mathbf{1}_V \psi_j)) d\mathbf{x} = + \int (\partial_\alpha \phi_i) m \mathbf{1}_V \psi_j d\mathbf{x}, \quad (17)$$

and the face term never arises. This is the step the thesis describes as transferring the derivatives onto the source and receiver coordinates of the propagator, and it is what converts a singular integro-differential equation into a nonsingular, compact integral equation.

### 5.2 The augmented dimension

After the transfer every term has the form  $(\partial^a \phi)^t \cdot \text{const} \cdot (\partial^b \psi)$  with  $a, b \in \{0, 1\}$ , so the form is *multiplicative* on the enlarged vector  $\mathcal{L}\psi = (\psi, \partial_1 \psi, \partial_2 \psi)$  — nominally  $6 + 6 + 6 = 18$  slots. The true dimension is smaller and is counted, not assumed.

**Result.** Exactly  $10 = 6 + 4$  slots are used. The four extra are  $\partial_1 v_1, \partial_2 v_1, \partial_1 v_2, \partial_2 v_2$ : the lateral gradients of the two *horizontal* velocity components, and nothing else.

Setting  $\partial_2 \rightarrow ik_y$ , as a  $2\frac{1}{2}$ -D formulation does, makes two of those four multiplicative, leaving  $6 + 2 = 8$ . That is precisely the thesis’s count for its augmented coupling matrix, reproduced here from an independent 3-D Cartesian derivation. The agreement is a corroboration of the thesis construction, not a restatement of it.

The transfer identity is verified three ways: the augmented matrix is constant (no derivative operators survive); the augmented form reproduces the transferred bilinear; and the untransferred and transferred bilinears differ by an explicitly exhibited lateral divergence  $\partial_\alpha P_\alpha$ .

## 6 The test space, derived

### 6.1 The argument

At this point there are ten augmented slots and, for an affine trial space, nine parameters. The temptation is to choose nine conditions. That temptation must be resisted: a subset chosen by scoring against the arbiter it is later judged by is a fitted parameter wearing a derivation's clothes.

There is no choice to make.  $\mathbf{J}_6$  is *antisymmetric*, and (4) applies to  $\Delta\mathcal{A}$  as well as to  $\mathcal{A}$  (it is linear in the operator), so

$$(\mathbf{J}_6 \Delta\mathcal{A})^t = \Delta\mathcal{A}^t \mathbf{J}_6^t = \Delta\mathcal{A}^t (-\mathbf{J}_6) = \mathbf{J}_6 \Delta\mathcal{A}. \quad (18)$$

$\mathbf{J}_6 \Delta\mathcal{A}$  is a **symmetric** matrix, so

$$b[\phi, \psi] = \phi^t \mathbf{J}_6 \Delta\mathcal{A} \psi \quad (19)$$

is a symmetric bilinear form. A symmetric form admits exactly one natural projection:

$$\boxed{\text{test space} = \mathbf{J}_6 \times \text{trial space.}} \quad (20)$$

Verified, together with the negative control: drop  $\mathbf{J}_6$  and the symmetry is lost, both at operator level and in the projected system.

### 6.2 The ten-versus-nine mismatch dissolves

The ten slots were never test conditions. They measure how many derivative channels the *form* touches — its representation — not how many independent functionals it supplies. Each trial function  $\psi_k$  gives exactly one test functional  $(\mathbf{J}_6 \psi_k)^t$ , so the number of conditions equals the trial dimension by construction.

### 6.3 The effective coupling matrix

For the affine trial space — nine parameters, three rigid translations and six uniform strains — the projected coupling

$$(\Delta\mathbf{C}_{\text{eff}})_{kl} = \int_V (\mathcal{L}\psi_k)^t \Delta\mathbf{C}(\mathcal{L}\psi_l) dV \quad (21)$$

is the 3-D Cartesian analogue of the thesis's effective coupling matrix. It is  $9 \times 9$ , symmetric, of full rank 9, vanishes with the contrast, and — because the trial functions are polynomials — requires **no Green's function at all**: the cube integral is exact.

The rigid-translation block is diagonal,

$$(\Delta\mathbf{C}_{\text{eff}})_{1:3,1:3} = i\omega^3 \Delta\rho V \mathbf{I}_3, \quad V = (2a)^3. \quad (22)$$

**The power of  $\omega$  is a basis statement, not a physical one.**  $\Delta\mathcal{A}$ 's density entry is  $i\omega\Delta\rho$  acting on the velocity slot; the trial carries  $\mathbf{v} = -i\omega\mathbf{u}$  there; and  $\mathbf{J}_6$  pulls the test function's

velocity slot into the traction rows, supplying a second  $-i\omega$ . Hence  $i\omega \cdot (-i\omega) \cdot (-i\omega) = i\omega^3$ . The displacement-basis answer  $i\omega\Delta\rho V$  is wrong by  $\omega^2$ . This is the displacement-versus-flux confusion in a new place, and the check asserts the correct factor *and* asserts the incorrect one is incorrect, so it cannot drift back.

## 7 What the first-order coupling actually buys

### 7.1 The trial spaces coincide

It is tempting to attribute any difference from the Navier formulation to the nonlinearity of §4.3, or to a different ansatz. Neither is right. Inside the scatterer the moduli are constant, so  $\boldsymbol{\tau}_3 = \mathbf{C}_{33}^{\text{tot}} \boldsymbol{\varepsilon}_{i3}$  is a constant *invertible* map. A uniform strain gives a uniform  $\boldsymbol{\tau}_3$  and an affine  $\mathbf{v}$ , and conversely: “ $\boldsymbol{\varepsilon}$  uniform” and “ $\boldsymbol{\tau}_3$  uniform plus in-plane  $\boldsymbol{\varepsilon}$  uniform” describe the same nine-parameter family.

The two formulations are therefore **different projections of the same residual onto the same trial space**. The difference lives in the test conditions.

### 7.2 The density channel cannot differ

With  $\Delta\lambda = \Delta\mu = 0$ ,  $\Delta\mathcal{A}$  is exactly  $i\omega\Delta\rho$  in rows 1–3, columns 4–6, and reaches *no* part of the stress-glut slot: a density contrast is a pure body force. The propagator’s body-force velocity block is  $-i\omega\mathbf{G}$ , so both routes return

$$A_u = \frac{1}{1 - \omega^2\Delta\rho\Gamma_0}, \quad \Gamma_0 = \int_V \mathbf{G} dV, \quad (23)$$

identically. The modulus channels are the only place the two can part.

### 7.3 The units the two couplings are written in

The statement just made is about the *amplification factor*, which is a ratio, and ratios are where this comparison is safe. The couplings themselves are not in the same units, and it is worth saying so plainly because every other check in this note is also a ratio, and none of them can see it.

The  $\mathbf{J}_6$  pairing tests against the field  $q$ , whose velocity slots carry  $v = -i\omega u$ . The Navier weak form, and this package’s local  $T$ -matrix, test against the *displacement*. So

$$\Delta\mathbf{C}_{\text{eff}} = i\omega \times [\text{the same object in displacement units}], \quad \Delta\mathbf{C}_{\text{eff}}|_{\text{rigid}} = i\omega^3\Delta\rho V \text{ against } \omega^2\Delta\rho V. \quad (24)$$

This is a change of units and not of physics. It cancels out of every amplification factor, every channel ratio and every Born limit expressed as a ratio; it does *not* cancel the moment  $\Delta\mathbf{C}_{\text{eff}}$  is combined with an object built in displacement units. The  $i\omega$  appearing in (70) is this same factor seen from the other side.

The check that establishes it is necessarily absolute rather than relative: after dividing by  $i\omega$ , the residual  $\|\Delta\mathbf{C}_{\text{eff}}/i\omega - \Delta\mathbf{C}_{\text{Nav}}\|/\|\Delta\mathbf{C}_{\text{Nav}}\|$  falls as  $4.19 \times 10^{-2}$ ,  $4.19 \times 10^{-3}$ ,  $4.19 \times 10^{-4}$ ,  $4.19 \times 10^{-5}$  under decade reductions of the contrast — the *first* power, so the residual is the genuine  $O(\Delta c^2)$

difference of §7 against an  $O(\Delta c)$  coupling, and not a second, smaller units error hiding beneath the first.

#### 7.4 Born agreement, and an $O(\Delta c^2)$ difference

Reading a proportionality constant off the rigid block alone gives  $i\omega$ , and testing it on the remaining eighty entries **refutes** the hypothesis  $\Delta \mathbf{C}_{\text{eff}} = i\omega \Delta \mathbf{C}_{\text{Navier}}$ . The two couplings agree exactly at  $O(\Delta c^0)$  and  $O(\Delta c^1)$  — identical Born limits, so nothing is broken — and differ from second order.

**Where.** Eleven of eighty-one entries differ. With the parameter order  $(c_1, c_2, c_3, e_{11}, e_{22}, e_{33}, e_{12}, e_{13}, e_{23})$  the nonzero block is  $\{e_{11}, e_{22}, e_{33}\}^2$  together with  $(e_{13}, e_{13})$  and  $(e_{23}, e_{23})$ .

That is exactly the set of strains  $\tau_3$  sees:

$$\tau_{13} = 2\mu e_{13}, \quad \tau_{23} = 2\mu e_{23}, \quad \tau_{33} = \lambda(e_{11} + e_{22} + e_{33}) + 2\mu e_{33}. \quad (25)$$

The in-plane shear  $e_{12}$  is **untouched**, as it must be:  $\tau_{12}$  was eliminated, not resummed. This was a falsifiable prediction about which entries may be nonzero, and it held. A representative entry is

$$(\Delta \mathbf{C}_{\text{eff}} - i\omega \Delta \mathbf{C}_{\text{Navier}})_{e_{11}e_{11}} = -\frac{i\omega V \Delta \lambda^2}{\lambda + 2\mu}, \quad (26)$$

pure second order, divided by  $K_c$ : the remainder of the geometric series for  $1/(K_c + \Delta K_c)$ .

#### 7.5 Half of it is a broken symmetry

An isotropic cube in an isotropic whole space has an  $O_h$ -symmetric  $T$ -matrix; no direction is special. But the first-order system is a recursion in *depth*: it eliminated  $\tau_1, \tau_2$  and kept  $\tau_3$ , so  $x_3$  is preferred by construction. The tell is already above —  $(e_{13}, e_{13})$  and  $(e_{23}, e_{23})$  differ while  $(e_{12}, e_{12})$  does not, and those three are  $O_h$ -equivalent.

Applying the cyclic element  $x \rightarrow y \rightarrow z \rightarrow x$  (positive controls:  $P^3 = \mathbf{I}$ , and the Navier coupling is invariant):

| contrast (GPa)              | invariant             | breaking              | ratio  |
|-----------------------------|-----------------------|-----------------------|--------|
| pure shear, $\Delta\mu = 1$ | $2.62 \times 10^{-2}$ | $1.94 \times 10^{-2}$ | 1.3520 |
| moderate, (2, 1, 0.1)       | $5.17 \times 10^{-4}$ | $2.93 \times 10^{-4}$ | 1.7669 |
| strong, (20, 10, 0.5)       | $1.05 \times 10^{-2}$ | $5.92 \times 10^{-3}$ | 1.7669 |

**Table 1.** Decomposition of  $\Delta \mathbf{C}_{\text{eff}} - i\omega \Delta \mathbf{C}_{\text{Navier}}$ , relative to  $\|\Delta \mathbf{C}_{\text{eff}}\|$ , at  $\lambda = 17.5$ ,  $\mu = 22.5$  GPa. The two mixed rows share a contrast *direction* ( $\Delta\lambda:\Delta\mu = 2:1$ ) at different magnitudes and return the identical ratio, so the split is a function of contrast direction alone — as a leading-order  $O(\Delta c^2)$  split must be.

So the difference is neither pure artefact nor pure physics. It carries both a genuine  $O_h$ -invariant resummation the Navier form lacks, and a spurious preference for  $x_3$ .

**A caveat on this test.** The group used is the cyclic subgroup of order 3, not all 48 elements of  $O_h$ . Invariance under it is necessary for  $O_h$  invariance, not sufficient.

## 7.6 What follows, and what does not

Two readings are available and the measurement above does not separate them.

*Against the route:* an approximation that breaks a symmetry the exact answer possesses has introduced an error that no propagator can remove.

*For the route:* the correction is not scattered. It lands exactly on the traction channel, treating  $\mathbf{C}_{33}$  exactly instead of leaving it to the Neumann series, and leaves the in-plane channel alone. Vertical  $P$  incidence excites  $e_{33}$ ; vertical  $S$  incidence excites  $e_{13}$  or  $e_{23}$ ; both are corrected. The in-plane shear, which the side-scattered field carries, is unchanged. If the response is dominated by reflection and transmission, the asymmetry is “z treated better” rather than “z treated worse”.

**Consequence for any comparison.** Do not score the raw coupling against an independent arbiter in a homogeneous background. Three effects are in play at comparable magnitude — the Galerkin-versus-collocation projection, the  $O_h$ -invariant resummation, and the artefact — and a single number separates none of them. Any comparison must be channel-resolved.

Whether the resummation is an *improvement* is not established here. The local inversion  $1/(K_c + \Delta K_c)$  is exact; whether that makes the closed  $T$ -matrix better depends on how the ansatz error and the series truncation interact, and those do not simply add.

## 8 The propagator, and the bridge to it

### 8.1 It already exists

Everything downstream needs the Green’s matrix of (1). It does not need to be built: the stratified  $6 \times 6$  Green’s function and its validated wrapper already return it, in the basis  $(u_z, u_x, u_y, T_{zz}, T_{xz}, T_{yz})$  — which is exactly the basis the thesis writes  $\mathbf{A}$  in. That the two are the same object had, however, been *assumed*.

### 8.2 The bridge identity

Away from the source plane the Green’s matrix solves the homogeneous system in its receiver coordinate, so for two receiver depths in one uniform layer

$$\mathbf{\Gamma}(z_2, z') = \exp(\mathbf{A}(z_2 - z_1)) \mathbf{\Gamma}(z_1, z'), \quad (27)$$

with no free parameters. The identity is *local*: it holds however complicated the medium between the source and  $z_1$  is, because the ODE knows only the layer the receiver sits in. The gate therefore places the source above a fast slab and the receivers below it, so transmission, reflection, conversion and interbed multiples are all present in  $\mathbf{\Gamma}$  and (27) must still hold exactly.

### 8.3 Reporting residuals against the arithmetic floor

A bare relative residual is not comparable across incidence angles. Beyond the critical slowness the propagator carries both a growing and a decaying evanescent branch, so  $\exp(\mathbf{A} dz)$  has dynamic range  $\sim e^{|k|dz}$  — about  $1.6 \times 10^7$  at the widest angle used. Forming  $\exp(\mathbf{A} dz) \Gamma(z_1)$  then amplifies the rounding already present by  $\|\exp(\mathbf{A} dz)\|$ , and no correct implementation can beat

$$\text{floor} = \epsilon \|\exp(\mathbf{A} dz)\| \frac{\|\Gamma(z_1)\|}{\|\Gamma(z_2)\|}. \quad (28)$$

Residuals below are quoted in units of that floor, so  $O(1)$  means the identity holds to the last bit available at that conditioning.

| $k_x$                                         | $k_y$ | $ k /\omega$ | raw                   | corrected             |
|-----------------------------------------------|-------|--------------|-----------------------|-----------------------|
| 2.00                                          | 1.30  | 0.0759       | $7.68 \times 10^{13}$ | $1.26 \times 10^{-2}$ |
| 0.40                                          | 0.05  | 0.0128       | $2.95 \times 10^{14}$ | $2.43 \times 10^{-2}$ |
| 14.00                                         | 9.00  | 0.5298       | $4.31 \times 10^{13}$ | $9.37 \times 10^{-3}$ |
| control, $\zeta \leftrightarrow \chi$ swapped |       |              | $7.68 \times 10^{13}$ | $1.36 \times 10^{12}$ |

**Table 2.** Residual of (27) in units of the arithmetic floor.  $\omega = 31.4159 \text{ rad s}^{-1}$ ,  $dz = 1 \text{ km}$ . The third row is evanescent for both branches.

**The corrected object is the physical state vector of  $\partial_3 \mathbf{q} = \mathcal{A} \mathbf{q}$ ;** the raw one misses by fourteen orders of magnitude.

The negative control swaps  $\zeta$  and  $\chi$  — two in-plane stiffnesses that differ by exactly  $2\mu$ , a plausible transcription slip that no symmetry check would catch. It is rejected by fourteen orders.

## 9 The single-site $T$ -matrix

This is the object the construction was for. On the nine-parameter affine trial space,

$$\mathbf{T} = \Delta \mathbf{C}_{\text{eff}} (\mathbf{I}_9 - \mathcal{G} \Delta \mathbf{C}_{\text{eff}})^{-1}, \quad (29)$$

with  $\Delta \mathbf{C}_{\text{eff}}$  the projected coupling of §6 and  $\mathcal{G}$  the self propagator moment of the cube.

### 9.1 What is new and what is borrowed

New is  $\Delta \mathbf{C}_{\text{eff}}$ . Borrowed is  $\mathcal{G}$ : the first-order route changes the *coupling*, not the Green's function. The two moments it needs are

$$\mathbf{G}_{ij} = \int_V G_{ij}^0 dV, \quad M_{in,pk} = \int_V \partial_p' \partial_k' G_{in}^0 dV. \quad (30)$$

These are *derived* in §9.2 below, not taken from a table. In the parameter basis

$$\mathcal{G} = \frac{1}{V} \begin{pmatrix} \Gamma_0 \mathbf{I}_3 & \mathbf{0} \\ \mathbf{0} & -\mathbf{S} \mathbf{W}^{-1} \end{pmatrix}, \quad \mathbf{W} = \text{diag}(1, 1, 1, 2, 2, 2), \quad (31)$$

where  $\mathbf{W}$  undoes the engineering doubling of the shear slots.

Note that (31) is asserted, not derived — which is exactly why §9.5 exists.

## 9.2 The self moments, derived

The integrand of  $M$  goes as  $1/r^3$ , which is not absolutely integrable in three dimensions. The customary remedy is to quote a tabulated Eshelby tensor and add a delta-function term at  $r = 0$  by hand; omit the delta and the depolarisation comes out with the *wrong sign*. Nothing below is tabulated and nothing is added by hand.

### 9.2.1 The static Green's tensor

In the static limit the elastodynamic tensor reduces to Kelvin's,

$$G_{in}^0(\mathbf{r}) = \frac{a_0 \delta_{in}}{r} + \frac{b_0 x_i x_n}{r^3}, \quad a_0 = \frac{1}{8\pi} \left( \frac{1}{\mu} + \frac{1}{\lambda + 2\mu} \right), \quad b_0 = \frac{1}{8\pi} \left( \frac{1}{\mu} - \frac{1}{\lambda + 2\mu} \right), \quad (32)$$

the coefficients being fixed by requiring that (32) satisfy the Navier operator away from the origin. An equivalent form makes the distributional content visible, since  $\partial_i \partial_n r = \delta_{in}/r - x_i x_n/r^3$ :

$$G_{in}^0 = \frac{1}{4\pi\mu} \left[ \frac{\delta_{in}}{r} + h_c \partial_i \partial_n r \right], \quad h_c = -\frac{\lambda + \mu}{2(\lambda + 2\mu)}, \quad (33)$$

and  $a_0 = (1 + h_c)/4\pi\mu$ ,  $b_0 = -h_c/4\pi\mu$  identically.

### 9.2.2 The moment as a distribution, and why no delta is inserted

Write the moment as a pairing against the test function  $w \mathbf{1}_V$  rather than as an integral of a function. With  $T = \partial_{\text{rest}} r^m$  and  $\partial_q \mathbf{1}_V = -n_q \delta_{\partial V}$ ,

$$\langle \partial_q T, w \mathbf{1}_V \rangle = \oint_{\partial V} n_q w T dA - \langle T, (\partial_q w) \mathbf{1}_V \rangle. \quad (34)$$

That is the entire construction. Two features make it exact rather than merely plausible:

- $T$  is only ever evaluated *on the boundary*, at distance  $\Delta/2$  from the origin, where every derivative of  $r^m$  is smooth and bounded. The singular pointwise structure of  $\partial^D r^m$  at the origin is never formed, so there is no principal value to take and no delta to insert.
- the recursion strictly lowers both the derivative count and the weight degree, terminating at  $\int_V x_w r^m dV$ , an ordinary convergent integral for  $m \geq -1$ .

The delta term is therefore neither added nor droppable — it is already inside the surface

integral. The sharp test is the Laplacian sum rule  $\sum_p E[-1; \{p, p\}] = -4\pi$ , which a spherical excision returns as 0.

### 9.2.3 The face integrals, and one of them by pure geometry

On the face  $x_3 = \Delta/2$ , scaling the in-face coordinates by  $\Delta/2$  reduces every surface term to one of three integrals over  $[-1, 1]^2$ :

$$j_1 = \iint \frac{du dv}{(1 + u^2 + v^2)^{3/2}} = \frac{2\pi}{3}, \quad j_2 = \iint \frac{u^2 du dv}{(1 + u^2 + v^2)^{5/2}} = -\frac{2(\sqrt{3} - \pi)}{9}, \quad k_1 = \iint \frac{du dv}{(1 + u^2 + v^2)^{5/2}} = \frac{2(2\sqrt{3} + \pi)}{9}. \quad (35)$$

All three are evaluated in closed form. The first needs no integration at all:  $j_1$  is the solid angle subtended by one face at the cube's centre, and the six faces subtend  $4\pi$ , so  $j_1 = 4\pi/6 = 2\pi/3$ . That identity is a check on the scaling and orientation of (35), obtained independently of any quadrature.

The zeroth moment needs one further integral, which also has a closed form:

$$\int_{[-1,1]^3} \frac{dV}{r} = -2(\pi + \log 64 - 12 \log(1 + \sqrt{3})) = \frac{4}{3} \log(70226 + 40545\sqrt{3}) - 2\pi = 9.520309455918214 \dots \quad (36)$$

The two closed forms are independent evaluations of the same constant and agree to the precision carried; the second is the one this package uses elsewhere.

### 9.2.4 The closed forms

Assembling (32) through the recursion (34) with the face integrals (35) gives, with

$$M_{in,pk} = A \delta_{in} \delta_{pk} + B (\delta_{ip} \delta_{nk} + \delta_{ik} \delta_{np}) + C E_{inpk}, \quad E_{inpk} = \begin{cases} 1 & i = n = p = k \\ 0 & \text{otherwise,} \end{cases} \quad (37)$$

the third structure being the one a cubic — as opposed to isotropic — geometry admits,

$$B = \frac{\sqrt{3}(\lambda + \mu)}{6\pi\mu(\lambda + 2\mu)}, \quad (38)$$

$$A = B - \frac{1}{3\mu}, \quad (39)$$

$$C = -\left(5 - \frac{2\pi}{\sqrt{3}}\right)B, \quad (40)$$

$$G_{ij} = \delta_{ij} \left(a_0 + \frac{b_0}{3}\right) \left(\frac{\Delta}{2}\right)^2 \int_{[-1,1]^3} \frac{dV}{r}. \quad (41)$$

$A$ ,  $B$  and  $C$  are scale-free;  $G$  carries  $\Delta^2$ . Note that  $A$ ,  $B$  and  $C$  are *not* three independent constants — the whole rank-four moment is carried by  $B$  together with the elementary  $1/3\mu$ , which is a structural statement the tabulated presentation obscures.  $G_{ij}$  must be a multiple of  $\delta_{ij}$ , because a rank-two cubic invariant has no other form: the cube cannot distinguish its own axes at that order.

### 9.2.5 The check that certifies the delta content

From (33) with  $\nabla^2(1/r) = -4\pi\delta^3$  and  $\nabla^2(\partial_i\partial_n r) = 2\partial_i\partial_n(1/r)$ , whose cube integral is  $-\frac{4\pi}{3}\delta_{in}$ ,

$$\sum_p M_{in,pp} = -\frac{\delta_{in}}{\mu} \left(1 + \frac{2h_c}{3}\right) = -\delta_{in} \frac{2\lambda + 5\mu}{3\mu(\lambda + 2\mu)}, \quad (42)$$

derived without reference to the recursion at all. The assembled moment satisfies (42) identically, and the export refuses to write its output otherwise. A spherically excised moment returns zero on the left-hand side, so (42) is exactly the statement that the  $r = 0$  content is present and correctly weighted.

## 9.3 The tabulated route agrees

This package also carries the conventional tabulated Eshelby route, in which  $j_1$ ,  $j_2$  and  $k_1$  are quoted rather than integrated. Evaluated alongside (38)–(41) at  $\lambda = 17.5$ ,  $\mu = 22.5$  GPa,  $\Delta = 1$ :

|   | moment                            | integrated                        | tabulated | relative              |
|---|-----------------------------------|-----------------------------------|-----------|-----------------------|
| A | $-1.220110745874 \times 10^{-11}$ | $-1.220110745874 \times 10^{-11}$ |           | 0                     |
| B | $2.613707356074 \times 10^{-12}$  | $2.613707356074 \times 10^{-12}$  |           | $3.1 \times 10^{-16}$ |
| C | $-3.587055298887 \times 10^{-12}$ | $-3.587055298887 \times 10^{-12}$ |           | $2.3 \times 10^{-16}$ |
| G | $6.622004020056 \times 10^{-12}$  | $6.622004020056 \times 10^{-12}$  |           | 0                     |

**Table 3.** The table reproduces the integration to machine precision. The direction of that statement matters: the integration *verifies* the table; the assembly reads the integrated values.

The dynamic (smooth radiation) part carries no table — it is already exact polynomial integration — and is taken from the existing route with its tabulated static piece removed and the integrated one substituted.

## 9.4 Two implementations of the coupling

$\Delta\mathbf{C}_{\text{eff}}$  is rebuilt in Python directly from  $\Delta\mathcal{A}$ , independently of the symbolic derivation, and compared against the values that derivation exports. Worst relative difference over three contrasts:  $1.4 \times 10^{-15}$  for  $\Delta\mathbf{C}_{\text{eff}}$  and exactly zero for the Navier coupling. Two implementations agreeing is the evidence standard; one agreeing with itself is not.

## 9.5 Calibration: every normalisation fixed against existing code

Run the *same* assembly with the ordinary Navier coupling in place of  $\Delta\mathbf{C}_{\text{eff}}$ . It must reproduce this package’s validated four amplification factors. That pins the volume factors, the sign of  $\mathcal{G}$  and the weight  $\mathbf{W}$  — the *bookkeeping* of (31) — with no freedom left to tune. It does not certify the moments themselves; that is what the integration of the previous subsection is for. Against  $\alpha = 5000$ ,  $\beta = 3000$ ,  $\rho = 2500$  (SI),  $\omega = 60$ ,  $a = 0.5$ :

| case                              | factor                                               | assembled vs existing (rel.)                         |
|-----------------------------------|------------------------------------------------------|------------------------------------------------------|
| pure shear ( $\Delta\mu = 1$ GPa) | $A_u, A_\theta, A_e^{\text{diag}}, A_e^{\text{off}}$ | 0, $2.2 \times 10^{-16}$ , 0, 0                      |
| moderate (2, 1, 100)              | same four                                            | 0, $2.3 \times 10^{-16}$ , $1.1 \times 10^{-16}$ , 0 |

**Table 4.** The assembly reproduces the existing amplification factors to machine precision. Only after this is the coupling swapped.

**One term had to be isolated to get there.** The weak form's  $\int_V \Delta\rho \omega^2 \phi \cdot \mathbf{u} dV$  contributes to the *strain*–strain block as well as the rigid one, because an affine trial displacement is not constant:  $\int_V x_a x_b dV = \delta_{ab} V a^2 / 3$ . This is a finite-size dynamic coupling between the density contrast and the strain channels, and the factorised amplification factors do not carry it — there  $A_u$  and the strain amplifications are independent. It shifts  $A_\theta$  by  $1.5 \times 10^{-7}$  at the moderate contrast and vanishes identically with  $\Delta\rho$ . It is a small correction that the factorisation drops, not a defect in either object, and it is reported rather than absorbed.

## 9.6 The first-order $T$ , channel by channel

A single norm would average a correction that lives in a few channels against others that do not move. The comparison is therefore resolved by channel.

| case       | channel                       | relative difference from the Navier route |
|------------|-------------------------------|-------------------------------------------|
| pure shear | $\theta$ (bulk)               | $1.13 \times 10^{-4}$                     |
|            | $e_{\text{diag}}$             | $1.1 \times 10^{-16}$                     |
|            | $e_{12}$ (in-plane shear)     | 0                                         |
|            | $e_{13}$ (out-of-plane shear) | <b><math>8.36 \times 10^{-4}</math></b>   |
|            | $u$ (density)                 | 0                                         |
| moderate   | $\theta$ (bulk)               | $1.74 \times 10^{-3}$                     |
|            | $e_{\text{diag}}$             | $1.1 \times 10^{-16}$                     |
|            | $e_{12}$                      | 0                                         |
|            | $e_{13}$                      | <b><math>8.36 \times 10^{-4}</math></b>   |
|            | $u$                           | 0                                         |

**Table 5.** Amplification  $(\mathbf{I} - \mathcal{G}\Delta\mathbf{C})^{-1}$  along each channel. The density channel is identical in both routes, as §7 requires; the in-plane shear is untouched; the bulk and out-of-plane shear channels move.

## 9.7 The broken symmetry, now visible in $T$ itself

$e_{12}$  and  $e_{13}$  are equivalent under  $O_h$ . For an isotropic cube in an isotropic background their amplifications *must* be equal. They are, in the Navier route, to exactly zero. In the first-order route they split:

|                                         | Navier | first-order           |
|-----------------------------------------|--------|-----------------------|
| $ A(e_{12}) - A(e_{13})  /  A(e_{12}) $ | 0      | $8.36 \times 10^{-4}$ |

**Table 6.** The  $O_h$  violation of §7.5, measured on the  $T$ -matrix rather than on the coupling. Identical at both contrasts, because both carry the same  $\Delta\mu$  and the shear channels see only that.

This is the sharpest form of the result. It is not a norm over a matrix: it is two numbers that a symmetry argument says must agree, and do not.

**What the split does and does not license.** It is tempting to read this as the first-order  $T$  failing an external test while the ordinary route passes. That reading claims more than the check can deliver, and the two halves must be kept apart.

*What it gives.* An **arbiter-free lower bound on the error of the first-order  $T$ .** If  $A(e_{12})$  and  $A(e_{13})$  must be equal and differ by  $\delta$ , then whichever lies closer to the truth, the larger of the two errors is at least  $\delta/2 \approx 4.2 \times 10^{-4}$ . An absolute error estimate obtained from a symmetry alone, with nothing external to compare against, is worth having — and it is the absolute companion this project requires alongside any ratio test.

*What it does not give.* Any bound whatever on the ordinary route’s error. A symmetric approximation can be symmetrically wrong, missing the truth in *both* channels by more than  $8.4 \times 10^{-4}$  while satisfying the symmetry exactly. Symmetry compliance is not accuracy. Setting a bound on one formulation against no bound on the other is not a comparison, and no ordering of the two follows from this table.

*What is nonetheless attributable.* The breaking is not manufactured by the projection mismatch of §9:  $\mathcal{G}$  is itself  $O_h$ -symmetric, and §7.5 measured the coupling’s asymmetry symbolically with no propagator present at all. The asymmetry is a property of the coupling, and it survives into  $\mathbf{T}$ .

**The channel set has to include  $e_{13}$ .** The coupling leaves  $e_{12}$  untouched by construction, so a comparison spanning only  $\theta$ ,  $e_{\text{diag}}$  and  $e_{12}$  reports that the shear channels agree — true, and uninformative. A channel-resolved comparison is only as good as the channels chosen, and the  $O_h$ -equivalent pair is the one that carries the information.

## 9.8 What this does and does not settle

The Born limits agree: at a contrast of  $10^{-6}$  the two couplings differ by  $4.2 \times 10^{-8}$ , i.e. at  $O(\Delta c)$  as they must.

**The caveat that no number above reveals.** The propagator moment used is the *single* (collocation) one, while the  $\mathbf{J}_6$  pairing makes the scheme Galerkin. The consistent pairing needs the *double* moment. This is the same class of mismatch the package has already measured as a genuine error elsewhere, so the table above should be read as the first-order coupling evaluated through a collocation propagator, not as the fully consistent first-order  $T$ . No winner is declared, and none should be read in.

## 10 The propagator moment from the first-order system

§9 took  $\mathcal{G}$  from the ordinary Kelvin tensor, so the  $T$ -matrix assembled there is first-order in its *coupling* only. This section replaces that half. The plan is docs/plans/2026-09-18-first-

`order-propagator-moment.md`; what follows is written as each step lands, and says plainly where the front is.

The object is the sandwiched form

$$[\Delta \mathbf{C}_{\text{eff}} \mathcal{G} \Delta \mathbf{C}_{\text{eff}}]_{kl} = \langle \mathbf{J}_6 \psi_k \mid \Delta \mathcal{A} \Gamma \Delta \mathcal{A} \mid \psi_l \rangle. \quad (43)$$

Plain Galerkin on  $\mathbf{q}$  is not available: its mass matrix  $\langle \mathbf{J}_6 \psi_k, \psi_l \rangle$  is antisymmetric, and an antisymmetric matrix in *odd* dimension — here nine — is singular.

### 10.1 Step 1: $\Gamma$ , built from $\mathcal{A}$ alone

For  $\partial_3 \mathbf{q} = \mathbf{A} \mathbf{q} + \mathbf{d} \delta(z - z')$  the radiating Green's matrix is

$$\Gamma(z, z') = \begin{cases} +\mathbf{P}_d e^{\mathbf{A}(z-z')}, & z > z', \\ -\mathbf{P}_u e^{\mathbf{A}(z-z')}, & z < z', \end{cases} \quad (44)$$

with  $\mathbf{P}_d, \mathbf{P}_u$  the spectral projectors onto the down- and up-going eigenspaces, so that  $\Gamma$  decays away from the source in both directions and jumps by  $\mathbf{P}_d + \mathbf{P}_u = \mathbf{I}$  across it.

**Projectors, not eigenvectors.** The thesis writes the eigen-matrix column by column with normalisation factors  $\epsilon_p, \epsilon_s, \epsilon_H$  left symbolic. Those factors are fixed by a flux normalisation that degenerates at normal incidence, where the horizontal direction no longer distinguishes the *SV* and *SH* polarisations — exactly the  $k_{\parallel} \rightarrow 0$  difficulty. The spectral projector

$$\mathbf{P}_d = \mathbf{V}_d (\mathbf{W}_d^t \mathbf{V}_d)^{-1} \mathbf{W}_d^t \quad (45)$$

is invariant under any rescaling of the columns of  $\mathbf{V}_d$ , so building  $\Gamma$  from projectors removes the ambiguity instead of carrying it. That invariance is verified directly, by rescaling every eigenvector by a random complex factor and confirming  $\mathbf{P}_d$  is unchanged.

**The matrix exponential must never be formed.** Equation (44) is a statement, not an algorithm. Evaluating it as written — projecting *after* exponentiating — builds the growing branch and then cancels it. At a separation of 640 km the growing entries reach  $e^{+28} \approx 10^{12}$ , so the cancellation leaves  $10^{-4}$  of absolute noise on a true answer of  $3 \times 10^{-8}$ . This is the Thomson–Haskell instability, the reason invariant imbedding exists, reappearing inside our own construction; it announces itself as a decay rate that turns *negative* at large separation. Applying the exponential to the projected eigenvalues never forms the growing branch at all,

$$\Gamma(z - z') = \pm \sum_{i \in \text{down/up}} e^{s_i(z-z')} \mathbf{v}_i \mathbf{w}_i^t, \quad \mathbf{w}_i \cdot \mathbf{v}_j = \delta_{ij}, \quad (46)$$

and is what is used. The moment integral needs only  $|z - z'| \leq 2a$ , where the instability cannot bite; a layered background with thick layers is where it would.

$\omega$  is complex from the outset. The vertical wavenumbers  $k_{z,c} = \sqrt{k_c^2 - k^2}$  vanish on the circles  $|k| = k_p$  and  $|k| = k_s$ , which the lateral integral must cross, and  $k_z$  sits in a denominator. At real  $\omega$  a quadrature straddles an integrable  $1/\sqrt{\cdot}$  singularity and does not converge; the directional-sweep work measured  $1.9 \times 10^{-1} \rightarrow 1.1 \times 10^{-1}$  over an eightfold refinement, non-monotone. With damping 0.02 the smallest  $|s|$  on the  $P$  circle is 0.240 rather than zero.

**Checks.** Thirteen, all passing: the spectrum is  $\{\pm ik_{z,p}\}$  and  $\{\pm ik_{z,s}\}$  twice;  $\mathbf{P}_d + \mathbf{P}_u = \mathbf{I}$ , both idempotent, both commuting with  $\mathbf{A}$ ;  $\mathbf{P}_d$  invariant under rescaling;  $\partial_z \Gamma = \mathbf{A} \Gamma$  on both sides; the jump exactly  $\mathbf{I}$ ; and the decay rate converging to the slowest eigenvalue,  $0.034990 \rightarrow 0.033901 \rightarrow 0.032857 \rightarrow 0.032436$  against 0.032393 — a rate, not a bar, because at any one finite separation the faster mode still contributes.

## 10.2 Step 2: the vertical integral, in closed form

The vertical integral is exact, but it is *not* an integral over the square. The two halves of  $\Gamma$  carry different matrices, so the double integral splits over the two triangles and they cannot be recombined: treating the kernel as a function of  $|z - z'|$  alone would symmetrise the propagator and destroy the up/down distinction the first-order system exists to carry.

With affine trial fields the weights are 1 and  $z$ , so four triangular integrals are needed,

$$T_{mn}(q) = \int_{-a}^a dz \int_{-a}^z dz' z^m z'^n e^{-q(z-z')}, \quad (47)$$

the lower triangle being  $T_{nm}$  by relabelling. All four are elementary:

$$T_{00} = \frac{2aq - 1 + e^{-2aq}}{q^2}, \quad (48)$$

$$T_{10} = -T_{01} = \frac{aq - 1 + (1 + aq)e^{-2aq}}{q^3}, \quad (49)$$

$$T_{11} = \frac{3 - 3(1 + aq)^2 e^{-2aq} + a^2 q^2 (2aq - 3)}{3q^4}. \quad (50)$$

Summing the two triangles with a common matrix recovers the familiar square integral  $2(2aq - 1 + e^{-2aq})/q^2$ , and the odd weight cancels between them by parity — both checked, as is the large- $q$  behaviour  $T_{00} \sim 2a/q$ , each triangle carrying half of the square's  $4a/q$ .

The decisive check of this step is not the closed forms themselves but the *assembly*: the vertical integration of  $\Gamma$  performed through (48)–(50), matrices and all, reproduces a brute-force double quadrature of  $\Gamma$  to  $2.5 \times 10^{-9}$  relative.

### 10.2.1 The same integrals as a series, and why they are needed

(48)–(50) are exact and, evaluated as written, lose most of their digits as  $q \rightarrow 0$ . Expanding (50) gives

$$3 - 3(1 + x)^2 e^{-2x} + x^2(2x - 3) = \frac{4}{5}x^5 - \frac{2}{3}x^6 + O(x^7), \quad x = aq, \quad (51)$$

so a numerator of order  $x^5$  is assembled from terms of order unity: at  $x = 0.025$  about eight digits are gone. That is not a corner of the parameter space — it is the whole propagating window  $k < \omega/\beta$ , which is precisely where the radiation lives.

The remedy is the series. Substituting  $u = z - z'$  in (47) and expanding the exponential,

$$T_{mn}(q) = \sum_{j \geq 0} \frac{(-q)^j}{j!} \int_0^{2a} u^j J_{mn}(u) du, \quad J_{mn}(u) = \int_{-a}^{a-u} (z' + u)^m z'^n dz', \quad (52)$$

and every  $u$ -integral is elementary. With  $b = 2a$ ,

$$J_{00} = b - u, \quad J_{10} = -J_{01} = \frac{b}{2}u - \frac{u^2}{2}, \quad J_{11} = \frac{b^3}{12} - \frac{b^2}{4}u + \frac{u^3}{6}, \quad (53)$$

giving coefficients of  $(-qb)^j/j!$  equal to

$$\frac{b^2}{(j+1)(j+2)}, \quad \frac{b^3}{2(j+2)(j+3)}, \quad b^4 \left[ \frac{1}{12(j+1)} - \frac{1}{4(j+2)} + \frac{1}{6(j+4)} \right] \quad (54)$$

for  $(0,0)$ ,  $(1,0)$  and  $(1,1)$  respectively. Two structural checks: at  $q = 0$  the first gives  $2a^2$  and the second  $2a^3/3$ , which are  $\iint_{z>z'} 1$  and  $\iint_{z>z'} z$ ; the third gives *zero*, and indeed  $\iint_{z>z'} zz' = 0$  by parity. Its  $j = 1$  term is  $\frac{4}{15}a^5q$ , which is exactly what (50) tends to,  $\frac{4}{5}a^5q^5/3q^4$ . The series is entire, so it is valid for the imaginary  $q$  that propagating modes carry; it is used below  $|2aq| = 3$ , where it is also the accurate branch.

### 10.3 Step 3: the lateral integral

Count powers as  $|k| \rightarrow \infty$  for the *untransferred* form, in which  $\partial_\alpha \rightarrow ik_\alpha$ : the box form factor gives  $k^{-2}$ , (48) gives  $q^{-1} \sim |k|^{-1}$ , the two contrast operators give  $k^2$ , and the measure gives  $k$ . The strain channel does not converge. (The transferred form of §10.3.6 does; that is one of the things the augmentation buys, and it is taken up there.)

The divergence is not an artefact to be trimmed. A delta function in real space is a constant in  $k$ -space, so the non-decaying tail of the integrand *is* the Eshelby term of §9.2. A bare cutoff discards precisely the piece whose omission reverses the sign of the depolarisation. Hence split-and-subtract,

$$I(\mathbf{k}) = [I(\mathbf{k}) - I_\infty(\mathbf{k})] + I_\infty(\mathbf{k}), \quad (55)$$

with the bracket integrated numerically — convergent, and its convergence measured as a rate — and  $\int I_\infty d^2k$  evaluated in closed form over the whole plane, where the delta content resides. The asymptote is taken of the  $z$ -integrated kernel (48)–(50), with the box form factors kept exact: they are what makes the remainder converge, so expanding them would defeat the subtraction.

#### 10.3.1 The scaling that makes the expansion uniform

$\mathbf{A}$  is not uniformly  $O(k)$ . Rows 1–3 carry  $k$ , but rows 5–6 carry  $k^2$ , because the basis mixes displacement with traction and the two have different dimensions; expanding  $\mathbf{A}$  in  $1/k$  as it stands mixes orders between blocks. The traction of an evanescent field scales as  $\mu k u$ , so the natural

balance is  $\mathbf{q} \rightarrow (\mathbf{u}, \boldsymbol{\tau}_3/k)$ . Under  $\mathbf{S} = \text{diag}(1, 1, 1, k, k, k)$ ,

$$\tilde{\mathbf{A}} = \mathbf{S}^{-1} \mathbf{A} \mathbf{S} \quad (56)$$

is uniformly  $O(k)$ : the upper-right block gains a factor  $k$ , the lower-left loses one, and the diagonal blocks are untouched. Measured,  $\|\mathbf{A}\|/k$  grows as  $5.3 \times 10^3 \rightarrow 5.3 \times 10^5$  over two decades in  $k$  while  $\|\tilde{\mathbf{A}}\|/k$  settles at  $53.0100 \rightarrow 53.0101$ . Being a similarity transform it changes no eigenvalue, and  $|s|/k \rightarrow 1$  for all six.

The scaling is also what makes the decomposition numerically usable. On the unscaled matrix the down- and up-going eigenvectors become nearly parallel as  $k$  grows — the evanescent limit degenerates — the eigen-decomposition is reported badly conditioned, and the kernel entries collapse into rounding error, from which no asymptotic power can be read. The decomposition is therefore done on  $\tilde{\mathbf{A}}$  and mapped back,  $\mathbf{K} = \mathbf{S} \tilde{\mathbf{K}} \mathbf{S}^{-1}$ , checked against the verified kernel of step 2 at moderate  $k$ .

### 10.3.2 Which entries fail to decay

The order in  $1/k$  that each entry of the  $z$ -integrated kernel carries is read off numerically, over decades in  $k$ , rather than derived: a measured power is harder to fool oneself with than one obtained from a  $6 \times 6$  eigen-perturbation. A dot marks an entry that vanishes, so no power can be read.

|             | $u_z$    | $u_x$    | $u_y$    | $\tau_{zz}$ | $\tau_{xz}$ | $\tau_{yz}$ |
|-------------|----------|----------|----------|-------------|-------------|-------------|
| $u_z$       | .        | -1       | -1       | -2          | .           | .           |
| $u_x$       | -1       | .        | .        | .           | -2          | -2          |
| $u_y$       | -1       | .        | .        | .           | -2          | -2          |
| $\tau_{zz}$ | <b>0</b> | .        | .        | .           | -1          | -1          |
| $\tau_{xz}$ | .        | <b>0</b> | <b>0</b> | -1          | .           | .           |
| $\tau_{yz}$ | .        | <b>0</b> | <b>0</b> | -1          | .           | .           |

**Table 7.** Power of  $k$  carried by each entry of the  $z$ -integrated kernel as  $|k| \rightarrow \infty$ . Exactly one block fails to decay.

**The entire ultraviolet problem sits in one block:** the traction rows against the displacement columns. Everything else decays at least as  $1/k$  and needs no subtraction.

That block is also the one the physics predicts. In the evanescent limit a short-wavelength displacement produces a traction  $\tau \sim \mu k u$ , while the vertical integral (48) supplies  $1/k$ ; the ratio is  $k$ -independent. A traction response to a displacement that does not depend on  $k$  is a local stiffness — which is the Eshelby delta of §9.2, seen from the first-order side. The two representations identify the same object, and  $I_\infty$  has to reproduce exactly this block.

### 10.3.3 The asymptote in closed form

The limit of the non-decaying block is not merely numerical. Writing  $\hat{\mathbf{k}}$  for the lateral direction,

$$C[\tau_{zz} \leftarrow u_z] = -\nu_1 = -\frac{4\mu(\lambda + \mu)}{\lambda + 2\mu}, \quad (57)$$

$$C[\tau_{\alpha z} \leftarrow u_\beta] = \mu(\hat{k}_\alpha \hat{k}_\beta - \delta_{\alpha\beta}). \quad (58)$$

**$\nu_1$  is not a new constant.** It is the same  $\nu_1$  that appears in  $\mathcal{A}_{12}$  — born in the elimination (7), and shown in §3.2 to be the thesis's  $\zeta$ . The subtraction is expressible entirely in the operator's own constants. A wrongly identified asymptote would have no reason to land on  $\nu_1$  and  $\mu$ .

Both are established as *rates*. The relative residuals fall  $1.50 \times 10^{-3} \rightarrow 1.50 \times 10^{-4} \rightarrow 1.50 \times 10^{-5}$  across decades in  $k$ , ratios 10.0006 and 10.011 — clean  $O(1/k)$ , which is the  $O(k_c/k)$  correction to the evanescent limit. Agreement at a single  $k$  would not distinguish a correct closed form from a coincidence.

The in-plane form is the sharper test. Two constants are fitted on *two* angles and then predict *six*:  $P = -22.4977$  and  $Q = +22.4948$ , both equal to  $\mu = 22.5$ , with predictions tracking measurements to five digits across  $\theta \in [0, 2.4]$ . The isotropic entry (57) is constant to seven digits over the same range.

**What this settles about the integration.** The asymptote *is* direction-dependent — the relative variation over  $\theta$  is 0.39 — so  $I_\infty$  is not a pure delta and the regulator cannot be avoided. But the structure is the most favourable available:  $\delta_{\alpha\beta}$  and  $\hat{k}_\alpha \hat{k}_\beta$  are the only two rank-two invariants a direction can build, and they transform to a local delta and a principal-value dipole kernel respectively, both standard. The  $\tau_{zz} \leftarrow u_z$  entry is isotropic and so is a pure delta of weight  $-\nu_1$ .

### 10.3.4 The lateral integral, and where it must be evaluated

The box form factor is separable — the lateral transform of the cube's indicator is  $(2a)^2 \text{sinc}(k_1 a) \text{sinc}(k_2 a)$  — so with (57)–(58) and one lateral derivative from each  $\Delta\mathcal{A}$ , the integrand carries  $k_c k_d$  against  $C$ , giving two families:  $\int d^2 k W k_c k_d$ , which is separable, and  $\int d^2 k W k_a k_b k_c k_d / k^2$ , which is not. The elementary one-dimensional integrals are

$$\int_{-\infty}^{\infty} \text{sinc}^2(ka) dk = \frac{\pi}{a}, \quad \int_{-\infty}^{\infty} k^2 \text{sinc}^2(ka) e^{-\epsilon|k|} dk = \frac{4}{4a^2\epsilon + \epsilon^3} = \frac{1}{a^2\epsilon} - \frac{\epsilon}{4a^4}. \quad (59)$$

**The second is a pure pole with no finite part.** So  $\int I_\infty d^2 k$  is not a number, and the split-and-subtract cannot be closed in  $k$ -space alone: adding back a pure pole restores exactly the divergence it was introduced to remove.

The method is not wrong; the tail is being evaluated in the wrong representation. Forming  $|\hat{\mathbf{S}}|^2$  first and then integrating term by term destroys the structure that makes the constant part

finite. By Parseval,

$$\iint f(\mathbf{x}) K(\mathbf{x} - \mathbf{x}') g(\mathbf{x}') = \int \frac{d^2 k}{(2\pi)^2} \hat{f}(-\mathbf{k}) \hat{K}(\mathbf{k}) \hat{g}(\mathbf{k}), \quad (60)$$

so a  $\hat{K}$  tending to a constant  $C$  contributes  $C \int_V f g d\mathbf{x}$  — a *local* term, finite, needing no regulator at all. This is verified directly in one dimension, both sides returning  $1/2a$ .

The decomposition therefore splits by representation rather than by magnitude:

| piece                                                                      | real-space object          | evaluation               |
|----------------------------------------------------------------------------|----------------------------|--------------------------|
| isotropic, $-\nu_1$ and $P\delta_{\alpha\beta}$                            | delta                      | local, elementary        |
| traceless $\hat{k}_\alpha \hat{k}_\beta - \frac{1}{2}\delta_{\alpha\beta}$ | 2-D principal-value dipole | conditionally convergent |
| remainder $I - I_\infty$                                                   | —                          | numerically, in $k$      |

**Table 8.** Split-and-subtract survives, but  $\int I_\infty$  is evaluated in real space. Only the remainder stays in  $k$ .

The traceless part is of degree  $-2$  with vanishing angular average — exactly the conditional convergence the real-space engine of §9.2 handles by peeling derivatives onto the cube faces. The two representations meet again here: the same distributional structure, in two dimensions rather than three.

### 10.3.5 The principal-value piece vanishes

In two dimensions  $\nabla^2 \log \rho = 2\pi\delta^2$ , and  $\partial_\alpha \partial_\beta \log \rho = (\delta_{\alpha\beta} - 2\hat{n}_\alpha \hat{n}_\beta)/\rho^2$  away from the origin; taking the trace fixes the delta weight, so that

$$\partial_\alpha \partial_\beta \log \rho = \text{PV} \left[ \frac{\delta_{\alpha\beta} - 2\hat{n}_\alpha \hat{n}_\beta}{\rho^2} \right] + \pi \delta_{\alpha\beta} \delta^2(\mathbf{x}). \quad (61)$$

By translation invariance the double integral over the square collapses onto the autocorrelation of its indicator,  $w(\mathbf{u}) = (2a - |u_1|)(2a - |u_2|)$ , exactly — removing two of the four dimensions — and both derivatives move onto  $w$ , which is legitimate because  $w$  has compact support.

$w$  is piecewise bilinear, so  $\partial_1 \partial_1 w$  is a sum of line deltas. It has three terms, not one:  $w$  vanishes at  $|u_1| = 2a$  but  $\partial_1 w$  does *not* — it jumps from  $\mp(2a - |u_2|)$  to zero — so there are edge deltas at  $u_1 = \pm 2a$  alongside the interior one at  $u_1 = 0$ . With them,

$$D_{11} = \underbrace{3}_{\text{interior}} + \underbrace{(\pi - 3)}_{\text{edges}} = \pi = \pi(2a)^2 \Big|_{a=1/2}, \quad D_{12} = 0, \quad (62)$$

and the trace identity  $\text{tr } D = \langle \nabla^2 \log \rho, w \rangle = 2\pi w(0) = 2\pi(2a)^2$  holds exactly. Dropping the edge deltas gives 6 in place of  $2\pi$ , so that identity is what detects them.

**Hence  $\langle \text{PV}_{\alpha\beta}, w \rangle = D_{\alpha\beta} - \pi \delta_{\alpha\beta} w(0) = 0$ : the traceless dipole contributes nothing.**

It could not have contributed.  $\langle \text{PV}_{\alpha\beta}, w \rangle$  is a symmetric two-tensor invariant under the square's  $D_4$  symmetry, and the only such tensor is a multiple of  $\delta_{\alpha\beta}$  — the traceless part has nowhere to live. The conditional convergence is real, but for this cross-section it integrates away.

**The whole of  $\int I_\infty$  is therefore the local delta term.** The piece that looked hardest — the conditionally convergent, shape-dependent dipole — is identically zero here, and what remains is elementary.

### 10.3.6 Why the transferred form converges

In the untransferred form  $\partial_\alpha \rightarrow ik_\alpha$ , each contrast operator contributes a factor of  $k$ , and the pole of (59) follows. After the by-parts transfer of (17) the derivatives act on the *polynomials, in real space*:  $\partial_\alpha$  of an affine trial function is a constant, whose lateral transform is again the box form factor. No factor of  $k$  survives, and the lateral integral converges.

**The augmentation therefore does two things at once:** it keeps the indicator from being differentiated, and it makes the lateral integral convergent. Both are the same integration by parts.

Convergence is nevertheless slow. With  $\hat{K} \rightarrow C$  the truncation error falls only as  $1/\Lambda$ , which is what the subtraction is for.

### 10.3.7 The tail, in closed form

The tail is elementary. For the in-plane entry  $C = \mu(\hat{k}_1^2 - 1) = -\mu k_2^2/k^2$ , so

$$\int d^2k W C = -\mu \int d^2k \operatorname{sinc}^2(k_1 a) \operatorname{sinc}^2(k_2 a) \frac{k_2^2}{k^2} = -\frac{\mu}{2} \left(\frac{\pi}{a}\right)^2, \quad (63)$$

because  $k_1^2/k^2 + k_2^2/k^2 = 1$  and the weight is symmetric under  $k_1 \leftrightarrow k_2$ , so each piece is exactly half of  $\int d^2k W = (\pi/a)^2$ . No quadrature is required; direct integration returns 19.7376 against  $(\pi/a)^2/2 = 19.7392$ , the difference being the quadrature's own error on a slowly converging integrand.

### 10.3.8 The remainder, and the test that closes the construction

Truncating the lateral integral at  $|k| < \Lambda$  for the  $\tau_{xz} \leftarrow u_x$  entry, with and without the subtraction:

The split is only worth making if both halves give the same answer. Truncating at  $|k| < \Lambda$  for the  $\tau_{xz} \leftarrow u_x$  entry:

| $\Lambda$ | unsubtracted | remainder | remainder + tail |
|-----------|--------------|-----------|------------------|
| 80        | −318.2214    | 118.83504 | −325.29716       |
| 160       | −321.6764    | 118.82727 | −325.30493       |
| 320       | −323.4539    | 118.82513 | −325.30707       |
| 640       | −324.3400    | 118.82466 | −325.30754       |

**Table 9.** The subtracted route is converged to six figures by  $\Lambda = 80$ ; the unsubtracted one is still 0.97 short of the limit at  $\Lambda = 640$ .

The remainder's successive differences fall  $7.8 \times 10^{-3} \rightarrow 2.1 \times 10^{-3} \rightarrow 4.7 \times 10^{-4}$ , a ratio near 0.27, against the unsubtracted 0.50 — the latter being the  $1/\Lambda$  decay the constant part imposes. Extrapolating both:

$$\underbrace{-325.226}_{\text{unsubtracted limit}} \quad \text{versus} \quad \underbrace{-325.308}_{\text{remainder+closed-form tail}}, \quad \text{relative difference } 2.5 \times 10^{-4}, \quad (64)$$

and that residue is dominated by the single-term extrapolation of the slowly converging unsubtracted series, not by the construction. At  $\Lambda = 640$  the unsubtracted error is 0.97 and the subtracted one  $5 \times 10^{-4}$ : a factor of about two thousand.

**Split-and-subtract is verified.** The two routes to the same number agree, the tail is closed in form by (63), and the remainder converges three orders of magnitude faster than the direct integral.

*What steps 1–3 establish.*  $\Gamma$  from spectral projectors; the triangular  $z$  integrals and their assembly; the scaling, the power table, and the asymptote  $-\nu_1$ ,  $\mu(\hat{k}\hat{k} - \delta)$  verified as an  $O(1/k)$  rate;  $\int I_\infty$  decomposed by representation, its principal-value part vanishing by  $D_4$  symmetry and its remaining part closed by (63); and the closure test (64). All of that is machinery for the lateral integral. It does not yet say what the integrand is, and the next subsection is about a question that has to be answered before it can.

#### 10.4 Step 4: how the augmentation composes with the Schwinger form

The augmentation of §6 turns the differential  $\Delta A$  into a multiplicative  $\Delta C_{\text{eff}}$  by moving each outer lateral derivative onto the test function. That is what keeps  $\partial_\alpha \mathbf{1}_V$  — the surface layer on the faces of the cube — from ever being formed. The Schwinger object

$$[\Delta C \mathcal{G} \Delta C]_{kl} = \langle J_6 \psi_k \mid \Delta A \Gamma \Delta A \mid \psi_l \rangle \quad (65)$$

carries *two* contrast operators, and the right-hand one has no test function next to it: what  $\Gamma$  needs is  $\Delta A \psi_l$  as an actual source, which is  $\partial_\alpha (\Delta c \mathbf{1}_V \partial_\beta \psi_l)$  — precisely the object the transfer exists to avoid. Read that way the two constructions look incompatible.

They are not, and the reason is that the Born term and the Schwinger term differ in a way that matters. In the Born term the indicator of the contrast and the test function sit at the *same* point, so  $\partial_\alpha (\mathbf{1}_V c)$  would have to be multiplied by another factor supported on  $\bar{V}$ ; a surface layer times an indicator is not defined, and the transfer is therefore not a convenience but the only well-posed form. In the Schwinger term the two indicators are separated by  $\Gamma$ . The source at  $\mathbf{x}'$  is paired, through a kernel, with a test object at  $\mathbf{x}$ ; nothing coincides, the surface layer is harmless — and in the lateral Fourier domain it need not be formed at all, because

$$\mathcal{F}[\partial_\alpha (\mathbf{1}_V c)](\mathbf{k}) = ik_\alpha \mathcal{F}[\mathbf{1}_V c](\mathbf{k}) \quad (66)$$

is exact. The derivative is carried by the *transform*; the indicator stays inside the form factor, undifferentiated.

The bookkeeping that follows is uniform. Every derivative in  $A$  is lateral — the depth derivative is the left-hand side of  $\partial_3 q = Aq + d$  — so in the lateral Fourier domain each is a factor  $ik_\alpha$  at the

single wavenumber the integral runs over. One entry of (65) has four derivative slots, and they are placed as

| leg   | slot  | where it goes                                               |
|-------|-------|-------------------------------------------------------------|
| left  | outer | onto the test polynomial (the augmentation)                 |
| left  | inner | onto the propagated field $\rightarrow ik_\beta$            |
| right | outer | onto the source transform $\rightarrow ik_\alpha$ , by (66) |
| right | inner | onto the trial polynomial (the augmentation)                |

The two outer slots are the ones in question. The left one goes exactly where  $\Delta C_{\text{eff}}$  already puts it; the right one goes onto the kernel. No surface term is evaluated on a face, and the augmentation is unchanged.

(66) is checked numerically rather than asserted. Paired with a smooth bump  $f$  it states  $\int f \partial_1 \mathbf{1}_S = - \int_S \partial_1 f$ , whose right side is an ordinary integral over the square; evaluating the left side in the Fourier domain as  $(2\pi)^{-2} \int \hat{f}(-\mathbf{k}) ik_1 F(\mathbf{k})$  reproduces it to  $1.0 \times 10^{-13}$ , with a zero imaginary residue and the opposite sign rejected by a factor of two.

### The propagator, checked against the Kelvin tensor

Before any of this can be assembled, one normalisation has to be settled that the bridge identity of §9 does not reach: the bridge compares  $\Gamma$  at two depths and so tests propagation, not the strength of the source jump. Inverting  $\Gamma$  laterally at a fixed vertical separation and reading the velocity–force block,

$$\frac{1}{-i\omega} (2\pi)^{-2} \int d^2k \Gamma_{3+i,j}(\mathbf{k}; \Delta z) \rightarrow G_{ij}(\Delta z), \quad (67)$$

the analytic Kelvin tensor comes back to  $7.3 \times 10^{-4}$  and still converging, with off-diagonal leakage  $2.9 \times 10^{-4}$ ; the residue is consistent with the  $O((k_s r)^2)$  the static comparison drops. This fixes both conventions at once — that a force  $f_i$  enters row  $i$  of  $\partial_3 q = Aq + s$ , and that  $q_{3+i} = v_i = -i\omega u_i$  comes out. The quadrature must be polar: the kernel depends on  $|\mathbf{k}|$  and on direction separately and so is not smooth at the origin, where a Cartesian tensor rule converges only as  $1/n$  (measured:  $4.0 \times 10^{-2} \rightarrow 1.2 \times 10^{-2}$  over a ninefold refinement, against  $5.4 \times 10^{-3} \rightarrow 7.3 \times 10^{-4}$  in polar). This is the opposite of the lateral moment itself, whose form factors are separable sinc products that polar cannot resolve.

### The density channel, assembled

With  $\Delta\lambda = \Delta\mu = 0$  only the multiplicative  $i\omega\Delta\rho$  term of  $\Delta A$  survives, both legs reduce to constants, and (65) can be predicted in closed form:

$$\langle J_6 \psi_0 \mid \Delta A \Gamma \Delta A \mid \psi_0 \rangle = i\omega^5 \Delta\rho^2 M, \quad M = \int_V \int_V G_{11} = \left(a_0 + \frac{b_0}{3}\right) D_0, \quad (68)$$

with  $a_0, b_0$  the Kelvin coefficients of §9.2 and  $D_0$  the purely geometric double-volume moment  $\int_V \int_V 1/r$ .  $D_0$  is evaluated by writing the double integral as a single one against the autocorrelation weight  $\prod_i (2a - |u_i|)$  — exact, and three dimensions cheaper — and removing the  $1/r$  outright with

a Duffy substitution  $\mathbf{u} = t(\xi, \eta, 1)$ , whose Jacobian  $t^2$  cancels it:

$$D_0 = 1.8823126444, \quad \int_V 1/r = 2.3800773640. \quad (69)$$

The lateral assembly with *both* contrast operators in place reproduces (68) to  $6.9 \times 10^{-5}$ , the two routes sharing nothing but the medium: one is a spectral construction on  $A$  integrated over lateral wavenumber, the other a real-space moment of the Kelvin tensor.

### What the nine-parameter $\mathcal{G}$ is, and is not

The same comparison locates an object that had been taken on trust. The  $\mathcal{G}$  used in §9 was *calibrated* against this package's  $T_9$  operator rather than derived, and it is not the moment that (65) calls for. The gap factorises exactly:

$$\frac{[\Delta C \mathcal{G} \Delta C]_{00}}{i\omega^5 \Delta \rho^2 M} = i\omega \frac{\int_V 1/r + i(\text{radiation})}{D_0}, \quad \text{residual } 3.6 \times 10^{-5}, \quad (70)$$

where the radiation term is the closed form  $V\omega(\alpha^{-3} + 2\beta^{-3})/12\pi\rho$  that the moment export is itself gated on. Two statements, both of which matter:

- The geometric factor  $\int_V 1/r/D_0 = 1.2644$  is the ratio of a *point*-volume moment to a *volume*-volume moment. The calibrated  $\mathcal{G}$  is a collocation object; (65) is a Galerkin one. That is consistent with this project's own finding that the single site is best posed by collocation, but the two cannot be mixed without the factor.
- The remaining  $i\omega$  is the normalisation of the  $q$ -field trial space:  $\Delta C \mathcal{G} \Delta C$  carries two  $J_6$  test factors while (65) carries one, so a  $\mathcal{G}$  that composes correctly must absorb the trial-function normalisation. How it should do so is not settled here; see §15.

### The derivative channels, and three defects the test exposed

The density channel settles the basis, the kernel, the  $z$  integrals and the quadrature, but its contrast operator is *multiplicative*: it does not exercise (66) at all. The channels that do are those carrying  $\Delta\lambda$  and  $\Delta\mu$ , and for them a sharp statement is available with no new arbiter.  $\mathbf{J}_6 \Delta \mathcal{A}$  is symmetric — that is what derived the test space in §6 — and  $\mathbf{\Gamma}$  is reciprocal,

$$\mathbf{\Gamma}(-\mathbf{k}; z', z)^\top = \mathbf{J}_6 \mathbf{\Gamma}(\mathbf{k}; z, z') \mathbf{J}_6, \quad (71)$$

which follows from  $\mathcal{A}(-\mathbf{k})^\top \mathbf{J}_6 = -\mathbf{J}_6 \mathcal{A}(\mathbf{k})$  and  $\mathbf{J}_6^{-1} = -\mathbf{J}_6$ . Hence the assembled matrix must be symmetric. It is not symmetric by construction: the assembly builds its two legs by opposite rules, transferring the left outer derivative onto the test polynomial and carrying the right one on the kernel as  $ik$ . Symmetry is therefore a test of exactly the bookkeeping in question, and on a quadrature grid symmetric under  $\mathbf{k} \rightarrow -\mathbf{k}$  it is an *algebraic* property, holding to round-off even where the grid is far too coarse to give the value accurately.

Applied, it found three defects, each in code that had already passed its own checks:

1. **The up/down split was decided by round-off inside the propagating window.** Classifying by  $\Re(\text{ev}) < 0$  is right for evanescent modes and meaningless for propagating ones: below  $k = \omega/\beta$  the vertical wavenumber is real,  $\Re(\text{ev})$  vanishes to arithmetic noise, and its sign was being decided independently at  $+\mathbf{k}$  and  $-\mathbf{k}$ . The reciprocity residual (71) at  $|k| = 1.1 \times 10^{-3}$  was 1.245 — a complete failure, in the region that carries all of the radiation. The limit  $\omega \rightarrow \omega(1 + i\epsilon)$  settles it without needing  $\epsilon$ : a propagating mode has  $\text{ev} = ik_z$  and the perturbation gives  $k_z$  a positive imaginary part, so  $\Re(\text{ev}) \rightarrow 0^-$  exactly when  $\Im(\text{ev}) > 0$ . Residual after:  $3.7 \times 10^{-13}$ .
2. **The triangular  $z$  integral cancels catastrophically at small  $aq$ .** The numerator of  $T^{11}$  is  $O((aq)^5)$  while its individual terms are  $O(1)$ , so by  $aq \approx 0.025$  it has lost about eight digits. That is not a corner case — it is again the propagating window. Expanding  $T^{mn}$  as a power series in  $q$  below  $|2aq| = 3$ , with the coefficients of §10.2.1, took that node from  $1.3 \times 10^{-6}$  to  $2.1 \times 10^{-13}$ .
3. **The scaling used to condition the eigenproblem is in the wrong direction.**  $\text{diag}(1, 1, 1, k, k, k)$  was introduced because the up- and down-going eigenvectors become parallel as  $k$  grows. Measured against *no scaling at all* it makes the eigenvector condition number *worse* at every  $k$  above the propagating window — by two orders at  $k = 60$ . Parlett–Reinsch balancing beats both by about twelve orders there ( $8.7 \times 10^7$  against  $1.6 \times 10^{20}$ ), because it is free to treat the third component differently from the first two, which no scalar rule can.

With all three repaired the assembled  $9 \times 9$  matrix is symmetric to  $4.7 \times 10^{-12}$ , while reversing the sign of the right leg's  $ik$  — the one step (66) licenses — breaks it at  $7.2 \times 10^{-2}$ . That is a separation of ten orders on precisely the disputed bookkeeping, and the derivative terms carry 80% of the matrix, so the test is not vacuous.

### The radiation reaction, from the first-order system

Repairing the branch selection changed the physics, not just the arithmetic. The previous classification suppressed the outgoing condition; with it restored, the propagator carries radiation damping, and both real-space arbiters above had to acquire the term they were missing. The sharpest form of the check is pointwise:

$$\text{IG}_{11}(\Delta z) \Big|_{\text{first-order}} = 5.2134 \times 10^{-14} \quad \text{against} \quad \frac{\omega}{12\pi\rho} \left( \frac{1}{\alpha^3} + \frac{2}{\beta^3} \right) = 5.2250 \times 10^{-14}, \quad (72)$$

agreeing to  $2.2 \times 10^{-3}$ . This quantity has *no static counterpart at all*: it is not a correction to the Kelvin tensor but a term the static theory does not contain, and it arrives here out of the spectral projectors of  $\mathcal{A}$  with nothing put in by hand.

Resolving it needs the quadrature to respect the two branch radii  $|k| = \omega/\alpha$  and  $\omega/\beta$ , where the vertical wavenumber vanishes. Those are *circles*: no Cartesian panel can follow one, and an unresolved square-root costs several per cent. The lateral integral is therefore split — Cartesian outside the propagating window, where the integrand is a separable product of sinc-type form factors that polar cannot resolve; polar inside it, where the branch radii are panel edges and

$ka < 0.02$  makes the form factors constant to  $10^{-4}$ . The central block of the Cartesian grid is exactly the inner square, so the two tile the box with no overlap and no gap. With that in place the density channel converges cleanly —  $9.858, 9.834, 9.828 \times 10^{-4}$  over a fourfold refinement — where before it drifted.

*Status of Task 1.* The design question is settled; the assembly computes a definite object; the density channel is verified against an independent real-space arbiter to  $9.8 \times 10^{-4}$ ; the derivative channels are verified by a symmetry with a ten-order separation from its control; and the radiation reaction comes out right. What remains is to read the shear moment  $B$  off the assembled matrix and compare it with  $\sqrt{3}(\lambda + \mu)/6\pi\mu(\lambda + 2\mu)$ , and to settle the  $i\omega$  of (70).

## 11 The thesis spectral representation, and what it removes

Three of the four numerical defects of §10.4 are not properties of the physics. They are consequences of asking a numerical eigensolver to do work that §2.1 of the thesis has already done in closed form, and they do not arise at all if  $\Gamma$  is built the way the thesis builds it.

### 11.1 What the thesis gives

Alongside (Akdef) the thesis carries the complete spectral representation

$$\mathbf{A}(k_x, k_y) = \mathbf{D}_z \mathbf{\Lambda} \mathbf{D}_z^{-1}, \quad \mathbf{\Lambda} = i \operatorname{diag}[k_{z,P}, k_{z,S}, k_{z,H}, -k_{z,P}, -k_{z,S}, -k_{z,H}], \quad (73)$$

with the columns of  $\mathbf{D}_z$  written out analytically —  $|\pm, P\rangle, |\pm, S\rangle, |\pm, H\rangle$ , the + columns downgoing — and three further things that matter more than they look:

1. **The branch is defined piecewise.**  $k_{z,c} = \sqrt{\hat{K}_c^2 - k_x^2}$  where  $|k_x| < |\hat{K}_c|$  and  $i\sqrt{k_x^2 - \hat{K}_c^2}$  otherwise: the positive root in both regimes, with the up/down assignment carried by the *column ordering* of (73) rather than by the sign of any computed quantity.
2. **The shear basis is chosen non-degenerate.** The thesis notes  $k_{z,S} = k_{z,H}$  and keeps the two labelled apart regardless, remarking that its quasi-SV and quasi-SH vectors — unlike the conventional seismological decomposition — remain independent at  $k_x = k_y = 0$ .
3. **The inverse requires no inversion.** From the quasi-Hamiltonian relation,

$$\mathbf{D}_z^{-1}(k_x, k_y) = -i \mathbf{J}_6 \mathbf{D}_z^t(-k_x, -k_y) \mathbf{J}_6. \quad (74)$$

All of it is verified here rather than taken on trust, and the checks are coupled: (73) cannot reproduce (Akdef) unless the eigenvectors, the ordering, the branch and (74) are all correct together. It does, to  $6.2 \times 10^{-14}$ . The normalisation factors  $\epsilon_c$ , which the thesis specifies only as “chosen so that  $\mathbf{D}_z^{-1}$  takes a simple form”, are *read off* rather than guessed: with  $\epsilon = 1$  the product (74)  $\cdot \mathbf{D}_z$  is diagonal — to  $1.1 \times 10^{-15}$ , itself a test of the eigenvectors — and  $\epsilon_c^2$  is the reciprocal of each diagonal entry.

## 11.2 What it removes

|                                                                                                                                                                 |                                                                                         |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------|
| found numerically, working from the operator form                                                                                                               | already settled in §2.1                                                                 |
| $\Re(\lambda)$ carries no information in the propagating window; its sign was decided independently at $\pm\mathbf{k}$ , giving a reciprocity residual of 1.245 | the branch of $k_{z,c}$ , defined piecewise, with the ordering of (73) carrying up/down |
| the doubly degenerate shear eigenspace returns an arbitrary <i>SV</i> / <i>SH</i> mixture; an intermittent factor of four in a mode-converted coefficient       | quasi- <i>SV</i> and quasi- <i>SH</i> chosen non-degenerate, and said to be so          |
| eigenvector conditioning of $1.6 \times 10^{20}$ ; a scalar scaling that proved to be in the wrong direction; then Parlett–Reinsch balancing                    | no eigenvectors are computed                                                            |
| $\mathbf{D}_z^{-1}$ by numerical inversion                                                                                                                      | (74)                                                                                    |

**Table 10.** The fourth defect, the cancellation in  $T^{11}$  of §10.2.1, is genuine: it concerns the  $z$  integral and not the spectral representation.

Rebuilt on (73),  $\Gamma$  is reciprocal to  $3.6 \times 10^{-12}$  with no classification performed anywhere, reproduces the Kelvin tensor to  $1.0 \times 10^{-3}$  and converging, and carries the radiation reaction to  $4.4 \times 10^{-3}$  — *with no  $\epsilon$  prescription*, since the outgoing condition is already in the branch of  $k_{z,c}$ . In the thesis basis the extraction is also plainer: a force  $f_j$  enters row  $3 + j$  with a minus sign and the displacement comes out in rows 0–2, so  $G_{ij} = -\Gamma_{i,3+j}$  with no velocity factor at all.

## 12 A layer in a whole space: the external arbiter

Every comparison so far in this note sets two formulations against each other. None of them says which is right. A single layer embedded in a whole space removes that limitation, because its reflection coefficient is known independently — from a validated reflectivity recursion, and from a closed form short enough to derive here — and because it exercises the operator in a way the cube cannot.

### 12.1 Why a layer is the right test

For a laterally invariant contrast the operator difference is simply

$$\Delta\mathcal{A}(z) = [\mathcal{A}_{\text{total}}(\mathbf{k}) - \mathcal{A}_{\text{background}}(\mathbf{k})] \mathbf{1}_{[0,h]}(z), \quad (75)$$

a piecewise-constant matrix. **There is no augmentation, no by-parts transfer, no surface term and no differentiated indicator**, because  $\mathcal{A}$  carries no  $\partial_3$ : a jump in the medium is simply a jump in the coefficients of an ODE, and is harmless. That is the structural advantage of the first-order form, and it is precisely the property the whole cube development of §6–§10.4 had

to work around. So this tests  $\mathcal{A}$ ,  $\Gamma$  and the Lippmann–Schwinger composition with none of the cube machinery in the way.

It also tests two things the cube checks could not. It is *nonlinear* in the contrast, so it reaches the resummation rather than the Born term alone; and  $\Delta\mathcal{A}$  carries  $\mathbf{k}$  explicitly through  $\mathcal{A}_{12} \sim v_1 k^2 / i\omega$ , so scanning slowness tests a  $k$  dependence that the density channel of §10.4, being multiplicative, left untouched.

## 12.2 The closed form, by continuity

At normal incidence the answer follows from matching. With  $e^{-i\omega t}$  and  $u$  the vertical displacement, each medium carries  $u = De^{ikz} + Ue^{-ikz}$  with  $k = \omega/c$ , and the traction is  $\tau = \rho c^2 \partial_z u = i\rho c \omega (De^{ikz} - Ue^{-ikz})$ . Continuity of  $u$  and  $\tau$  at  $z = 0$  and  $z = h$ , with no up-going wave below, gives

$$R = \frac{r_{12} + r_{23} e^{2i\omega h/\alpha_2}}{1 + r_{12} r_{23} e^{2i\omega h/\alpha_2}}, \quad r_{12} = \frac{Z_1 - Z_2}{Z_1 + Z_2} = -r_{23}, \quad Z = \rho c. \quad (76)$$

The point of deriving it rather than quoting it is that the *sign* is then fixed by the derivation. Throughout this section

$$R = \frac{u_3 \text{ of the up-going field at } z = 0^-}{u_3 \text{ of the incident field at } z = 0^-}, \quad (77)$$

on both sides of every comparison, which leaves no convention free.

## 12.3 The first-order route

The laterally invariant Lippmann–Schwinger equation

$$q(z) = q_{\text{inc}}(z) + \int_0^h \Gamma(\mathbf{k}; z, z') \Delta\mathcal{A} q(z') dz' \quad (78)$$

is solved by piecewise-constant collocation, the cell integrals taken in closed form from the spectral representation of  $\Gamma$  and *split at*  $z' = z$ : the jump across the diagonal is the source delta, and quadrature across it would smear exactly the object that defines the propagator.

One structural simplification is used, and it is exact. With  $\mathbf{k}$  along  $x_1$  the system is block diagonal:  $\{-\tau_{23}, v_2\}$  is *SH* and does not couple to  $\{-\tau_{13}, -\tau_{33}, v_1, v_3\}$ . Working in the  $4 \times 4$  *P-SV* block is therefore no approximation — and it removes the double degeneracy of  $q_S$ , which is a genuine hazard rather than a nicety: a degenerate eigenspace has no preferred basis, so `eig` returns an arbitrary mixture of *SV* and *SH*. Before the block restriction the *SS* coefficient was correct at three slownesses and wrong *by a factor of four* at the fourth, which reads as a critical angle and is linear algebra.

## 12.4 What it gives

| contrast | (76)                    | first order             | relative             |
|----------|-------------------------|-------------------------|----------------------|
| +2%      | $-0.019937 - 0.019792i$ | $-0.019937 - 0.019792i$ | $2.7 \times 10^{-5}$ |
| +10%     | $-0.127846 - 0.087946i$ | $-0.127849 - 0.087948i$ | $2.4 \times 10^{-5}$ |
| +40%     | $-0.578990 - 0.067730i$ | $-0.578999 - 0.067736i$ | $1.7 \times 10^{-5}$ |

**Table 11.** Normal incidence,  $h = 200$  m,  $\omega h/\alpha = 2.4$ . The +40% row is the one that matters: it is far from Born, so it tests the *resummation*. Refinement in the cell count gives ratios 4.00, 4.00, 4.00 — clean second order.

Against the reflectivity recursion the comparison runs to  $\sin i = 0.75$ :

$$\frac{R_{PP}^{\text{first order}}}{R_{PP}^{\text{Kennett}}} = -1.00004, -1.00004, -1.00003, -1.00002, \quad (79)$$

and the same for SS to  $1.0 \times 10^{-4}$ . Two remarks make this a statement rather than a fitted factor. First, the recursion’s  $R$  proved *independent of the thickness of the layer above* — identical for  $d = 50, 100, 200, 400$  m — so it is already referenced at the target layer, and the two-way phase correction that seemed obviously required in fact manufactures a spurious slowness-dependent offset. Second, what is checked is not the value  $-1$  but that the factor is *constant*: a convention cannot depend on angle, and a physics error would not be constant. The sign itself is pinned by (76), which owes nothing to either code.

The mode-converted coefficients are covered by an identity that needs no knowledge of the convention at all. If the two differ by a diagonal rescaling  $\text{diag}(d_p, d_s)$  — the displacement-versus-flux factor, invisible to any per-mode ratio test — then the off-diagonal ratios are  $d_p/d_s$  and  $d_s/d_p$ , whose product is unity whatever  $\mathbf{D}$  is. Measured: 1.00008, 1.00010, 1.00008.

Finally, a control isolates what this test uniquely reaches. Building  $\Delta\mathcal{A}$  at  $\mathbf{k} = 0$  while leaving  $\mathbf{\Gamma}$  at the true  $\mathbf{k}$  changes *nothing* at normal incidence (relative difference 0, exactly) and breaks the answer by  $6.8 \times 10^{-2}$  off it. The  $k$  dependence of the contrast operator is therefore load-bearing, and correct.

**What this establishes, and what it does not.** The system matrix  $\mathcal{A}$ , the propagator  $\mathbf{\Gamma}$ , their composition, and the resummation reproduce an independently validated reflectivity code across contrast and across angle. That is the first external validation in this note. It does *not* validate the cube: the augmentation, the form factors and the moments of §9.2 play no part in a laterally invariant problem. Those remain checked only against other formulations.

## 13 Which formulation is more accurate

§12 validates the operator and the propagator but cannot compare the two formulations, because a laterally invariant problem is one the cube moment method does not address at all. The geometry that does compare them is the one this package already uses to score itself: a periodic patch of space-filling cubes, which for identical cubes *is* a homogeneous layer, and whose reflection

coefficient the reflectivity recursion gives exactly. Both routes then build a single-site  $T$  for the same cube, both pass through the same lattice sum, the same contact correction and the same Foldy–Lax solve, and both are scored against the same reference. Only  $T_0$  differs.

### 13.1 Substituting at the one place both routes speak the same language

$\mathbf{T} = \Delta\mathbf{C}(\mathbf{I} - \mathcal{G}\Delta\mathbf{C})^{-1}$  is not in the package’s  $9 \times 9$  convention — the strain block differs in sign and by the engineering doubling, and  $\Delta\mathbf{C}_{\text{eff}}$  is in velocity units besides (§7.3). Mapping all of that would risk measuring a convention as though it were physics. The substitution is therefore made through the four *dimensionless* amplification factors  $A_u$ ,  $A_\theta$ ,  $A_{e,\text{diag}}$ ,  $A_{e,\text{off}}$ , which are ratios and carry no units. The package’s own mapping to renormalised contrasts and its own  $T_0$  construction are then used unchanged for both routes, so nothing distinguishes them but those four numbers. What licenses this is that the Navier route reproduces the package’s own factors to  $1.5 \times 10^{-6}$ , re-checked rather than assumed.

### 13.2 The measurement

| contrast                            | ordinary              | first order           | ratio |
|-------------------------------------|-----------------------|-----------------------|-------|
| weak, $\Delta\mu/\mu \approx 0.4\%$ | $2.08 \times 10^{-3}$ | $2.09 \times 10^{-3}$ | 1.007 |
| moderate, $\approx 4\%$             | $2.00 \times 10^{-2}$ | $1.94 \times 10^{-2}$ | 0.966 |
| strong, $\approx 13\%$              | $5.59 \times 10^{-2}$ | $5.02 \times 10^{-2}$ | 0.899 |

**Table 12.** Relative error in  $R_{pp}$  against the reflectivity reference,  $M = 8$ ,  $N_z = 16$ ,  $a = 0.5$  m,  $\omega = 60$ ,  $p = 0$ . At weak contrast the two agree, as they must — they agree at Born order — and the advantage appears where the  $O(\Delta c^2)$  difference lives.

### 13.3 Why this is not cancellation

At weak contrast the two  $T_0$  differ by  $1.8 \times 10^{-5}$  and yet both miss the reference by  $2.1 \times 10^{-3}$ : **most of the error is shared**, and belongs to the lattice rather than to  $T_0$ . A difference in a mostly-shared error is exactly the shape an accidental cancellation would take — the first-order  $O(\Delta c^2)$  term happening to offset the lattice’s own. Refining the lattice discriminates between the two readings:

| $M$ | shared error at weak contrast | ratio at strong contrast |
|-----|-------------------------------|--------------------------|
| 8   | $2.0995 \times 10^{-3}$       | 0.8987                   |
| 12  | $1.7157 \times 10^{-3}$       | 0.8766                   |
| 16  | $1.4690 \times 10^{-3}$       | 0.8570                   |

The shared error falls and the advantage *grows*. Cancellation would erode as the thing being cancelled against is removed; instead the advantage was being *masked* at coarse  $M$ . So 14.3% is a **lower bound** on the part attributable to  $T_0$ , not an estimate of it — it is still improving at the finest lattice tested.

A remark on how that check was posed, since the first version of it was wrong: it demanded that the ratio be *stable* under refinement, and it failed. Stability is not what should be expected.

As the shared error goes to zero the ratio should approach the  $T_0$ -attributable value, and monotone movement towards a limit — not constancy — is the correct signature.

### 13.4 Which channel carries it

The improvement can be attributed, because the substitution runs through four separable factors. Replacing one at a time, at the strongest contrast and the finest lattice:

| substituted              | error                   | ratio         |
|--------------------------|-------------------------|---------------|
| none (all ordinary)      | $3.7826 \times 10^{-2}$ | 1.0000        |
| $A_u$ only               | $3.7826 \times 10^{-2}$ | 1.0000        |
| $A_\theta$ only          | $3.2418 \times 10^{-2}$ | <b>0.8570</b> |
| $A_{e,\text{diag}}$ only | $3.7826 \times 10^{-2}$ | 1.0000        |
| $A_{e,\text{off}}$ only  | $3.7826 \times 10^{-2}$ | 1.0000        |
| all four                 | $3.2418 \times 10^{-2}$ | 0.8570        |

**The whole of it is in  $A_\theta$** , the trace — which is the  $A_{1g}$  channel, invariant under the full cubic group. So the advantage measured at normal incidence is the genuine  $O_h$ -invariant resummation, and the  $x_3$ - preferring part of the  $O(\Delta c^2)$  difference contributes nothing. The reason is elementary once seen: a normally incident  $P$  wave is pure compression along  $z$  and never probes a shear channel at all. At  $p = 0$  the measurement is blind to precisely the part of the formulation that singles out  $x_3$ .

### 13.5 What the $x_3$ preference is worth, bracketed

Two facts have to be put together. First, a structural one: `effective_stiffness_voigt` takes a *single*  $\Delta\mu_{\text{off}}^*$  and writes it into all three shear slots, so the package's  $T_0$  is cubic-symmetric by construction and **cannot represent an  $e_{12}/e_{13}$  split**. The symmetry breaking was therefore excluded from §13 not by choice but for want of a slot. Second,  $SH$  incidence at  $p \neq 0$  does probe the shear channels, and there:

| $SH$ at $p = 10^{-4}$           | error in $R_{SH}$       | ratio         |
|---------------------------------|-------------------------|---------------|
| all ordinary                    | $4.7933 \times 10^{-3}$ | 1.0000        |
| first order, $e_{12}$ amplitude | $4.7933 \times 10^{-3}$ | 1.0000        |
| first order, $e_{13}$ amplitude | $1.8432 \times 10^{-3}$ | <b>0.3845</b> |

The  $e_{12}$  row reproducing the ordinary route *exactly* is §7 confirmed end to end: the two differ only in the traction channel, and the in-plane shear is untouched. The  $e_{13}$  row is the interesting one. An  $SH$  wave reflecting off a layer near normal incidence is dominated by  $\partial_z u_y$  — a  $yz$  shear, the  $z$ -involving one — and the  $x_3$ -distinguished amplitude gives an error 2.6 times smaller. The first-order route separates the two by  $7.2 \times 10^{-3}$ ; the ordinary route does not separate them at all.

*This brackets the effect rather than realising it.* Feeding the  $e_{13}$  amplitude through a cubic-symmetric  $T_0$  writes it into the  $e_{12}$  slot as well, which is also wrong; it is the other *uniform* choice, not the split. A faithful first-order  $T$  needs a  $T_0$  of tetragonal rather than cubic symmetry,

$e_{12} \neq e_{13} = e_{23}$  — three shear slots where there is currently one. That is a small and well-defined piece of work, and until it is done the  $x_3$  benefit is bounded, not collected.

**The answer, with its scope.** On the question as posed — same geometry, same lattice machinery, same external reference, only  $T_0$  swapped — the first-order  $T$  is the more accurate of the two, and the margin grows with contrast. The part of that measured at normal incidence is entirely the  $O_h$ -invariant resummation; the  $x_3$ -preferring part is separately shown to be worth a further factor of 2.6 in the one channel that probes it, but cannot be carried by the present  $T_0$ . This is one geometry, one frequency, normal incidence, three contrasts, and an unconverged lattice. It scores the *single-site*  $T$  only, so an error the two share is invisible by construction. And it does not touch §7.5: the  $O_h$  symmetry breaking remains a genuine  $\approx 4.2 \times 10^{-4}$  debit against the first-order route in a homogeneous whole space, where there is no preferred direction for it to be right about.

## 14 Summary of checks

| object                          | what is established                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|---------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| $\mathcal{A}$                   | built from $c_{ijkl}$ ; matches the published isotropic blocks<br>matches the thesis $\mathbf{A}(k_x, k_y)$ under one similarity transform<br>symmetry relations, spectrum, normal incidence<br>stiffness matrices, $\mathbf{U}_{al}$ , $\nu_1$ , $\nu_2$<br>source vector, negative control                                                                                                                                                                                                                                                                                                                                                                                                              |
| $\Delta\mathcal{A}$             | exact contrast operator; Navier body force on a generic field<br>nonlinearity in the contrast, quantified<br>augmentation: constant, reproduces the form, divergence exhibited<br>$O_h$ /trial-space structure, density channel                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| $\Delta\mathbf{C}_{\text{eff}}$ | $\mathbf{J}_6\Delta\mathcal{A}$ symmetric; test space derived; controls<br>$9 \times 9$ symmetric, rank 9, rigid block, basis trap<br>Born agreement, $O(\Delta c^2)$ difference, localisation<br>$O_h$ decomposition, positive and negative controls                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| $\mathbf{\Gamma}$               | propagates by $\exp(\mathbf{A} dz)$ with all multiples present                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| moments                         | face integrals evaluated; $j_1$ also by solid angle<br>$\int dV/r$ in two independent closed forms<br>$B$ , and $A = B - 1/3\mu$ , $C = -(5 - 2\pi/\sqrt{3})B$ , $G$<br>Kelvin coefficients against the $h_c$ form<br>the trace identity (42), derived without the recursion<br>the tabulated route reproduces all four                                                                                                                                                                                                                                                                                                                                                                                   |
| $\mathbf{T}$                    | coupling rebuilt independently, matches the symbolic derivation<br>calibration: reproduces the existing amplification factors<br>the finite-size density-strain term, isolated and sized<br>density channel identical; $O_h$ split on the $e_{12}/e_{13}$ pair<br>Born limits agree                                                                                                                                                                                                                                                                                                                                                                                                                       |
| Schwinger                       | $\mathcal{F}[\partial_\alpha(\mathbf{1}_V c)] = ik_\alpha \mathcal{F}[\mathbf{1}_V c]$ , with a sign control<br>$\mathcal{A}$ in the paper basis = $S^{-1}\mathcal{A}_{\text{thesis}}S$ ; spectrum; $\mathcal{A}(-\mathbf{k})^\top \mathbf{J}_6 = -\mathbf{J}_6 \mathcal{A}(\mathbf{k})$ , with a control<br>$\mathbf{\Gamma}$ inverted laterally is the Kelvin tensor<br>the radiation reaction (72), no static counterpart<br>reciprocity (71), pointwise and z-integrated<br>the two- $\Delta\mathcal{A}$ assembly is the Galerkin double-volume moment<br>the gap to the calibrated $\mathcal{Z}$ , factorised<br>the $9 \times 9$ matrix is symmetric, with a sign control and a not-vacuous control |
| thesis                          | $\mathbf{D}_z^{-1}$ from (74) is diagonal on the thesis eigenvectors, and exact after normalisation<br>$\mathbf{D}_z \mathbf{A} \mathbf{D}_z^{-1}$ reproduces (Akdef)<br>the quasi-Hamiltonian relation as the thesis states it<br>quasi-SV and quasi-SH independent at $k = 0$<br>$\mathbf{\Gamma}$ reciprocal, and the Kelvin tensor with the radiation reaction, from the analytic modes alone                                                                                                                                                                                                                                                                                                         |
| slab                            | the Navier route reproduces the package's amplification factors<br>the advantage grows as the shared lattice error falls<br>the first-order $T_0$ is the more accurate of the two<br>the improvement survives either choice of the split channel<br>SH incidence does probe the shear <sup>38</sup> channels<br>the one-pass assembly reproduces the density channel                                                                                                                                                                                                                                                                                                                                      |
| layer                           | normal incidence vs (76), three contrasts                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |

**Table 13.** All checks pass. The symbolic scripts report 25/25 and 44/44, the  $T$ -matrix assembly 14/14, the Schwinger settlement 17/17, the layer comparison 7/7, the thesis spectral representation 8/8, the slab comparison 5/5, and the propagator gate exits clean.

## 15 What is not established

In keeping with the project’s standard that an unvalidated formula be labelled as such, the following are *open*:

1. **The  $T$ -matrix uses the collocation propagator, not the consistent one.**  $\mathbf{T}$  is assembled in §9 and calibrated exactly, but  $\mathcal{Z}$  there is the *single* cube moment while the  $\mathbf{J}_6$  pairing makes the scheme Galerkin. This is now measured rather than suspected: (70) separates the gap into a geometric factor  $\int_V 1/r/D_0 = 1.2644$  and a factor  $i\omega$ , and the double moment  $D_0$  that the consistent pairing needs is extracted in §10.4. The  $i\omega$  in (70) is *not* open — see §7.3 — so what remains is the **geometric** factor alone. Until that is resolved, everything in §9 is the first-order coupling seen through a collocation propagator.
2. **The external arbiter reaches the operator, not the cube.** §12 closes this item for  $\mathcal{A}$ ,  $\mathbf{\Gamma}$ , their composition and the resummation: a layer in a whole space reproduces an independently validated reflectivity recursion across contrast and angle. But a laterally invariant problem does not touch the augmentation, the form factors or the moments, so the quantities in §7 and §9 still compare two *formulations* with each other rather than either against truth. The  $O_h$  split remains the only partial exception there, and it is one-sided: it lower-bounds the first-order error at  $\approx 4.2 \times 10^{-4}$  and bounds the ordinary route not at all. §13 then answers the accuracy question for the cube by the one geometry in which both formulations apply and an external reference exists — a space-filling patch — and the first-order  $T$  wins there by a margin that grows with contrast. What remains genuinely open is the *scope* of that: one frequency, normal incidence, an unconverged lattice, and no measurement at all away from space-filling.
3. **Whether the resummation improves accuracy is unknown.** It is exact as a local inversion; that is not the same claim.
4.  $O_h$  **invariance was tested on a subgroup of order 3**, not the full group.
5. **The derivative count for a collocation scheme differs from the Galerkin one.** In the Galerkin pairing the transferred derivative lands on polynomials, so the third-derivative moment cannot arise; under collocation it transfers onto  $\mathbf{\Gamma}$  and can. Do not quote the Galerkin count for a collocation build. (66) is how that transfer is carried out in practice, and it is verified; what is not yet done is the count itself.
6. **The derivative channels are verified only by symmetry.** That is a strong test — ten orders of separation from its control, and the derivative terms carry 80% of the matrix — but symmetry is a *structural* property. A bookkeeping error that is symmetric would pass it. The value of those channels has not been compared with an independent arbiter, which would need the Galerkin double-volume moments of  $\partial_i \partial_j G_{kl}$ ; those follow from the same autocorrelation and Duffy construction as  $D_0$ , with  $\partial_i \partial_j$  moved onto the weight, but they have not been computed.

7. **The present  $T_0$  cannot carry the  $x_3$  benefit.** `effective_stiffness_voigt` has one shear slot and writes it into all three, so it is cubic-symmetric by construction. §13.4 shows the  $x_3$ -distinguished shear amplitude is worth a factor of 2.6 in the channel that probes it, and that this is *bracketed* rather than realised: passing  $e_{13}$ 's amplitude through a cubic  $T_0$  writes it into the  $e_{12}$  slot as well. Collecting it needs a tetragonal  $T_0$ ,  $e_{12} \neq e_{13} = e_{23}$ . Small, well defined, not done.
8. **The shear moment  $B$  has not been read off.** This is the comparison Task 1 exists for —  $B$  against  $\sqrt{3}(\lambda + \mu)/6\pi\mu(\lambda + 2\mu)$  — and it is now the one remaining step of the task rather than a matter of method.

## 16 Where the route belongs

The  $x_3$  preference of §7.5 is an artefact only because the problem posed there — an isotropic cube in an isotropic whole space — has no preferred direction. A layered background does. There the formulation's directional structure matches the physics rather than contradicting it, the  $O_h$  argument does not apply, and the propagator is already available and already bridged to  $\mathcal{A}$  by (27).

That was written as an expectation and is now a measurement. In the stratified geometry of §13, the  $x_3$ -distinguished shear amplitude is the one the physics wants, by a factor of 2.6 in the channel that probes it (§13.4). The same feature that is demonstrably wrong in a whole space —  $\geq 4.2 \times 10^{-4}$  of error, by §7.5 — is demonstrably right once there is a stack of interfaces for it to be right about. Both statements are measured, and they are not in tension: they are about different problems.

That is also the case in which the ordinary route has no competitor: the cube moment method rests on volume moments of the whole-space Green's tensor, and in a stratified background that object does not exist. The first-order form supplies what the moment method cannot.