

# The First-Order Contrast Operator

The elastodynamic matrix–vector wave equation as a route to the single-site  $T$ -matrix: derivation, validation, and what it actually buys

17 September 2026

## Abstract

The elastodynamic matrix–vector wave equation  $\partial_3 \mathbf{q} = \mathcal{A} \mathbf{q} + \mathbf{d}$  expresses the wave field as a first-order recursion in depth. This note asks whether the cube moment method can be applied to it directly, to obtain a single-site  $T$ -matrix without passing through the Navier equation. The answer is yes, and the construction is not new: the thesis already carries it in  $2\frac{1}{2}$ -D. What is new here is the 3-D Cartesian operator, and a measurement of what the route is worth.

Three results. *First*, the operator matrix  $\mathcal{A}$  is derived from the stiffness tensor and reproduces the thesis system matrix (Akdef) under one constant similarity transform; the paper’s  $\nu_1, \nu_2$  are the thesis’s  $\zeta, \chi$ . A fourth, arrived at late and by the wrong road, is that the thesis’s own spectral representation of  $\mathbf{A}$  — analytic eigenvectors, a piecewise branch, and an inverse obtained from the symplectic structure rather than by inversion — removes three numerical hazards outright that the operator route had to discover and repair one at a time. *Second*, the test space is *derived* rather than chosen:  $\mathbf{J}_6$  is antisymmetric and  $\Delta \mathcal{A}^t \mathbf{J}_6 = -\mathbf{J}_6 \Delta \mathcal{A}$ , so  $\mathbf{J}_6 \Delta \mathcal{A}$  is symmetric and the only natural projection is test =  $\mathbf{J}_6 \times$  trial. This dissolves an apparent mismatch of ten conditions against nine unknowns. *Third*, and least expected, the resulting coupling agrees with the ordinary Navier weak form at Born order but differs at  $O(\Delta c^2)$  entirely within the traction channel — and that difference splits into a genuine  $O_h$ -invariant resummation *and* a spurious preference for the  $x_3$  axis, in comparable amounts. A single-site score against an independent arbiter in a homogeneous background would mix the two and settle neither. The directional preference ceases to be an artefact once the background is layered, which is where the route was wanted.

The  $T$ -matrix itself is assembled in §9 and the comparison is reported channel by channel. The sharpest statement is not a norm:  $e_{12}$  and  $e_{13}$  are equivalent under  $O_h$  and so must amplify identically, and in the first-order  $T$  they split by  $8.4 \times 10^{-4}$  while the ordinary route splits them by exactly zero. That split is an arbiter-free *lower bound* on the first-order error,  $\geq \delta/2$ ; it bounds nothing about the ordinary route, which can be symmetrically wrong, so no ordering of the two follows from it.

The ordering is settled elsewhere, in the one geometry where both formulations apply and an external reference exists. A periodic patch of space-filling cubes *is* a homogeneous layer, so a reflectivity recursion gives its reflection coefficient exactly; running both routes through the same lattice sum and the same solve, with only  $T_0$  substituted, the first-order  $T$  is the more accurate of the two, by a margin that grows with contrast and that *grows further as the shared lattice error is refined away* — which is what distinguishes a real advantage from an accidental cancellation (§14).

Every claim below is backed by an executable check; the scripts report 25/25, 44/44, 14/14, 17/17, 9/9, 7/7, 7/7 and 11/11, and a separate gate binds the numerical propagator to the derived  $\mathcal{A}$ . Section 18 states plainly what is *not* yet established.

## Contents

|          |                                                                                             |           |
|----------|---------------------------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Conventions and objects</b>                                                              | <b>5</b>  |
| 1.1      | Indices, time, and units . . . . .                                                          | 5         |
| 1.2      | The field vector . . . . .                                                                  | 5         |
| 1.3      | The metric $\mathbf{J}_6$ . . . . .                                                         | 5         |
| <b>2</b> | <b>The first-order system for an isotropic medium</b>                                       | <b>6</b>  |
| 2.1      | From the stiffness tensor to the stiffness matrices . . . . .                               | 6         |
| 2.2      | Eliminating the in-plane tractions . . . . .                                                | 6         |
| 2.3      | The operator matrices . . . . .                                                             | 7         |
| <b>3</b> | <b><math>\mathcal{A}</math> against the thesis, and against the published operator form</b> | <b>7</b>  |
| 3.1      | Arbiter A: the published isotropic operator matrices . . . . .                              | 7         |
| 3.2      | The anchor: the thesis system matrix (Akdef) . . . . .                                      | 8         |
| 3.3      | Structural checks . . . . .                                                                 | 8         |
| <b>4</b> | <b>The contrast operator</b>                                                                | <b>9</b>  |
| 4.1      | Definition and the Lippmann–Schwinger equation . . . . .                                    | 9         |
| 4.2      | The foundation check . . . . .                                                              | 9         |
| 4.3      | $\Delta\mathcal{A}$ is not linear in the contrast . . . . .                                 | 9         |
| <b>5</b> | <b>The augmentation</b>                                                                     | <b>10</b> |
| 5.1      | Why a derivative must be moved . . . . .                                                    | 10        |
| 5.2      | The augmented dimension . . . . .                                                           | 10        |
| <b>6</b> | <b>The test space, derived</b>                                                              | <b>10</b> |
| 6.1      | The argument . . . . .                                                                      | 10        |
| 6.2      | The ten-versus-nine mismatch dissolves . . . . .                                            | 11        |
| 6.3      | The effective coupling matrix . . . . .                                                     | 11        |
| <b>7</b> | <b>What the first-order coupling actually buys</b>                                          | <b>11</b> |
| 7.1      | The trial spaces coincide . . . . .                                                         | 11        |
| 7.2      | The density channel cannot differ . . . . .                                                 | 12        |
| 7.3      | The units the two couplings are written in . . . . .                                        | 12        |
| 7.4      | Born agreement, and an $O(\Delta c^2)$ difference . . . . .                                 | 12        |
| 7.5      | Half of it is a broken symmetry . . . . .                                                   | 13        |
| 7.6      | What follows, and what does not . . . . .                                                   | 13        |

|           |                                                                         |           |
|-----------|-------------------------------------------------------------------------|-----------|
| <b>8</b>  | <b>The propagator, and the bridge to it</b>                             | <b>14</b> |
| 8.1       | It already exists . . . . .                                             | 14        |
| 8.2       | The bridge identity . . . . .                                           | 14        |
| 8.3       | Reporting residuals against the arithmetic floor . . . . .              | 14        |
| <b>9</b>  | <b>The single-site <math>T</math>-matrix</b>                            | <b>15</b> |
| 9.1       | What is new and what is borrowed . . . . .                              | 15        |
| 9.2       | The self moments, derived . . . . .                                     | 15        |
| 9.2.1     | The static Green's tensor . . . . .                                     | 16        |
| 9.2.2     | The moment as a distribution, and why no delta is inserted . . . . .    | 16        |
| 9.2.3     | The face integrals, and one of them by pure geometry . . . . .          | 16        |
| 9.2.4     | The closed forms . . . . .                                              | 17        |
| 9.2.5     | The check that certifies the delta content . . . . .                    | 17        |
| 9.3       | The tabulated route agrees . . . . .                                    | 18        |
| 9.4       | Two implementations of the coupling . . . . .                           | 18        |
| 9.5       | Calibration: every normalisation fixed against existing code . . . . .  | 18        |
| 9.6       | The first-order $T$ , channel by channel . . . . .                      | 19        |
| 9.7       | The broken symmetry, now visible in $T$ itself . . . . .                | 19        |
| 9.8       | What this does and does not settle . . . . .                            | 20        |
| <b>10</b> | <b>The propagator moment from the first-order system</b>                | <b>20</b> |
| 10.1      | Step 1: $\Gamma$ , built from $\mathcal{A}$ alone . . . . .             | 21        |
| 10.2      | Step 2: the vertical integral, in closed form . . . . .                 | 22        |
| 10.2.1    | The same integrals as a series, and why they are needed . . . . .       | 22        |
| 10.3      | Step 3: the lateral integral . . . . .                                  | 23        |
| 10.3.1    | The scaling that makes the expansion uniform . . . . .                  | 23        |
| 10.3.2    | Which entries fail to decay . . . . .                                   | 24        |
| 10.3.3    | The asymptote in closed form . . . . .                                  | 24        |
| 10.3.4    | The lateral integral, and where it must be evaluated . . . . .          | 25        |
| 10.3.5    | The principal-value piece vanishes . . . . .                            | 26        |
| 10.3.6    | Why the transferred form converges . . . . .                            | 26        |
| 10.3.7    | The tail, in closed form . . . . .                                      | 27        |
| 10.3.8    | The remainder, and the test that closes the construction . . . . .      | 27        |
| 10.4      | Step 4: how the augmentation composes with the Schwinger form . . . . . | 28        |
| <b>11</b> | <b>The thesis spectral representation, and what it removes</b>          | <b>32</b> |

|           |                                                                                       |           |
|-----------|---------------------------------------------------------------------------------------|-----------|
| 11.1      | What the thesis gives . . . . .                                                       | 32        |
| 11.2      | What it removes . . . . .                                                             | 33        |
| <b>12</b> | <b>A layer in a whole space: the external arbiter</b>                                 | <b>34</b> |
| 12.1      | Why a layer is the right test . . . . .                                               | 34        |
| 12.2      | The closed form, by continuity . . . . .                                              | 34        |
| 12.3      | The first-order route . . . . .                                                       | 35        |
| 12.4      | What it gives . . . . .                                                               | 35        |
| <b>13</b> | <b>The static moments <math>A</math>, <math>B</math> and <math>C</math>, read off</b> | <b>36</b> |
| 13.1      | Which realisation, and why it is the easy one . . . . .                               | 36        |
| 13.2      | The sign, pinned independently . . . . .                                              | 36        |
| 13.3      | The result . . . . .                                                                  | 37        |
| 13.4      | $A$ and $C$ , and the delta they carry . . . . .                                      | 37        |
| 13.5      | Two things that nearly went wrong . . . . .                                           | 38        |
| <b>14</b> | <b>Which formulation is more accurate</b>                                             | <b>39</b> |
| 14.1      | Substituting at the one place both routes speak the same language . . . . .           | 39        |
| 14.2      | The measurement . . . . .                                                             | 39        |
| 14.3      | Why this is not cancellation . . . . .                                                | 39        |
| 14.4      | Which channel carries it . . . . .                                                    | 40        |
| 14.5      | What the $x_3$ preference is worth, bracketed . . . . .                               | 40        |
| 14.6      | The tetragonal $T_0$ . . . . .                                                        | 41        |
| <b>15</b> | <b>The layered background</b>                                                         | <b>42</b> |
| 15.1      | Nothing here is new machinery . . . . .                                               | 42        |
| 15.2      | The source normalisation, which the bridge could not see . . . . .                    | 42        |
| 15.3      | The uniform limit, which cannot be what was planned . . . . .                         | 43        |
| <b>16</b> | <b>From the integral equation to invariant imbedding</b>                              | <b>43</b> |
| 16.1      | The kernel is one-sided, so the integral equation splits . . . . .                    | 44        |
| 16.2      | One-sidedness makes the discretised operator tridiagonal . . . . .                    | 44        |
| 16.3      | The imbedding equation . . . . .                                                      | 45        |
| 16.4      | What this replaces, and what it does not . . . . .                                    | 45        |
| 16.5      | Building the coefficient matrices . . . . .                                           | 46        |
| 16.6      | Is the march stiff? . . . . .                                                         | 46        |
| 16.7      | The real obstruction is non-normality . . . . .                                       | 46        |

|                                                            |           |
|------------------------------------------------------------|-----------|
| 16.8 The impedance form, and the ceiling cleared . . . . . | 47        |
| 16.9 Marching it through a slab, and scoring it . . . . .  | 48        |
| 16.10 What the convention actually is . . . . .            | 49        |
| 16.11 Where the lateral coupling enters . . . . .          | 50        |
| <b>17 Summary of checks</b>                                | <b>50</b> |
| <b>18 What is not established</b>                          | <b>52</b> |
| <b>19 Where the route belongs</b>                          | <b>53</b> |

## 1 Conventions and objects

### 1.1 Indices, time, and units

The positive  $x_3$  axis points downward;  $x_1, x_2$  are lateral. Greek indices  $\alpha, \beta$  run over the lateral pair  $\{1, 2\}$  only, Latin indices over  $\{1, 2, 3\}$ . The time convention is  $e^{-i\omega t}$  throughout, so  $\partial_t \rightarrow -i\omega$  and the particle velocity is  $\mathbf{v} = -i\omega \mathbf{u}$ . The lateral transform is  $e^{+i(k_1 x_1 + k_2 x_2)}$ , hence  $\partial_\alpha \rightarrow ik_\alpha$ .

Numerical work uses seismic units: km,  $\text{km s}^{-1}$ ,  $\text{g cm}^{-3}$ , GPa. These are consistent, since  $\text{GPa} = \text{g cm}^{-3} \cdot (\text{km s}^{-1})^2$ .

Two orderings appear in the sources and must not be conflated. The paper orders components  $(1, 2, 3)$ ; the thesis orders them  $(z, x, y) = (x_3, x_1, x_2)$ . The map between them is written explicitly in §3.2.

### 1.2 The field vector

The matrix–vector wave equation is

$$\partial_3 \mathbf{q} = \mathcal{A} \mathbf{q} + \mathbf{d}, \quad \mathbf{q} = \begin{pmatrix} \mathbf{q}_1 \\ \mathbf{q}_2 \end{pmatrix}, \quad \mathcal{A} = \begin{pmatrix} \mathcal{A}_{11} & \mathcal{A}_{12} \\ \mathcal{A}_{21} & \mathcal{A}_{22} \end{pmatrix}, \quad (1)$$

with the elastodynamic choice

$$\mathbf{q}_1 = -\boldsymbol{\tau}_3, \quad \mathbf{q}_2 = \mathbf{v}, \quad (\boldsymbol{\tau}_j)_i = \tau_{ij}, \quad (2)$$

so that  $\mathbf{q} = (-\tau_{13}, -\tau_{23}, -\tau_{33}, v_1, v_2, v_3)^t$ . This is not an arbitrary pairing:  $-\boldsymbol{\tau}_3$  and  $\mathbf{v}$  are conjugate in the vertical power flux,  $j = \frac{1}{4}(-\boldsymbol{\tau}_3^\dagger \mathbf{v} - \mathbf{v}^\dagger \boldsymbol{\tau}_3)$ , and that is what makes the metric of §1.3 the natural one.

### 1.3 The metric $\mathbf{J}_6$

The symplectic metric is

$$\mathbf{J}_6 = \begin{pmatrix} \mathbf{0} & \mathbf{I}_3 \\ -\mathbf{I}_3 & \mathbf{0} \end{pmatrix}, \quad \mathbf{J}_6^t = -\mathbf{J}_6, \quad \mathbf{J}_6^2 = -\mathbf{I}_6. \quad (3)$$

**A naming warning.** This project, following the thesis, calls this matrix  $\mathbf{J}_6$ ; it is already defined as such in the package. The paper calls the *same* matrix  $\mathbf{N}$ , and uses  $\mathbf{J}$  for  $\text{diag}(\mathbf{I}_3, -\mathbf{I}_3)$ , a different object. Translate on the way in. Every occurrence of  $\mathbf{N}$  in the paper's equations (16)–(30) corresponds to  $\mathbf{J}_6$  here.

The operator matrix is *quasi-Hamiltonian* with respect to it,

$$\mathcal{A}^t \mathbf{J}_6 = -\mathbf{J}_6 \mathcal{A}, \quad (4)$$

which is the paper's eq. (27) and the thesis's ATdef. Everything in §6 follows from (4) together with the antisymmetry of  $\mathbf{J}_6$ , and from nothing else.

## 2 The first-order system for an isotropic medium

### 2.1 From the stiffness tensor to the stiffness matrices

For an isotropic solid,

$$c_{ijkl} = \lambda \delta_{ij} \delta_{kl} + \mu (\delta_{ik} \delta_{jl} + \delta_{il} \delta_{jk}), \quad \rho_{ij} = \rho \delta_{ij}, \quad (5)$$

with  $\lambda = \lambda(\mathbf{x}, \omega)$ ,  $\mu = \mu(\mathbf{x}, \omega)$ ,  $\rho = \rho(\mathbf{x}, \omega)$  and  $K_c = \lambda + 2\mu$ . The stiffness matrices are  $(\mathbf{C}_{jl})_{ik} = c_{ijkl}$ . In particular

$$\mathbf{C}_{33} = \text{diag}(\mu, \mu, K_c), \quad \mathbf{C}_{33}^{-1} = \text{diag}(\mu^{-1}, \mu^{-1}, K_c^{-1}), \quad (6)$$

and  $\mathbf{C}_{jl}^t = \mathbf{C}_{lj}$ , which follows from  $c_{ijkl} = c_{klij}$ . All nine blocks are reproduced from (5) and checked entry by entry.

### 2.2 Eliminating the in-plane tractions

The components  $\boldsymbol{\tau}_1, \boldsymbol{\tau}_2$  do not appear in (2) and must be removed. Doing so introduces the reduced stiffness

$$\mathbf{U}_{\alpha l} = \mathbf{C}_{\alpha l} - \mathbf{C}_{\alpha 3} \mathbf{C}_{33}^{-1} \mathbf{C}_{3l}, \quad (7)$$

with the two structural properties  $\mathbf{U}_{\alpha 3} = \mathbf{0}$  and  $\mathbf{U}_{\alpha\beta}^t = \mathbf{U}_{\beta\alpha}$ . Both are checked rather than assumed, because it is they — not the algebra that produced them — that make the symmetry relations of §3.3 hold.

The isotropic  $\mathbf{U}_{\alpha\beta}$  is where two constants are *born*:

$$v_1 \equiv (\mathbf{U}_{11})_{11} = \frac{4\mu(\lambda + \mu)}{\lambda + 2\mu}, \quad v_2 \equiv (\mathbf{U}_{12})_{12} = \frac{2\mu\lambda}{\lambda + 2\mu}, \quad v_1 - v_2 = 2\mu. \quad (8)$$

They are read off  $\mathbf{U}$  and only then compared with the published definitions; they are not independent quantities. Every  $\mathbf{U}_{\alpha\beta}$  has a null third row and column, which is the statement that  $\tau_{33}$  has been fully eliminated.

### 2.3 The operator matrices

The general construction gives

$$\mathcal{A}_{11} = -\partial_\alpha \mathbf{C}_{\alpha 3} \mathbf{C}_{33}^{-1}, \quad \mathcal{A}_{12} = i\omega\rho - \frac{1}{i\omega} \partial_\alpha \mathbf{U}_{\alpha\beta} \partial_\beta, \quad \mathcal{A}_{21} = i\omega \mathbf{C}_{33}^{-1}, \quad \mathcal{A}_{22} = -\mathbf{C}_{33}^{-1} \mathbf{C}_{3\beta} \partial_\beta, \quad (9)$$

and for the isotropic medium these evaluate to

$$\mathcal{A}_{11} = \begin{pmatrix} 0 & 0 & -\partial_1 \frac{\lambda}{\lambda+2\mu} \\ 0 & 0 & -\partial_2 \frac{\lambda}{\lambda+2\mu} \\ -\partial_1 & -\partial_2 & 0 \end{pmatrix}, \quad \mathcal{A}_{22} = \begin{pmatrix} 0 & 0 & -\partial_1 \\ 0 & 0 & -\partial_2 \\ -\frac{\lambda}{\lambda+2\mu} \partial_1 & -\frac{\lambda}{\lambda+2\mu} \partial_2 & 0 \end{pmatrix}, \quad (10)$$

$$\mathcal{A}_{12} = \begin{pmatrix} i\omega\rho - \frac{1}{i\omega}(\partial_1 \nu_1 \partial_1 + \partial_2 \mu \partial_2) & -\frac{1}{i\omega}(\partial_2 \mu \partial_1 + \partial_1 \nu_2 \partial_2) & 0 \\ -\frac{1}{i\omega}(\partial_2 \nu_2 \partial_1 + \partial_1 \mu \partial_2) & i\omega\rho - \frac{1}{i\omega}(\partial_1 \mu \partial_1 + \partial_2 \nu_1 \partial_2) & 0 \\ 0 & 0 & i\omega\rho \end{pmatrix}, \quad (11)$$

with  $\mathcal{A}_{21} = i\omega \text{diag}(\mu^{-1}, \mu^{-1}, K_c^{-1})$ .

**The ordering is the content.** Compare the (1, 3) entry of  $\mathcal{A}_{11}$  with the (3, 1) entry of  $\mathcal{A}_{22}$ . In the first the derivative stands to the *left* of the modulus and differentiates the product; in the second it stands to the *right* and differentiates the field alone. For a laterally invariant medium the distinction collapses and is invisible. For a laterally varying one it is the physics. This is why the checks in §3.1 compare *operators* — by their action on generic fields — rather than expressions, and why they are accompanied by a negative control that corrupts precisely this ordering.

## 3 $\mathcal{A}$ against the thesis, and against the published operator form

**A note on where this starts.** An earlier version of this note presented the published operator matrices as the foundation and the thesis system matrix as an arbiter for them. That is the wrong way round. The thesis carries  $\mathbf{A}(k_x, k_y)$  in §2.1, equation (Akdef), and the purpose of this work is to extend and test the thesis, not to check the thesis against a later paper. What the operator form genuinely contributes is narrower: the thesis has the system only in the laterally *invariant* case, where  $\partial_\alpha \rightarrow ik_\alpha$  has already been taken, and a cube requires the coefficients to vary laterally. So the operator form is a *generalisation* of (Akdef), and it is that generalisation — not the system — that is new here. The consequences of getting this backwards were not only presentational; they are collected in §11.

### 3.1 Arbiter A: the published isotropic operator matrices

The blocks (10)–(11) are transcribed separately and compared against those built from (9). The comparison is made by acting on a vector of generic symbolic functions, so it tests the derivative ordering as well as the coefficients. All four match.

A check that cannot fail is worth nothing, so a corrupted  $\mathcal{A}_{22}$  — with the derivative wrongly placed outside the modulus — is also tested, and is rejected.

### 3.2 The anchor: the thesis system matrix (Akdef)

The thesis carries a  $6 \times 6$  Fourier-domain system matrix  $\mathbf{A}(k_x, k_y)$  for a laterally invariant medium, in a *different* field vector

$$\mathbf{b} = (u_3, u_1, u_2, \tau_{33}, \tau_{31}, \tau_{32})^t, \quad (12)$$

i.e. displacement–stress rather than velocity–traction, ordered  $(z, x, y)$  rather than  $(1, 2, 3)$ , and normalised so that  $\rho\omega^2$  appears where the paper has  $i\omega\rho$ . It uses the shorthand

$$\gamma = \frac{\lambda}{\lambda + 2\mu}, \quad a = \frac{1}{\lambda + 2\mu}, \quad b = \frac{1}{\mu}, \quad \zeta = \frac{4\mu(\lambda + \mu)}{\lambda + 2\mu}, \quad \chi = \frac{2\mu\lambda}{\lambda + 2\mu}. \quad (13)$$

The thesis attaches a warning to these that must travel with them: the shorthand  $\gamma, a, b, \zeta, \chi$  is *local to* (Akdef), and is *not* the  $\gamma, \zeta, \chi$  of the List of Symbols, which denote the Chapter 5 material-contrast ratios. Since this note uses  $\zeta$  and  $\chi$  throughout as the in-plane stiffnesses that survive the elimination of  $\epsilon_{33}$ , the collision is worth stating once rather than inheriting silently. With that said,  $\zeta - \chi \equiv 2\mu$ , which is the thesis’s own footnote and the cleanest check on the pair.

**Comparing (8) with these:**  $\zeta \equiv \nu_1$  and  $\chi \equiv \nu_2$ . The thesis’s in-plane stiffnesses and the paper’s auxiliary moduli are the same two quantities, born in the same elimination (7). This is verified, not asserted.

Setting  $\lambda, \mu, \rho$  constant and acting on  $e^{i(k_1 x_1 + k_2 x_2)}$  reduces  $\mathcal{A}$  to a matrix. With the constant map  $\mathbf{b} = \mathbf{S}\mathbf{q}$  built from  $\mathbf{u} = \mathbf{v}/(-i\omega)$ ,  $\boldsymbol{\tau}_3 = -\mathbf{q}_{1:3}$  and the permutation  $(1, 2, 3) \rightarrow (3, 1, 2)$ , the identity

$$\mathbf{A}_{\text{thesis}} = \mathbf{S} \mathbf{A}_{\text{paper}} \mathbf{S}^{-1} \quad (14)$$

holds entry by entry. Since the two matrices were derived independently and published separately, this is evidence rather than bookkeeping.

### 3.3 Structural checks

Three further properties are verified on the reduced  $6 \times 6$ :

1. the operator symmetry relations  $\mathcal{A}_{11}^t = -\mathcal{A}_{22}$ ,  $\mathcal{A}_{12}^t = \mathcal{A}_{12}$ ,  $\mathcal{A}_{21}^t = \mathcal{A}_{21}$ , where the transpose is the formal adjoint under integration by parts ( $\partial_\alpha^t = -\partial_\alpha$ );
2. the quasi-Hamiltonian relation (4), which in the wavenumber domain reads  $\mathbf{A}^t(-\mathbf{k})\mathbf{J}_6 = -\mathbf{J}_6\mathbf{A}(\mathbf{k})$ ;

3. the characteristic polynomial factorises as  $(s^2 + k_{z,p}^2)(s^2 + k_{z,s}^2)^2$ , recovering  $\alpha^2 = (\lambda + 2\mu)/\rho$  and  $\beta^2 = \mu/\rho$  from a construction in which no wave speed was ever written down.

Item 3 is the one that binds the algebra to physics: nothing upstream knows the wave speeds.

## 4 The contrast operator

### 4.1 Definition and the Lippmann–Schwinger equation

Inside the scatterer the medium is  $(\lambda + \Delta\lambda, \mu + \Delta\mu, \rho + \Delta\rho)$ . Writing  $\mathcal{A}^{\text{tot}} = \mathcal{A} + \Delta\mathcal{A}$ , the total field solves the background system driven by its own contrast,

$$(\partial_3 - \mathcal{A}) \mathbf{q} = \Delta\mathcal{A} \mathbf{q} + \mathbf{d}, \quad (15)$$

which is the Lippmann–Schwinger equation of the first-order system.

### 4.2 The foundation check

Before (15) means anything,  $\mathbf{q}$  must actually satisfy (1). Building  $\mathbf{q}$  from a *generic* displacement field in a *laterally varying* medium and forming  $\partial_3 \mathbf{q} - \mathcal{A} \mathbf{q}$  gives two halves with different strengths:

- rows 4–6 vanish **identically**, for any  $\mathbf{u}$  whatever. They are the constitutive relation rearranged; no equation of motion is used. Failure here would mean the elimination (7) is wrong.
- rows 1–3 equal **exactly** the Navier body force  $f_i = -\rho\omega^2 u_i - \partial_j \tau_{ij}$ . Anything left over would be a defect in  $\mathcal{A}_{11}$  or  $\mathcal{A}_{12}$ .

Both hold. This is a considerably stronger statement than the plane-wave and eigenvalue checks of §3.3, because  $\mathbf{u}$  is generic and the moduli vary laterally, so every ordering term in (10) is exercised.

### 4.3 $\Delta\mathcal{A}$ is not linear in the contrast

The Navier potential is exactly linear in  $(\Delta\lambda, \Delta\mu, \Delta\rho)$ .  $\Delta\mathcal{A}$  is not, because the elimination inverted  $\mathbf{C}_{33}$ :

$$(\Delta\mathcal{A}_{21})_{11} = i\omega \left( \frac{1}{\mu + \Delta\mu} - \frac{1}{\mu} \right) = -\frac{i\omega \Delta\mu}{\mu(\mu + \Delta\mu)}, \quad (16)$$

whose ratio to its linearisation  $-i\omega\Delta\mu/\mu^2$  is exactly  $\mu/(\mu + \Delta\mu)$ . At the project’s validated moderate contrast ( $\alpha = 5$ ,  $\beta = 3$ ,  $\rho = 2.5$ , so  $\lambda = 17.5$ ,  $\mu = 22.5$  GPa;  $\Delta\lambda = 2$ ,  $\Delta\mu = 1$  GPa):

| quantity                                      | value   |
|-----------------------------------------------|---------|
| $\Delta\mu/\mu$                               | 0.04444 |
| exact / linear for $\Delta b = \Delta(1/\mu)$ | 0.95745 |
| exact / linear for $\Delta a = \Delta(1/K_c)$ | 0.93985 |

The first-order form has already summed part of the local response exactly. §7 measures what that is worth, and finds the answer is not the obvious one.

## 5 The augmentation

### 5.1 Why a derivative must be moved

Inside the scatterer the contrasts are constant, so every coefficient of  $\Delta\mathcal{A}$  carries the indicator  $\mathbf{1}_V$ . A derivative standing to the *left* of such a coefficient differentiates the indicator and produces a surface delta on the cube faces. Integration by parts over all space moves that derivative onto the test function instead, where it meets something smooth:

$$\int \phi_i (-\partial_\alpha (m \mathbf{1}_V \psi_j)) d\mathbf{x} = + \int (\partial_\alpha \phi_i) m \mathbf{1}_V \psi_j d\mathbf{x}, \quad (17)$$

and the face term never arises. This is the step the thesis describes as transferring the derivatives onto the source and receiver coordinates of the propagator, and it is what converts a singular integro-differential equation into a nonsingular, compact integral equation.

### 5.2 The augmented dimension

After the transfer every term has the form  $(\partial^a \phi)^t \cdot \text{const} \cdot (\partial^b \psi)$  with  $a, b \in \{0, 1\}$ , so the form is *multiplicative* on the enlarged vector  $\mathcal{L}\psi = (\psi, \partial_1 \psi, \partial_2 \psi)$  — nominally  $6 + 6 + 6 = 18$  slots. The true dimension is smaller and is counted, not assumed.

**Result.** Exactly  $10 = 6 + 4$  slots are used. The four extra are  $\partial_1 v_1, \partial_2 v_1, \partial_1 v_2, \partial_2 v_2$ : the lateral gradients of the two *horizontal* velocity components, and nothing else.

Setting  $\partial_2 \rightarrow ik_y$ , as a  $2\frac{1}{2}$ -D formulation does, makes two of those four multiplicative, leaving  $6 + 2 = 8$ . That is precisely the thesis's count for its augmented coupling matrix, reproduced here from an independent 3-D Cartesian derivation. The agreement is a corroboration of the thesis construction, not a restatement of it.

The transfer identity is verified three ways: the augmented matrix is constant (no derivative operators survive); the augmented form reproduces the transferred bilinear; and the untransferred and transferred bilinears differ by an explicitly exhibited lateral divergence  $\partial_\alpha P_\alpha$ .

## 6 The test space, derived

### 6.1 The argument

At this point there are ten augmented slots and, for an affine trial space, nine parameters. The temptation is to choose nine conditions. That temptation must be resisted: a subset chosen by scoring against the arbiter it is later judged by is a fitted parameter wearing a derivation's clothes.

There is no choice to make.  $\mathbf{J}_6$  is *antisymmetric*, and (4) applies to  $\Delta\mathcal{A}$  as well as to  $\mathcal{A}$  (it is linear in the operator), so

$$(\mathbf{J}_6 \Delta\mathcal{A})^t = \Delta\mathcal{A}^t \mathbf{J}_6^t = \Delta\mathcal{A}^t (-\mathbf{J}_6) = \mathbf{J}_6 \Delta\mathcal{A}. \quad (18)$$

$\mathbf{J}_6 \Delta \mathcal{A}$  is a **symmetric** matrix, so

$$b[\phi, \psi] = \phi^t \mathbf{J}_6 \Delta \mathcal{A} \psi \quad (19)$$

is a symmetric bilinear form. A symmetric form admits exactly one natural projection:

$$\boxed{\text{test space} = \mathbf{J}_6 \times \text{trial space.}} \quad (20)$$

Verified, together with the negative control: drop  $\mathbf{J}_6$  and the symmetry is lost, both at operator level and in the projected system.

## 6.2 The ten-versus-nine mismatch dissolves

The ten slots were never test conditions. They measure how many derivative channels the *form* touches — its representation — not how many independent functionals it supplies. Each trial function  $\psi_k$  gives exactly one test functional  $(\mathbf{J}_6 \psi_k)^t$ , so the number of conditions equals the trial dimension by construction.

## 6.3 The effective coupling matrix

For the affine trial space — nine parameters, three rigid translations and six uniform strains — the projected coupling

$$(\Delta \mathbf{C}_{\text{eff}})_{kl} = \int_V (\mathcal{L} \psi_k)^t \Delta \mathbf{C} (\mathcal{L} \psi_l) dV \quad (21)$$

is the 3-D Cartesian analogue of the thesis's effective coupling matrix. It is  $9 \times 9$ , symmetric, of full rank 9, vanishes with the contrast, and — because the trial functions are polynomials — requires **no Green's function at all**: the cube integral is exact.

The rigid-translation block is diagonal,

$$(\Delta \mathbf{C}_{\text{eff}})_{1:3,1:3} = i\omega^3 \Delta \rho V \mathbf{I}_3, \quad V = (2a)^3. \quad (22)$$

**The power of  $\omega$  is a basis statement, not a physical one.**  $\Delta \mathcal{A}$ 's density entry is  $i\omega \Delta \rho$  acting on the velocity slot; the trial carries  $\mathbf{v} = -i\omega \mathbf{u}$  there; and  $\mathbf{J}_6$  pulls the test function's velocity slot into the traction rows, supplying a second  $-i\omega$ . Hence  $i\omega \cdot (-i\omega) \cdot (-i\omega) = i\omega^3$ . The displacement-basis answer  $i\omega \Delta \rho V$  is wrong by  $\omega^2$ . This is the displacement-versus-flux confusion in a new place, and the check asserts the correct factor *and* asserts the incorrect one is incorrect, so it cannot drift back.

## 7 What the first-order coupling actually buys

### 7.1 The trial spaces coincide

It is tempting to attribute any difference from the Navier formulation to the nonlinearity of §4.3, or to a different ansatz. Neither is right. Inside the scatterer the moduli are constant, so  $\boldsymbol{\tau}_3 = \mathbf{C}_{33}^{\text{tot}} \boldsymbol{\varepsilon}_{i3}$

is a constant *invertible* map. A uniform strain gives a uniform  $\boldsymbol{\tau}_3$  and an affine  $\mathbf{v}$ , and conversely: “ $\boldsymbol{\varepsilon}$  uniform” and “ $\boldsymbol{\tau}_3$  uniform plus in-plane  $\boldsymbol{\varepsilon}$  uniform” describe the same nine-parameter family.

The two formulations are therefore **different projections of the same residual onto the same trial space**. The difference lives in the test conditions.

## 7.2 The density channel cannot differ

With  $\Delta\lambda = \Delta\mu = 0$ ,  $\Delta\mathcal{A}$  is exactly  $i\omega\Delta\rho$  in rows 1–3, columns 4–6, and reaches *no* part of the stress-glut slot: a density contrast is a pure body force. The propagator’s body-force velocity block is  $-i\omega\mathbf{G}$ , so both routes return

$$A_u = \frac{1}{1 - \omega^2\Delta\rho\Gamma_0}, \quad \Gamma_0 = \int_V \mathbf{G} dV, \quad (23)$$

identically. The modulus channels are the only place the two can part.

## 7.3 The units the two couplings are written in

The statement just made is about the *amplification factor*, which is a ratio, and ratios are where this comparison is safe. The couplings themselves are not in the same units, and it is worth saying so plainly because every other check in this note is also a ratio, and none of them can see it.

The  $\mathbf{J}_6$  pairing tests against the field  $q$ , whose velocity slots carry  $v = -i\omega u$ . The Navier weak form, and this package’s local  $T$ -matrix, test against the *displacement*. So

$$\Delta\mathbf{C}_{\text{eff}} = i\omega \times [\text{the same object in displacement units}], \quad \Delta\mathbf{C}_{\text{eff}}|_{\text{rigid}} = i\omega^3\Delta\rho V \text{ against } \omega^2\Delta\rho V. \quad (24)$$

This is a change of units and not of physics. It cancels out of every amplification factor, every channel ratio and every Born limit expressed as a ratio; it does *not* cancel the moment  $\Delta\mathbf{C}_{\text{eff}}$  is combined with an object built in displacement units. The  $i\omega$  appearing in (70) is this same factor seen from the other side.

The check that establishes it is necessarily absolute rather than relative: after dividing by  $i\omega$ , the residual  $\|\Delta\mathbf{C}_{\text{eff}}/i\omega - \Delta\mathbf{C}_{\text{Nav}}\|/\|\Delta\mathbf{C}_{\text{Nav}}\|$  falls as  $4.19 \times 10^{-2}$ ,  $4.19 \times 10^{-3}$ ,  $4.19 \times 10^{-4}$ ,  $4.19 \times 10^{-5}$  under decade reductions of the contrast — the *first* power, so the residual is the genuine  $O(\Delta c^2)$  difference of §7 against an  $O(\Delta c)$  coupling, and not a second, smaller units error hiding beneath the first.

## 7.4 Born agreement, and an $O(\Delta c^2)$ difference

Reading a proportionality constant off the rigid block alone gives  $i\omega$ , and testing it on the remaining eighty entries **refutes** the hypothesis  $\Delta\mathbf{C}_{\text{eff}} = i\omega\Delta\mathbf{C}_{\text{Navier}}$ . The two couplings agree exactly at  $O(\Delta c^0)$  and  $O(\Delta c^1)$  — identical Born limits, so nothing is broken — and differ from second order.

**Where.** Eleven of eighty-one entries differ. With the parameter order  $(c_1, c_2, c_3, e_{11}, e_{22}, e_{33}, e_{12}, e_{13}, e_{23})$  the nonzero block is  $\{e_{11}, e_{22}, e_{33}\}^2$  together with  $(e_{13}, e_{13})$  and  $(e_{23}, e_{23})$ .

That is exactly the set of strains  $\tau_3$  sees:

$$\tau_{13} = 2\mu e_{13}, \quad \tau_{23} = 2\mu e_{23}, \quad \tau_{33} = \lambda(e_{11} + e_{22} + e_{33}) + 2\mu e_{33}. \quad (25)$$

The in-plane shear  $e_{12}$  is **untouched**, as it must be:  $\tau_{12}$  was eliminated, not resummed. This was a falsifiable prediction about which entries may be nonzero, and it held. A representative entry is

$$(\Delta \mathbf{C}_{\text{eff}} - i\omega \Delta \mathbf{C}_{\text{Navier}})_{e_{11}e_{11}} = -\frac{i\omega V \Delta \lambda^2}{\lambda + 2\mu}, \quad (26)$$

pure second order, divided by  $K_c$ : the remainder of the geometric series for  $1/(K_c + \Delta K_c)$ .

## 7.5 Half of it is a broken symmetry

An isotropic cube in an isotropic whole space has an  $O_h$ -symmetric  $T$ -matrix; no direction is special. But the first-order system is a recursion in *depth*: it eliminated  $\tau_1, \tau_2$  and kept  $\tau_3$ , so  $x_3$  is preferred by construction. The tell is already above —  $(e_{13}, e_{13})$  and  $(e_{23}, e_{23})$  differ while  $(e_{12}, e_{12})$  does not, and those three are  $O_h$ -equivalent.

Applying the cyclic element  $x \rightarrow y \rightarrow z \rightarrow x$  (positive controls:  $P^3 = \mathbf{I}$ , and the Navier coupling *is* invariant):

| contrast (GPa)              | invariant             | breaking              | ratio  |
|-----------------------------|-----------------------|-----------------------|--------|
| pure shear, $\Delta\mu = 1$ | $2.62 \times 10^{-2}$ | $1.94 \times 10^{-2}$ | 1.3520 |
| moderate, (2, 1, 0.1)       | $5.17 \times 10^{-4}$ | $2.93 \times 10^{-4}$ | 1.7669 |
| strong, (20, 10, 0.5)       | $1.05 \times 10^{-2}$ | $5.92 \times 10^{-3}$ | 1.7669 |

**Table 1.** Decomposition of  $\Delta \mathbf{C}_{\text{eff}} - i\omega \Delta \mathbf{C}_{\text{Navier}}$ , relative to  $\|\Delta \mathbf{C}_{\text{eff}}\|$ , at  $\lambda = 17.5$ ,  $\mu = 22.5$  GPa. The two mixed rows share a contrast *direction* ( $\Delta\lambda:\Delta\mu = 2:1$ ) at different magnitudes and return the identical ratio, so the split is a function of contrast direction alone — as a leading-order  $O(\Delta c^2)$  split must be.

So the difference is neither pure artefact nor pure physics. It carries both a genuine  $O_h$ -invariant resummation the Navier form lacks, and a spurious preference for  $x_3$ .

**A caveat on this test.** The group used is the cyclic subgroup of order 3, not all 48 elements of  $O_h$ . Invariance under it is necessary for  $O_h$  invariance, not sufficient.

## 7.6 What follows, and what does not

Two readings are available and the measurement above does not separate them.

*Against the route:* an approximation that breaks a symmetry the exact answer possesses has introduced an error that no propagator can remove.

*For the route:* the correction is not scattered. It lands exactly on the traction channel, treating  $\mathbf{C}_{33}$  exactly instead of leaving it to the Neumann series, and leaves the in-plane channel alone. Vertical  $P$  incidence excites  $e_{33}$ ; vertical  $S$  incidence excites  $e_{13}$  or  $e_{23}$ ; both are corrected. The in-plane shear, which the side-scattered field carries, is unchanged. If the response is dominated by reflection and transmission, the asymmetry is “z treated better” rather than “z treated worse”.

**Consequence for any comparison.** Do not score the raw coupling against an independent arbiter in a homogeneous background. Three effects are in play at comparable magnitude — the Galerkin-versus-collocation projection, the  $O_h$ -invariant resummation, and the artefact — and a single number separates none of them. Any comparison must be channel-resolved.

Whether the resummation is an *improvement* is not established here. The local inversion  $1/(K_c + \Delta K_c)$  is exact; whether that makes the closed  $T$ -matrix better depends on how the ansatz error and the series truncation interact, and those do not simply add.

## 8 The propagator, and the bridge to it

### 8.1 It already exists

Everything downstream needs the Green’s matrix of (1). It does not need to be built: the stratified  $6 \times 6$  Green’s function and its validated wrapper already return it, in the basis  $(u_z, u_x, u_y, T_{zz}, T_{xz}, T_{yz})$  — which is exactly the basis the thesis writes  $\mathbf{A}$  in. That the two are the same object had, however, been *assumed*.

### 8.2 The bridge identity

Away from the source plane the Green’s matrix solves the homogeneous system in its receiver coordinate, so for two receiver depths in one uniform layer

$$\mathbf{\Gamma}(z_2, z') = \exp(\mathbf{A}(z_2 - z_1)) \mathbf{\Gamma}(z_1, z'), \quad (27)$$

with no free parameters. The identity is *local*: it holds however complicated the medium between the source and  $z_1$  is, because the ODE knows only the layer the receiver sits in. The gate therefore places the source above a fast slab and the receivers below it, so transmission, reflection, conversion and interbed multiples are all present in  $\mathbf{\Gamma}$  and (27) must still hold exactly.

### 8.3 Reporting residuals against the arithmetic floor

A bare relative residual is not comparable across incidence angles. Beyond the critical slowness the propagator carries both a growing and a decaying evanescent branch, so  $\exp(\mathbf{A} dz)$  has dynamic range  $\sim e^{|k|dz}$  — about  $1.6 \times 10^7$  at the widest angle used. Forming  $\exp(\mathbf{A} dz) \mathbf{\Gamma}(z_1)$  then amplifies the rounding already present by  $\|\exp(\mathbf{A} dz)\|$ , and no correct implementation can beat

$$\text{floor} = \epsilon \|\exp(\mathbf{A} dz)\| \frac{\|\mathbf{\Gamma}(z_1)\|}{\|\mathbf{\Gamma}(z_2)\|}. \quad (28)$$

Residuals below are quoted in units of that floor, so  $O(1)$  means the identity holds to the last bit available at that conditioning.

| $k_x$                                         | $k_y$ | $ k /\omega$ | raw                   | corrected             |
|-----------------------------------------------|-------|--------------|-----------------------|-----------------------|
| 2.00                                          | 1.30  | 0.0759       | $7.68 \times 10^{13}$ | $1.26 \times 10^{-2}$ |
| 0.40                                          | 0.05  | 0.0128       | $2.95 \times 10^{14}$ | $2.43 \times 10^{-2}$ |
| 14.00                                         | 9.00  | 0.5298       | $4.31 \times 10^{13}$ | $9.37 \times 10^{-3}$ |
| control, $\zeta \leftrightarrow \chi$ swapped |       |              | $7.68 \times 10^{13}$ | $1.36 \times 10^{12}$ |

**Table 2.** Residual of (27) in units of the arithmetic floor.  $\omega = 31.4159 \text{ rad s}^{-1}$ ,  $dz = 1 \text{ km}$ . The third row is evanescent for both branches.

**The corrected object is the physical state vector of  $\partial_3 \mathbf{q} = \mathcal{A} \mathbf{q}$ ;** the raw one misses by fourteen orders of magnitude.

The negative control swaps  $\zeta$  and  $\chi$  — two in-plane stiffnesses that differ by exactly  $2\mu$ , a plausible transcription slip that no symmetry check would catch. It is rejected by fourteen orders.

## 9 The single-site $T$ -matrix

This is the object the construction was for. On the nine-parameter affine trial space,

$$\mathbf{T} = \Delta \mathbf{C}_{\text{eff}} (\mathbf{I}_9 - \mathcal{G} \Delta \mathbf{C}_{\text{eff}})^{-1}, \quad (29)$$

with  $\Delta \mathbf{C}_{\text{eff}}$  the projected coupling of §6 and  $\mathcal{G}$  the self propagator moment of the cube.

### 9.1 What is new and what is borrowed

New is  $\Delta \mathbf{C}_{\text{eff}}$ . Borrowed is  $\mathcal{G}$ : the first-order route changes the *coupling*, not the Green's function. The two moments it needs are

$$\mathbf{G}_{ij} = \int_V G_{ij}^0 dV, \quad M_{in,pk} = \int_V \partial'_p \partial'_k G_{in}^0 dV. \quad (30)$$

These are *derived* in §9.2 below, not taken from a table. In the parameter basis

$$\mathcal{G} = \frac{1}{V} \begin{pmatrix} \Gamma_0 \mathbf{I}_3 & \mathbf{0} \\ \mathbf{0} & -\mathbf{S} \mathbf{W}^{-1} \end{pmatrix}, \quad \mathbf{W} = \text{diag}(1, 1, 1, 2, 2, 2), \quad (31)$$

where  $\mathbf{W}$  undoes the engineering doubling of the shear slots.

Note that (31) is asserted, not derived — which is exactly why §9.5 exists.

### 9.2 The self moments, derived

The integrand of  $M$  goes as  $1/r^3$ , which is not absolutely integrable in three dimensions. The customary remedy is to quote a tabulated Eshelby tensor and add a delta-function term at  $r = 0$  by hand; omit the delta and the depolarisation comes out with the *wrong sign*. Nothing below is tabulated and nothing is added by hand.

### 9.2.1 The static Green's tensor

In the static limit the elastodynamic tensor reduces to Kelvin's,

$$G_{in}^0(\mathbf{r}) = \frac{a_0 \delta_{in}}{r} + \frac{b_0 x_i x_n}{r^3}, \quad a_0 = \frac{1}{8\pi} \left( \frac{1}{\mu} + \frac{1}{\lambda + 2\mu} \right), \quad b_0 = \frac{1}{8\pi} \left( \frac{1}{\mu} - \frac{1}{\lambda + 2\mu} \right), \quad (32)$$

the coefficients being fixed by requiring that (32) satisfy the Navier operator away from the origin. An equivalent form makes the distributional content visible, since  $\partial_i \partial_n r = \delta_{in}/r - x_i x_n/r^3$ :

$$G_{in}^0 = \frac{1}{4\pi\mu} \left[ \frac{\delta_{in}}{r} + h_c \partial_i \partial_n r \right], \quad h_c = -\frac{\lambda + \mu}{2(\lambda + 2\mu)}, \quad (33)$$

and  $a_0 = (1 + h_c)/4\pi\mu$ ,  $b_0 = -h_c/4\pi\mu$  identically.

### 9.2.2 The moment as a distribution, and why no delta is inserted

Write the moment as a pairing against the test function  $w \mathbf{1}_V$  rather than as an integral of a function. With  $T = \partial_{\text{rest}} r^m$  and  $\partial_q \mathbf{1}_V = -n_q \delta_{\partial V}$ ,

$$\langle \partial_q T, w \mathbf{1}_V \rangle = \oint_{\partial V} n_q w T dA - \langle T, (\partial_q w) \mathbf{1}_V \rangle. \quad (34)$$

That is the entire construction. Two features make it exact rather than merely plausible:

- $T$  is only ever evaluated *on the boundary*, at distance  $\Delta/2$  from the origin, where every derivative of  $r^m$  is smooth and bounded. The singular pointwise structure of  $\partial^D r^m$  at the origin is never formed, so there is no principal value to take and no delta to insert.
- the recursion strictly lowers both the derivative count and the weight degree, terminating at  $\int_V x_w r^m dV$ , an ordinary convergent integral for  $m \geq -1$ .

The delta term is therefore neither added nor droppable — it is already inside the surface integral. The sharp test is the Laplacian sum rule  $\sum_p E[-1; \{p, p\}] = -4\pi$ , which a spherical excision returns as 0.

### 9.2.3 The face integrals, and one of them by pure geometry

On the face  $x_3 = \Delta/2$ , scaling the in-face coordinates by  $\Delta/2$  reduces every surface term to one of three integrals over  $[-1, 1]^2$ :

$$j_1 = \iint \frac{du dv}{(1 + u^2 + v^2)^{3/2}} = \frac{2\pi}{3}, \quad j_2 = \iint \frac{u^2 du dv}{(1 + u^2 + v^2)^{5/2}} = -\frac{2(\sqrt{3} - \pi)}{9}, \quad k_1 = \iint \frac{du dv}{(1 + u^2 + v^2)^{5/2}} = \frac{2(2\sqrt{3} + \pi)}{9}. \quad (35)$$

All three are evaluated in closed form. The first needs no integration at all:  $j_1$  is the solid angle subtended by one face at the cube's centre, and the six faces subtend  $4\pi$ , so  $j_1 = 4\pi/6 = 2\pi/3$ . That identity is a check on the scaling and orientation of (35), obtained independently of any quadrature.

The zeroth moment needs one further integral, which also has a closed form:

$$\int_{[-1,1]^3} \frac{dV}{r} = -2(\pi + \log 64 - 12 \log(1 + \sqrt{3})) = \frac{4}{3} \log(70226 + 40545\sqrt{3}) - 2\pi = 9.520309455918214 \dots \quad (36)$$

The two closed forms are independent evaluations of the same constant and agree to the precision carried; the second is the one this package uses elsewhere.

#### 9.2.4 The closed forms

Assembling (32) through the recursion (34) with the face integrals (35) gives, with

$$M_{in,pk} = A \delta_{in} \delta_{pk} + B (\delta_{ip} \delta_{nk} + \delta_{ik} \delta_{np}) + C E_{inpk}, \quad E_{inpk} = \begin{cases} 1 & i = n = p = k \\ 0 & \text{otherwise,} \end{cases} \quad (37)$$

the third structure being the one a cubic — as opposed to isotropic — geometry admits,

$$B = \frac{\sqrt{3}(\lambda + \mu)}{6\pi\mu(\lambda + 2\mu)}, \quad (38)$$

$$A = B - \frac{1}{3\mu}, \quad (39)$$

$$C = -\left(5 - \frac{2\pi}{\sqrt{3}}\right)B, \quad (40)$$

$$G_{ij} = \delta_{ij} \left(a_0 + \frac{b_0}{3}\right) \left(\frac{\Delta}{2}\right)^2 \int_{[-1,1]^3} \frac{dV}{r}. \quad (41)$$

$A$ ,  $B$  and  $C$  are scale-free;  $G$  carries  $\Delta^2$ . Note that  $A$ ,  $B$  and  $C$  are *not* three independent constants — the whole rank-four moment is carried by  $B$  together with the elementary  $1/3\mu$ , which is a structural statement the tabulated presentation obscures.  $G_{ij}$  must be a multiple of  $\delta_{ij}$ , because a rank-two cubic invariant has no other form: the cube cannot distinguish its own axes at that order.

#### 9.2.5 The check that certifies the delta content

From (33) with  $\nabla^2(1/r) = -4\pi\delta^3$  and  $\nabla^2(\partial_i \partial_n r) = 2\partial_i \partial_n(1/r)$ , whose cube integral is  $-\frac{4\pi}{3}\delta_{in}$ ,

$$\sum_p M_{in,pp} = -\frac{\delta_{in}}{\mu} \left(1 + \frac{2h_c}{3}\right) = -\delta_{in} \frac{2\lambda + 5\mu}{3\mu(\lambda + 2\mu)}, \quad (42)$$

derived without reference to the recursion at all. The assembled moment satisfies (42) identically, and the export refuses to write its output otherwise. A spherically excised moment returns zero on the left-hand side, so (42) is exactly the statement that the  $r = 0$  content is present and correctly weighted.

### 9.3 The tabulated route agrees

This package also carries the conventional tabulated Eshelby route, in which  $j_1$ ,  $j_2$  and  $k_1$  are quoted rather than integrated. Evaluated alongside (38)–(41) at  $\lambda = 17.5$ ,  $\mu = 22.5$  GPa,  $\Delta = 1$ :

| moment | integrated                        | tabulated                         | relative              |
|--------|-----------------------------------|-----------------------------------|-----------------------|
| $A$    | $-1.220110745874 \times 10^{-11}$ | $-1.220110745874 \times 10^{-11}$ | 0                     |
| $B$    | $2.613707356074 \times 10^{-12}$  | $2.613707356074 \times 10^{-12}$  | $3.1 \times 10^{-16}$ |
| $C$    | $-3.587055298887 \times 10^{-12}$ | $-3.587055298887 \times 10^{-12}$ | $2.3 \times 10^{-16}$ |
| $G$    | $6.622004020056 \times 10^{-12}$  | $6.622004020056 \times 10^{-12}$  | 0                     |

**Table 3.** The table reproduces the integration to machine precision. The direction of that statement matters: the integration *verifies* the table; the assembly reads the integrated values.

The dynamic (smooth radiation) part carries no table — it is already exact polynomial integration — and is taken from the existing route with its tabulated static piece removed and the integrated one substituted.

### 9.4 Two implementations of the coupling

$\Delta\mathbf{C}_{\text{eff}}$  is rebuilt in Python directly from  $\Delta\mathcal{A}$ , independently of the symbolic derivation, and compared against the values that derivation exports. Worst relative difference over three contrasts:  $1.4 \times 10^{-15}$  for  $\Delta\mathbf{C}_{\text{eff}}$  and exactly zero for the Navier coupling. Two implementations agreeing is the evidence standard; one agreeing with itself is not.

### 9.5 Calibration: every normalisation fixed against existing code

Run the *same* assembly with the ordinary Navier coupling in place of  $\Delta\mathbf{C}_{\text{eff}}$ . It must reproduce this package’s validated four amplification factors. That pins the volume factors, the sign of  $\mathcal{G}$  and the weight  $\mathbf{W}$  — the *bookkeeping* of (31) — with no freedom left to tune. It does not certify the moments themselves; that is what the integration of the previous subsection is for. Against  $\alpha = 5000$ ,  $\beta = 3000$ ,  $\rho = 2500$  (SI),  $\omega = 60$ ,  $a = 0.5$ :

| case                              | factor                                               | assembled vs existing (rel.)                         |
|-----------------------------------|------------------------------------------------------|------------------------------------------------------|
| pure shear ( $\Delta\mu = 1$ GPa) | $A_u, A_\theta, A_e^{\text{diag}}, A_e^{\text{off}}$ | 0, $2.2 \times 10^{-16}$ , 0, 0                      |
| moderate (2, 1, 100)              | same four                                            | 0, $2.3 \times 10^{-16}$ , $1.1 \times 10^{-16}$ , 0 |

**Table 4.** The assembly reproduces the existing amplification factors to machine precision. Only after this is the coupling swapped.

**One term had to be isolated to get there.** The weak form’s  $\int_V \Delta\rho \omega^2 \boldsymbol{\phi} \cdot \mathbf{u} dV$  contributes to the *strain*–strain block as well as the rigid one, because an affine trial displacement is not constant:  $\int_V x_a x_b dV = \delta_{ab} V a^2 / 3$ . This is a finite-size dynamic coupling between the density contrast and the strain channels, and the factorised amplification factors do not carry it — there  $A_u$  and the strain amplifications are independent. It shifts  $A_\theta$  by  $1.5 \times 10^{-7}$  at the moderate contrast and

vanishes identically with  $\Delta\rho$ . It is a small correction that the factorisation drops, not a defect in either object, and it is reported rather than absorbed.

## 9.6 The first-order $T$ , channel by channel

A single norm would average a correction that lives in a few channels against others that do not move. The comparison is therefore resolved by channel.

| case       | channel                       | relative difference from the Navier route |
|------------|-------------------------------|-------------------------------------------|
| pure shear | $\theta$ (bulk)               | $1.13 \times 10^{-4}$                     |
|            | $e_{\text{diag}}$             | $1.1 \times 10^{-16}$                     |
|            | $e_{12}$ (in-plane shear)     | 0                                         |
|            | $e_{13}$ (out-of-plane shear) | $8.36 \times 10^{-4}$                     |
|            | $u$ (density)                 | 0                                         |
| moderate   | $\theta$ (bulk)               | $1.74 \times 10^{-3}$                     |
|            | $e_{\text{diag}}$             | $1.1 \times 10^{-16}$                     |
|            | $e_{12}$                      | 0                                         |
|            | $e_{13}$                      | $8.36 \times 10^{-4}$                     |
|            | $u$                           | 0                                         |

**Table 5.** Amplification  $(\mathbf{I} - \mathcal{G}\Delta\mathbf{C})^{-1}$  along each channel. The density channel is identical in both routes, as §7 requires; the in-plane shear is untouched; the bulk and out-of-plane shear channels move.

## 9.7 The broken symmetry, now visible in $T$ itself

$e_{12}$  and  $e_{13}$  are equivalent under  $O_h$ . For an isotropic cube in an isotropic background their amplifications *must* be equal. They are, in the Navier route, to exactly zero. In the first-order route they split:

|                                         | Navier | first-order           |
|-----------------------------------------|--------|-----------------------|
| $ A(e_{12}) - A(e_{13})  /  A(e_{12}) $ | 0      | $8.36 \times 10^{-4}$ |

**Table 6.** The  $O_h$  violation of §7.5, measured on the  $T$ -matrix rather than on the coupling. Identical at both contrasts, because both carry the same  $\Delta\mu$  and the shear channels see only that.

This is the sharpest form of the result. It is not a norm over a matrix: it is two numbers that a symmetry argument says must agree, and do not.

**What the split does and does not license.** It is tempting to read this as the first-order  $T$  failing an external test while the ordinary route passes. That reading claims more than the check can deliver, and the two halves must be kept apart.

*What it gives.* An **arbiter-free lower bound on the error of the first-order  $T$** . If  $A(e_{12})$  and  $A(e_{13})$  must be equal and differ by  $\delta$ , then whichever lies closer to the truth, the larger of the two errors is at least  $\delta/2 \approx 4.2 \times 10^{-4}$ . An absolute error estimate obtained from a symmetry alone,

with nothing external to compare against, is worth having — and it is the absolute companion this project requires alongside any ratio test.

*What it does not give.* Any bound whatever on the ordinary route’s error. A symmetric approximation can be symmetrically wrong, missing the truth in *both* channels by more than  $8.4 \times 10^{-4}$  while satisfying the symmetry exactly. Symmetry compliance is not accuracy. Setting a bound on one formulation against no bound on the other is not a comparison, and no ordering of the two follows from this table.

*What is nonetheless attributable.* The breaking is not manufactured by the projection mismatch of §9:  $\mathcal{G}$  is itself  $O_h$ -symmetric, and §7.5 measured the coupling’s asymmetry symbolically with no propagator present at all. The asymmetry is a property of the coupling, and it survives into  $\mathbf{T}$ .

**The channel set has to include  $e_{13}$ .** The coupling leaves  $e_{12}$  untouched by construction, so a comparison spanning only  $\theta$ ,  $e_{\text{diag}}$  and  $e_{12}$  reports that the shear channels agree — true, and uninformative. A channel-resolved comparison is only as good as the channels chosen, and the  $O_h$ -equivalent pair is the one that carries the information.

## 9.8 What this does and does not settle

The Born limits agree: at a contrast of  $10^{-6}$  the two couplings differ by  $4.2 \times 10^{-8}$ , i.e. at  $O(\Delta c)$  as they must.

**The caveat that no number above reveals.** The propagator moment used is the *single* (collocation) one, while the  $\mathbf{J}_6$  pairing makes the scheme Galerkin. The consistent pairing needs the *double* moment. This is the same class of mismatch the package has already measured as a genuine error elsewhere, so the table above should be read as the first-order coupling evaluated through a collocation propagator, not as the fully consistent first-order  $T$ . No winner is declared, and none should be read in.

## 10 The propagator moment from the first-order system

§9 took  $\mathcal{G}$  from the ordinary Kelvin tensor, so the  $T$ -matrix assembled there is first-order in its *coupling* only. This section replaces that half. The plan is `docs/plans/2026-09-18-first-order-propagator-moment.md`; what follows is written as each step lands, and says plainly where the front is.

The object is the sandwiched form

$$[\Delta \mathbf{C}_{\text{eff}} \mathcal{G} \Delta \mathbf{C}_{\text{eff}}]_{kl} = \langle \mathbf{J}_6 \psi_k \mid \Delta \mathcal{A} \Gamma \Delta \mathcal{A} \mid \psi_l \rangle. \quad (43)$$

Plain Galerkin on  $\mathbf{q}$  is not available: its mass matrix  $\langle \mathbf{J}_6 \psi_k, \psi_l \rangle$  is antisymmetric, and an antisymmetric matrix in *odd* dimension — here nine — is singular.

### 10.1 Step 1: $\Gamma$ , built from $\mathcal{A}$ alone

For  $\partial_3 \mathbf{q} = \mathbf{A}\mathbf{q} + \mathbf{d}\delta(z - z')$  the radiating Green's matrix is

$$\Gamma(z, z') = \begin{cases} +\mathbf{P}_d e^{\mathbf{A}(z-z')}, & z > z', \\ -\mathbf{P}_u e^{\mathbf{A}(z-z')}, & z < z', \end{cases} \quad (44)$$

with  $\mathbf{P}_d, \mathbf{P}_u$  the spectral projectors onto the down- and up-going eigenspaces, so that  $\Gamma$  decays away from the source in both directions and jumps by  $\mathbf{P}_d + \mathbf{P}_u = \mathbf{I}$  across it.

**Projectors, not eigenvectors.** The thesis writes the eigen-matrix column by column with normalisation factors  $\epsilon_P, \epsilon_S, \epsilon_H$  left symbolic. Those factors are fixed by a flux normalisation that degenerates at normal incidence, where the horizontal direction no longer distinguishes the *SV* and *SH* polarisations — exactly the  $k_{\parallel} \rightarrow 0$  difficulty. The spectral projector

$$\mathbf{P}_d = \mathbf{V}_d (\mathbf{W}_d^t \mathbf{V}_d)^{-1} \mathbf{W}_d^t \quad (45)$$

is invariant under any rescaling of the columns of  $\mathbf{V}_d$ , so building  $\Gamma$  from projectors removes the ambiguity instead of carrying it. That invariance is verified directly, by rescaling every eigenvector by a random complex factor and confirming  $\mathbf{P}_d$  is unchanged.

**The matrix exponential must never be formed.** Equation (44) is a statement, not an algorithm. Evaluating it as written — projecting *after* exponentiating — builds the growing branch and then cancels it. At a separation of 640 km the growing entries reach  $e^{+28} \approx 10^{12}$ , so the cancellation leaves  $10^{-4}$  of absolute noise on a true answer of  $3 \times 10^{-8}$ . This is the Thomson–Haskell instability, the reason invariant imbedding exists, reappearing inside our own construction; it announces itself as a decay rate that turns *negative* at large separation. Applying the exponential to the projected eigenvalues never forms the growing branch at all,

$$\Gamma(z - z') = \pm \sum_{i \in \text{down/up}} e^{s_i(z-z')} \mathbf{v}_i \mathbf{w}_i^t, \quad \mathbf{w}_i \cdot \mathbf{v}_j = \delta_{ij}, \quad (46)$$

and is what is used. The moment integral needs only  $|z - z'| \leq 2a$ , where the instability cannot bite; a layered background with thick layers is where it would.

**$\omega$  is complex from the outset.** The vertical wavenumbers  $k_{z,c} = \sqrt{k_c^2 - k^2}$  vanish on the circles  $|k| = k_P$  and  $|k| = k_S$ , which the lateral integral must cross, and  $k_z$  sits in a denominator. At real  $\omega$  a quadrature straddles an integrable  $1/\sqrt{\cdot}$  singularity and does not converge; the directional-sweep work measured  $1.9 \times 10^{-1} \rightarrow 1.1 \times 10^{-1}$  over an eightfold refinement, non-monotone. With damping 0.02 the smallest  $|s|$  on the *P* circle is 0.240 rather than zero.

**Checks.** Thirteen, all passing: the spectrum is  $\{\pm ik_{z,P}\}$  and  $\{\pm ik_{z,S}\}$  twice;  $\mathbf{P}_d + \mathbf{P}_u = \mathbf{I}$ , both idempotent, both commuting with  $\mathbf{A}$ ;  $\mathbf{P}_d$  invariant under rescaling;  $\partial_z \Gamma = \mathbf{A}\Gamma$  on both sides; the jump exactly  $\mathbf{I}$ ; and the decay rate converging to the slowest eigenvalue,  $0.034990 \rightarrow 0.033901 \rightarrow$

0.032857  $\rightarrow$  0.032436 against 0.032393 — a rate, not a bar, because at any one finite separation the faster mode still contributes.

## 10.2 Step 2: the vertical integral, in closed form

The vertical integral is exact, but it is *not* an integral over the square. The two halves of  $\Gamma$  carry different matrices, so the double integral splits over the two triangles and they cannot be recombined: treating the kernel as a function of  $|z - z'|$  alone would symmetrise the propagator and destroy the up/down distinction the first-order system exists to carry.

With affine trial fields the weights are 1 and  $z$ , so four triangular integrals are needed,

$$T_{mn}(q) = \int_{-a}^a dz \int_{-a}^z dz' z^m z'^n e^{-q(z-z')}, \quad (47)$$

the lower triangle being  $T_{nm}$  by relabelling. All four are elementary:

$$T_{00} = \frac{2aq - 1 + e^{-2aq}}{q^2}, \quad (48)$$

$$T_{10} = -T_{01} = \frac{aq - 1 + (1 + aq)e^{-2aq}}{q^3}, \quad (49)$$

$$T_{11} = \frac{3 - 3(1 + aq)^2 e^{-2aq} + a^2 q^2 (2aq - 3)}{3q^4}. \quad (50)$$

Summing the two triangles with a common matrix recovers the familiar square integral  $2(2aq - 1 + e^{-2aq})/q^2$ , and the odd weight cancels between them by parity — both checked, as is the large- $q$  behaviour  $T_{00} \sim 2a/q$ , each triangle carrying half of the square's  $4a/q$ .

The decisive check of this step is not the closed forms themselves but the *assembly*: the vertical integration of  $\Gamma$  performed through (48)–(50), matrices and all, reproduces a brute-force double quadrature of  $\Gamma$  to  $2.5 \times 10^{-9}$  relative.

### 10.2.1 The same integrals as a series, and why they are needed

(48)–(50) are exact and, evaluated as written, lose most of their digits as  $q \rightarrow 0$ . Expanding (50) gives

$$3 - 3(1 + x)^2 e^{-2x} + x^2(2x - 3) = \frac{4}{5}x^5 - \frac{2}{3}x^6 + O(x^7), \quad x = aq, \quad (51)$$

so a numerator of order  $x^5$  is assembled from terms of order unity: at  $x = 0.025$  about eight digits are gone. That is not a corner of the parameter space — it is the whole propagating window  $k < \omega/\beta$ , which is precisely where the radiation lives.

The remedy is the series. Substituting  $u = z - z'$  in (47) and expanding the exponential,

$$T_{mn}(q) = \sum_{j \geq 0} \frac{(-q)^j}{j!} \int_0^{2a} u^j J_{mn}(u) du, \quad J_{mn}(u) = \int_{-a}^{a-u} (z' + u)^m z'^n dz', \quad (52)$$

and every  $u$ -integral is elementary. With  $b = 2a$ ,

$$J_{00} = b - u, \quad J_{10} = -J_{01} = \frac{b}{2}u - \frac{u^2}{2}, \quad J_{11} = \frac{b^3}{12} - \frac{b^2}{4}u + \frac{u^3}{6}, \quad (53)$$

giving coefficients of  $(-qb)^j/j!$  equal to

$$\frac{b^2}{(j+1)(j+2)}, \quad \frac{b^3}{2(j+2)(j+3)}, \quad b^4 \left[ \frac{1}{12(j+1)} - \frac{1}{4(j+2)} + \frac{1}{6(j+4)} \right] \quad (54)$$

for  $(0,0)$ ,  $(1,0)$  and  $(1,1)$  respectively. Two structural checks: at  $q = 0$  the first gives  $2a^2$  and the second  $2a^3/3$ , which are  $\iint_{z>z'} 1$  and  $\iint_{z>z'} z$ ; the third gives *zero*, and indeed  $\iint_{z>z'} zz' = 0$  by parity. Its  $j = 1$  term is  $\frac{4}{15}a^5q$ , which is exactly what (50) tends to,  $\frac{4}{5}a^5q^5/3q^4$ . The series is entire, so it is valid for the imaginary  $q$  that propagating modes carry; it is used below  $|2aq| = 3$ , where it is also the accurate branch.

### 10.3 Step 3: the lateral integral

Count powers as  $|k| \rightarrow \infty$  for the *untransferred* form, in which  $\partial_\alpha \rightarrow ik_\alpha$ : the box form factor gives  $k^{-2}$ , (48) gives  $q^{-1} \sim |k|^{-1}$ , the two contrast operators give  $k^2$ , and the measure gives  $k$ . The strain channel does not converge. (The transferred form of §10.3.6 does; that is one of the things the augmentation buys, and it is taken up there.)

The divergence is not an artefact to be trimmed. A delta function in real space is a constant in  $k$ -space, so the non-decaying tail of the integrand *is* the Eshelby term of §9.2. A bare cutoff discards precisely the piece whose omission reverses the sign of the depolarisation. Hence split-and-subtract,

$$I(\mathbf{k}) = [I(\mathbf{k}) - I_\infty(\mathbf{k})] + I_\infty(\mathbf{k}), \quad (55)$$

with the bracket integrated numerically — convergent, and its convergence measured as a rate — and  $\int I_\infty d^2k$  evaluated in closed form over the whole plane, where the delta content resides. The asymptote is taken of the  $z$ -integrated kernel (48)–(50), with the box form factors kept exact: they are what makes the remainder converge, so expanding them would defeat the subtraction.

#### 10.3.1 The scaling that makes the expansion uniform

$\mathbf{A}$  is not uniformly  $O(k)$ . Rows 1–3 carry  $k$ , but rows 5–6 carry  $k^2$ , because the basis mixes displacement with traction and the two have different dimensions; expanding  $\mathbf{A}$  in  $1/k$  as it stands mixes orders between blocks. The traction of an evanescent field scales as  $\mu ku$ , so the natural balance is  $\mathbf{q} \rightarrow (\mathbf{u}, \boldsymbol{\tau}_3/k)$ . Under  $\mathbf{S} = \text{diag}(1, 1, 1, k, k, k)$ ,

$$\tilde{\mathbf{A}} = \mathbf{S}^{-1} \mathbf{A} \mathbf{S} \quad (56)$$

is uniformly  $O(k)$ : the upper-right block gains a factor  $k$ , the lower-left loses one, and the diagonal blocks are untouched. Measured,  $\|\mathbf{A}\|/k$  grows as  $5.3 \times 10^3 \rightarrow 5.3 \times 10^5$  over two decades in  $k$  while  $\|\tilde{\mathbf{A}}\|/k$  settles at  $53.0100 \rightarrow 53.0101$ . Being a similarity transform it changes no eigenvalue, and  $|s|/k \rightarrow 1$  for all six.

The scaling is also what makes the decomposition numerically usable. On the unscaled matrix

the down- and up-going eigenvectors become nearly parallel as  $k$  grows — the evanescent limit degenerates — the eigen-decomposition is reported badly conditioned, and the kernel entries collapse into rounding error, from which no asymptotic power can be read. The decomposition is therefore done on  $\tilde{\mathbf{A}}$  and mapped back,  $\mathbf{K} = \mathbf{S} \tilde{\mathbf{K}} \mathbf{S}^{-1}$ , checked against the verified kernel of step 2 at moderate  $k$ .

### 10.3.2 Which entries fail to decay

The order in  $1/k$  that each entry of the  $z$ -integrated kernel carries is read off numerically, over decades in  $k$ , rather than derived: a measured power is harder to fool oneself with than one obtained from a  $6 \times 6$  eigen-perturbation. A dot marks an entry that vanishes, so no power can be read.

|             | $u_z$    | $u_x$    | $u_y$    | $\tau_{zz}$ | $\tau_{xz}$ | $\tau_{yz}$ |
|-------------|----------|----------|----------|-------------|-------------|-------------|
| $u_z$       | .        | -1       | -1       | -2          | .           | .           |
| $u_x$       | -1       | .        | .        | .           | -2          | -2          |
| $u_y$       | -1       | .        | .        | .           | -2          | -2          |
| $\tau_{zz}$ | <b>0</b> | .        | .        | .           | -1          | -1          |
| $\tau_{xz}$ | .        | <b>0</b> | <b>0</b> | -1          | .           | .           |
| $\tau_{yz}$ | .        | <b>0</b> | <b>0</b> | -1          | .           | .           |

**Table 7.** Power of  $k$  carried by each entry of the  $z$ -integrated kernel as  $|k| \rightarrow \infty$ . Exactly one block fails to decay.

**The entire ultraviolet problem sits in one block:** the traction rows against the displacement columns. Everything else decays at least as  $1/k$  and needs no subtraction.

That block is also the one the physics predicts. In the evanescent limit a short-wavelength displacement produces a traction  $\tau \sim \mu k u$ , while the vertical integral (48) supplies  $1/k$ ; the ratio is  $k$ -independent. A traction response to a displacement that does not depend on  $k$  is a local stiffness — which is the Eshelby delta of §9.2, seen from the first-order side. The two representations identify the same object, and  $I_\infty$  has to reproduce exactly this block.

### 10.3.3 The asymptote in closed form

The limit of the non-decaying block is not merely numerical. Writing  $\hat{\mathbf{k}}$  for the lateral direction,

$$C[\tau_{zz} \leftarrow u_z] = -\nu_1 = -\frac{4\mu(\lambda + \mu)}{\lambda + 2\mu}, \quad (57)$$

$$C[\tau_{\alpha z} \leftarrow u_\beta] = \mu(\hat{k}_\alpha \hat{k}_\beta - \delta_{\alpha\beta}). \quad (58)$$

**$\nu_1$  is not a new constant.** It is the same  $\nu_1$  that appears in  $\mathcal{A}_{12}$  — born in the elimination (7), and shown in §3.2 to be the thesis’s  $\zeta$ . The subtraction is expressible entirely in the operator’s own constants. A wrongly identified asymptote would have no reason to land on  $\nu_1$  and  $\mu$ .

Both are established as *rates*. The relative residuals fall  $1.50 \times 10^{-3} \rightarrow 1.50 \times 10^{-4} \rightarrow 1.50 \times 10^{-5}$  across decades in  $k$ , ratios 10.0006 and 10.011 — clean  $O(1/k)$ , which is the  $O(k_c/k)$  correction to the evanescent limit. Agreement at a single  $k$  would not distinguish a correct closed form from a coincidence.

The in-plane form is the sharper test. Two constants are fitted on *two* angles and then predict *six*:  $P = -22.4977$  and  $Q = +22.4948$ , both equal to  $\mu = 22.5$ , with predictions tracking measurements to five digits across  $\theta \in [0, 2.4]$ . The isotropic entry (57) is constant to seven digits over the same range.

**What this settles about the integration.** The asymptote *is* direction-dependent — the relative variation over  $\theta$  is 0.39 — so  $I_\infty$  is not a pure delta and the regulator cannot be avoided. But the structure is the most favourable available:  $\delta_{\alpha\beta}$  and  $\hat{k}_\alpha \hat{k}_\beta$  are the only two rank-two invariants a direction can build, and they transform to a local delta and a principal-value dipole kernel respectively, both standard. The  $\tau_{zz} \leftarrow u_z$  entry is isotropic and so is a pure delta of weight  $-\nu_1$ .

#### 10.3.4 The lateral integral, and where it must be evaluated

The box form factor is separable — the lateral transform of the cube’s indicator is  $(2a)^2 \text{sinc}(k_1 a) \text{sinc}(k_2 a)$  — so with (57)–(58) and one lateral derivative from each  $\Delta \mathcal{A}$ , the integrand carries  $k_c k_d$  against  $C$ , giving two families:  $\int d^2 k W k_c k_d$ , which is separable, and  $\int d^2 k W k_a k_b k_c k_d / k^2$ , which is not. The elementary one-dimensional integrals are

$$\int_{-\infty}^{\infty} \text{sinc}^2(ka) dk = \frac{\pi}{a}, \quad \int_{-\infty}^{\infty} k^2 \text{sinc}^2(ka) e^{-\epsilon|k|} dk = \frac{4}{4a^2\epsilon + \epsilon^3} = \frac{1}{a^2\epsilon} - \frac{\epsilon}{4a^4}. \quad (59)$$

**The second is a pure pole with no finite part.** So  $\int I_\infty d^2 k$  is not a number, and the split-and-subtract cannot be closed in  $k$ -space alone: adding back a pure pole restores exactly the divergence it was introduced to remove.

The method is not wrong; the tail is being evaluated in the wrong representation. Forming  $|\hat{S}|^2$  first and then integrating term by term destroys the structure that makes the constant part finite. By Parseval,

$$\iint f(\mathbf{x}) K(\mathbf{x} - \mathbf{x}') g(\mathbf{x}') = \int \frac{d^2 k}{(2\pi)^2} \hat{f}(-\mathbf{k}) \hat{K}(\mathbf{k}) \hat{g}(\mathbf{k}), \quad (60)$$

so a  $\hat{K}$  tending to a constant  $C$  contributes  $C \int_V f g d\mathbf{x}$  — a *local* term, finite, needing no regulator at all. This is verified directly in one dimension, both sides returning  $1/2a$ .

The decomposition therefore splits by representation rather than by magnitude:

| piece                                                                      | real-space object          | evaluation               |
|----------------------------------------------------------------------------|----------------------------|--------------------------|
| isotropic, $-\nu_1$ and $P\delta_{\alpha\beta}$                            | delta                      | local, elementary        |
| traceless $\hat{k}_\alpha \hat{k}_\beta - \frac{1}{2}\delta_{\alpha\beta}$ | 2-D principal-value dipole | conditionally convergent |
| remainder $I - I_\infty$                                                   | —                          | numerically, in $k$      |

**Table 8.** Split-and-subtract survives, but  $\int I_\infty$  is evaluated in real space. Only the remainder stays in  $k$ .

The traceless part is of degree  $-2$  with vanishing angular average — exactly the conditional convergence the real-space engine of §9.2 handles by peeling derivatives onto the cube faces. The two representations meet again here: the same distributional structure, in two dimensions rather than three.

### 10.3.5 The principal-value piece vanishes

In two dimensions  $\nabla^2 \log \rho = 2\pi\delta^2$ , and  $\partial_\alpha \partial_\beta \log \rho = (\delta_{\alpha\beta} - 2\hat{n}_\alpha \hat{n}_\beta)/\rho^2$  away from the origin; taking the trace fixes the delta weight, so that

$$\partial_\alpha \partial_\beta \log \rho = \text{PV} \left[ \frac{\delta_{\alpha\beta} - 2\hat{n}_\alpha \hat{n}_\beta}{\rho^2} \right] + \pi \delta_{\alpha\beta} \delta^2(\mathbf{x}). \quad (61)$$

By translation invariance the double integral over the square collapses onto the autocorrelation of its indicator,  $w(\mathbf{u}) = (2a - |u_1|)(2a - |u_2|)$ , exactly — removing two of the four dimensions — and both derivatives move onto  $w$ , which is legitimate because  $w$  has compact support.

$w$  is piecewise bilinear, so  $\partial_1 \partial_1 w$  is a sum of line deltas. It has three terms, not one:  $w$  vanishes at  $|u_1| = 2a$  but  $\partial_1 w$  does *not* — it jumps from  $\mp(2a - |u_2|)$  to zero — so there are edge deltas at  $u_1 = \pm 2a$  alongside the interior one at  $u_1 = 0$ . With them,

$$D_{11} = \underbrace{3}_{\text{interior}} + \underbrace{(\pi - 3)}_{\text{edges}} = \pi = \pi(2a)^2 \Big|_{a=1/2}, \quad D_{12} = 0, \quad (62)$$

and the trace identity  $\text{tr } D = \langle \nabla^2 \log \rho, w \rangle = 2\pi w(0) = 2\pi(2a)^2$  holds exactly. Dropping the edge deltas gives 6 in place of  $2\pi$ , so that identity is what detects them.

**Hence  $\langle \text{PV}_{\alpha\beta}, w \rangle = D_{\alpha\beta} - \pi \delta_{\alpha\beta} w(0) = 0$ : the traceless dipole contributes nothing.**

It could not have contributed.  $\langle \text{PV}_{\alpha\beta}, w \rangle$  is a symmetric two-tensor invariant under the square's  $D_4$  symmetry, and the only such tensor is a multiple of  $\delta_{\alpha\beta}$  — the traceless part has nowhere to live. The conditional convergence is real, but for this cross-section it integrates away.

**The whole of  $\int I_\infty$  is therefore the local delta term.** The piece that looked hardest — the conditionally convergent, shape-dependent dipole — is identically zero here, and what remains is elementary.

### 10.3.6 Why the transferred form converges

In the untransferred form  $\partial_\alpha \rightarrow ik_\alpha$ , each contrast operator contributes a factor of  $k$ , and the pole of (59) follows. After the by-parts transfer of (17) the derivatives act on the *polynomials, in real space*:  $\partial_\alpha$  of an affine trial function is a constant, whose lateral transform is again the box form factor. No factor of  $k$  survives, and the lateral integral converges.

**The augmentation therefore does two things at once:** it keeps the indicator from being differentiated, and it makes the lateral integral convergent. Both are the same integration by parts.

Convergence is nevertheless slow. With  $\hat{K} \rightarrow C$  the truncation error falls only as  $1/\Lambda$ , which is what the subtraction is for.

### 10.3.7 The tail, in closed form

The tail is elementary. For the in-plane entry  $C = \mu(\hat{k}_1^2 - 1) = -\mu k_2^2/k^2$ , so

$$\int d^2k W C = -\mu \int d^2k \operatorname{sinc}^2(k_1 a) \operatorname{sinc}^2(k_2 a) \frac{k_2^2}{k^2} = -\frac{\mu}{2} \left(\frac{\pi}{a}\right)^2, \quad (63)$$

because  $k_1^2/k^2 + k_2^2/k^2 = 1$  and the weight is symmetric under  $k_1 \leftrightarrow k_2$ , so each piece is exactly half of  $\int d^2k W = (\pi/a)^2$ . No quadrature is required; direct integration returns 19.7376 against  $(\pi/a)^2/2 = 19.7392$ , the difference being the quadrature's own error on a slowly converging integrand.

### 10.3.8 The remainder, and the test that closes the construction

Truncating the lateral integral at  $|k| < \Lambda$  for the  $\tau_{xz} \leftarrow u_x$  entry, with and without the subtraction:

The split is only worth making if both halves give the same answer. Truncating at  $|k| < \Lambda$  for the  $\tau_{xz} \leftarrow u_x$  entry:

| $\Lambda$ | unsubtracted | remainder | remainder + tail |
|-----------|--------------|-----------|------------------|
| 80        | -318.2214    | 118.83504 | -325.29716       |
| 160       | -321.6764    | 118.82727 | -325.30493       |
| 320       | -323.4539    | 118.82513 | -325.30707       |
| 640       | -324.3400    | 118.82466 | -325.30754       |

**Table 9.** The subtracted route is converged to six figures by  $\Lambda = 80$ ; the unsubtracted one is still 0.97 short of the limit at  $\Lambda = 640$ .

The remainder's successive differences fall  $7.8 \times 10^{-3} \rightarrow 2.1 \times 10^{-3} \rightarrow 4.7 \times 10^{-4}$ , a ratio near 0.27, against the unsubtracted 0.50 — the latter being the  $1/\Lambda$  decay the constant part imposes. Extrapolating both:

$$\underbrace{-325.226}_{\text{unsubtracted limit}} \quad \text{versus} \quad \underbrace{-325.308}_{\text{remainder+closed-form tail}}, \quad \text{relative difference } 2.5 \times 10^{-4}, \quad (64)$$

and that residue is dominated by the single-term extrapolation of the slowly converging unsubtracted series, not by the construction. At  $\Lambda = 640$  the unsubtracted error is 0.97 and the subtracted one  $5 \times 10^{-4}$ : a factor of about two thousand.

**Split-and-subtract is verified.** The two routes to the same number agree, the tail

is closed in form by (63), and the remainder converges three orders of magnitude faster than the direct integral.

**This is the method, and §13.4 is where it pays.** The finding above — that  $\int I_\infty$  is a pure pole with no finite part, so the split *cannot* be closed in  $k$ -space and must be taken to real space — reads here as a difficulty encountered and worked around. It is better than that: it is the general statement, and the moments  $A$  and  $C$  are obtained by applying it verbatim. There the asymptotic integrand is again a function of direction alone, its  $n = 0$  part is again a delta, and its  $n \neq 0$  part is again a principal-value dipole over the square. The only change is which angular harmonics survive, which the symmetry of the trial function decides.

Worth stating plainly, since the cost was paid twice: four separate attempts were made to close the  $A$  and  $C$  tail in  $k$ -space — cosine-integral tails, graded angular meshes for the logarithmic singularities that then appear, a square-versus-disc geometry correction — before this paragraph was reread. Every one of them was a way of doing what it says cannot be done.

*What steps 1–3 establish.*  $\Gamma$  from spectral projectors; the triangular  $z$  integrals and their assembly; the scaling, the power table, and the asymptote  $-\nu_1$ ,  $\mu(\hat{k}\hat{k} - \delta)$  verified as an  $O(1/k)$  rate;  $\int I_\infty$  decomposed by representation, its principal-value part vanishing by  $D_4$  symmetry and its remaining part closed by (63); and the closure test (64). All of that is machinery for the lateral integral. It does not yet say what the integrand is, and the next subsection is about a question that has to be answered before it can.

#### 10.4 Step 4: how the augmentation composes with the Schwinger form

The augmentation of §6 turns the differential  $\Delta A$  into a multiplicative  $\Delta C_{\text{eff}}$  by moving each outer lateral derivative onto the test function. That is what keeps  $\partial_\alpha \mathbf{1}_V$  — the surface layer on the faces of the cube — from ever being formed. The Schwinger object

$$[\Delta C \mathcal{G} \Delta C]_{kl} = \langle J_6 \psi_k \mid \Delta A \Gamma \Delta A \mid \psi_l \rangle \quad (65)$$

carries *two* contrast operators, and the right-hand one has no test function next to it: what  $\Gamma$  needs is  $\Delta A \psi_l$  as an actual source, which is  $\partial_\alpha (\Delta c \mathbf{1}_V \partial_\beta \psi_l)$  — precisely the object the transfer exists to avoid. Read that way the two constructions look incompatible.

They are not, and the reason is that the Born term and the Schwinger term differ in a way that matters. In the Born term the indicator of the contrast and the test function sit at the *same* point, so  $\partial_\alpha (\mathbf{1}_V c)$  would have to be multiplied by another factor supported on  $\bar{V}$ ; a surface layer times an indicator is not defined, and the transfer is therefore not a convenience but the only well-posed form. In the Schwinger term the two indicators are separated by  $\Gamma$ . The source at  $\mathbf{x}'$  is paired, through a kernel, with a test object at  $\mathbf{x}$ ; nothing coincides, the surface layer is harmless — and in the lateral Fourier domain it need not be formed at all, because

$$\mathcal{F}[\partial_\alpha (\mathbf{1}_V c)](\mathbf{k}) = ik_\alpha \mathcal{F}[\mathbf{1}_V c](\mathbf{k}) \quad (66)$$

is exact. The derivative is carried by the *transform*; the indicator stays inside the form factor, undifferentiated.

The bookkeeping that follows is uniform. Every derivative in  $A$  is lateral — the depth derivative is the left-hand side of  $\partial_3 q = Aq + d$  — so in the lateral Fourier domain each is a factor  $ik_\alpha$  at the single wavenumber the integral runs over. One entry of (65) has four derivative slots, and they are placed as

| leg   | slot  | where it goes                                               |
|-------|-------|-------------------------------------------------------------|
| left  | outer | onto the test polynomial (the augmentation)                 |
| left  | inner | onto the propagated field $\rightarrow ik_\beta$            |
| right | outer | onto the source transform $\rightarrow ik_\alpha$ , by (66) |
| right | inner | onto the trial polynomial (the augmentation)                |

The two outer slots are the ones in question. The left one goes exactly where  $\Delta C_{\text{eff}}$  already puts it; the right one goes onto the kernel. No surface term is evaluated on a face, and the augmentation is unchanged.

(66) is checked numerically rather than asserted. Paired with a smooth bump  $f$  it states  $\int f \partial_1 \mathbf{1}_S = - \int_S \partial_1 f$ , whose right side is an ordinary integral over the square; evaluating the left side in the Fourier domain as  $(2\pi)^{-2} \int \hat{f}(-\mathbf{k}) ik_1 F(\mathbf{k})$  reproduces it to  $1.0 \times 10^{-13}$ , with a zero imaginary residue and the opposite sign rejected by a factor of two.

### The propagator, checked against the Kelvin tensor

Before any of this can be assembled, one normalisation has to be settled that the bridge identity of §9 does not reach: the bridge compares  $\Gamma$  at two depths and so tests propagation, not the strength of the source jump. Inverting  $\Gamma$  laterally at a fixed vertical separation and reading the velocity–force block,

$$\frac{1}{-i\omega} (2\pi)^{-2} \int d^2k \Gamma_{3+i,j}(\mathbf{k}; \Delta z) \rightarrow G_{ij}(\Delta z), \quad (67)$$

the analytic Kelvin tensor comes back to  $7.3 \times 10^{-4}$  and still converging, with off-diagonal leakage  $2.9 \times 10^{-4}$ ; the residue is consistent with the  $O((k_{sr})^2)$  the static comparison drops. This fixes both conventions at once — that a force  $f_i$  enters row  $i$  of  $\partial_3 q = Aq + s$ , and that  $q_{3+i} = v_i = -i\omega u_i$  comes out. The quadrature must be polar: the kernel depends on  $|\mathbf{k}|$  and on direction separately and so is not smooth at the origin, where a Cartesian tensor rule converges only as  $1/n$  (measured:  $4.0 \times 10^{-2} \rightarrow 1.2 \times 10^{-2}$  over a ninefold refinement, against  $5.4 \times 10^{-3} \rightarrow 7.3 \times 10^{-4}$  in polar). This is the opposite of the lateral moment itself, whose form factors are separable sinc products that polar cannot resolve.

### The density channel, assembled

With  $\Delta\lambda = \Delta\mu = 0$  only the multiplicative  $i\omega\Delta\rho$  term of  $\Delta A$  survives, both legs reduce to constants, and (65) can be predicted in closed form:

$$\langle J_6 \psi_0 \mid \Delta A \Gamma \Delta A \mid \psi_0 \rangle = i\omega^5 \Delta\rho^2 M, \quad M = \int_V \int_V G_{11} = \left(a_0 + \frac{b_0}{3}\right) D_0, \quad (68)$$

with  $a_0, b_0$  the Kelvin coefficients of §9.2 and  $D_0$  the purely geometric double-volume moment  $\int_V \int_V 1/r$ .  $D_0$  is evaluated by writing the double integral as a single one against the autocorrelation weight  $\prod_i (2a - |u_i|)$  — exact, and three dimensions cheaper — and removing the  $1/r$  outright with a Duffy substitution  $\mathbf{u} = t(\xi, \eta, 1)$ , whose Jacobian  $t^2$  cancels it:

$$D_0 = 1.8823126444, \quad \int_V 1/r = 2.3800773640. \quad (69)$$

The lateral assembly with *both* contrast operators in place reproduces (68) to  $6.9 \times 10^{-5}$ , the two routes sharing nothing but the medium: one is a spectral construction on  $A$  integrated over lateral wavenumber, the other a real-space moment of the Kelvin tensor.

### What the nine-parameter $\mathcal{G}$ is, and is not

The same comparison locates an object that had been taken on trust. The  $\mathcal{G}$  used in §9 was *calibrated* against this package's  $T_9$  operator rather than derived, and it is not the moment that (65) calls for. The gap factorises exactly:

$$\frac{[\Delta C \mathcal{G} \Delta C]_{00}}{i\omega^5 \Delta \rho^2 M} = i\omega \frac{\int_V 1/r + i(\text{radiation})}{D_0}, \quad \text{residual } 3.6 \times 10^{-5}, \quad (70)$$

where the radiation term is the closed form  $V\omega(\alpha^{-3} + 2\beta^{-3})/12\pi\rho$  that the moment export is itself gated on. Two statements, both of which matter:

- The geometric factor  $\int_V 1/r/D_0 = 1.2644$  is the ratio of a *point*-volume moment to a *volume*-volume moment. The calibrated  $\mathcal{G}$  is a collocation object; (65) is a Galerkin one. That is consistent with this project's own finding that the single site is best posed by collocation, but the two cannot be mixed without the factor.
- The remaining  $i\omega$  is the normalisation of the  $q$ -field trial space:  $\Delta C \mathcal{G} \Delta C$  carries two  $J_6$  test factors while (65) carries one, so a  $\mathcal{G}$  that composes correctly must absorb the trial-function normalisation. How it should do so is not settled here; see §18.

### The derivative channels, and three defects the test exposed

The density channel settles the basis, the kernel, the  $z$  integrals and the quadrature, but its contrast operator is *multiplicative*: it does not exercise (66) at all. The channels that do are those carrying  $\Delta\lambda$  and  $\Delta\mu$ , and for them a sharp statement is available with no new arbiter.  $\mathbf{J}_6 \Delta \mathcal{A}$  is symmetric — that is what derived the test space in §6 — and  $\mathbf{\Gamma}$  is reciprocal,

$$\mathbf{\Gamma}(-\mathbf{k}; z', z)^\top = \mathbf{J}_6 \mathbf{\Gamma}(\mathbf{k}; z, z') \mathbf{J}_6, \quad (71)$$

which follows from  $\mathcal{A}(-\mathbf{k})^\top \mathbf{J}_6 = -\mathbf{J}_6 \mathcal{A}(\mathbf{k})$  and  $\mathbf{J}_6^{-1} = -\mathbf{J}_6$ . Hence the assembled matrix must be symmetric. It is not symmetric by construction: the assembly builds its two legs by opposite rules, transferring the left outer derivative onto the test polynomial and carrying the right one on the kernel as  $ik$ . Symmetry is therefore a test of exactly the bookkeeping in question, and on a

quadrature grid symmetric under  $\mathbf{k} \rightarrow -\mathbf{k}$  it is an *algebraic* property, holding to round-off even where the grid is far too coarse to give the value accurately.

Applied, it found three defects, each in code that had already passed its own checks:

1. **The up/down split was decided by round-off inside the propagating window.** Classifying by  $\Re(\text{ev}) < 0$  is right for evanescent modes and meaningless for propagating ones: below  $k = \omega/\beta$  the vertical wavenumber is real,  $\Re(\text{ev})$  vanishes to arithmetic noise, and its sign was being decided independently at  $+\mathbf{k}$  and  $-\mathbf{k}$ . The reciprocity residual (71) at  $|k| = 1.1 \times 10^{-3}$  was 1.245 — a complete failure, in the region that carries all of the radiation. The limit  $\omega \rightarrow \omega(1 + i\epsilon)$  settles it without needing  $\epsilon$ : a propagating mode has  $\text{ev} = ik_z$  and the perturbation gives  $k_z$  a positive imaginary part, so  $\Re(\text{ev}) \rightarrow 0^-$  exactly when  $\text{I}(\text{ev}) > 0$ . Residual after:  $3.7 \times 10^{-13}$ .
2. **The triangular  $z$  integral cancels catastrophically at small  $aq$ .** The numerator of  $T^{11}$  is  $O((aq)^5)$  while its individual terms are  $O(1)$ , so by  $aq \approx 0.025$  it has lost about eight digits. That is not a corner case — it is again the propagating window. Expanding  $T^{mn}$  as a power series in  $q$  below  $|2aq| = 3$ , with the coefficients of §10.2.1, took that node from  $1.3 \times 10^{-6}$  to  $2.1 \times 10^{-13}$ .
3. **The scaling used to condition the eigenproblem is in the wrong direction.**  $\text{diag}(1, 1, 1, k, k, k)$  was introduced because the up- and down-going eigenvectors become parallel as  $k$  grows. Measured against *no scaling at all* it makes the eigenvector condition number *worse* at every  $k$  above the propagating window — by two orders at  $k = 60$ . Parlett–Reinsch balancing beats both by about twelve orders there ( $8.7 \times 10^7$  against  $1.6 \times 10^{20}$ ), because it is free to treat the third component differently from the first two, which no scalar rule can.

With all three repaired the assembled  $9 \times 9$  matrix is symmetric to  $4.7 \times 10^{-12}$ , while reversing the sign of the right leg's  $ik$  — the one step (66) licenses — breaks it at  $7.2 \times 10^{-2}$ . That is a separation of ten orders on precisely the disputed bookkeeping, and the derivative terms carry 80% of the matrix, so the test is not vacuous.

### The radiation reaction, from the first-order system

Repairing the branch selection changed the physics, not just the arithmetic. The previous classification suppressed the outgoing condition; with it restored, the propagator carries radiation damping, and both real-space arbiters above had to acquire the term they were missing. The sharpest form of the check is pointwise:

$$\text{IG}_{11}(\Delta z) \Big|_{\text{first-order}} = 5.2134 \times 10^{-14} \quad \text{against} \quad \frac{\omega}{12\pi\rho} \left( \frac{1}{\alpha^3} + \frac{2}{\beta^3} \right) = 5.2250 \times 10^{-14}, \quad (72)$$

agreeing to  $2.2 \times 10^{-3}$ . This quantity has *no static counterpart at all*: it is not a correction to the Kelvin tensor but a term the static theory does not contain, and it arrives here out of the spectral projectors of  $\mathcal{A}$  with nothing put in by hand.

Resolving it needs the quadrature to respect the two branch radii  $|k| = \omega/\alpha$  and  $\omega/\beta$ , where the vertical wavenumber vanishes. Those are *circles*: no Cartesian panel can follow one, and an

unresolved square-root costs several per cent. The lateral integral is therefore split — Cartesian outside the propagating window, where the integrand is a separable product of sinc-type form factors that polar cannot resolve; polar inside it, where the branch radii are panel edges and  $ka < 0.02$  makes the form factors constant to  $10^{-4}$ . The central block of the Cartesian grid is exactly the inner square, so the two tile the box with no overlap and no gap. With that in place the density channel converges cleanly —  $9.858, 9.834, 9.828 \times 10^{-4}$  over a fourfold refinement — where before it drifted.

*Status of Task 1.* The design question is settled; the assembly computes a definite object; the density channel is verified against an independent real-space arbiter to  $9.8 \times 10^{-4}$ ; the derivative channels are verified by a symmetry with a ten-order separation from its control; and the radiation reaction comes out right. What remains is to read the shear moment  $B$  off the assembled matrix and compare it with  $\sqrt{3}(\lambda + \mu)/6\pi\mu(\lambda + 2\mu)$ , and to settle the  $i\omega$  of (70).

## 11 The thesis spectral representation, and what it removes

Three of the four numerical defects of §10.4 are not properties of the physics. They are consequences of asking a numerical eigensolver to do work that §2.1 of the thesis has already done in closed form, and they do not arise at all if  $\mathbf{\Gamma}$  is built the way the thesis builds it.

### 11.1 What the thesis gives

Alongside (Akdef) the thesis carries the complete spectral representation

$$\mathbf{A}(k_x, k_y) = \mathbf{D}_z \mathbf{\Lambda} \mathbf{D}_z^{-1}, \quad \mathbf{\Lambda} = i \operatorname{diag}[k_{z,P}, k_{z,S}, k_{z,H}, -k_{z,P}, -k_{z,S}, -k_{z,H}], \quad (73)$$

with the columns of  $\mathbf{D}_z$  written out analytically —  $|\pm, P\rangle, |\pm, S\rangle, |\pm, H\rangle$ , the  $+$  columns downgoing — and three further things that matter more than they look:

1. **The branch is defined piecewise.**  $k_{z,c} = \sqrt{\hat{K}_c^2 - k_x^2}$  where  $|k_x| < |\hat{K}_c|$  and  $i\sqrt{k_x^2 - \hat{K}_c^2}$  otherwise: the positive root in both regimes, with the up/down assignment carried by the *column ordering* of (73) rather than by the sign of any computed quantity.
2. **The shear basis is chosen non-degenerate.** The thesis notes  $k_{z,S} = k_{z,H}$  and keeps the two labelled apart regardless, remarking that its quasi-SV and quasi-SH vectors — unlike the conventional seismological decomposition — remain independent at  $k_x = k_y = 0$ .
3. **The inverse requires no inversion.** From the quasi-Hamiltonian relation,

$$\mathbf{D}_z^{-1}(k_x, k_y) = -i \mathbf{J}_6 \mathbf{D}_z^t(-k_x, -k_y) \mathbf{J}_6. \quad (74)$$

All of it is verified here rather than taken on trust, and the checks are coupled: (73) cannot reproduce (Akdef) unless the eigenvectors, the ordering, the branch and (74) are all correct together. It does, to  $6.2 \times 10^{-14}$ . The normalisation factors  $\epsilon_c$ , which the thesis specifies only as “chosen so that  $\mathbf{D}_z^{-1}$  takes a simple form”, are *read off* rather than guessed: with  $\epsilon = 1$  the product (74)  $\cdot \mathbf{D}_z$  is diagonal — to  $1.1 \times 10^{-15}$ , itself a test of the eigenvectors — and  $\epsilon_c^2$  is the reciprocal of each diagonal entry.

**The one place  $\mathbf{D}_z$  is genuinely singular, and why it is not a difficulty.** Reading  $\epsilon_c$  off the diagonal fails where a diagonal entry vanishes, and that is not round-off: it is the branch point  $k_{z,c} = 0$ , at  $|k| = \hat{K}_p$  and  $|k| = \hat{K}_s$  exactly, where that wave type’s up- and down-going columns coalesce and  $\mathbf{D}_z$  has rank five. It must be tested on  $k_{z,c}$  itself and not on the spread of the diagonal — the entries of  $\mathbf{D}_z$  span many orders at moderate  $k$ , so a smallest-entry threshold measures dynamic range and fires where nothing is wrong.

The degeneracy exists only in the lossless idealisation. Since  $k_{z,c}^2 = \hat{K}_c^2 - k^2$  and  $\hat{K}_c = \omega/c$ , it cannot vanish at real  $k$  once  $\omega$  carries an imaginary part, so the attenuation that §10 already adopts for the quadrature — “ $\omega$  is complex from the outset” — removes the coalescence outright rather than dodging it. At  $Q = 200$  the closest approach of  $|k_{z,s}|$  across its own circle is  $1.4 \times 10^{-3}$  against a scale of  $2 \times 10^{-2}$ ;  $\mathbf{D}_z$  inverts *on* the circle to  $1.2 \times 10^{-15}$ , where the lossless evaluation is singular; and (73) is rebuilt there to  $3.7 \times 10^{-16}$ . The device is the one the package already uses: putting the loss in a complex slowness rather than a complex  $\omega$  is the same statement, since  $k_{z,c}$  depends on the two only through  $\omega/c$ , and the two agree to  $6.3 \times 10^{-6}$  at  $Q = 200$  — the  $O(Q^{-2})$  difference. Nudging a node off the point remains the fallback for a deliberately lossless calculation, and placing a panel edge there, as the lateral gates do, is the right treatment of the integrable  $1/\sqrt{\cdot}$  that survives in that limit.

## 11.2 What it removes

|                                                                                                                                                                  |                                                                                         |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------|
| found numerically, working from the operator form                                                                                                                | already settled in §2.1                                                                 |
| $\Re(\lambda)$ carries no information in the propagating window; its sign was decided independently at $\pm \mathbf{k}$ , giving a reciprocity residual of 1.245 | the branch of $k_{z,c}$ , defined piecewise, with the ordering of (73) carrying up/down |
| the doubly degenerate shear eigenspace returns an arbitrary <i>SV</i> / <i>SH</i> mixture; an intermittent factor of four in a mode-converted coefficient        | quasi- <i>SV</i> and quasi- <i>SH</i> chosen non-degenerate, and said to be so          |
| eigenvector conditioning of $1.6 \times 10^{20}$ ; a scalar scaling that proved to be in the wrong direction; then Parlett–Reinsch balancing                     | no eigenvectors are computed                                                            |
| $\mathbf{D}_z^{-1}$ by numerical inversion                                                                                                                       | (74)                                                                                    |

**Table 10.** The fourth defect, the cancellation in  $T^{11}$  of §10.2.1, is genuine: it concerns the  $z$  integral and not the spectral representation.

Rebuilt on (73),  $\Gamma$  is reciprocal to  $3.6 \times 10^{-12}$  with no classification performed anywhere, reproduces the Kelvin tensor to  $1.0 \times 10^{-3}$  and converging, and carries the radiation reaction to  $4.4 \times 10^{-3}$  — *with no  $\epsilon$  prescription*, since the outgoing condition is already in the branch of  $k_{z,c}$ . In the thesis basis the extraction is also plainer: a force  $f_j$  enters row  $3 + j$  with a minus sign and the displacement comes out in rows  $0-2$ , so  $G_{ij} = -\Gamma_{i,3+j}$  with no velocity factor at all.

## 12 A layer in a whole space: the external arbiter

Every comparison so far in this note sets two formulations against each other. None of them says which is right. A single layer embedded in a whole space removes that limitation, because its reflection coefficient is known independently — from a validated reflectivity recursion, and from a closed form short enough to derive here — and because it exercises the operator in a way the cube cannot.

### 12.1 Why a layer is the right test

For a laterally invariant contrast the operator difference is simply

$$\Delta\mathcal{A}(z) = [\mathcal{A}_{\text{total}}(\mathbf{k}) - \mathcal{A}_{\text{background}}(\mathbf{k})] \mathbf{1}_{[0,h]}(z), \quad (75)$$

a piecewise-constant matrix. **There is no augmentation, no by-parts transfer, no surface term and no differentiated indicator**, because  $\mathcal{A}$  carries no  $\partial_z$ : a jump in the medium is simply a jump in the coefficients of an ODE, and is harmless. That is the structural advantage of the first-order form, and it is precisely the property the whole cube development of §6–§10.4 had to work around. So this tests  $\mathcal{A}$ ,  $\Gamma$  and the Lippmann–Schwinger composition with none of the cube machinery in the way.

It also tests two things the cube checks could not. It is *nonlinear* in the contrast, so it reaches the resummation rather than the Born term alone; and  $\Delta\mathcal{A}$  carries  $\mathbf{k}$  explicitly through  $\mathcal{A}_{12} \sim v_1 k^2 / i\omega$ , so scanning slowness tests a  $k$  dependence that the density channel of §10.4, being multiplicative, left untouched.

### 12.2 The closed form, by continuity

At normal incidence the answer follows from matching. With  $e^{-i\omega t}$  and  $u$  the vertical displacement, each medium carries  $u = De^{ikz} + Ue^{-ikz}$  with  $k = \omega/c$ , and the traction is  $\tau = \rho c^2 \partial_z u = i\rho c \omega (De^{ikz} - Ue^{-ikz})$ . Continuity of  $u$  and  $\tau$  at  $z = 0$  and  $z = h$ , with no up-going wave below, gives

$$R = \frac{r_{12} + r_{23} e^{2i\omega h/\alpha_2}}{1 + r_{12} r_{23} e^{2i\omega h/\alpha_2}}, \quad r_{12} = \frac{Z_1 - Z_2}{Z_1 + Z_2} = -r_{23}, \quad Z = \rho c. \quad (76)$$

The point of deriving it rather than quoting it is that the *sign* is then fixed by the derivation. Throughout this section

$$R = \frac{u_3 \text{ of the up-going field at } z = 0^-}{u_3 \text{ of the incident field at } z = 0^-}, \quad (77)$$

on both sides of every comparison, which leaves no convention free.

### 12.3 The first-order route

The laterally invariant Lippmann–Schwinger equation

$$q(z) = q_{\text{inc}}(z) + \int_0^h \mathbf{\Gamma}(\mathbf{k}; z, z') \Delta \mathcal{A} q(z') dz' \quad (78)$$

is solved by piecewise-constant collocation, the cell integrals taken in closed form from the spectral representation of  $\mathbf{\Gamma}$  and *split at*  $z' = z$ : the jump across the diagonal is the source delta, and quadrature across it would smear exactly the object that defines the propagator.

One structural simplification is used, and it is exact. With  $\mathbf{k}$  along  $x_1$  the system is block diagonal:  $\{-\tau_{23}, v_2\}$  is  $SH$  and does not couple to  $\{-\tau_{13}, -\tau_{33}, v_1, v_3\}$ . Working in the  $4 \times 4$   $P$ -SV block is therefore no approximation — and it removes the double degeneracy of  $q_S$ , which is a genuine hazard rather than a nicety: a degenerate eigenspace has no preferred basis, so `eig` returns an arbitrary mixture of  $SV$  and  $SH$ . Before the block restriction the  $SS$  coefficient was correct at three slownesses and wrong *by a factor of four* at the fourth, which reads as a critical angle and is linear algebra.

### 12.4 What it gives

| contrast | (76)                    | first order             | relative             |
|----------|-------------------------|-------------------------|----------------------|
| +2%      | $-0.019937 - 0.019792i$ | $-0.019937 - 0.019792i$ | $2.7 \times 10^{-5}$ |
| +10%     | $-0.127846 - 0.087946i$ | $-0.127849 - 0.087948i$ | $2.4 \times 10^{-5}$ |
| +40%     | $-0.578990 - 0.067730i$ | $-0.578999 - 0.067736i$ | $1.7 \times 10^{-5}$ |

**Table 11.** Normal incidence,  $h = 200$  m,  $\omega h/\alpha = 2.4$ . The +40% row is the one that matters: it is far from Born, so it tests the *resummation*. Refinement in the cell count gives ratios 4.00, 4.00, 4.00 — clean second order.

Against the reflectivity recursion the comparison runs to  $\sin i = 0.75$ :

$$\frac{R_{PP}^{\text{first order}}}{R_{PP}^{\text{Kennett}}} = -1.00004, -1.00004, -1.00003, -1.00002, \quad (79)$$

and the same for  $SS$  to  $1.0 \times 10^{-4}$ . Two remarks make this a statement rather than a fitted factor. First, the recursion’s  $R$  proved *independent of the thickness of the layer above* — identical for  $d = 50, 100, 200, 400$  m — so it is already referenced at the target layer, and the two-way phase correction that seemed obviously required in fact manufactures a spurious slowness-dependent offset. Second, what is checked is not the value  $-1$  but that the factor is *constant*: a convention cannot depend on angle, and a physics error would not be constant. The sign itself is pinned by (76), which owes nothing to either code.

The mode-converted coefficients are covered by an identity that needs no knowledge of the convention at all. If the two differ by a diagonal rescaling  $\text{diag}(d_P, d_S)$  — the displacement-versus-flux factor, invisible to any per-mode ratio test — then the off-diagonal ratios are  $d_P/d_S$  and  $d_S/d_P$ , whose product is unity whatever  $\mathbf{D}$  is. Measured: 1.00008, 1.00010, 1.00008.

Finally, a control isolates what this test uniquely reaches. Building  $\Delta \mathcal{A}$  at  $\mathbf{k} = 0$  while leaving

$\Gamma$  at the true  $\mathbf{k}$  changes *nothing* at normal incidence (relative difference 0, exactly) and breaks the answer by  $6.8 \times 10^{-2}$  off it. The  $k$  dependence of the contrast operator is therefore load-bearing, and correct.

**What this establishes, and what it does not.** The system matrix  $\mathcal{A}$ , the propagator  $\Gamma$ , their composition, and the resummation reproduce an independently validated reflectivity code across contrast and across angle. That is the first external validation in this note. It does *not* validate the cube: the augmentation, the form factors and the moments of §9.2 play no part in a laterally invariant problem. Those remain checked only against other formulations.

## 13 The static moments $A$ , $B$ and $C$ , read off

This is what Task 1 was for: the cube's propagator moment obtained from  $\partial_3 \mathbf{q} = \mathcal{A} \mathbf{q} + \mathbf{d}$  rather than from the Kelvin tensor, and compared with a closed form the first-order machinery knows nothing about.

### 13.1 Which realisation, and why it is the easy one

$B$  multiplies  $(\delta_{ik}\delta_{jm} + \delta_{im}\delta_{jk})$  in the moment tensor, so

$$B = I_{xyxy} = \int_V \partial_y \partial_x G_{xy} dV, \quad (80)$$

and *that* realisation carries only lateral derivatives. In the lateral Fourier domain they are factors of  $ik$ , with no  $\partial_z$  and hence no  $\mathcal{A}\Gamma$  product to form:

$$B = \frac{1}{(2\pi)^2} \int d^2k (ik_x)(ik_y) F(\mathbf{k}) \int_{-a}^a G_{xy}(\mathbf{k}; -z') dz'. \quad (81)$$

Every other realisation of  $B$  involves  $\partial_z$ . The  $z$  integral is closed form from (73); the lateral one is quadrature.

### 13.2 The sign, pinned independently

The two factors of  $i$  give  $(ik_x)(ik_y) = -k_x k_y$ . Dropping them costs exactly a sign, which is invisible against  $|B|$  and would be reported as success. So the sign is fixed by an evaluation owing nothing to the first-order machinery: by the divergence theorem,

$$\int_V \partial_y \partial_x G_{xy} dV = \oint_{\partial V} n_y (\partial_x G_{xy}) dA, \quad (82)$$

which is legitimate — and, more to the point, *complete* — because the divergence theorem applied to a locally integrable  $F$  returns the full *distributional* integral. Any delta at the origin is therefore included automatically, which is what makes this a safe way to evaluate a quantity whose volume form has an Eshelby delta in it. Here  $\partial_x G_{xy} \sim 1/r^2$ , integrable in three dimensions

and non-singular on faces standing at distance  $a$  from the centre; only the  $y = \pm a$  faces contribute and their integrands add, being odd in  $y$ . Evaluated,

$$\oint_{\partial V} n_y \partial_x G_{xy} dA = +2.6137073561 \times 10^{-12} = B \quad \text{to } 1.7 \times 10^{-13}, \quad (83)$$

so (80) is the right realisation and the sign is positive.

### 13.3 The result

| $\Lambda$ | nodes per half-period | $B$ from (81)                  | relative             |
|-----------|-----------------------|--------------------------------|----------------------|
| 400       | 8                     | $2.6144253494 \times 10^{-12}$ | $2.7 \times 10^{-4}$ |
| 400       | 12                    | $2.6138081029 \times 10^{-12}$ | $3.9 \times 10^{-5}$ |
| 800       | 12                    | $2.6137960462 \times 10^{-12}$ | $3.4 \times 10^{-5}$ |
| 1200      | 14                    | $2.6137524215 \times 10^{-12}$ | $1.7 \times 10^{-5}$ |

**Table 12.** Against  $B = \sqrt{3}(\lambda + \mu)/6\pi\mu(\lambda + 2\mu) = 2.6137073561 \times 10^{-12}$ , at  $\omega = 6$  so that  $(k_{sa})^2 = 10^{-6}$  and the static form is the right comparison. The residual falls under refinement with no sign of a floor, so it bounds the quadrature and is not a discrepancy.

Dropping the two factors of  $i$  returns  $-1.000275 \times B$ : the control flips the sign exactly, as it must.

### 13.4 $A$ and $C$ , and the delta they carry

The same index map gives the other two with lateral derivatives alone:

$$I_{xxyy} = A, \quad I_{xxxx} = A + 2B + C, \quad (84)$$

so nothing here needs  $\partial_z$  either. What they do have, and  $B$  does not, is the Eshelby delta. It lives in the  $\delta_{ij}$  channels;  $B$  has  $i \neq j$  and so carries none, which is why its integral converges outright while these two do not. *Convergence tracking the delta's presence is a check on the setup in itself.*

### The split closes by representation, not by magnitude

Since  $\hat{g} \rightarrow C(\theta)/k^2$ , the asymptotic integrand is

$$-k_x^2 F(\mathbf{k}) \frac{C(\theta)}{k^2} = -\cos^2 \theta C(\theta) F(\mathbf{k}), \quad (85)$$

a function of *direction* only. Its  $k$ -space integral is a pure pole with no finite part, so adding it back restores exactly the divergence the subtraction was introduced to remove: **the split cannot be closed in  $k$ -space at all**, as §10 already found for the tail there. It closes by representation instead.

*The remainder stays in  $k$* , and is regular at the origin — which is the point, because  $-k_x^2$  times  $-C/k^2$  is  $+\cos^2 \theta C$ , bounded. No taper, no cutoff, no matching of one region to another: it is

independent of the box half-width to seven digits across a fourfold change, where the unsubtracted integral drifted by per cent.

*The local part goes to real space.* Writing  $h(\theta) = -\cos^2\theta C(\theta) = \sum_n h_n e^{in\theta}$ , the  $n = 0$  term is a delta contributing  $h_0 \mathbf{1}_S(0) = h_0$ . For  $n \neq 0$  the inverse transform of  $e^{in\theta}$  is a homogeneous distribution of degree  $-2$ ,

$$\check{f}_n(r, \varphi) = \frac{(-i)^n |n|}{2\pi} \frac{e^{in\varphi}}{r^2}, \quad \int_0^\infty k J_n(kr) dk = \frac{|n|}{r^2}, \quad (86)$$

whose principal-value integral over the square is elementary, the  $\ln \epsilon$  of  $\int_\epsilon^R dr/r$  cancelling because  $e^{in\varphi}$  has zero angular mean:

$$\int_S \frac{e^{in\varphi}}{r^2} d^2x = \int_0^{2\pi} e^{in\varphi} \ln R(\varphi) d\varphi, \quad R(\varphi) = \frac{a}{\max(|\cos \varphi|, |\sin \varphi|)}. \quad (87)$$

$D_4$  symmetry admits  $n = 0, \pm 4, \pm 8, \dots$ , and the measurement says the series stops at the first of them:  $n = 4$  carries the whole correction and everything above it contributes  $4 \times 10^{-9}$  relative. Taking the local term as the delta alone is wrong by 5.5%, so the dipole is load-bearing and not decorative — both are gated.

### The result, and the relations tested

|   | first order                  | closed form                  | relative             |
|---|------------------------------|------------------------------|----------------------|
| A | $-1.2205739 \times 10^{-11}$ | $-1.2201107 \times 10^{-11}$ | $3.8 \times 10^{-4}$ |
| B | $+2.6137524 \times 10^{-12}$ | $+2.6137074 \times 10^{-12}$ | $1.7 \times 10^{-5}$ |
| C | $-3.5777391 \times 10^{-12}$ | $-3.5870553 \times 10^{-12}$ | $2.6 \times 10^{-3}$ |

Because all three come out of  $\partial_3 \mathbf{q} = \mathcal{A} \mathbf{q} + \mathbf{d}$  separately, the structural relations of §9.2 become *tests* rather than definitions. On first-order numbers only, with nothing substituted from the closed forms:

$$A \text{ vs } B - \frac{1}{3\mu} : 3.8 \times 10^{-4}, \quad C \text{ vs } -(5 - 2\pi/\sqrt{3}) B : 2.6 \times 10^{-3}. \quad (88)$$

### 13.5 Two things that nearly went wrong

*The grid.* The integrand is  $4 \sin(k_x a) \sin(k_y a)$  times the kernel, a product of sines separable in the *Cartesian* components. On a polar grid it does not converge even at fixed cutoff, and that drift, read as cutoff dependence, looks exactly like a logarithmic divergence — which invites a hunt for a missing Eshelby delta. There is none. On Cartesian panels cut at the sinc half-period  $\pi/a$  the integral converges cleanly. The resolution study at *fixed* cutoff is what separates the two readings, and nothing else does.

*A near miss.* On the polar grid at  $\Lambda = 50$  the value was  $2.6147 \times 10^{-12}$  — agreeing with  $B$  to  $3.8 \times 10^{-4}$ , and entirely an artefact of where the truncation happened to land.

Finally, the normalisation factors  $\epsilon_c$  play no part:  $\epsilon$  scales column  $i$  of  $\mathbf{D}_z$  and divides row  $i$  of its inverse, so it cancels out of every projector. Dropping it removes the one step in the

construction that cancels — the step of §11 that limits the useful range in  $k$  — and lets the whole assembly batch.

## 14 Which formulation is more accurate

§12 validates the operator and the propagator but cannot compare the two formulations, because a laterally invariant problem is one the cube moment method does not address at all. The geometry that does compare them is the one this package already uses to score itself: a periodic patch of space-filling cubes, which for identical cubes *is* a homogeneous layer, and whose reflection coefficient the reflectivity recursion gives exactly. Both routes then build a single-site  $T$  for the same cube, both pass through the same lattice sum, the same contact correction and the same Foldy–Lax solve, and both are scored against the same reference. Only  $T_0$  differs.

### 14.1 Substituting at the one place both routes speak the same language

$\mathbf{T} = \Delta\mathbf{C}(\mathbf{I} - \mathcal{G}\Delta\mathbf{C})^{-1}$  is not in the package’s  $9 \times 9$  convention — the strain block differs in sign and by the engineering doubling, and  $\Delta\mathbf{C}_{\text{eff}}$  is in velocity units besides (§7.3). Mapping all of that would risk measuring a convention as though it were physics. The substitution is therefore made through the four *dimensionless* amplification factors  $A_u$ ,  $A_\theta$ ,  $A_{e,\text{diag}}$ ,  $A_{e,\text{off}}$ , which are ratios and carry no units. The package’s own mapping to renormalised contrasts and its own  $T_0$  construction are then used unchanged for both routes, so nothing distinguishes them but those four numbers. What licenses this is that the Navier route reproduces the package’s own factors to  $1.5 \times 10^{-6}$ , re-checked rather than assumed.

### 14.2 The measurement

| contrast                            | ordinary              | first order           | ratio |
|-------------------------------------|-----------------------|-----------------------|-------|
| weak, $\Delta\mu/\mu \approx 0.4\%$ | $2.08 \times 10^{-3}$ | $2.09 \times 10^{-3}$ | 1.007 |
| moderate, $\approx 4\%$             | $2.00 \times 10^{-2}$ | $1.94 \times 10^{-2}$ | 0.966 |
| strong, $\approx 13\%$              | $5.59 \times 10^{-2}$ | $5.02 \times 10^{-2}$ | 0.899 |

**Table 13.** Relative error in  $R_{pp}$  against the reflectivity reference,  $M = 8$ ,  $N_z = 16$ ,  $a = 0.5$  m,  $\omega = 60$ ,  $p = 0$ . At weak contrast the two agree, as they must — they agree at Born order — and the advantage appears where the  $O(\Delta c^2)$  difference lives.

### 14.3 Why this is not cancellation

At weak contrast the two  $T_0$  differ by  $1.8 \times 10^{-5}$  and yet both miss the reference by  $2.1 \times 10^{-3}$ : **most of the error is shared**, and belongs to the lattice rather than to  $T_0$ . A difference in a mostly-shared error is exactly the shape an accidental cancellation would take — the first-order  $O(\Delta c^2)$  term happening to offset the lattice’s own. Refining the lattice discriminates between the two readings:

| $M$ | shared error at weak contrast | ratio at strong contrast |
|-----|-------------------------------|--------------------------|
| 8   | $2.0995 \times 10^{-3}$       | 0.8987                   |
| 12  | $1.7157 \times 10^{-3}$       | 0.8766                   |
| 16  | $1.4690 \times 10^{-3}$       | 0.8570                   |

The shared error falls and the advantage *grows*. Cancellation would erode as the thing being cancelled against is removed; instead the advantage was being *masked* at coarse  $M$ . So 14.3% is a **lower bound** on the part attributable to  $T_0$ , not an estimate of it — it is still improving at the finest lattice tested.

A remark on how that check was posed, since the first version of it was wrong: it demanded that the ratio be *stable* under refinement, and it failed. Stability is not what should be expected. As the shared error goes to zero the ratio should approach the  $T_0$ -attributable value, and monotone movement towards a limit — not constancy — is the correct signature.

#### 14.4 Which channel carries it

The improvement can be attributed, because the substitution runs through four separable factors. Replacing one at a time, at the strongest contrast and the finest lattice:

| substituted              | error                   | ratio         |
|--------------------------|-------------------------|---------------|
| none (all ordinary)      | $3.7826 \times 10^{-2}$ | 1.0000        |
| $A_u$ only               | $3.7826 \times 10^{-2}$ | 1.0000        |
| $A_\theta$ only          | $3.2418 \times 10^{-2}$ | <b>0.8570</b> |
| $A_{e,\text{diag}}$ only | $3.7826 \times 10^{-2}$ | 1.0000        |
| $A_{e,\text{off}}$ only  | $3.7826 \times 10^{-2}$ | 1.0000        |
| all four                 | $3.2418 \times 10^{-2}$ | 0.8570        |

**The whole of it is in  $A_\theta$** , the trace — which is the  $A_{1g}$  channel, invariant under the full cubic group. So the advantage measured at normal incidence is the genuine  $O_h$ -invariant resummation, and the  $x_3$ - preferring part of the  $O(\Delta c^2)$  difference contributes nothing. The reason is elementary once seen: a normally incident  $P$  wave is pure compression along  $z$  and never probes a shear channel at all. At  $p = 0$  the measurement is blind to precisely the part of the formulation that singles out  $x_3$ .

#### 14.5 What the $x_3$ preference is worth, bracketed

Two facts have to be put together. First, a structural one: `effective_stiffness_voigt` takes a *single*  $\Delta\mu_{\text{off}}^*$  and writes it into all three shear slots, so the package's  $T_0$  is cubic-symmetric by construction and **cannot represent an  $e_{12}/e_{13}$  split**. The symmetry breaking was therefore excluded from §14 not by choice but for want of a slot. Second,  $SH$  incidence at  $p \neq 0$  does probe the shear channels, and there:

| $SH$ at $p = 10^{-4}$           | error in $R_{SH}$       | ratio         |
|---------------------------------|-------------------------|---------------|
| all ordinary                    | $4.7933 \times 10^{-3}$ | 1.0000        |
| first order, $e_{12}$ amplitude | $4.7933 \times 10^{-3}$ | 1.0000        |
| first order, $e_{13}$ amplitude | $1.8432 \times 10^{-3}$ | <b>0.3845</b> |

The  $e_{12}$  row reproducing the ordinary route *exactly* is §7 confirmed end to end: the two differ only in the traction channel, and the in-plane shear is untouched. The  $e_{13}$  row is the interesting one. An  $SH$  wave reflecting off a layer near normal incidence is dominated by  $\partial_z u_y$  — a  $yz$  shear, the  $z$ -involving one — and the  $x_3$ -distinguished amplitude gives an error 2.6 times smaller. The first-order route separates the two by  $7.2 \times 10^{-3}$ ; the ordinary route does not separate them at all.

*That brackets the effect rather than realising it.* Feeding the  $e_{13}$  amplitude through a cubic-symmetric  $T_0$  writes it into the  $e_{12}$  slot as well, which is also wrong; it is the other *uniform* choice, not the split.

## 14.6 The tetragonal $T_0$

What the split needs is a  $T_0$  of tetragonal rather than cubic symmetry —  $e_{12} \neq e_{13} = e_{23}$ , one extra shear parameter, which is the minimum that can carry it. The in-plane shear keeps its own amplitude and the two depth-involving shears share the other.

The whole of the correctness is the slot assignment, and it is measured rather than inherited. The package's Voigt ordering is  $(z, x, y)$  with  $1 \rightarrow z$ ,  $2 \rightarrow x$ ,  $3 \rightarrow y$ , so slot 3 is  $(x, y)$ , slot 4 is  $(z, y)$  and slot 5 is  $(z, x)$ ; the first-order gate's paper indices put  $e_{12}$  in slot 3 and  $e_{13}, e_{23}$  in slots 4 and 5. The two agree — but two docstrings agreeing is not evidence, so each slot is perturbed on its own and the  $SH$  channel is asked which one it reads:

| +10% on                   | moves the $SH$ error by |
|---------------------------|-------------------------|
| slot 3, in-plane $(x, y)$ | $9.390 \times 10^{-1}$  |
| slot 4, depth $(z, y)$    | $9.537 \times 10^0$     |
| slot 5, depth $(z, x)$    | 0                       |

which is exactly the physics of an  $SH$  wave in the  $(x, z)$  plane polarised along  $y$ :  $\partial_z u_y$  dominates,  $\partial_x u_y$  contributes near normal incidence, and a  $zx$  shear does not enter an  $SH$  problem at all. Slot 5 reading *identically* zero is the sharpest part of that.

With the slots so assigned:

| $SH$ at $p = 10^{-4}$         | error in $R_{SH}$       | ratio         |
|-------------------------------|-------------------------|---------------|
| ordinary                      | $4.7933 \times 10^{-3}$ | 1.0000        |
| first order, cubic $T_0$      | $4.7933 \times 10^{-3}$ | 1.0000        |
| first order, tetragonal $T_0$ | $1.4905 \times 10^{-3}$ | <b>0.3110</b> |

The tetragonal form collects *more* than the bracket of the previous subsection predicted —  $3.2\times$  against  $2.6\times$  — and for a reason: the uniform- $e_{13}$  substitution also wrote that amplitude into slot 3, where it is wrong and where slot 3 demonstrably contributes. Nothing is traded away for it: the normal-incidence channel is unchanged at 0.8570.

**The answer, with its scope.** On the question as posed — same geometry, same lattice machinery, same external reference, only  $T_0$  swapped — the first-order  $T$  is the more accurate of the two, and the margin grows with contrast. The part of that measured at normal incidence is entirely the  $O_h$ -invariant resummation; the  $x_3$ -preferring part is separately shown to be worth a further factor of 2.6 in the one channel that probes it, but cannot be carried by the present  $T_0$ . This is one geometry, one frequency, normal incidence, three contrasts, and an unconverged lattice. It scores the *single-site*  $T$  only, so an error the two share is invisible by construction. And it does not touch §7.5: the  $O_h$  symmetry breaking remains a genuine  $\approx 4.2 \times 10^{-4}$  debit against the first-order route in a homogeneous whole space, where there is no preferred direction for it to be right about.

## 15 The layered background

This is the payoff, and it is also where the route stops having a competitor. Everything up to here has had a closed-form whole-space  $\mathbf{G}$  standing behind it, so the moment method could be scored against something. Over a stack there is no such object, and  $\mathbf{\Gamma}$  must simply *be* the stratified propagator.

### 15.1 Nothing here is new machinery

The stratified  $6 \times 6$  already exists and is validated, and §8 has already measured that it is the Green's matrix of *this*  $\mathbf{A}$ . What it is not is a function of depth: the block-Riccati sweep is indexed by *interface*, while the moment integral needs  $\mathbf{\Gamma}(z, z')$  at continuously varying  $z$  and  $z'$ .

The tempting repair — propagate away from the nearest interface with  $\exp(\mathbf{A}\delta)$  — is the Thomson–Haskell instability reintroduced deliberately, and an earlier attempt did exactly that: a propagator-matrix traversal that reduced correctly in the no-contrast limit, to  $7.4 \times 10^{-15}$ , and was unusable,  $\text{cond}(\mathbf{Y}_{\text{up}})$  reaching  $9.0 \times 10^{15}$ , because forming  $\exp(\mathbf{A}\delta)$  forms the growing evanescent branch explicitly. It was discarded rather than patched.

The stable construction does not propagate at all. The model is *cut* at the depths wanted, with zero contrast across the cut; the sweep is unchanged, is still a Riccati recursion, and still never forms the growing branch. Cutting must then be harmless, and is: between two *fixed* physical depths, two cuts and six cuts agree to  $2.5 \times 10^{-15}$  across the propagating and evanescent regimes alike. The bridge identity, previously verified only at the interfaces of a fixed stack, is then re-measured at depths lying on no lattice and holds to a fraction of one arithmetic floor.

### 15.2 The source normalisation, which the bridge could not see

§8 tests a *ratio* of two receiver depths. The source-side operator cancels out of a ratio, so that gate says nothing about how  $\mathbf{\Gamma}$  is normalised at the source — and the normalisation is what makes it a Green's matrix rather than merely a solution. Measuring it needs both receivers carried back onto the source plane,

$$e^{-\mathbf{A}\delta} \mathbf{\Gamma}(z'+\delta, z') - e^{+\mathbf{A}\delta} \mathbf{\Gamma}(z'-\delta, z') = \mathbf{I}_6, \quad (89)$$

and (89) holds to  $7 \times 10^{-14}$ , flat in  $\delta$  over a factor of eight. So  $[\partial_z - \mathbf{A}]\mathbf{\Gamma} = \mathbf{I}_6 \delta(z - z')$  with *no* calibration constant anywhere, and the  $6 \times 6$  half of the source-convention question is closed. Reciprocity in the thesis basis,  $\mathbf{\Gamma}(-\mathbf{k}; z', z)^t = \mathbf{J}_6 \mathbf{\Gamma}(\mathbf{k}; z, z') \mathbf{J}_6$  with source and receiver exchanged, holds to  $1.2 \times 10^{-15}$ .

### 15.3 The uniform limit, which cannot be what was planned

The plan called for the uniform limit of the layered assembly to reduce to the whole space. It cannot, and the reason is structural rather than numerical: layer 0 of the stratified solver is acoustic, so a uniform *elastic* whole space is not a model the solver can be given. The ocean floor is always there and always reflects.

What replaces it is sharper, and has no free parameter. Whatever the ocean floor returns is, at *every* receiver, downgoing — it left the source upward, turned once, and is coming back down — and the half-space below radiates, so nothing else can contribute. Therefore

$$\mathbf{\Gamma}_{\text{layered}}(z, z') - \mathbf{\Gamma}_{\text{whole space}}(z - z') = \mathbf{P}_d [\cdot], \quad \mathbf{P}_d = \mathbf{D}_z \text{diag}(1, 1, 1, 0, 0, 0) \mathbf{D}_z^{-1}, \quad (90)$$

with the split taken from the column ordering of (73) and nothing classified numerically. Measured,  $\|\mathbf{P}_u \Delta\|/\|\Delta\|$  is  $2.6 \times 10^{-16}$  to  $5.2 \times 10^{-14}$  over five wavenumbers and three receiver depths, *including* receivers above the source, where the whole-space term is upgoing and the difference is still not. The difference itself is 54%–74% of  $\|\mathbf{\Gamma}\|$ , so the projection is not being satisfied by a vanishing numerator.

Beyond the critical wavenumber the returned field is evanescent over the whole round trip and the difference vanishes outright: at  $k_x = 6, 7, 9$  against  $\omega/\beta = 5.66$ ,  $\|\Delta\|/\|\mathbf{\Gamma}\|$  is  $3 \times 10^{-15}$ – $3 \times 10^{-14}$ . There the layered  $\mathbf{\Gamma}$  is the whole-space  $\mathbf{\Gamma}$ , to the arithmetic floor — which is the reduction the plan wanted, holding exactly where the boundary cannot reach and failing elsewhere only by the boundary term it should fail by.

The control matters here, because (90) is a statement about a projection and a projection can be satisfied for uninteresting reasons. Moving the contrast *below* the receiver instead sends an upgoing field back, and  $\|\mathbf{P}_u \Delta\|/\|\Delta\|$  rises from  $10^{-16}$  to  $1.9 \times 10^{-1}$ – $4.8 \times 10^{-1}$ : twelve orders, on the same geometry.

## 16 From the integral equation to invariant imbedding

*Status.* The equivalence of §16.1 is a derivation, and its pivot — that the jump is exactly  $\mathbf{I}_6$  — is measured, at  $7 \times 10^{-14}$ . The Riccati equation (96) is derived here and the march is not yet implemented in this route: the existing implementation is the sibling repository's layer recursion, which §16.6 identifies as the same object discretised. Its coefficient matrices, however, are built and gated — §16.5 to §16.8 report measurements, not expectations, including the impedance form and the wavenumber range it reaches. What remains unbuilt is the march itself and the lateral generalisation of §16.11, which is written down to fix what would have to be built.

### 16.1 The kernel is one-sided, so the integral equation splits

The propagator of §11 is one-sided by construction: downgoing projectors below the source, upgoing above,

$$\Gamma(z - z') = \begin{cases} + \sum_{i \in \text{down}} e^{\lambda_i(z-z')} \mathbf{P}_i, & z > z', \\ - \sum_{i \in \text{up}} e^{\lambda_i(z-z')} \mathbf{P}_i, & z < z', \end{cases} \quad \mathbf{P}_d + \mathbf{P}_u = \mathbf{I}_6. \quad (91)$$

The Lippmann–Schwinger equation therefore splits at the *receiver* — down from above, up from below:

$$\mathbf{q}(z) = \mathbf{q}^0(z) + \int_{-\infty}^z \Gamma_d(z - z') \Delta \mathcal{A}(z') \mathbf{q}(z') dz' + \int_z^{\infty} \Gamma_u(z - z') \Delta \mathcal{A}(z') \mathbf{q}(z') dz'. \quad (92)$$

Both limits of integration move with  $z$ , so differentiating (92) produces two Leibniz boundary terms in addition to the interior derivative:

$$\partial_z \mathbf{q} = \underbrace{\frac{\mathbf{A} \mathbf{q}^0}{\partial_z \mathbf{q}^0}}_{\partial_z \Gamma = \mathbf{A} \Gamma} + \underbrace{\frac{\mathbf{A} (\mathbf{q} - \mathbf{q}^0)}{\partial_z \Gamma = \mathbf{A} \Gamma}}_{\text{upper limit}} + \underbrace{\frac{\Gamma_d(0^+) \Delta \mathcal{A} \mathbf{q}}{\text{upper limit}}}_{\text{lower limit}} - \underbrace{\frac{\Gamma_u(0^-) \Delta \mathcal{A} \mathbf{q}}{\text{lower limit}}}_{\text{lower limit}}. \quad (93)$$

By (91) the two boundary terms are  $\mathbf{P}_d \Delta \mathcal{A} \mathbf{q}$  and  $+\mathbf{P}_u \Delta \mathcal{A} \mathbf{q}$ , and they sum to  $(\mathbf{P}_d + \mathbf{P}_u) \Delta \mathcal{A} \mathbf{q} = \mathbf{I}_6 \Delta \mathcal{A} \mathbf{q}$ . Hence

$$\boxed{\partial_z \mathbf{q} = [\mathbf{A} + \Delta \mathcal{A}(z)] \mathbf{q}} \quad (94)$$

— the integral equation and the local first-order system with the *total* operator are the same statement.

Two things are worth naming about (94).

**The jump is load-bearing, not bookkeeping.** The coefficient of  $\Delta \mathcal{A}$  in (94) is exactly one *because*  $\mathbf{P}_d + \mathbf{P}_u = \mathbf{I}_6$ . Any other normalisation of  $\Gamma$  would leave a constant standing there and the equivalence would be false. That is the quantity (89) measures on the stratified object, at  $7 \times 10^{-14}$  and flat in  $\delta$ ; §15 presents it as closing a convention question, and this is the reason it was worth closing.

**What the integral form carries that the ODE does not.** The ODE needs boundary conditions; (92) already has them, in the one-sidedness of its kernel — radiation at  $\pm\infty$  is built into which projectors appear on which side. The integral equation is the ODE *plus* its boundary conditions, which is exactly why it is the useful form to discretise.

### 16.2 One-sidedness makes the discretised operator tridiagonal

Because the kernel is one-sided, a depth discretisation of (92) couples only neighbouring depths. The thesis states this directly: in Chapter 5 the array  $(\mathbf{I} - \mathcal{S}_{\text{int}} \mathcal{E})$  — the discretised Lippmann–Schwinger operator — is *block tridiagonal* (`GstratRep.tex`, eq. `Amat`), and block-tridiagonal elimination *is* the Riccati recursion, given there as Algorithm 5.1. The Introduction names the connection: Riccati factorisation for discretely layered media is “implicit in the invariant embedding

schemes of Kennett (1984), Tromp & Snieder (1988)”, after Chin, Hedstrom & Thigpen (1984) and Ha (1985).

So the three descriptions — the split integral equation, the two-point boundary-value ODE, and the invariant imbedding — are one object seen three ways, and the sibling repository’s sweep is its discretisation.

### 16.3 The imbedding equation

Write (94) in mode coordinates,  $\mathbf{q} = \mathbf{D}_z \mathbf{a}$  with  $\mathbf{a} = (\mathbf{a}_d, \mathbf{a}_u)$  split by the column ordering of (73), and let  $\mathbf{V} = \mathbf{D}_z^{-1} \Delta \mathcal{A} \mathbf{D}_z$  carry the same block structure:

$$\mathbf{a}'_d = \Lambda_d \mathbf{a}_d + \mathbf{V}_{dd} \mathbf{a}_d + \mathbf{V}_{du} \mathbf{a}_u, \quad \mathbf{a}'_u = \Lambda_u \mathbf{a}_u + \mathbf{V}_{ud} \mathbf{a}_d + \mathbf{V}_{uu} \mathbf{a}_u. \quad (95)$$

Define  $\mathbf{R}(z)$  as the operator returning the upgoing amplitude produced by everything *below*  $z$  in response to the downgoing amplitude there,  $\mathbf{a}_u(z) = \mathbf{R}(z) \mathbf{a}_d(z)$ . Substituting into (95) and requiring the result to hold for every  $\mathbf{a}_d$  gives the matrix Riccati equation

$$\mathbf{R}' = \mathbf{V}_{ud} + \Lambda_u \mathbf{R} - \mathbf{R} \Lambda_d + \mathbf{V}_{uu} \mathbf{R} - \mathbf{R} \mathbf{V}_{dd} - \mathbf{R} \mathbf{V}_{du} \mathbf{R}, \quad \mathbf{R} = \mathbf{0} \text{ below the heterogeneity.} \quad (96)$$

The initial condition is the radiation condition, and it is literally the sibling repository’s  $\mathbf{Y}[\mathbf{M}+1] = \mathbf{0}$ : nothing comes back from the radiating half-space, so the sweep starts from zero reflection and marches up.

**Why this is the stable form.**  $\mathbf{R}$  is a reflection operator and is bounded — in an energy norm,  $\|\mathbf{R}\| \leq 1$  for a passive medium — so integrating (96) never forms a growing exponential. A propagator matrix does form one, which is the Thomson–Haskell instability and is what the discarded build of §15 ran into at  $\text{cond} = 9.0 \times 10^{15}$ . The nonlinearity  $-\mathbf{R} \mathbf{V}_{du} \mathbf{R}$  is not a complication to be removed: it is the resummation of the two-way multiple scattering, and it is what keeps the marched quantity bounded.

### 16.4 What this replaces, and what it does not

For a laterally invariant contrast, (96) supersedes the moment assembly this note has been building toward. There is no moment integral, no split-and-subtract, and — the substantive point — **no ultraviolet treatment at all**. The delta has been absorbed *analytically*, by the Leibniz boundary terms of (93), rather than subtracted numerically. Everything in §13.4 exists to handle a singularity that, in this formulation, never has to be represented.  $\mathbf{\Gamma}$  stops being an object to build and becomes an object to factorise.

That is a statement about the *background*, and it is worth being exact about how far it reaches. It does not retire the single-site  $T$ : a compact scatterer in a Foldy–Lax solve needs a  $T_0$ , and a  $T_0$  is not what a depth march produces. The two decompositions answer different questions, and the moment route remains the one that answers the second.

## 16.5 Building the coefficient matrices

Every piece of  $\mathbf{V} = \mathbf{D}_z^{-1} \Delta \mathcal{A} \mathbf{D}_z$  already exists.  $\Delta \mathcal{A}$  is the difference of two (Akdef) matrices;  $\mathbf{D}_z$  is analytic, from (Peigen), (SVeigen) and (SHeigen); and  $\mathbf{D}_z^{-1}$  comes from (74), with the balanced numerical inverse taking over where §11 says it must. The row scaling cancels out of  $\mathbf{V}$  identically —  $(\mathbf{R}\mathbf{D})^{-1}(\mathbf{R}\Delta \mathcal{A}\mathbf{R}^{-1})(\mathbf{R}\mathbf{D}) = \mathbf{D}^{-1}\Delta \mathcal{A}\mathbf{D}$  — so the two routes must agree, and they do, to  $2 \times 10^{-16}$ , which checks both.

$\Delta \mathcal{A}$  is sparse for a structural reason: every entry of (Akdef) built from  $\mathbf{k}$  alone is kinematic and cancels in the difference, leaving 12 of 36. The compliance block is exactly diagonal,  $\text{diag}(\Delta a, \Delta b, \Delta b)$ , so  $\mathbf{V}_{du}$  has a three-term source; the inertia block carries  $\Delta \rho \omega^2$  and the  $\Delta \zeta, \Delta \chi, \Delta \mu$  pair; and the two remaining blocks are nilpotent and are each other's negative transpose.

That last is not a coincidence.  $\Delta \mathcal{A}$  inherits (ATdef) from the two operators it differences, and (74) carries the relation into the mode basis:

$$\mathbf{J}_6 \mathbf{V}^t(-\mathbf{k}) + \mathbf{V}(\mathbf{k}) \mathbf{J}_6 = 0 \iff \begin{cases} \mathbf{V}_{du}(\mathbf{k}) = \mathbf{V}_{du}^t(-\mathbf{k}), \\ \mathbf{V}_{ud}(\mathbf{k}) = \mathbf{V}_{ud}^t(-\mathbf{k}), \\ \mathbf{V}_{dd}(\mathbf{k}) = -\mathbf{V}_{uu}^t(-\mathbf{k}), \end{cases} \quad (97)$$

verified to  $3 \times 10^{-14}$ . Only half the blocks are independent, and (96) preserves the symmetry of  $\mathbf{R}$  — which is reciprocity of the reflection operator, and a gate rather than an assumption.

## 16.6 Is the march stiff?

Yes, and the answer is more useful than that, because of *where* the stiffness sits. The linear part of (96) is diagonal on the entries of  $\mathbf{R}$ : entry  $(c, c')$  sees

$$\mu_{cc'} = -(\lambda_{u,c} - \lambda_{d,c'}) = -i(k_{z,c} + k_{z,c'}). \quad (98)$$

Inside the propagating window  $k_z$  is real, so (98) is purely imaginary: the march is *oscillatory*, not stiff, and the step is set by resolving the phase. Beyond it  $k_z = i|k_z|$  and  $\mu_{cc'}$  is real and negative marching upward — decaying, at a rate up to  $2|\mathbf{k}|$ . That is stiffness, and the ratio grows as  $|\mathbf{k}|\beta/\omega$ : measured 1.0, 1.3, 3.5, 10.6, 35.5, 106.4 at  $|\mathbf{k}| = 0.5, 8, 20, 60, 200, 600$  against  $\omega/\beta = 5.66$ .

But (98) is *diagonal*, so an integrating factor removes it exactly and for nothing: march  $\tilde{\mathbf{R}} = e^{-\Lambda_u z} \mathbf{R} e^{\Lambda_d z}$  and no stiff linear term remains. No implicit solver is needed, and a general-purpose stiff integrator would be the wrong tool.

This is not a new prescription. It is what the sibling repository's layer recursion already does: in a homogeneous layer  $\mathbf{V} = 0$  and the exact update is  $\mathbf{R} \leftarrow e^{\Lambda_u h} \mathbf{R} e^{-\Lambda_d h}$ , which is precisely the phase factor it applies. The recursion is an exponential integrator with the layer as its step, which is why it is stable at any layer thickness.

## 16.7 The real obstruction is non-normality

Stiffness turns out not to be the difficulty. Physics requires the spectrum of  $\mathbf{V}$  to scale with  $\Lambda \sim |\mathbf{k}|$  — a fixed material contrast cannot strengthen with lateral wavenumber — and it does:

$\max |\text{eig } \mathbf{V}|/2|\mathbf{k}|$  is flat at 0.0656 over four decades,  $|\mathbf{k}| = 2$  to 2000. The *norm* does not.  $\|\mathbf{V}\|/2|\mathbf{k}|$  runs 0.07, 1.8, 16.7, 186, 1671, 18566 over the same range, growing as  $|\mathbf{k}|^2$ . Nor is it a scaling choice: Parlett–Reinsch balancing of  $\mathbf{V}$  returns the identity, so the non-normality is structural — (97) pins  $\mathbf{V}_{du}$  and  $\mathbf{V}_{ud}$  to comparable size and leaves a diagonal rescaling nothing to do.

The mechanism is measurable and is physical. As  $|\mathbf{k}|$  grows, both  $k_{z,P}$  and  $k_{z,S}$  tend to  $i|\mathbf{k}|$ , so the  $P$  and  $S$  columns of  $\mathbf{D}_z$  become parallel:  $\cos(+P, +S)$  reaches 1.000000, and  $\text{cond}(\mathbf{D}_z)$  grows as  $|\mathbf{k}|^{2.001}$  after the symplectic row scaling has been applied. This is the elastostatic degeneracy — at short lateral scale there is no distinction between  $P$  and  $S$ , both being the same decaying static field — and unlike the unit mismatch of §11, which the row scaling does remove, it cannot be scaled away. Up- and downgoing modes are untouched:  $\cos(+P, -P)$  falls to  $2 \times 10^{-6}$ .

**The ceiling this imposes.** At  $|\mathbf{k}| = 600$ ,  $\mathbf{V}_{du}$  and  $\mathbf{V}_{ud}$  are each  $2.0 \times 10^6$  while the eigenvalues of their product are only  $79^2$  — a nine-order cancellation. The similarity  $\text{eig}(\mathbf{V}) = \text{eig}(\Delta\mathcal{A})$ , which is exact in theory, degrades accordingly:  $1.5 \times 10^{-15}$ ,  $4.2 \times 10^{-13}$ ,  $3.5 \times 10^{-8}$ , 2.0 at  $|\mathbf{k}| = 2, 20, 200, 600$ . **The mode basis is sound to  $|\mathbf{k}| \approx 200$  here, about 35 times the  $S$  wavenumber, and not beyond.** Whether that is a real limit depends on the lateral resolution the scatterer demands,  $|\mathbf{k}|_{\max} \sim \pi N/a$ , and so it is recorded as a range rather than dismissed.

**What it suggests.** Where the modes coalesce, the mode basis is the wrong basis, and the remedy is not to condition  $\mathbf{V}$  but to avoid forming it: march an impedance in the physical  $(\mathbf{u}, \boldsymbol{\tau}_3)$  variables, where  $\mathbf{D}_z$  never appears and the coefficient matrices are the blocks of  $\mathbf{A} + \Delta\mathcal{A}$  themselves, polynomial in  $\mathbf{k}$ . §16.8 builds it and finds that it does clear the ceiling.

## 16.8 The impedance form, and the ceiling cleared

*A correction. An earlier draft of the paragraph above asserted that marching an impedance “is what the existing Riccati solver and the classical imbedding codes do”. The survey that preceded this subsection found otherwise: no matrix impedance exists in this package, in the sibling repository, or in the thesis. What exists is a scalar SH impedance used as an eigenvector normalisation convention, and two Riccati sweeps — both of which march a reflection matrix in mode amplitudes. The claim has been withdrawn and the construction below is new.*

Split  $\mathbf{q} = (\mathbf{u}, \boldsymbol{\tau}_3)$  and  $\mathbf{A}$  into  $3 \times 3$  blocks, and let  $\boldsymbol{\tau}_3 = \mathbf{Y}\mathbf{u}$ . Then  $\mathbf{u}' = (\mathbf{A}_{11} + \mathbf{A}_{12}\mathbf{Y})\mathbf{u}$  and  $\boldsymbol{\tau}_3' = (\mathbf{A}_{21} + \mathbf{A}_{22}\mathbf{Y})\mathbf{u}$ , while  $\boldsymbol{\tau}_3' = \mathbf{Y}'\mathbf{u} + \mathbf{Y}\mathbf{u}'$ , so

$$\mathbf{Y}' = \mathbf{A}_{21} + \mathbf{A}_{22}\mathbf{Y} - \mathbf{Y}\mathbf{A}_{11} - \mathbf{Y}\mathbf{A}_{12}\mathbf{Y}. \quad (99)$$

The coefficients are the blocks of  $\mathbf{A}$  itself, polynomial in  $\mathbf{k}$ , with no eigenvectors anywhere. In a homogeneous medium  $\mathbf{Y}' = 0$  and (99) becomes an algebraic Riccati equation whose stabilising solution is the downgoing half-space impedance; the equation is confirmed by checking that the impedance read off  $\mathbf{D}_z$  satisfies it, to  $10^{-14}$ , where  $\mathbf{D}_z$  is still trustworthy.

**Why it behaves differently, in one line.**  $\mathbf{V}$  needs the *full* inverse  $\mathbf{D}_z^{-1}$ , which mixes up- and downgoing and is where the cancellation lives.  $\mathbf{Y} = \mathbf{T}_d \mathbf{U}_d^{-1}$  needs only the three *downgoing* columns. Measured side by side over  $|\mathbf{k}| = 2$  to 6000: the impedance residual tracks  $\text{cond}(\mathbf{U}_d)\epsilon$  linearly,

reaching  $7.2 \times 10^{-10}$ , while  $\mathbf{V}$  returns 2.000 from  $|\mathbf{k}| = 600$  on. That 2.000 is  $|-ik_z - ik_z|/|k_z|$ : a lost branch, not a lost digit.

**Refining with  $\mathbf{A}$  alone.** Newton’s method on (99) — the correction solving a Sylvester equation built from  $\mathbf{A}$  — holds the residual at  $10^{-16}$  at *every* wavenumber tested. The division of labour is the natural one:  $\mathbf{D}_z$  is analytic, so it names the right branch however imprecise it has become, and  $\mathbf{A}$  supplies the precision. Continuation in  $|\mathbf{k}|$  was tried instead of a  $\mathbf{D}_z$  start and *flipped branch* at  $\omega/\alpha$ , returning the upgoing root; the arbiter caught it at exactly 2.000. The branch is then confirmed without reference to  $\mathbf{D}_z$  at all, by requiring every eigenvalue of  $\mathbf{A}_{11} + \mathbf{A}_{12}\mathbf{Y}$  to have negative real part.

**$\mathbf{Y}$  is bounded, where  $\mathbf{V}$  was not.** This is the whole difference.  $\|\mathbf{Y}\|/\mu|\mathbf{k}|$  converges to 1.529022 and  $\text{cond}(\mathbf{Y})$  to 2.0000 across four decades — the elastostatic half-space stiffness,  $O(1)$  conditioning, no growth relative to its own spectrum. The limit is also  $\omega$ -independent, as a static limit must be: two frequencies a factor of four apart agree to  $7.9 \times 10^{-6}$  at  $|\mathbf{k}| = 6000$ , converging as  $|\mathbf{k}|^{-2}$ .

**The arbiter, and the verdict.** Self-consistency would not settle this, so  $\mathbf{Y}$  is scored against an independently known analytic quantity:  $\text{eig}(\mathbf{A}_{11} + \mathbf{A}_{12}\mathbf{Y})$  must be  $\{ik_{z,P}, ik_{z,S}, ik_{z,c}\}$  with  $k_{z,c}$  from (kzcDef) in closed form. The impedance holds it to  $3.4 \times 10^{-10}$  at  $|\mathbf{k}| = 6000$ , drifting as about  $|\mathbf{k}|^{1.6}$  from  $2 \times 10^{-13}$  at  $|\mathbf{k}| = 60$ . It is not exact. What it does not do is fail. **The usable range goes from  $|\mathbf{k}| \approx 200$  to beyond 6000 — a factor of thirty — and the failure mode changes from losing the branch to losing digits slowly.**

## 16.9 Marching it through a slab, and scoring it

§16.8 only ever solved the *algebraic* Riccati equation, in a homogeneous medium. This is the march: integrate (99) upward from the base of a slab, starting from the downgoing impedance of the background half-space — the radiation condition, and the continuum form of  $\mathbf{Y}_{M+1} = 0$  — then read a reflection matrix off the result using continuity of  $\mathbf{u}$  and  $\boldsymbol{\tau}_3$  and nothing else,

$$\mathbf{R} = -(\mathbf{T}_u - \mathbf{Y}\mathbf{U}_u)^{-1}(\mathbf{T}_d - \mathbf{Y}\mathbf{U}_d). \quad (100)$$

Two integrators are carried so that each checks the other. For a piecewise-constant medium the step is a *linear-fractional transform* and is exact,

$$\mathbf{Y}_{\text{top}} = (\mathbf{M}_{21} + \mathbf{M}_{22}\mathbf{Y}_{\text{bot}})(\mathbf{M}_{11} + \mathbf{M}_{12}\mathbf{Y}_{\text{bot}})^{-1}, \quad \mathbf{M} = e^{-\mathbf{A}h}, \quad (101)$$

so one step settles a uniform slab outright; RK4 applied to (99) directly never forms  $e^{-\mathbf{A}h}$  at all, and is the control.

**Absolute, before anything relative.** At normal incidence one step of (101) reproduces the closed form to  $1.4 \times 10^{-16}$ ,  $4.5 \times 10^{-16}$  and  $3.8 \times 10^{-15}$  at contrasts of +2%, +10% and +40% — with no discretisation in it anywhere. RK4 then converges on the same answer at ratios 11.4, 15.3, 15.9, 15.9: fourth order, confirming that (99) is what is being integrated.

**Against Kennett.** The thesis and Kennett amplitude conventions differ, and §16.10 determines exactly how; for the moment what is compared is the three *diagonal* ratios, which must be constant because a convention cannot depend on angle. Across  $\sin i = 0$  to 0.75 they are constant to  $4.6 \times 10^{-14}$ , and equal  $(-1, -1, +1)$  to  $4.8 \times 10^{-14}$ . **The SH sign does not flip where P–SV does**; the split is real, and lumping the three together is what made this gate’s first run report a spurious 2.000. The collocation route of §12 reaches the same three constants,  $(-1.000042, -1.000166, +1.000166)$  at 120 cells, by an entirely different path — so this is a cross-check between two independent routes rather than a fitted offset, and the march reaches it ten orders more sharply. The constants survive past both critical slownesses, where the march has to carry evanescence.

**The graded slab, which is the point.** A uniform slab never needed a march. Grade the contrast with depth and no closed form exists — while Kennett, cut at the *same* depths, is solving the identical piecewise-constant problem. The two must therefore agree to the arithmetic floor at finite step count, not merely converge together, and they do:  $3.6 \times 10^{-15}$ ,  $2.2 \times 10^{-15}$  and  $1.1 \times 10^{-14}$  at 4, 16 and 64 steps, at normal and oblique incidence alike. Reversing the depth profile — same layers, same thicknesses, reordered — changes  $\mathbf{R}$  by 50%, so the agreement is not the insensitivity of a march that cannot see where the contrast sits.

## 16.10 What the convention actually is

*A correction, and a instructive one. Both this section and §12 previously said the two conventions “differ by a fixed matrix”. That was asserted on the strength of the three diagonal ratios, which cannot establish it — and it is wrong twice: the relation is two-sided rather than a single matrix, and it is not fixed.*

A convention is a renormalisation of mode amplitudes, and the incident and reflected sets need not be renormalised the same way. If  $\mathbf{a}_d^{\text{th}} = \mathbf{S}_d \mathbf{a}_d^{\text{ken}}$  and  $\mathbf{a}_u^{\text{th}} = \mathbf{S}_u \mathbf{a}_u^{\text{ken}}$  then

$$\mathbf{R}_{\text{th}} = \mathbf{S}_u \mathbf{R}_{\text{ken}} \mathbf{S}_d^{-1}, \quad (102)$$

which is a *two-sided* transform. If both are diagonal, the elementwise ratio  $Q_{ij} = (\mathbf{R}_{\text{th}})_{ij}/(\mathbf{R}_{\text{ken}})_{ij}$  satisfies  $Q_{ij} = (\mathbf{S}_u)_i/(\mathbf{S}_d)_j$  — rank one in log space, so

$$Q_{PP} Q_{SS} = Q_{PS} Q_{SP} \quad (103)$$

exactly. That is a sharp test of diagonality, and the diagonal-only measurement never performed it. Measured, (103) holds to  $1.9 \times 10^{-14}$ : both renormalisations *are* diagonal.

Writing  $Q_{ij} = -s_i/s_j$  with  $s = (1, ic)$ , the off-diagonals give  $c$ , and

$$c = \sqrt{k_{z,S}/k_{z,P}}, \quad (104)$$

to  $5.6 \times 10^{-15}$  across the sweep. This is the displacement-versus-flux difference: Kennett normalises modes to unit vertical energy flux, which carries  $\sqrt{k_z}$  per mode, while the columns here are normalised by their own displacement component.  **$c$  runs from 1.293 to 2.466 over  $p = 2 \times 10^{-5}$  to  $1.95 \times 10^{-4}$  and diverges as  $k_{z,P} \rightarrow 0$  at the  $P$  critical slowness, so the convention**

**is not fixed.** What is fixed is its diagonal —  $\mathbf{S}_u = -\mathbf{S}_d$  on  $P$ - $SV$  and  $\mathbf{S}_u = +\mathbf{S}_d$  on  $SH$ , giving  $(-1, -1, +1)$  — which is all §16.9 and §12 ever saw.

Two further things are worth stating precisely. The overall scale is not observable at all: (102) is unchanged by  $\mathbf{S} \rightarrow \lambda \mathbf{S}$  on both sides, so only  $\mathbf{S}_u \mathbf{S}_d^{-1}$  and the  $P/S$  ratio within each set are determined — reporting a single convention matrix would be overclaiming a third time. And the standing signature of a displacement-versus-flux difference is that it is *invisible to any per-mode ratio test*, which is precisely the test that was being run.

### 16.11 Where the lateral coupling enters

A cube is not laterally invariant, so  $\Delta \mathcal{A}$  depends on  $x$  and  $y$  and couples lateral wavenumbers. Then  $\mathbf{V}$  in (95) is an operator in  $\mathbf{k}_\perp$  as well as in mode index, and  $\mathbf{R}$  becomes a matrix on  $(3 \text{ modes}) \times (\text{lateral grid})$ .

The structure survives exactly: (96) is unchanged in form, and the tridiagonality in depth is untouched, because one-sidedness in  $z$  is a property of the kernel and not of the contrast. What changes is that each block goes dense in  $\mathbf{k}_\perp$ , and  $\mathbf{R} \mathbf{V}_{du} \mathbf{R}$  — two dense products per step — becomes the cost. That term is also where a near-field preconditioner would act, and it is the reason such a preconditioner is effective on an imbedding scheme specifically: it conditions the quadratic term that carries the strong scattering, which is the axis a directional-sweep architecture does not touch at all.

## 17 Summary of checks

| object                          | what is established                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|---------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| $\mathcal{A}$                   | built from $c_{ijkl}$ ; matches the published isotropic blocks<br>matches the thesis $\mathbf{A}(k_x, k_y)$ under one similarity transform<br>symmetry relations, spectrum, normal incidence<br>stiffness matrices, $\mathbf{U}_{al}$ , $\nu_1$ , $\nu_2$<br>source vector, negative control                                                                                                                                                                                                                                                                                                                                                                                                                       |
| $\Delta\mathcal{A}$             | exact contrast operator; Navier body force on a generic field<br>nonlinearity in the contrast, quantified<br>augmentation: constant, reproduces the form, divergence exhibited<br>$O_h$ /trial-space structure, density channel                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
| $\Delta\mathbf{C}_{\text{eff}}$ | $\mathbf{J}_6\Delta\mathcal{A}$ symmetric; test space derived; controls<br>$9 \times 9$ symmetric, rank 9, rigid block, basis trap<br>Born agreement, $O(\Delta c^2)$ difference, localisation<br>$O_h$ decomposition, positive and negative controls                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| $\mathbf{\Gamma}$               | propagates by $\exp(\mathbf{A} dz)$ with all multiples present                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| moments                         | face integrals evaluated; $j_1$ also by solid angle<br>$\int dV/r$ in two independent closed forms<br>$B$ , and $A = B - 1/3\mu$ , $C = -(5 - 2\pi/\sqrt{3})B$ , $G$<br>Kelvin coefficients against the $h_c$ form<br>the trace identity (42), derived without the recursion<br>the tabulated route reproduces all four                                                                                                                                                                                                                                                                                                                                                                                            |
| $\mathbf{T}$                    | coupling rebuilt independently, matches the symbolic derivation<br>calibration: reproduces the existing amplification factors<br>the finite-size density–strain term, isolated and sized<br>density channel identical; $O_h$ split on the $e_{12}/e_{13}$ pair<br>Born limits agree                                                                                                                                                                                                                                                                                                                                                                                                                                |
| Schwinger                       | $\mathcal{F}[\partial_\alpha(\mathbf{1}_V c)] = ik_\alpha \mathcal{F}[\mathbf{1}_V c]$ , with a sign control<br>$\mathcal{A}$ in the paper basis $= S^{-1} \mathcal{A}_{\text{thesis}} S$ ; spectrum; $\mathcal{A}(-\mathbf{k})^\top \mathbf{J}_6 = -\mathbf{J}_6 \mathcal{A}(\mathbf{k})$ , with a control<br>$\mathbf{\Gamma}$ inverted laterally <i>is</i> the Kelvin tensor<br>the radiation reaction (72), no static counterpart<br>reciprocity (71), pointwise and z-integrated<br>the two- $\Delta\mathcal{A}$ assembly is the Galerkin double-volume moment<br>the gap to the calibrated $\mathcal{G}$ , factorised<br>the $9 \times 9$ matrix is symmetric, with a sign control and a not-vacuous control |
| thesis                          | $\mathbf{D}_z^{-1}$ from (74) is diagonal on the thesis eigenvectors, and exact after normalisation<br>$\mathbf{D}_z \mathbf{A} \mathbf{D}_z^{-1}$ reproduces (Akdef)<br>the quasi-Hamiltonian relation as the thesis states it<br>quasi-SV and quasi-SH independent at $k = 0$<br>$\mathbf{\Gamma}$ reciprocal, and the Kelvin tensor with the radiation reaction, from the analytic modes alone<br>the valid range in $k$ , recorded<br>the branch point is real, and attenuation removes it; complex $\omega$ is the complex slowness                                                                                                                                                                           |
| $A, B, C$                       | the batched eigenvectors are the scalar ones<br>(82) pins the realisation and the sign<br>$B$ from the first-order system matches the closed form<br>the two factors of $i$ are load-bearing<br>the subtracted remainder is cutoff-independent<br>the local term stops at $n = 4$ ; the dipole is load-bearing                                                                                                                                                                                                                                                                                                                                                                                                     |

**Table 14.** All checks pass. The symbolic scripts report 25/25 and 44/44, the  $T$ -matrix assembly 14/14, the Schwinger settlement 17/17, the layer comparison 7/7, the thesis spectral representation 13/13, the slab comparison 7/7, the static moments 11/11, the Riccati blocks 11/11, the impedance 9/9, its march 11/11, the layered  $\Gamma$  8/8, and the propagator gate exits clean.

## 18 What is not established

In keeping with the project’s standard that an unvalidated formula be labelled as such, the following are *open*:

1. **The  $T$ -matrix uses the collocation propagator, not the consistent one.**  $\mathbf{T}$  is assembled in §9 and calibrated exactly, but  $\mathcal{G}$  there is the *single* cube moment while the  $\mathbf{J}_6$  pairing makes the scheme Galerkin. This is now measured rather than suspected: (70) separates the gap into a geometric factor  $\int_V 1/r/D_0 = 1.2644$  and a factor  $i\omega$ , and the double moment  $D_0$  that the consistent pairing needs is extracted in §10.4. The  $i\omega$  in (70) is *not* open — see §7.3 — so what remains is the **geometric** factor alone. Until that is resolved, everything in §9 is the first-order coupling seen through a collocation propagator.
2. **The external arbiter reaches the operator, not the cube.** §12 closes this item for  $\mathcal{A}$ ,  $\Gamma$ , their composition and the resummation: a layer in a whole space reproduces an independently validated reflectivity recursion across contrast and angle. But a laterally invariant problem does not touch the augmentation, the form factors or the moments, so the quantities in §7 and §9 still compare two *formulations* with each other rather than either against truth. The  $O_h$  split remains the only partial exception there, and it is one-sided: it lower-bounds the first-order error at  $\approx 4.2 \times 10^{-4}$  and bounds the ordinary route not at all. §14 then answers the accuracy question for the cube by the one geometry in which both formulations apply and an external reference exists — a space-filling patch — and the first-order  $T$  wins there by a margin that grows with contrast. What remains genuinely open is the *scope* of that: one frequency, normal incidence, an unconverged lattice, and no measurement at all away from space-filling.
3. **Whether the resummation improves accuracy is unknown.** It is exact as a local inversion; that is not the same claim.
4.  **$O_h$  invariance was tested on a subgroup of order 3, not the full group.**
5. **The derivative count for a collocation scheme differs from the Galerkin one.** In the Galerkin pairing the transferred derivative lands on polynomials, so the third-derivative moment cannot arise; under collocation it transfers onto  $\Gamma$  and can. Do not quote the Galerkin count for a collocation build. (66) is how that transfer is carried out in practice, and it is verified; what is not yet done is the count itself.
6. **The derivative channels are verified only by symmetry.** That is a strong test — ten orders of separation from its control, and the derivative terms carry 80% of the matrix — but symmetry is a *structural* property. A bookkeeping error that is symmetric would pass it. The value of those channels has not been compared with an independent arbiter, which would need the Galerkin double-volume moments of  $\partial_i \partial_j G_{kl}$ ; those follow from the

same autocorrelation and Duffy construction as  $D_0$ , with  $\partial_i \partial_j$  moved onto the weight, but they have not been computed.

7. **The tetragonal  $T_0$  is measured but not shipped.** §14.6 builds it and shows it collects the  $x_3$  benefit, 3.2× in the channel that probes it. But it is built in the gate, not in the package: `effective_stiffness_voigt` still takes one shear parameter and writes it into all three slots. Widening it is a small backwards-compatible change, and the right moment for it is when the first-order route is wired into the package proper rather than now, when nothing upstream in the package produces two shear amplitudes.

## 19 Where the route belongs

The  $x_3$  preference of §7.5 is an artefact only because the problem posed there — an isotropic cube in an isotropic whole space — has no preferred direction. A layered background does. There the formulation’s directional structure matches the physics rather than contradicting it, the  $O_h$  argument does not apply, and the propagator is already available and already bridged to  $\mathcal{A}$  by (27).

That was written as an expectation and is now a measurement. In the stratified geometry of §14, the  $x_3$ -distinguished shear amplitude is the one the physics wants, by a factor of 2.6 in the channel that probes it (§14.4). The same feature that is demonstrably wrong in a whole space —  $\geq 4.2 \times 10^{-4}$  of error, by §7.5 — is demonstrably right once there is a stack of interfaces for it to be right about. Both statements are measured, and they are not in tension: they are about different problems.

That is also the case in which the ordinary route has no competitor: the cube moment method rests on volume moments of the whole-space Green’s tensor, and in a stratified background that object does not exist. The first-order form supplies what the moment method cannot.