

Graded Voxels in Three Dimensions

A Galerkin First-Moment Cube with a Contrast Linear in the Cell, Fourth Order in its Discretisation

T. M. Nestor

Independent researcher, Australia

29 September 2026

Abstract

Voxel schemes represent a heterogeneous elastic medium as cells of constant properties, and in a smoothly varying medium that representation caps them at second order in the cell size whatever the cell carries. In one dimension a cell that carries the Legendre moments of its field to degree p and a contrast expanded to degree r converges at order $\min(2p + 2, 2r + 2)$. We build the three-dimensional cell with $p = r = 1$: a Galerkin cube carrying the mean and first moments of the nine-component state, displacement and an independent strain, with a contrast linear in the cell. The singular integrals of its coupling blocks are derived twice, by moving derivatives onto the cells' autocorrelation and by integrating by parts onto their polynomials, and agree to 10^{-10} for representative moment pairs; the blocks reproduce the Eshelby and biharmonic sum rules, the cube's Coulomb energy and an exact subdivision identity. On a sphere whose contrast falls smoothly to zero at its surface, against its exact solution, the discretisation converges at fourth order (apparent orders 4.01 and 4.13) where uniform-field voxels converge at second. At full contrast the error is 42 to 66 times below the collocation voxel at the same cell size. Solved with an FFT matrix-vector product to 16 cells across, the full solve's apparent order rises to 3.80 at both frequencies, approaching four from below. The slow approach does not depend on the frequency, nor on the smoothness of the contrast profile; it depends on how many cells span the graded shell, and with the shell widened from four to six cells the apparent order reaches 3.92. The far-field extraction must treat the engineering shear stress of the source consistently with the coupling kernel; a factor of two there is invisible at Born order along a cell axis and appears as an error in the second-order term that no refinement removes.

1 Introduction

A voxel scheme for elastic scattering replaces a heterogeneous body by cubic cells, each carrying a single-site T -matrix and coupled through the background Green's tensor. In a layer of such cells the scheme converges at second order when each cell carries a uniform internal field, and a cell that also carries the first moment of its field converges at fourth order [1]. That paper also showed that where the medium varies within a cell, a contrast sampled as constant caps every voxel at second order, and that the rule for a cell whose field has Legendre degree p and whose

contrast has degree r is $\min(2p + 2, 2r + 2)$. All of that is one-dimensional, or three-dimensional only for a laterally uniform layer, where the lateral moments are never excited.

This paper builds the three-dimensional graded cell ($p = r = 1$) and asks whether it keeps the fourth order on a body with genuine three-dimensional structure: a sphere whose contrast falls smoothly to zero at its surface, which has an exact solution and no staircase. The cell and its single-site T -matrix are described in §2 and §3, the coupling between cells in §4, and the sphere in §5.

2 The graded voxel

A voxel is a cube $V = [-h, h]^3$ of half-width h centred at $\mathbf{x}_c$, with local coordinates $\boldsymbol{\xi} = (\mathbf{x} - \mathbf{x}_c)/h \in [-1, 1]^3$ in the axis order (z, x, y) . The field in it is the nine-component state $\boldsymbol{\psi} = (\mathbf{u}, \boldsymbol{\varepsilon})$: the displacement $\mathbf{u}$ and the engineering strain $\boldsymbol{\varepsilon}$ in Voigt order $(zz, xx, yy, 2xy, 2zy, 2zx)$. The strain is an *independent* unknown, not the gradient of $\mathbf{u}$. A basis whose strain is derived from its displacement leaves a spurious source at every face between cells, because each cell's projected displacement jumps across the face; in a layer that scheme converges at second or third order however rich the basis [1]. The quadratic 27-mode cube T -matrix is of that kind and is not considered further.

The state is expanded in the Legendre polynomials of degree at most one,

$$\boldsymbol{\psi}(\mathbf{x}) = \sum_{a=0}^3 \boldsymbol{\psi}_a L_a(\boldsymbol{\xi}), \quad L_0 = 1, \quad L_i = \xi_{i-1} \quad (i = 1, 2, 3), \quad (1)$$

36 weights per cell. The contrast operator $\boldsymbol{\Delta} = \text{diag}(\omega^2 \Delta \rho \mathbf{I}_3, \Delta \mathbf{C})$, with ω the angular frequency, $\Delta \rho$ the density contrast and $\Delta \mathbf{C}$ the 6×6 Voigt stiffness contrast, is linear in the cell,

$$\boldsymbol{\Delta}(\boldsymbol{\xi}) = \boldsymbol{\Delta}_0 + \sum_{i=1}^3 \boldsymbol{\Delta}_i \xi_{i-1}, \quad (2)$$

its coefficients the L^2 projection of the medium onto $\{1, \xi_i\}$. The source density $\boldsymbol{\Delta} \boldsymbol{\psi}$ is then a polynomial of degree at most two. Writing m_c ($c = 0, \dots, 9$) for the ten monomials of degree at most two,

$$\boldsymbol{\Delta}(\boldsymbol{\xi}) L_b(\boldsymbol{\xi}) = \sum_{c=0}^9 \mathbf{E}_{cb} m_c(\boldsymbol{\xi}), \quad (3)$$

where $\mathbf{E}_{cb}$ collects the 9×9 coefficients: the Legendre linearisation of the one-dimensional construction [1], with no new kernel object. The Lippmann–Schwinger equation $\boldsymbol{\psi} = \boldsymbol{\psi}^0 + \int \mathbf{P}(\mathbf{x} - \mathbf{x}') \boldsymbol{\Delta} \boldsymbol{\psi} d^3 x'$, with $\boldsymbol{\psi}^0$ the incident state and $\mathbf{P}$ the 9×9 point propagator $[[\mathbf{G}, \mathbf{C}], [\mathbf{H}, \mathbf{S}]]$ built from the elastodynamic Green's tensor and its derivatives, is tested with each L_a over each cell (Galerkin). Cells m and n couple through the blocks

$$\mathbf{K}_{ac}(\mathbf{R}) = \int_{V_m} \int_{V_n} L_a(\boldsymbol{\xi}) \mathbf{P}(\mathbf{x} - \mathbf{x}') m_c(\boldsymbol{\xi}') d^3 x d^3 x', \quad \mathbf{R} = \mathbf{x}_m - \mathbf{x}_n, \quad (4)$$

and the discrete system is

$$\sum_b \mathbf{M}_{ab} \psi_{m,b} - \sum_n \sum_{b,c} \mathbf{K}_{ac}(\mathbf{R}_{mn}) \mathbf{E}_{cb}^{(n)} \psi_{n,b} = \langle L_a, \psi^0 \rangle_m, \quad (5)$$

with $\mathbf{M}_{ab} = \langle L_a, L_b \rangle \mathbf{I}_9$ the Gram matrix, $\text{diag}(8h^3, 8h^3/3, 8h^3/3, 8h^3/3)$.

3 The single site

For one isolated cell the system closes on the self block $\mathbf{K}(0)$. With the source moments $\sigma_a = \langle L_a, \Delta\psi \rangle = \sum_b \mathbf{F}_{ab} \psi_b$, $\mathbf{F}_{ab} = \sum_c \langle L_a, m_c \rangle \mathbf{E}_{cb}$, the graded single site is

$$\mathbf{T}_{36} = \mathbf{F} (\mathbf{M} - \mathbf{K}(0) \mathbf{E})^{-1} \mathbf{M}, \quad (6)$$

which maps the incident field's Legendre coefficients to the source moments: force and stress means and their first moments. The first moment of the force, which a nine-component composite cannot hold, is part of the object.

Parity. Under inversion $\mathbf{x} \rightarrow -\mathbf{x}$ the displacement is odd, the strain even and ξ_i odd. With $g = 0$ (no gradient) $\mathbf{T}_{36}$ therefore splits exactly into a 15-dimensional set (strain $\times 1$, displacement $\times \xi$) and a 21-dimensional set (displacement $\times 1$, strain $\times \xi$); the coupling between them is 5×10^{-18} of the largest entry. A gradient is odd about the centre and couples the two sets linearly: $0.325 gh$ of the largest entry, for a gradient g along one axis, at $gh = 0.05, 0.1$ and 0.2 . An axis-aligned gradient leaves the symmetry of rotations by a quarter turn about its axis, which $\mathbf{T}_{36}$ satisfies to 10^{-11} .

Reciprocity. The propagator satisfies $\mathbf{W}\mathbf{P}(\mathbf{r})$ symmetric at every $\mathbf{r}$, with $\mathbf{W} = \text{diag}(\mathbf{I}_6, \frac{1}{2}\mathbf{I}_3)$, and $\mathbf{P}(-\mathbf{r}) = \mathbf{\Pi}\mathbf{P}(\mathbf{r})\mathbf{\Pi}$ with $\mathbf{\Pi} = \text{diag}(-\mathbf{I}_3, \mathbf{I}_6)$. Every block between test functions is therefore $\mathbf{W}$ -symmetric, $(\mathbf{W}\mathbf{K}_{ab})^\top = \mathbf{W}\mathbf{K}_{ab}$, and the self block obeys $\mathbf{K}_{ba}(0) = \mathbf{\Pi}\mathbf{K}_{ab}(0)\mathbf{\Pi}$ (exchange the two points of the cell). Both hold to 10^{-11} .

The Born limit. Without a gradient $\mathbf{T}_{36}$ is exact at first order in the contrast for any incident field: the source moments are $\Delta_0 \langle L_a, \psi^0 \rangle$. With a gradient the first-order source $\langle L_a, (\mathbf{g} \cdot \boldsymbol{\xi}) \psi^0 \rangle$ needs the incident field's *quadratic* moments, which the degree-one input drops. The error is $O(\Delta gh (kh)^2)$, with k the incident wavenumber: measured slopes 2.0001 and 2.0004 for $kh = 0.025 \rightarrow 0.05 \rightarrow 0.1$, where it reaches 8.9×10^{-5} at $gh = 0.2$ (these and the parity numbers: `scripts/measure_graded_voxel_site.py`). It is the single site's view of the rule $\min(2p+2, 2r+2)$ for a field of degree p and a contrast of degree r [1].

An isolated cube has an error floor that refinement in k does not remove. In the static limit a cube of contrast $\Delta\mathbf{C}$ under a uniform remote strain has the internal strain $\boldsymbol{\varepsilon}(\mathbf{x}) = \boldsymbol{\varepsilon}^0 + \int_V \partial\partial\mathbf{G}(\mathbf{x} - \mathbf{x}') \Delta\mathbf{C} \boldsymbol{\varepsilon}(\mathbf{x}') d^3x'$. The self-interaction $\mathbf{S}(\mathbf{x}) = \int_V \partial\partial\mathbf{G}(\mathbf{x} - \mathbf{x}') d^3x'$ is constant inside an ellipsoid and only there [2]; in a cube it concentrates towards the edges and

corners. A closure that keeps averages is exact through second order in the contrast, since the second-order source needs only $\int_V \mathbf{S}$, and first errs at third order, by $\Delta \mathbf{C}^3 V \text{Var}(\mathbf{S})$; a collocation closure, which uses $\mathbf{S}(0)$, errs already at second order. Neither error depends on k . It is not what a tiled medium sees: between identical cells the faces are not physical, the exact field is smooth there, and the concentration does not arise. The order of convergence is a property of the tiled medium, the closure error a property of the isolated cell; this paper measures the former (§5).

4 The coupling between cells

The point propagator. $\mathbf{P}$ is built from $\mathbf{G}_{ij} = \frac{1}{4\pi\mu} [\delta_{ij}g_S(r) + k_S^{-2}\partial_i\partial_j(g_S - g_P)(r)]$, $g_k = e^{ikr}/r$, with μ the background shear modulus and k_P, k_S the P and S wavenumbers. Every derivative of a radial function f is a combination of $F_q = (r^{-1}d/dr)^q f$. Near the origin the closed form cancels catastrophically (the k_S^{-2} difference), so for $k_S r \leq \frac{1}{2}$ the F_q are summed term by term from the power series of g_k . This kernel equals the package's closed form to 10^{-11} where both are accurate, and its static part is the Kelvin tensor to 10^{-13} . Its static (Kelvin) part is the $t = 0$ term of the δ_{ij} series and the $t = 2$ term of the derivative series, $\frac{1}{4\pi\mu}[\delta_{ij}r^{-1} + b_2\partial_i\partial_j r]$ with $b_2 = -\frac{1}{2}(1 - \beta^2/\alpha^2)$; the dynamic remainder is at most $1/r$.

The s-form. With $\mathbf{x} = \mathbf{x}_m + \mathbf{u}$ and $\mathbf{x}' = \mathbf{x}_n + \mathbf{u}'$ the kernel sees $\mathbf{R} + \mathbf{s}$, $\mathbf{s} = \mathbf{u} - \mathbf{u}' \in [-2h, 2h]^3$, and Eq. (4) becomes

$$\mathbf{K}_{ac}(\mathbf{R}) = \int W_{ac}(\mathbf{s}) \mathbf{P}(\mathbf{R} + \mathbf{s}) d^3s, \quad W_{ac}(\mathbf{s}) = \prod_i w_{e_{a,i}e_{c,i}}(s_i), \quad (7)$$

with the one-dimensional cross-correlation $w_{\alpha\beta}(s) = h \int v^\alpha (v - s/h)^\beta dv$ over $v \in [\max(-1, s/h - 1), \min(1, s/h + 1)]$: continuous, a polynomial on $[-2h, 0]$ and on $[0, 2h]$, zero at $\pm 2h$, with kinks at 0 and $\pm 2h$. For cells that do not touch the kernel is smooth on each of the eight pieces and a three-dimensional Gauss rule suffices ($8n^3$ kernel evaluations against n^6 for the double integral). It agrees with six-dimensional Gauss quadrature of Eq. (4) to 10^{-9} one cell beyond contact.

The self cell and its 26 neighbours. The static part is a sum of terms $\partial^l r^m$ ($m = -1$ with $|\iota| \leq 2$, $m = 1$ with $|\iota| \leq 4$), distributions at the origin. Moving $k = \max(0, |\iota| - (m + 1)) \leq 2$ derivatives onto W leaves a kernel at most $1/r$: $\int W \partial^l f = (-1)^k \int (\partial^{\text{moved}} W) \partial^{\text{rest}} f$. Two derivatives on one axis turn the kinks of w into plane deltas $J_b \delta(s - bh)$, $b \in \{-2, 0, 2\}$, J_b the jump of w' at b . Every piece is then a box or a rectangle whose corners lie on the breakpoints, so the singular point $\mathbf{s} = -\mathbf{R}$ is either a vertex, integrated with Duffy's pyramids, or at distance $2h$ or more. The Eshelby delta is not added by hand; it is carried by the plane deltas. The construction reproduces the Eshelby sum rule $\sum_p \iint_{V \times V} \partial_p \partial_p r^{-1} = -4\pi V$ and the cube's Coulomb self-energy

$$\iint_{V \times V} \frac{d^3x d^3x'}{|\mathbf{x} - \mathbf{x}'|} = 2 \left[\frac{1}{5}(1 + \sqrt{2} - 2\sqrt{3}) - \frac{\pi}{3} + \ln((1 + \sqrt{2})(2 + \sqrt{3})) \right] (2h)^5 \quad (8)$$

to 10^{-11} , converges in the Gauss order to 10^{-11} , and satisfies the reciprocity and exchange laws of §3. The biharmonic rule $\sum_{p,q} \partial_p \partial_p \partial_q \partial_q r = -8\pi\delta$ holds to the same accuracy, which checks the delta carried by the r terms. And every block equals the sum of the 64 blocks between the two cells' half-size sub-cells, with the polynomials re-expressed on the sub-cells, to 2×10^{-15} for all (a, c) at the self cell, its touching neighbours and beyond (`scripts/check_graded_voxel_subdivision.py`).

An independent construction. Integration by parts can move the derivatives onto the cells' polynomials instead of onto W : one derivative onto each cell, each move either differentiating the polynomial or producing a face term, which leaves $1/r$ kernels over cell×cell, face×cell and face×face, all absolutely convergent. Built in Mathematica on its own (`GradedVoxel_WeakKernel.wl`), this reproduces the sum rule and the Coulomb energy, and, as the source cell shrinks to a point, the distributional single average of the moment engine [1]; all to 10^{-14} or better.

5 The smoothly graded sphere

The body and its exact solution. A sphere of radius $a = 10$ m has a homogeneous core $r < a/2$ with the contrast $\Delta\lambda = 2$ GPa, $\Delta\mu = 1$ GPa, $\Delta\rho = 100$ kg m⁻³ in a background $\alpha = 5$ km s⁻¹, $\beta = 3$ km s⁻¹, $\rho = 2500$ kg m⁻³. Across the shell $a/2 < r < a$ the contrast is multiplied by the smoothstep $10x^3 - 15x^4 + 6x^5$, $x = (a-r)/(a/2)$, so it vanishes like $(a-r)^3$ at the surface: there is no staircase. Its per-order T -matrices are exact (Takeuchi–Saito integration, two implementations agreeing to 2×10^{-11}) [1]. The incident wave is a P wave along z , at $k_s a = 0.5$ and 1.

Arms and measurement. At the same cell size, on a grid of n^3 cells across the bounding cube: the collocation uniform-field voxel with each cell at the profile's centre value (T_9), which keeps the cells whose centre lies inside the sphere; the Galerkin uniform-field voxel ($p = r = 0$); and the graded first-moment voxel ($p = r = 1$). A Galerkin cell carries the L^2 projection of its contrast (six-point Gauss rule per axis), so the Galerkin arms keep every cell overlapping the sphere (408 cells at $n = 8$, against 280): keeping only cells whose centre lies inside drops contrast from partly filled shell cells, and the uniform Galerkin voxel then converges at order 1.07 instead of 2. The error is the largest far-field difference from the exact solution over nine angles, relative to its peak, at 5×10^8 radii; at 5×10^4 radii the near-field terms that the asymptotic far field omits are 1.4×10^{-4} , larger than the graded voxel's own error.

The solve. A block $\mathbf{K}_{ac}(\mathbf{o})$ depends only on the offset $\mathbf{o}$ between two cells of the grid. The background is isotropic and the cube is mapped to itself by its 48 signed permutations $\mathbf{Q}$, so $\mathbf{P}(\mathbf{Q}\mathbf{r}) = \mathbf{S}\mathbf{P}(\mathbf{r})\mathbf{S}^\top$ with $\mathbf{S} = \text{diag}(\mathbf{Q}, \mathbf{R}_6(\mathbf{Q}))$, where $\mathbf{R}_6$ is the representation of $\mathbf{e} \mapsto \mathbf{Q}\mathbf{e}\mathbf{Q}^\top$ on Voigt vectors (for a signed permutation it is the same for tensor and engineering components). With the moments of the test and source polynomials transforming as $L_a(\mathbf{Q}\boldsymbol{\eta}) = \sum_b T_{ab} L_b(\boldsymbol{\eta})$ and $m_c(\mathbf{Q}\boldsymbol{\eta}) = \sum_d U_{cd} m_d(\boldsymbol{\eta})$,

$$\mathbf{K}_{ac}(\mathbf{Q}\mathbf{o}) = \sum_{b,d} T_{ab} U_{cd} \mathbf{S} \mathbf{K}_{bd}(\mathbf{o}) \mathbf{S}^\top, \quad (9)$$

so one block is computed per orbit of offsets and the others follow exactly (to 10^{-11} against direct computation). The coupling is then a convolution over the grid, applied by FFT on a $(2n - 1)^3$ circular embedding, and the system is solved by GMRES preconditioned with each cell's own block. On the grids where both run, this solve and the dense one give far-field errors that agree to 5×10^{-9} relative (`cubic_scattering/graded_voxel/fft.py`); at $n = 16$ (99,936 unknowns) it takes 105 s and 3.9 GB.

The discretisation alone. At a contrast 10^{-6} of the above, multiple scattering is negligible and the error is that of the discretisation (Table 1). The graded voxel converges at fourth order, with apparent orders 4.01 and 4.13 on the finest pair; both uniform-field voxels are of second order. At $n = 4$ a cell is as wide as the graded shell and the orders there are not asymptotic.

n	$k_s a = 0.5$			$k_s a = 1$		
	T_9	$p = r = 0$	$p = r = 1$	T_9	$p = r = 0$	$p = r = 1$
4	1.04×10^{-1}	8.85×10^{-3}	7.53×10^{-5}	1.16×10^{-1}	3.45×10^{-2}	2.32×10^{-4}
6	4.71×10^{-3}	3.61×10^{-3}	5.03×10^{-6}	8.26×10^{-3}	1.42×10^{-2}	9.14×10^{-5}
8	2.71×10^{-3}	2.06×10^{-3}	1.59×10^{-6}	5.17×10^{-3}	8.15×10^{-3}	2.78×10^{-5}
order $6 \rightarrow 8$	1.92	1.96	4.01	1.63	1.93	4.13

Table 1. Weak contrast (10^{-6}): the discretisation error of the far field.

Full contrast. With multiple scattering (Table 4) the graded voxel's error falls to 6.3×10^{-5} and 7.6×10^{-5} at $n = 8, 42$ and 66 times below the collocation voxel at the same cell size; at $k_s a = 1$ it is below the collocation voxel at $n = 24$ (4.2×10^{-4} with 64,872 unknowns, against 7.6×10^{-5} with 14,688). Its apparent order between $n = 6$ and 8 is 3.35 and 3.40: below four, and rising. The part of the error that multiple scattering adds is carried by the second-order Born term T_2 (one application of the coupling). Computed on its own, which needs no solve and so reaches finer grids, its apparent order at $k_s a = 0.5$ rises through 2.85, 3.40, 3.52 and 3.65 from $n = 4$ to 12 (errors 4.1×10^{-2} to 1.1×10^{-3} of T_2 ; `scripts/measure_graded_voxel_t2.py`). At $k_s a = 1$ the sequence is almost the same: apparent orders 2.88, 3.42, 3.54, 3.67 (errors 4.2×10^{-2} to 1.1×10^{-3}). The slow approach does not depend on the frequency. The Born term contributes little: its error, 7.5×10^{-5} of T_1 at $n = 4$ and 1.3×10^{-7} at $n = 16$, is below T_2 's share of the full error at every n .

The full solve, taken to $n = 16$ (Table 3), follows the same law. Its apparent order rises through 3.54, 3.67, 3.75 and 3.80 from $n = 8$ to 16 at $k_s a = 0.5$, and through 3.56, 3.67, 3.75, 3.80 at $k_s a = 1$, within 0.02 of the apparent orders of T_2 on the same grids. The correction to fourth order is large, but its own order is not fixed by these grids: on the even grids from $n = 10$, with $h = 1/n$, $Ah^4 + Bh^5$ and $Ah^4 + Ch^6$ fit T_2 at $k_s a = 0.5$ to 0.21% and 0.23%, and odd n , for which the sphere's centre falls at a cell centre rather than a vertex, scatter by 1 to 2% about the curve of the even ones.

What sets the approach. Three causes are separable (Table 2; `scripts/measure_graded_voxel_resolution.py`). The profile $s(x)$ is only C^2 : its third derivative jumps at $r = a/2$ and

at $r = a$. Replaced by the C^∞ step $s(x) = e^{-1/x}/(e^{-1/x} + e^{-1/(1-x)})$, whose derivatives all vanish at both ends, the approach is not faster but slower (apparent orders 3.31 to 3.77, irregular), so the jumps are not its cause; that step is steeper, with a maximum slope of 2 against 1.875 and much larger higher derivatives. Widening the shell instead, with the core at $a/4$, so that $n = 16$ puts six cells across it rather than four, speeds the approach for both profiles: the apparent order reaches 3.92 with s and 3.88 with the C^∞ step. The full solve follows T_2 throughout. The approach to fourth order is therefore governed by how finely the cells resolve the variation of the contrast; the data are consistent with an expansion in the ratio of the cell to the profile's length scale whose higher terms are still large when four cells span the shell, though they do not determine its terms.

n	core $a/2$ (4 cells across the shell at $n = 16$)		core $a/4$ (6 cells)	
	C^2 step	C^∞ step	C^2 step	C^∞ step
8	6.28×10^{-5}	8.33×10^{-5}	2.91×10^{-5}	4.01×10^{-5}
10	2.85×10^{-5}	3.98×10^{-5}	1.26×10^{-5}	1.78×10^{-5}
12	1.46×10^{-5}	2.10×10^{-5}	6.25×10^{-6}	9.38×10^{-6}
14	8.18×10^{-6}	1.18×10^{-5}	3.44×10^{-6}	5.26×10^{-6}
16	4.93×10^{-6}	7.35×10^{-6}	2.04×10^{-6}	3.13×10^{-6}
order $14 \rightarrow 16$	3.80	3.52	3.92	3.88

Table 2. Full contrast, $k_S a = 0.5$: the graded voxel's far-field error for two shell profiles and two shell widths.

n	unknowns	$k_S a = 0.5$		$k_S a = 1$	
		error	order	error	order
4	2,304	5.47×10^{-4}		6.34×10^{-4}	
6	7,488	1.65×10^{-4}	2.96	2.02×10^{-4}	2.82
8	14,688	6.28×10^{-5}	3.35	7.60×10^{-5}	3.40
10	27,648	2.85×10^{-5}	3.54	3.43×10^{-5}	3.56
12	46,080	1.46×10^{-5}	3.67	1.76×10^{-5}	3.67
14	68,832	8.18×10^{-6}	3.75	9.85×10^{-6}	3.75
16	99,936	4.93×10^{-6}	3.80	5.93×10^{-6}	3.80

Table 3. Full contrast: the graded voxel ($p = r = 1$) solved by FFT to $n = 16$; the order is the apparent order from the previous row.

n	$k_S a = 0.5$			$k_S a = 1$		
	T_9	$p = r = 0$	$p = r = 1$	T_9	$p = r = 0$	$p = r = 1$
4	1.03×10^{-1}	7.70×10^{-3}	5.47×10^{-4}	1.15×10^{-1}	3.24×10^{-2}	6.34×10^{-4}
6	4.66×10^{-3}	3.04×10^{-3}	1.65×10^{-4}	8.15×10^{-3}	1.32×10^{-2}	2.02×10^{-4}
8	2.64×10^{-3}	1.68×10^{-3}	6.28×10^{-5}	5.01×10^{-3}	7.52×10^{-3}	7.60×10^{-5}
order $6 \rightarrow 8$	1.98	2.05	3.35	1.70	1.95	3.40

Table 4. Full contrast: the far-field error of the three voxels.

A convention the Born limit cannot see. The shear entries of the source are engineering stress $2\sigma_{pq}$: the contrast operator maps the engineering strain γ_{pq} to $2\Delta\mu\gamma_{pq}$, and the kernel's




Figure 1. Far-field error against the number of cells across the sphere, for the collocation and Galerkin uniform-field voxels and the graded first-moment voxel; dashed guides show h^2 and h^4 . Top: weak contrast (the discretisation alone). Bottom: full contrast, with the graded voxel solved by FFT to $n = 16$ (Table 3).

shear columns carry a factor $\frac{1}{2}$. The far field must take the tensor stress σ_{pq} from them, half the entry. Radiating the entry itself doubles the shear stress, and a P wave along a cell axis has no shear stress at Born order, so the weak-contrast test cannot see it. It appears at second order, through the density–shear cross terms, as an error in T_2 that does not fall with refinement (4.65% and 4.05% at $n = 4$ and 6, against 4.22% and 1.31% when the far field agrees with the kernel for every source component).

6 Discussion

What three dimensions change. Two things differ from the layer. The lateral first moments are now excited, and they couple through the near field of neighbouring cells; and the discrete source, projected cell by cell, jumps by $O(h^2)$ across every face between cells, where in the layer

only the faces normal to the stratification exist. The Born term does not see either: its error is the quadrature error of the source, and it is of fourth order. The multiple-scattering term sees both, but neither costs an order: its apparent order rises towards four, and the full solve follows it (§5). The approach is slow where few cells span the variation of the contrast, and does not depend on the frequency or on the smoothness of the profile. Why the multiple-scattering term is so much more sensitive to that resolution than the Born term, whose relative error is three orders of magnitude smaller, has not been established.

Galerkin and collocation. With the lattice couplings then available, a uniform-field Galerkin cell was measured several times worse than the collocation cell [1]. Here, with every block computed exactly, the uniform-field Galerkin voxel is comparable to the collocation voxel at second order, and the graded Galerkin voxel is far better. The double average is the route to fourth order when the coupling is exact.

Cost. The graded cell has 36 unknowns against 9. On the sphere at full contrast it reaches 7.6×10^{-5} with 14,688 unknowns; the collocation voxel reaches only 4.2×10^{-4} with 64,872, and would need more to reach 7.6×10^{-5} ; at $n = 16$ the graded voxel reaches 5.9×10^{-6} with 99,936. The blocks between non-touching cells cost a three-dimensional Gauss rule over the eight pieces of the autocorrelation; the self block and its 26 neighbours are computed once per cell size.

7 Conclusions

A Galerkin cube that carries the mean and first moments of displacement and an independent strain, with a contrast linear in the cell, converges at fourth order in its discretisation (at Born order) on a three-dimensional body with an exact solution, where cells of constant properties converge at second. At full contrast it is between one and two orders of magnitude more accurate than a uniform-field voxel at the same cell size. Taken by an FFT solve to 16 cells across, the full solve's apparent order rises to 3.80, and to 3.92 when six cells rather than four span the graded shell: it approaches fourth order at a rate set by the resolution of the contrast profile, not by the frequency or the profile's smoothness. The singular integrals of its coupling blocks follow from two independent constructions that agree to 10^{-10} for representative moment pairs. The far field of such a scheme must take the source's engineering shear stress as the coupling kernel does; the Born limit along a cell axis cannot reveal a mismatch there.

References

- [1] T. M. Nestor. The continuum limit of a space-filling layer of cubes: A tiling identity, an exact lattice kernel, and the closed-form error of the discrete layer. Submitted, 2026.
- [2] J. D. Eshelby. The determination of the elastic field of an ellipsoidal inclusion, and related problems. *Proc. R. Soc. Lond. A*, 241:376–396, 1957. doi: 10.1098/rspa.1957.0133.